

# **The R Reference Manual**

## **Base Package**

### **Volume 2**

The R Development Core Team

Version 1.8.1

A catalogue record for this book is available from the British Library.

First printing, January 2004 (24/1/2004).

Published by Network Theory Limited.

15 Royal Park  
Bristol  
BS8 3AL  
United Kingdom

Email: [info@network-theory.co.uk](mailto:info@network-theory.co.uk)

ISBN 0-9546120-1-9

Further information about this book is available from  
<http://www.network-theory.co.uk/R/reference/>

Copyright (C) 1999–2003 R Development Core Team.

Permission is granted to make and distribute verbatim copies of this manual provided the copyright notice and this permission notice are preserved on all copies.

Permission is granted to copy and distribute modified versions of this manual under the conditions for verbatim copying, provided that the entire resulting derived work is distributed under the terms of a permission notice identical to this one.

Permission is granted to copy and distribute translations of this manual into another language, under the above conditions for modified versions, except that this permission notice may be stated in a translation approved by the R Development Core Team.

# Contents

<b>Publisher's Preface</b>	<b>1</b>
<b>1 Base package — graphics</b>	<b>5</b>
abline . . . . .	6
arrows . . . . .	8
assocplot . . . . .	10
axis . . . . .	12
axis.POSIXct . . . . .	15
axTicks . . . . .	17
barplot . . . . .	19
box . . . . .	23
boxplot . . . . .	24
boxplot.stats . . . . .	28
bxp . . . . .	30
chull . . . . .	33
col2rgb . . . . .	35
colors . . . . .	37
contour . . . . .	38
coplot . . . . .	42
curve . . . . .	46
dev.xxx . . . . .	48
dev2 . . . . .	50
dev2bitmap . . . . .	53
Devices . . . . .	55
dotchart . . . . .	57
filled.contour . . . . .	59
fourfoldplot . . . . .	62
frame . . . . .	65
Gnome . . . . .	66
gray . . . . .	68
grid . . . . .	69
gtk . . . . .	71

Hershey . . . . .	72
hist . . . . .	76
hsv . . . . .	80
identify . . . . .	82
image . . . . .	84
interaction.plot . . . . .	87
Japanese . . . . .	90
jitter . . . . .	91
layout . . . . .	93
legend . . . . .	96
lines . . . . .	102
locator . . . . .	104
matplot . . . . .	106
mosaicplot . . . . .	110
mtext . . . . .	114
n2mfrow . . . . .	117
pairs . . . . .	118
palette . . . . .	121
Palettes . . . . .	123
panel.smooth . . . . .	125
par . . . . .	126
pdf . . . . .	135
persp . . . . .	137
pictex . . . . .	141
pie . . . . .	143
plot . . . . .	145
plot.data.frame . . . . .	147
plot.default . . . . .	148
plot.density . . . . .	152
plot.design . . . . .	153
plot.factor . . . . .	156
plot.formula . . . . .	157
plot.histogram . . . . .	159
plot.lm . . . . .	161
plot.table . . . . .	164
plot.ts . . . . .	166
plot.window . . . . .	168
plot.xy . . . . .	170
plotmath . . . . .	171
png . . . . .	175
points . . . . .	178
polygon . . . . .	181
postscript . . . . .	184

ppoints . . . . .	190
preplot . . . . .	191
pretty . . . . .	192
qqnorm . . . . .	194
quartz . . . . .	196
recordPlot . . . . .	197
rect . . . . .	198
rgb . . . . .	201
rug . . . . .	202
screen . . . . .	204
segments . . . . .	207
stars . . . . .	209
stripchart . . . . .	213
strwidth . . . . .	215
sunflowerplot . . . . .	217
symbols . . . . .	220
termplot . . . . .	223
text . . . . .	225
title . . . . .	228
units . . . . .	230
x11 . . . . .	231
xfig . . . . .	233
xy.coords . . . . .	235
xyz.coords . . . . .	237
<b>2 Base package — math</b>	<b>239</b>
abs . . . . .	240
all.equal . . . . .	241
approxfun . . . . .	243
Arithmetic . . . . .	246
backsolve . . . . .	248
Bessel . . . . .	250
chol . . . . .	253
chol2inv . . . . .	256
colSums . . . . .	258
convolve . . . . .	260
crossprod . . . . .	262
cumsum . . . . .	263
deriv . . . . .	264
eigen . . . . .	267
Extremes . . . . .	270
fft . . . . .	272
findInterval . . . . .	274

gl . . . . .	276
Hyperbolic . . . . .	277
integrate . . . . .	278
kappa . . . . .	281
log . . . . .	283
matmult . . . . .	285
matrix . . . . .	286
nextn . . . . .	288
poly . . . . .	289
polyroot . . . . .	291
prod . . . . .	293
qr . . . . .	294
QR.Auxiliaries . . . . .	297
range . . . . .	299
Round . . . . .	301
sign . . . . .	303
solve . . . . .	304
sort . . . . .	306
Special . . . . .	309
splinefun . . . . .	311
sum . . . . .	314
svd . . . . .	315
tabulate . . . . .	317
Trig . . . . .	318
<b>3 Base package — distributions and random numbers</b>	<b>319</b>
bandwidth . . . . .	320
Beta . . . . .	322
Binomial . . . . .	324
birthday . . . . .	326
Cauchy . . . . .	328
Chisquare . . . . .	330
density . . . . .	333
Exponential . . . . .	338
FDist . . . . .	340
GammaDist . . . . .	342
Geometric . . . . .	344
Hypergeometric . . . . .	346
Logistic . . . . .	348
Lognormal . . . . .	350
Multinomial . . . . .	352
NegBinomial . . . . .	354
Normal . . . . .	357

Poisson . . . . .	359
r2dtable . . . . .	361
Random . . . . .	363
Random.user . . . . .	368
sample . . . . .	370
SignRank . . . . .	372
TDist . . . . .	374
Tukey . . . . .	376
Uniform . . . . .	378
Weibull . . . . .	380
Wilcoxon . . . . .	382
<b>4 Base package — models</b>	<b>385</b>
add1 . . . . .	386
AIC . . . . .	389
alias . . . . .	391
anova . . . . .	394
anova.glm . . . . .	395
anova.lm . . . . .	397
aov . . . . .	399
AsIs . . . . .	401
C . . . . .	403
case/variable.names . . . . .	405
coef . . . . .	407
confint . . . . .	408
constrOptim . . . . .	410
contrast . . . . .	413
contrasts . . . . .	415
deviance . . . . .	417
df.residual . . . . .	418
dummy.coef . . . . .	419
eff.aovlist . . . . .	421
effects . . . . .	423
expand.grid . . . . .	425
expand.model.frame . . . . .	426
extractAIC . . . . .	428
factor.scope . . . . .	430
family . . . . .	432
fitted . . . . .	435
formula . . . . .	436
glm . . . . .	439
glm.control . . . . .	445
glm.summaries . . . . .	447

influence.measures . . . . .	449
is.empty.model . . . . .	453
labels . . . . .	454
lm . . . . .	455
lm.fit . . . . .	459
lm.influence . . . . .	461
lm.summaries . . . . .	464
logLik . . . . .	466
logLik.glm . . . . .	468
logLik.lm . . . . .	469
loglin . . . . .	471
ls.diag . . . . .	474
ls.print . . . . .	476
lsfit . . . . .	477
make.link . . . . .	479
makepredictcall . . . . .	480
manova . . . . .	482
model.extract . . . . .	483
model.frame . . . . .	485
model.matrix . . . . .	487
model.tables . . . . .	489
naprint . . . . .	491
naresid . . . . .	492
nlm . . . . .	493
offset . . . . .	497
optim . . . . .	498
optimize . . . . .	505
power . . . . .	508
predict.glm . . . . .	509
predict.lm . . . . .	511
profile . . . . .	513
proj . . . . .	514
relevel . . . . .	517
replications . . . . .	518
residuals . . . . .	520
se.contrast . . . . .	521
stat.anova . . . . .	523
step . . . . .	525
summary.aov . . . . .	528
summary.glm . . . . .	530
summary.lm . . . . .	533
summary.manova . . . . .	536
terms . . . . .	538



terms.formula . . . . .	539
terms.object . . . . .	541
TukeyHSD . . . . .	543
uniroot . . . . .	545
update . . . . .	547
update.formula . . . . .	549
vcov . . . . .	550
weighted.residuals . . . . .	551
<b>5 Base package — dates, time and time-series</b>	<b>553</b>
as.POSIX* . . . . .	554
cut.POSIXt . . . . .	556
DateTimeClasses . . . . .	558
diff . . . . .	561
difftime . . . . .	563
hist.POSIXt . . . . .	565
print.ts . . . . .	567
rep . . . . .	568
round.POSIXt . . . . .	571
seq.POSIXt . . . . .	572
start . . . . .	574
strptime . . . . .	575
Sys.time . . . . .	579
time . . . . .	580
ts . . . . .	582
ts-methods . . . . .	585
tsp . . . . .	586
weekdays . . . . .	587
window . . . . .	589
<b>6 Base package — datasets</b>	<b>591</b>
airmiles . . . . .	592
airquality . . . . .	593
anscombe . . . . .	595
attenu . . . . .	597
attitude . . . . .	599
cars . . . . .	601
chickwts . . . . .	603
co2 . . . . .	604
data . . . . .	605
discoveries . . . . .	608
esoph . . . . .	609
euro . . . . .	611
eurodist . . . . .	613

faithful . . . . .	614
Formaldehyde . . . . .	616
freeny . . . . .	617
HairEyeColor . . . . .	618
infert . . . . .	620
InsectSprays . . . . .	622
iris . . . . .	623
islands . . . . .	625
LifeCycleSavings . . . . .	626
longley . . . . .	628
morley . . . . .	630
mtcars . . . . .	632
nhtemp . . . . .	633
OrchardSprays . . . . .	634
phones . . . . .	636
PlantGrowth . . . . .	637
precip . . . . .	638
presidents . . . . .	639
pressure . . . . .	640
quakes . . . . .	641
randu . . . . .	642
rivers . . . . .	643
sleep . . . . .	644
stackloss . . . . .	645
state . . . . .	647
sunspots . . . . .	649
swiss . . . . .	650
Titanic . . . . .	652
ToothGrowth . . . . .	654
trees . . . . .	655
UCBAdmissions . . . . .	656
USArrests . . . . .	658
USJudgeRatings . . . . .	659
USPersonalExpenditure . . . . .	660
uspop . . . . .	661
VADeaths . . . . .	662
volcano . . . . .	664
warpbreaks . . . . .	665
women . . . . .	667

## Publisher's Preface

This reference manual documents the use of R, an environment for statistical computing and graphics.

R is *free software*. The term “free software” refers to your freedom to run, copy, distribute, study, change and improve the software. With R you have all these freedoms.

R is part of the GNU Project. The GNU Project was launched in 1984 to develop a complete Unix-like operating system which is free software: the GNU system. It was conceived as a way of bringing back the cooperative spirit that prevailed in the computing community in earlier days, by removing the obstacles to cooperation imposed by the owners of proprietary software.

You can support the GNU Project by becoming an associate member of the Free Software Foundation. The Free Software Foundation is a tax-exempt charity dedicated to promoting computer users' right to use, study, copy, modify, and redistribute computer programs. It also helps to spread awareness of the ethical and political issues of freedom in the use of software. For more information visit the website [www.fsf.org](http://www.fsf.org).

The development of R itself is guided by the R Foundation, a not for profit organization working in the public interest. Individuals and organisations using R can support its continued development by becoming members of the R Foundation. Further information is available at the website [www.r-project.org](http://www.r-project.org).

Brian Gough  
Publisher  
November 2003



# Introduction

This is the second volume of the R Reference Manual. This volume documents the commands for graphics, mathematical functions, random distributions, models, date-time calculations, time-series and example datasets in the *base* package of R. For ease of use, the commands have been grouped into chapters by topic, and then sorted alphabetically within each chapter.

The documentation for the other base package commands can be found in the first volume of this series. The first volume documents the core commands: programming constructs, fundamental numerical routines and system functions. The base package commands are automatically available when the R environment is started. Documentation for additional packages is available in further volumes of this series.

R is available from many commercial distributors, including the Free Software Foundation on their source code CD-ROMs. Information about R itself can be found online at [www.r-project.org](http://www.r-project.org).

To start the program once it is installed simply use the command `R` on Unix-like systems,

```
$ R
R : Copyright 2003, The R Development Core Team
Type 'demo()' for some demos, 'help()' for on-line help,
or 'help.start()' for a HTML browser interface to help.
Type 'q()' to quit R.
>
```

Commands can then be typed at the R prompt (`>`). The commands given in this manual should generally be entered with arguments in parentheses, e.g. `help()`, as shown in the examples. Typing the name of the function without parentheses, such as `help`, displays its internal definition and does not execute the command.

To obtain online help for any command type `help(command)`. A tutorial for new users of R is available in the book “*An Introduction to R*” (ISBN 0-9541617-4-2).



# Chapter 1

## Base package — graphics

---

**abline**     *Add a Straight Line to a Plot*


---

**Description**

This function adds one or more straight lines through the current plot.

**Usage**

```
abline(a, b, untf = FALSE, ...)
abline(h=, untf = FALSE, ...)
abline(v=, untf = FALSE, ...)
abline(coef=, untf = FALSE, ...)
abline(reg=, untf = FALSE, ...)
```

**Arguments**

<b>a,b</b>	the intercept and slope.
<b>untf</b>	logical asking to <i>untransform</i> . See Details.
<b>h</b>	the y-value for a horizontal line.
<b>v</b>	the x-value for a vertical line.
<b>coef</b>	a vector of length two giving the intercept and slope.
<b>reg</b>	an object with a <b>coef</b> component. See Details.
<b>...</b>	graphical parameters.

**Details**

The first form specifies the line in intercept/slope form (alternatively **a** can be specified on its own and is taken to contain the slope and intercept in vector form).

The **h=** and **v=** forms draw horizontal and vertical lines at the specified coordinates.

The **coef** form specifies the line by a vector containing the slope and intercept.

**reg** is a regression object which contains **reg\$coef**. If it is of length 1 then the value is taken to be the slope of a line through the origin, otherwise, the first 2 values are taken to be the intercept and slope.

If **untf** is true, and one or both axes are log-transformed, then a curve is drawn corresponding to a line in original coordinates, otherwise a line



is drawn in the transformed coordinate system. The `h` and `v` parameters always refer to original coordinates.

The graphical parameters `col` and `lty` can be specified as arguments to `abline`; see `par` for details.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`lines` and `segments` for connected and arbitrary lines given by their *endpoints*. `par`.

## Examples

```
data(cars)
z <- lm(dist ~ speed, data = cars)
plot(cars)
abline(z)
```

---

**arrows**     *Add Arrows to a Plot*

---

**Description**

Draw arrows between pairs of points.

**Usage**

```
arrows(x0, y0, x1, y1, length = 0.25, angle = 30, code = 2,  
       col = par("fg"), lty = NULL, lwd = par("lwd"),  
       xpd = NULL)
```

**Arguments**

<code>x0, y0</code>	coordinates of points <b>from</b> which to draw.
<code>x1, y1</code>	coordinates of points <b>to</b> which to draw.
<code>length</code>	length of the edges of the arrow head (in inches).
<code>angle</code>	angle from the shaft of the arrow to the edge of the arrow head.
<code>code</code>	integer code, determining <i>kind</i> of arrows to be drawn.
<code>col, lty, lwd, xpd</code>	usual graphical parameters as in <code>par</code> .

**Details**

For each `i`, an arrow is drawn between the point `(x0[i], y0[i])` and the point `(x1[i], y1[i])`.

If `code=2` an arrowhead is drawn at `(x0[i],y0[i])` and if `code=1` an arrowhead is drawn at `(x1[i],y1[i])`. If `code=3` a head is drawn at both ends of the arrow. Unless `length = 0`, when no head is drawn.

The graphical parameters `col` and `lty` can be used to specify a color and line texture for the line segments which make up the arrows (`col` may be a vector).

The direction of a zero-length arrow is indeterminate, and hence so is the direction of the arrowheads. To allow for rounding error, arrowheads are omitted (with a warning) on any arrow of length less than 1/1000 inch.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`segments` to draw segments.

## Examples

```
x <- runif(12); y <- rnorm(12)
i <- order(x,y); x <- x[i]; y <- y[i]
plot(x,y, main="arrows(.) and segments(.)")
## draw arrows from point to point :
s <- seq(length(x)-1) # one shorter than data
arrows(x[s], y[s], x[s+1], y[s+1], col= 1:3)
s <- s[-length(s)]
segments(x[s], y[s], x[s+2], y[s+2], col= 'pink')
```

---

**assocplot**      *Association Plots*


---

**Description**

Produce a Cohen-Friendly association plot indicating deviations from independence of rows and columns in a 2-dimensional contingency table.

**Usage**

```
assocplot(x, col = c("black", "red"), space = 0.3,
          main = NULL, xlab = NULL, ylab = NULL)
```

**Arguments**

<b>x</b>	a two-dimensional contingency table in matrix form.
<b>col</b>	a character vector of length two giving the colors used for drawing positive and negative Pearson residuals, respectively.
<b>space</b>	the amount of space (as a fraction of the average rectangle width and height) left between each rectangle.
<b>main</b>	overall title for the plot.
<b>xlab</b>	a label for the x axis. Defaults to the name of the row variable in <b>x</b> if non-NULL.
<b>ylab</b>	a label for the y axis. Defaults to the column names of the column variable in <b>x</b> if non-NULL.

**Details**

For a two-way contingency table, the signed contribution to Pearson's  $\chi^2$  for cell  $i, j$  is  $d_{ij} = (f_{ij} - e_{ij}) / \sqrt{e_{ij}}$ , where  $f_{ij}$  and  $e_{ij}$  are the observed and expected counts corresponding to the cell. In the Cohen-Friendly association plot, each cell is represented by a rectangle that has (signed) height proportional to  $d_{ij}$  and width proportional to  $\sqrt{e_{ij}}$ , so that the area of the box is proportional to the difference in observed and expected frequencies. The rectangles in each row are positioned relative to a baseline indicating independence ( $d_{ij} = 0$ ). If the observed frequency of a cell is greater than the expected one, the box rises above the baseline and is shaded in the color specified by the first element of **col**, which defaults to black; otherwise, the box falls below the baseline and is shaded in the color specified by the second element of **col**, which defaults to red.

## References

Cohen, A. (1980), On the graphical display of the significant components in a two-way contingency table. *Communications in Statistics—Theory and Methods*, **A9**, 1025–1041.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Conference Proceedings*, **17**, 190–200. <http://www.math.yorku.ca/SCS/sugi/sugi17-paper.html>

## See Also

`mosaicplot`; `chisq.test`.

## Examples

```
data(HairEyeColor)
## Aggregate over sex:
x <- margin.table(HairEyeColor, c(1, 2))
x
assocplot(x, main = "Relation between hair and eye color")
```

---

**axis**     *Add an Axis to a Plot*


---

**Description**

Adds an axis to the current plot, allowing the specification of the side, position, labels, and other options.

**Usage**

```
axis(side, at = NULL, labels = TRUE, tick = TRUE,
      line = NA, pos = NA, outer = FALSE, font = NA,
      vfont = NULL, lty = "solid", lwd = 1, col = NULL, ...)
```

**Arguments**

<b>side</b>	an integer specifying which side of the plot the axis is to be drawn on. The axis is placed as follows: 1=below, 2=left, 3=above and 4=right.
<b>at</b>	the points at which tick-marks are to be drawn. Non-finite (infinite, NaN or NA) values are omitted. By default, when NULL, tickmark locations are computed, see Details below.
<b>labels</b>	this can either be a logical value specifying whether (numerical) annotations are to be made at the tick-marks, or a vector of character strings to be placed at the tickpoints.
<b>tick</b>	a logical value specifying whether tickmarks should be drawn
<b>line</b>	the number of lines into the margin which the axis will be drawn. This overrides the value of the graphical parameter <code>mgp[3]</code> . The relative placing of tickmarks and tick labels is unchanged.
<b>pos</b>	the coordinate at which the axis line is to be drawn. this overrides the value of both <code>line</code> and <code>mgp[3]</code> .
<b>outer</b>	a logical value indicating whether the axis should be drawn in the outer plot margin, rather than the standard plot margin.
<b>font</b>	font for text.
<b>vfont</b>	vector font for text.

<code>lty</code> , <code>lwd</code>	line type, width for the axis line and the tick marks.
<code>col</code>	color for the axis line and the tick marks. The default <code>NULL</code> means to use <code>par("fg")</code> .
<code>...</code>	other graphical parameters may also be passed as arguments to this function, e.g., <code>las</code> for vertical/horizontal label orientation, or <code>fg</code> instead of <code>col</code> , see <code>par</code> on these.

## Details

The axis line is drawn from the lowest to the highest value of `at`, but will be clipped at the plot region. Only ticks which are drawn from points within the plot region (up to a tolerance for rounding error) are plotted, but the ticks and their labels may well extend outside the plot region.

When `at = NULL`, pretty tick mark locations are computed internally, the same `axTicks(side)` would, from `par("usr", "lab")`, and `par("xlog")` (or `ylog` respectively).

## Value

This function is invoked for its side effect, which is to add an axis to an already existing plot.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`axTicks` returns the axis tick locations corresponding to `at=NULL`; `pretty` is more flexible for computing pretty tick coordinates and does *not* depend on (nor adapt to) the coordinate system in use.

## Examples

```
plot(1:4, rnorm(4), axes=FALSE)
axis(1, 1:4, LETTERS[1:4])
axis(2)
box() # to make it look "as usual"

plot(1:7, rnorm(7), main = "axis() examples",
     type = "s", xaxt="n", frame = FALSE, col = "red")
```

```
axis(1, 1:7, LETTERS[1:7], col.axis = "blue")
# unusual options:
axis(4, col = "violet", col.axis="dark violet",lwd = 2)
axis(3, col = "gold", lty = 2, lwd = 0.5)
```



---

**axis.POSIXct**      *Date-time Plotting Functions*

---

**Description**

Functions to plot objects of classes "POSIXlt" and "POSIXct" representing calendar dates and times.

**Usage**

```
axis.POSIXct(side, x, at, format, ...)
```

```
## S3 method for class 'POSIXct':  
plot(x, y, xlab = "", ...)
```

```
## S3 method for class 'POSIXlt':  
plot(x, y, xlab = "", ...)
```

**Arguments**

<b>x, at</b>	A date-time object.
<b>y</b>	numeric values to be plotted against <b>x</b> .
<b>xlab</b>	a character string giving the label for the <b>x</b> axis.
<b>side</b>	See <b>axis</b> .
<b>format</b>	See <b>strptime</b> .
<b>...</b>	Further arguments to be passed from or to other methods, typically graphical parameters or arguments of <b>plot.default</b> .

**Details**

The functions plot against an x-axis of date-times. **axis.POSIXct** works quite hard to choose suitable time units (years, months, days, hours, minutes or seconds) and a sensible output format, but this can be overridden by supplying a **format** specification.

If **at** is supplied for **axis.POSIXct** it specifies the locations of the ticks and labels: if **x** is specified a suitable grid of labels is chosen.

**See Also**

DateTimeClasses for details of the classes.

## Examples

```
res <- try(data(beav1, package = "MASS"))
if(!inherits(res, "try-error")) {
  attach(beav1)
  time <- strptime(
    paste(1990, day, time %% 100, time %% 100),
    "%Y %j %H %M"
  )
  plot(time, temp, type="l") # axis at 4-hour intervals.
  # now label every hour on the time axis
  plot(time, temp, type="l", xaxt="n")
  r <- as.POSIXct(round(range(time), "hours"))
  axis.POSIXct(1, at=seq(r[1], r[2], by="hour"), format="%H")
  rm(time)
  detach(beav1)
}

plot(.leap.seconds, 1:22, type="n", yaxt="n",
     xlab="leap seconds", ylab="", bty="n")
rug(.leap.seconds)
```

---

**axTicks**      *Compute Axis Tickmark Locations*

---

**Description**

Compute tickmark locations, the same way as R does internally. This is only non-trivial when **log** coordinates are active. By default, gives the **at** values which **axis(side)** would use.

**Usage**

```
axTicks(side, axp = NULL, usr = NULL, log = NULL)
```

**Arguments**

<b>side</b>	integer in 1:4, as for <b>axis</b> .
<b>axp</b>	numeric vector of length three, defaulting to <b>par("Zaxp")</b> where “Z” is “x” or “y” depending on the <b>side</b> argument.
<b>usr</b>	numeric vector of length four, defaulting to <b>par("usr")</b> giving horizontal (‘x’) and vertical (‘y’) user coordinate limits.
<b>log</b>	logical indicating if log coordinates are active; defaults to <b>par("Zlog")</b> where ‘Z’ is as for the <b>axp</b> argument above.

**Details**

The **axp**, **usr**, and **log** arguments must be consistent as their default values (the **par(...)** results) are. Note that the meaning of **axp** alters very much when **log** is **TRUE**, see the documentation on **par(xaxp=.)**.

**axTicks()** can be regarded as an R implementation of the C function **CreateAtVector()** in ‘.../src/main/graphics.c’ which is called by **axis(side,\*)** when no argument **at** is specified.

**Value**

numeric vector of coordinate values at which axis tickmarks can be drawn. By default, when only the first argument is specified, these values should be identical to those that **axis(side)** would use or has used.

## See Also

`axis`, `par`. `pretty` uses the same algorithm but is independent of the graphics environment and has more options.

## Examples

```
plot(1:7, 10*21:27)
axTicks(1)
axTicks(2)
stopifnot(identical(axTicks(1), axTicks(3)),
           identical(axTicks(2), axTicks(4)))

## Show how axTicks() and axis() correspond :
op <- par(mfrow = c(3,1))
for(x in 9999*c(1,2,8)) {
  plot(x,9, log = "x")
  cat(formatC(par("xaxp"),wid=5),";",T<-axTicks(1),"\n")
  rug(T, col="red")
}
par(op)
```

---

**barplot**     *Bar Plots*

---

**Description**

Creates a bar plot with vertical or horizontal bars.

**Usage**

```
## Default S3 method:
barplot(height, width = 1, space = NULL,
  names.arg = NULL, legend.text = NULL, beside = FALSE,
  horiz = FALSE, density = NULL, angle = 45,
  col = heat.colors(NR), border = par("fg"),
  main = NULL, sub = NULL, xlab = NULL, ylab = NULL,
  xlim = NULL, ylim = NULL, xpd = TRUE,
  axes = TRUE, axisnames = TRUE,
  cex.axis = par("cex.axis"), cex.names = par("cex.axis"),
  inside = TRUE, plot = TRUE, axis.lty = 0, ...)
```

**Arguments**

- |               |  |
|---------------|--|
| <b>height</b> | either a vector or matrix of values describing the bars which make up the plot. If <b>height</b> is a vector, the plot consists of a sequence of rectangular bars with heights given by the values in the vector. If <b>height</b> is a matrix and <b>beside</b> is <b>FALSE</b> then each bar of the plot corresponds to a column of <b>height</b> , with the values in the column giving the heights of stacked “sub-bars” making up the bar. If <b>height</b> is a matrix and <b>beside</b> is <b>TRUE</b> , then the values in each column are juxtaposed rather than stacked. |
| <b>width</b>  | optional vector of bar widths. Re-cycled to length the number of bars drawn. Specifying a single value will have no visible effect unless <b>xlim</b> is specified.  |
| <b>space</b>  | the amount of space (as a fraction of the average bar width) left before each bar. May be given as a single number or one number per bar. If <b>height</b> is a matrix and <b>beside</b> is <b>TRUE</b> , <b>space</b> may be specified by two numbers, where the first is the space between bars in the same group, and the second the space between the groups. If not given explicitly, it defaults to <code>c(0,1)</code>  |

	if <b>height</b> is a matrix and <b>beside</b> is TRUE, and to 0.2 otherwise.
<b>names.arg</b>	a vector of names to be plotted below each bar or group of bars. If this argument is omitted, then the names are taken from the <b>names</b> attribute of <b>height</b> if this is a vector, or the column names if it is a matrix.
<b>legend.text</b>	a vector of text used to construct a legend for the plot, or a logical indicating whether a legend should be included. This is only useful when <b>height</b> is a matrix. In that case given legend labels should correspond to the rows of <b>height</b> ; if <b>legend.text</b> is true, the row names of <b>height</b> will be used as labels if they are non-null.
<b>beside</b>	a logical value. If FALSE, the columns of <b>height</b> are portrayed as stacked bars, and if TRUE the columns are portrayed as juxtaposed bars.
<b>horiz</b>	a logical value. If FALSE, the bars are drawn vertically with the first bar to the left. If TRUE, the bars are drawn horizontally with the first at the bottom.
<b>density</b>	a vector giving the density of shading lines, in lines per inch, for the bars or bar components. The default value of NULL means that no shading lines are drawn. Non-positive values of <b>density</b> also inhibit the drawing of shading lines.
<b>angle</b>	the slope of shading lines, given as an angle in degrees (counter-clockwise), for the bars or bar components.
<b>col</b>	a vector of colors for the bars or bar components.
<b>border</b>	the color to be used for the border of the bars.
<b>main,sub</b>	overall and sub title for the plot.
<b>xlab</b>	a label for the x axis.
<b>ylab</b>	a label for the y axis.
<b>xlim</b>	limits for the x axis.
<b>ylim</b>	limits for the y axis.
<b>xpd</b>	logical. Should bars be allowed to go outside region?
<b>axes</b>	logical. If TRUE, a vertical (or horizontal, if <b>horiz</b> is true) axis is drawn.
<b>axisnames</b>	logical. If TRUE, and if there are <b>names.arg</b> (see above), the other axis is drawn (with <b>lty=0</b> ) and labeled.

<code>cex.axis</code>	expansion factor for numeric axis labels.
<code>cex.names</code>	expansion factor for axis names (bar labels).
<code>inside</code>	logical. If <code>TRUE</code> , the lines which divide adjacent (non-stacked!) bars will be drawn. Only applies when <code>space = 0</code> (which it partly is when <code>beside = TRUE</code> ).
<code>plot</code>	logical. If <code>FALSE</code> , nothing is plotted.
<code>axis.lty</code>	the graphics parameter <code>lty</code> applied to the axis and tick marks of the categorical (default horizontal) axis. Note that by default the axis is suppressed.
<code>...</code>	further graphical parameters ( <code>par</code> ) are passed to <code>plot.window()</code> , <code>title()</code> and <code>axis</code> .

## Details

This is a generic function, it currently only has a default method. A formula interface may be added eventually.

## Value

A numeric vector (or matrix, when `beside = TRUE`), say `mp`, giving the coordinates of *all* the bar midpoints drawn, useful for adding to the graph.

If `beside` is true, use `colMeans(mp)` for the midpoints of each *group* of bars, see example.

## Note

Prior to R 1.6.0, `barplot` behaved as if `axis.lty = 1`, unintentionally.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`plot(..., type="h")`, `dotchart`, `hist`.

## Examples

```
tN <- table(Ni <- rpois(100, lambda=5))
r <- barplot(tN, col='gray')
# type = "h" plotting is 'bar'plot
```

```

lines(r, tN, type='h', col='red', lwd=2)

barplot(tN, space = 1.5, axisnames=FALSE,
  sub = "barplot(..., space= 1.5, axisnames = FALSE)")

data(VADeaths, package = "base")
barplot(VADeaths, plot = FALSE)
barplot(VADeaths, plot = FALSE, beside = TRUE)

mp <- barplot(VADeaths) # default
tot <- colMeans(VADeaths)
text(mp, tot + 3, format(tot), xpd = TRUE, col = "blue")
barplot(VADeaths, beside = TRUE,
  col = c("lightblue", "mistyrose", "lightcyan",
    "lavender", "cornsilk"),
  legend = rownames(VADeaths), ylim = c(0, 100))
title(main = "Death Rates in Virginia", font.main = 4)

hh <- t(VADeaths)[, 5:1]
mybarcol <- "gray20"
mp <- barplot(hh, beside = TRUE,
  col = c("lightblue", "mistyrose", "lightcyan", "lavender"),
  legend = colnames(VADeaths), ylim = c(0, 100),
  main = "Death Rates in Virginia", font.main = 4,
  sub = "Faked upper 2*sigma error bars",
  col.sub = mybarcol,
  cex.names = 1.5)
segments(mp, hh, mp, hh + 2*sqrt(1000*hh/100),
  col = mybarcol, lwd = 1.5)
stopifnot(dim(mp) == dim(hh)) # corresponding matrices
mtext(side = 1, at = colMeans(mp), line = -2,
  text = paste("Mean", formatC(colMeans(hh))), col = "red")

# Bar shading example
barplot(VADeaths, angle = 15+10*1:5, density = 20,
  col = "black", legend = rownames(VADeaths))
title(main = list("Death Rates in Virginia", font = 4))

# border :
barplot(VADeaths, border = "dark blue")

```



---

**box**     *Draw a Box around a Plot*

---

## Description

This function draws a box around the current plot in the given color and linetype. The **bty** parameter determines the type of box drawn. See **par** for details.

## Usage

```
box(which="plot", lty="solid", ...)
```

## Arguments

<b>which</b>	character, one of "plot", "figure", "inner" and "outer".
<b>lty</b>	line type of the box.
<b>...</b>	further graphical parameters, such as <b>bty</b> , <b>col</b> , or <b>lwd</b> , see <b>par</b> .

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**rect** for drawing of arbitrary rectangles.

## Examples

```
plot(1:7,abs(rnorm(7)), type='h', axes = FALSE)
axis(1, labels = letters[1:7])
box(lty='1373', col = 'red')
```

---

**boxplot**     *Box Plots*


---

**Description**

Produce box-and-whisker plot(s) of the given (grouped) values.

**Usage**

```
boxplot(x, ...)

## S3 method for class 'formula':
boxplot(formula, data = NULL, ..., subset)

## Default S3 method:
boxplot(x, ..., range=1.5, width=NULL, varwidth=FALSE,
        notch=FALSE, outline=TRUE, names, boxwex=0.8, plot=TRUE,
        border=par("fg"), col=NULL, log="", pars=NULL,
        horizontal=FALSE, add=FALSE, at=NULL)
```

**Arguments**

<b>formula</b>	a formula, such as <code>y ~ x</code> .
<b>data</b>	a <code>data.frame</code> (or list) from which the variables in <b>formula</b> should be taken.
<b>subset</b>	an optional vector specifying a subset of observations to be used for plotting.
<b>x</b>	for specifying data from which the boxplots are to be produced as well as for giving graphical parameters. Additional unnamed arguments specify further data, either as separate vectors (each corresponding to a component boxplot) or as a single list containing such vectors. NAs are allowed in the data.
<b>...</b>	For the <b>formula</b> method, arguments to the default method and graphical parameters.  For the default method, unnamed arguments are additional data vectors, and named arguments are graphical parameters in addition to the ones given by argument <b>pars</b> .

<b>range</b>	this determines how far the plot whiskers extend out from the box. If <b>range</b> is positive, the whiskers extend to the most extreme data point which is no more than <b>range</b> times the interquartile range from the box. A value of zero causes the whiskers to extend to the data extremes.
<b>width</b>	a vector giving the relative widths of the boxes making up the plot.
<b>varwidth</b>	if <b>varwidth</b> is TRUE, the boxes are drawn with widths proportional to the square-roots of the number of observations in the groups.
<b>notch</b>	if <b>notch</b> is TRUE, a notch is drawn in each side of the boxes. If the notches of two plots do not overlap then the medians are significantly different at the 5 percent level.
<b>outline</b>	if <b>outline</b> is not true, the boxplot lines are not drawn.
<b>names</b>	group labels which will be printed under each boxplot.
<b>boxwex</b>	a scale factor to be applied to all boxes. When there are only a few groups, the appearance of the plot can be improved by making the boxes narrower.
<b>plot</b>	if TRUE (the default) then a boxplot is produced. If not, the summaries which the boxplots are based on are returned.
<b>border</b>	an optional vector of colors for the outlines of the boxplots. The values in <b>border</b> are recycled if the length of <b>border</b> is less than the number of plots.
<b>col</b>	if <b>col</b> is non-null it is assumed to contain colors to be used to color the bodies of the box plots.
<b>log</b>	character indicating if x or y or both coordinates should be plotted in log scale.
<b>pars</b>	a list of graphical parameters; these are passed to <b>bxp</b> (if <b>plot</b> is true).
<b>horizontal</b>	logical indicating if the boxplots should be horizontal; default FALSE means vertical boxes.
<b>add</b>	logical, if true <i>add</i> boxplot to current plot.
<b>at</b>	numeric vector giving the locations where the boxplots should be drawn, particularly when <b>add</b> = TRUE; defaults to 1:n where <b>n</b> is the number of boxes.

## Details

The generic function `boxplot` currently has a default method (`boxplot.default`) and a formula interface (`boxplot.formula`).

## Value

List with the following components:

<code>stats</code>	a matrix, each column contains the extreme of the lower whisker, the lower hinge, the median, the upper hinge and the extreme of the upper whisker for one group/plot.
<code>n</code>	a vector with the number of observations in each group.
<code>conf</code>	a matrix where each column contains the lower and upper extremes of the notch.
<code>out</code>	the values of any data points which lie beyond the extremes of the whiskers.
<code>group</code>	a vector of the same length as <code>out</code> whose elements indicate which group the outlier belongs to
<code>names</code>	a vector of names for the groups

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

See also `boxplot.stats`.

## See Also

`boxplot.stats` which does the computation, `bxp` for the plotting; and `stripchart` for an alternative (with small data sets).

## Examples

```
## boxplot on a formula:
data(InsectSprays)
boxplot(count ~ spray, data = InsectSprays,
        col = "lightgray")
# add notches (somewhat funny here):
boxplot(count ~ spray, data = InsectSprays,
        notch = TRUE, add = TRUE, col = "blue")
```

```

data(OrchardSprays)
boxplot(decrease ~ treatment, data = OrchardSprays,
        log = "y", col="bisque")

rb <- boxplot(decrease ~ treatment, data = OrchardSprays,
              col="bisque")
title("Comparing boxplot()s and non-robust mean +/- SD")

mn.t <- tapply(OrchardSprays$decrease,
               OrchardSprays$treatment, mean)
sd.t <- tapply(OrchardSprays$decrease,
               OrchardSprays$treatment, sd)
xi <- 0.3 + seq(rb$n)
points(xi, mn.t, col = "orange", pch = 18)
arrows(xi, mn.t - sd.t, xi, mn.t + sd.t,
       code = 3, col = "pink", angle = 75, length = .1)

## boxplot on a matrix:
mat <- cbind(Uni05 = (1:100)/21, Norm = rnorm(100),
             T5 = rt(100, df = 5), Gam2 = rgamma(100, shape = 2))
boxplot(data.frame(mat),
        main = "boxplot(data.frame(mat), main = ...)")
par(las=1) # all axis labels horizontal
boxplot(data.frame(mat),
        main = "boxplot(*, horizontal = TRUE)",
        horizontal = TRUE)

## Using 'at = ' and adding boxplots -- example idea by
## Roger Bivand :
data(ToothGrowth)
boxplot(len ~ dose, data = ToothGrowth,
        boxwex = 0.25, at = 1:3 - 0.2,
        subset= supp == "VC", col="yellow",
        main="Guinea Pigs' Tooth Growth",
        xlab="Vitamin C dose mg",
        ylab="tooth length", ylim=c(0,35))
boxplot(len ~ dose, data = ToothGrowth, add = TRUE,
        boxwex = 0.25, at = 1:3 + 0.2,
        subset= supp == "OJ", col="orange")
legend(2, 9, c("Ascorbic acid", "Orange juice"),
      fill = c("yellow", "orange"))

```

---

**boxplot.stats**      *Box Plot Statistics*


---

**Description**

This function is typically called by `boxplot` to gather the statistics necessary for producing box plots, but may be invoked separately.

**Usage**

```
boxplot.stats(x, coef = 1.5, do.conf=TRUE, do.out=TRUE)
```

**Arguments**

- |                        |   |
|------------------------|---|
| <b>x</b>               | a numeric vector for which the boxplot will be constructed (NAs and NaNs are allowed and omitted).  |
| <b>coef</b>            | this determines how far the plot “whiskers” extend out from the box. If <b>coef</b> is positive, the whiskers extend to the most extreme data point which is no more than <b>coef</b> times the length of the box away from the box. A value of zero causes the whiskers to extend to the data extremes (with this setting no outliers will be returned). |
| <b>do.conf, do.out</b> | logicals; if <b>FALSE</b> , the <b>conf</b> or <b>out</b> component respectively will be empty in the result.   |

**Details**

The two “hinges” are versions of the first and third quartile, i.e., close to `quantile(x, c(1,3)/4)`. The hinges equal the quartiles for odd  $n$  (where  $n <- \text{length}(x)$ ) and differ for even  $n$ . Where the quartiles only equal observations for  $n \% 4 == 1$  ( $n \equiv 1 \pmod{4}$ ), the hinges do so *additionally* for  $n \% 4 == 2$  ( $n \equiv 2 \pmod{4}$ ), and are in the middle of two observations otherwise.

**Value**

List with named components as follows:

- |              |   |
|--------------|---|
| <b>stats</b> | a vector of length 5, containing the extreme of the lower whisker, the lower “hinge”, the median, the upper “hinge” and the extreme of the upper whisker. |
|--------------|---|

<code>n</code>	the number of non-NA observations in the sample.
<code>conf</code>	the lower and upper extremes of the “notch” ( <code>if(do.conf)</code> ).
<code>out</code>	the values of any data points which lie beyond the extremes of the whiskers ( <code>if(do.out)</code> ).

Note that `$stats` and `$conf` are sorted in *increasing* order, unlike `S`, and that `$n` and `$out` include any  $\pm \text{Inf}$  values.

## References

- Tukey, J. W. (1977) *Exploratory Data Analysis*. Section 2C.
- McGill, R., Tukey, J. W. and Larsen, W. A. (1978) Variations of box plots. *The American Statistician* **32**, 12–16.
- Velleman, P. F. and Hoaglin, D. C. (1981) *Applications, Basics and Computing of Exploratory Data Analysis*. Duxbury Press.
- Emerson, J. D and Strenio, J. (1983). Boxplots and batch comparison. Chapter 3 of *Understanding Robust and Exploratory Data Analysis*, eds. D. C. Hoaglin, F. Mosteller and J. W. Tukey. Wiley.

## See Also

`fivenum`, `boxplot`, `bxp`.

## Examples

```
x <- c(1:100, 1000)
str(b1 <- boxplot.stats(x))
str(b2 <- boxplot.stats(x, do.conf=FALSE, do.out=FALSE))
# do.out=F is still robust
stopifnot(b1 $ stats == b2 $ stats)
str(boxplot.stats(x, coef = 3, do.conf=FALSE))
## no outlier treatment:
str(boxplot.stats(x, coef = 0))

str(boxplot.stats(c(x, NA))) # slight change : n + 1
str(r <- boxplot.stats(c(x, -1:1/0)))
stopifnot(r$out == c(1000, -Inf, Inf))
```

---

**bxp**     *Box Plots from Summaries*


---

**Description**

**bxp** draws box plots based on the given summaries in **z**. It is usually called from within **boxplot**, but can be invoked directly.

**Usage**

```
bxp(z, notch = FALSE, width = NULL, varwidth = FALSE,
     outline = TRUE, notch.frac = 0.5, boxwex = 0.8,
     border = par("fg"), col = NULL, log="", pars = NULL,
     frame.plot = axes, horizontal = FALSE,
     add = FALSE, at = NULL, show.names = NULL, ...)
```

**Arguments**

<b>z</b>	a list containing data summaries to be used in constructing the plots. These are usually the result of a call to <b>boxplot</b> , but can be generated in any fashion.
<b>notch</b>	if <b>notch</b> is <b>TRUE</b> , a notch is drawn in each side of the boxes. If the notches of two plots do not overlap then the medians are significantly different at the 5 percent level.
<b>width</b>	a vector giving the relative widths of the boxes making up the plot.
<b>varwidth</b>	if <b>varwidth</b> is <b>TRUE</b> , the boxes are drawn with widths proportional to the square-roots of the number of observations in the groups.
<b>outline</b>	if <b>outline</b> is not true, the boxplot lines are not drawn.
<b>boxwex</b>	a scale factor to be applied to all boxes. When there are only a few groups, the appearance of the plot can be improved by making the boxes narrower.
<b>notch.frac</b>	numeric in (0,1). When <b>notch=TRUE</b> , the fraction of the box width that the notches should use.
<b>border</b>	character, the color of the box borders. Is recycled for multiple boxes.
<b>col</b>	character; the color within the box. Is recycled for multiple boxes



<code>log</code>	character, indicating if any axis should be drawn in logarithmic scale, as in <code>plot.default</code> .
<code>frame.plot</code>	logical, indicating if a “frame” (box) should be drawn; defaults to <code>TRUE</code> , unless <code>axes = FALSE</code> is specified.
<code>horizontal</code>	logical indicating if the boxplots should be horizontal; default <code>FALSE</code> means vertical boxes.
<code>add</code>	logical, if true <i>add</i> boxplot to current plot.
<code>at</code>	numeric vector giving the locations where the boxplots should be drawn, particularly when <code>add = TRUE</code> ; defaults to <code>1:n</code> where <code>n</code> is the number of boxes.
<code>show.names</code>	Set to <code>TRUE</code> or <code>FALSE</code> to override the defaults on whether an x-axis label is printed for each group.
<code>pars, ...</code>	graphical parameters can be passed as arguments to this function, either as a list ( <code>pars</code> ) or normally( <code>...</code> ). Currently, <code>pch</code> , <code>cex</code> , and <code>bg</code> are passed to <code>points</code> , <code>ylim</code> and <code>axes</code> to the main plot ( <code>plot.default</code> ), <code>xaxt</code> , <code>yaxt</code> , <code>las</code> to <code>axis</code> and the others to <code>title</code> .

## Value

An invisible vector, actually identical to the `at` argument, with the coordinates (“x” if horizontal is false, “y” otherwise) of box centers, useful for adding to the plot.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## Examples

```
set.seed(753)
str(bx.p <- boxplot(split(rt(100, 4), gl(5,20))))
op <- par(mfrow= c(2,2))
bxp(bx.p, xaxt = "n")
bxp(bx.p, notch = TRUE, axes = FALSE, pch = 4)
bxp(bx.p, notch = TRUE, col= "lightblue", frame= FALSE,
    outl= FALSE, main = "bxp(*, frame=FALSE, outl=FALSE)")
bxp(bx.p, notch = TRUE, col= "lightblue", border="red",
    ylim = c(-4,4), pch = 22, bg = "green", log = "x",
    main = "... log='x', ylim=*")
par(op)
```

```
op <- par(mfrow= c(1,2))
data(PlantGrowth)
## single group -- no label
boxplot(weight ~ group, data = PlantGrowth,
        subset = group == "ctrl")
bx<-boxplot(weight ~ group, data = PlantGrowth,
            subset = group == "ctrl", plot = FALSE)
## with label
bxp(bx, show.names=TRUE)
par(op)
```

---

**chull**      *Compute Convex Hull of a Set of Points*

---

**Description**

Computes the subset of points which lie on the convex hull of the set of points specified.

**Usage**

```
chull(x, y=NULL)
```

**Arguments**

**x**, **y**                      coordinate vectors of points. This can be specified as two vectors **x** and **y**, a 2-column matrix **x**, a list **x** with two components, etc, see **xy.coords**.

**Details**

**xy.coords** is used to interpret the specification of the points. The algorithm is that given by Eddy (1977).

‘Peeling’ as used in the S function **chull** can be implemented by calling **chull** recursively.

**Value**

An integer vector giving the indices of the points lying on the convex hull, in clockwise order.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Eddy, W. F. (1977) A new convex hull algorithm for planar sets. *ACM Transactions on Mathematical Software*, **3**, 398–403.

Eddy, W. F. (1977) Algorithm 523. CONVEX, A new convex hull algorithm for planar sets[Z]. *ACM Transactions on Mathematical Software*, **3**, 411–412.

**See Also**

**xy.coords**, **polygon**

**Examples**

```
X <- matrix(rnorm(2000), ncol=2)
plot(X, cex=0.5)
hpts <- chull(X)
hpts <- c(hpts, hpts[1])
lines(X[hpts, ])
```

---

**col2rgb**     *Color to RGB Conversion*

---

**Description**

“Any R color” to RGB (red/green/blue) conversion.

**Usage**

```
col2rgb(col)
```

**Arguments**

<code>col</code>	vector of any of the three kind of R colors, i.e., either a color name (an element of <code>colors()</code> ), a hexadecimal string of the form <code>"#rrggbb"</code> , or an integer <code>i</code> meaning <code>palette()[i]</code> .
------------------	---

**Details**

For integer colors, 0 is shorthand for the current `par("bg")`, and `NA` means “nothing” which effectively does not draw the corresponding item.

For character colors, `"NA"` is equivalent to `NA` above.

**Value**

an integer matrix with three rows and number of columns the length (and names if any) as `col`.

**Author(s)**

Martin Maechler

**See Also**

`rgb`, `colors`, `palette`, etc.

## Examples

```

col2rgb("peachpuff")
# names kept
col2rgb(c(blu = "royalblue", reddish = "tomato"))

col2rgb(1:8) # the ones from the palette() :
col2rgb(paste("gold", 1:4, sep=""))
col2rgb("#08a0ff")
## all three kind of colors mixed :
col2rgb(c(red="red", palette= 1:3, hex="#abcdef"))

## Non-introductory examples
grC <- col2rgb(paste("gray",0:100,sep=""))
# '2' or '3': almost equidistant
table(print(diff(grC["red",])))
## The 'named' grays are in between {"slate gray" is not
## gray, strictly}
col2rgb(c(g66="gray66", darkg= "dark gray", g67="gray67",
          g74="gray74", gray =      "gray", g75="gray75",
          g82="gray82", light="light gray", g83="gray83"))

crgb <- col2rgb(cc <- colors())
colnames(crgb) <- cc
t(crgb) ## The whole table

ccodes <- c(256^(2:0) %*% crgb) ## = internal codes
## How many names are 'aliases' of each other:
table(tcc <- table(ccodes))
length(uc <- unique(sort(ccodes))) # 502
## All the multiply named colors:
mult <- uc[tcc >= 2]
cl <- lapply(mult, function(m) cc[ccodes == m])
names(cl) <- apply(col2rgb(sapply(cl, function(x)x[1])),
                  2, function(n)paste(n, collapse=","))
str(cl)

if(require(xgobi)) {
  ## Look at the color cube dynamically :
  tc <- t(crgb[, !duplicated(ccodes)])
  # (397, 105)
  table(is.gray <- tc[,1] == tc[,2] & tc[,2] == tc[,3])
  xgobi(tc, color = c("gold", "gray")[1 + is.gray])
}

```

---

**colors**     *Color Names*

---

**Description**

Returns the built-in color names which R knows about.

**Usage**

```
colors()
```

**Details**

These color names can be used with a `col=` specification in graphics functions.

An even wider variety of colors can be created with primitives `rgb` and `hsv` or the derived `rainbow`, `heat.colors`, etc.

**Value**

A character vector containing all the built-in color names.

**See Also**

`palette` for setting the “palette” of colors for `par(col=num)`; `rgb`, `hsv`, `gray`; `rainbow` for a nice example; and `heat.colors`, `topo.colors` for images.

`col2rgb` for translating to RGB numbers and extended examples.

**Examples**

```
str(colors())
```

---

**contour**     *Display Contours*


---

**Description**

Create a contour plot, or add contour lines to an existing plot.

**Usage**

```
contour(x, ...)
## Default S3 method:
contour(x = seq(0, 1, len = nrow(z)),
        y = seq(0, 1, len = ncol(z)),
        z,
        nlevels = 10, levels = pretty(zlim, nlevels),
        labels = NULL,
        xlim = range(x, finite = TRUE),
        ylim = range(y, finite = TRUE),
        zlim = range(z, finite = TRUE),
        labcex = 0.6, drawlabels = TRUE,
        method = "flattest",
        vfont = c("sans serif", "plain"),
        axes = TRUE, frame.plot = axes,
        col = par("fg"), lty = par("lty"), lwd = par("lwd"),
        add = FALSE, ...)
```

**Arguments**

<b>x,y</b>	locations of grid lines at which the values in <b>z</b> are measured. These must be in ascending order. By default, equally spaced values from 0 to 1 are used. If <b>x</b> is a <i>list</i> , its components <b>x\$x</b> and <b>x\$y</b> are used for <b>x</b> and <b>y</b> , respectively. If the list has component <b>z</b> this is used for <b>z</b> .
<b>z</b>	a matrix containing the values to be plotted (NAs are allowed). Note that <b>x</b> can be used instead of <b>z</b> for convenience.
<b>nlevels</b>	number of contour levels desired <b>iff</b> <b>levels</b> is not supplied.
<b>levels</b>	numeric vector of levels at which to draw contour lines.
<b>labels</b>	a vector giving the labels for the contour lines. If <b>NULL</b> then the levels are used as labels.



<code>labcex</code>	<code>cex</code> for contour labelling.
<code>drawlabels</code>	logical. Contours are labelled if <code>TRUE</code> .
<code>method</code>	character string specifying where the labels will be located. Possible values are <code>"simple"</code> , <code>"edge"</code> and <code>"flattest"</code> (the default). See the Details section.
<code>vfont</code>	if a character vector of length 2 is specified, then Hershey vector fonts are used for the contour labels. The first element of the vector selects a typeface and the second element selects a fontindex (see <code>text</code> for more information).
<code>xlim, ylim, zlim</code>	x-, y- and z-limits for the plot.
<code>axes, frame.plot</code>	logical indicating whether axes or a box should be drawn, see <code>plot.default</code> .
<code>col</code>	color for the lines drawn.
<code>lty</code>	line type for the lines drawn.
<code>lwd</code>	line width for the lines drawn.
<code>add</code>	logical. If <code>TRUE</code> , add to a current plot.
<code>...</code>	additional graphical parameters (see <code>par</code> ) and the arguments to <code>title</code> may also be supplied.

## Details

`contour` is a generic function with only a default method in base R.

There is currently no documentation about the algorithm. The source code is in `'$R_HOME/src/main/plot3d.c'`.

The methods for positioning the labels on contours are `"simple"` (draw at the edge of the plot, overlaying the contour line), `"edge"` (draw at the edge of the plot, embedded in the contour line, with no labels overlapping) and `"flattest"` (draw on the flattest section of the contour, embedded in the contour line, with no labels overlapping). The second and third may not draw a label on every contour line.

For information about vector fonts, see the help for `text` and `Hershey`.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`filled.contour` for “color-filled” contours, `image` and the `graphics` demo which can be invoked as `demo(graphics)`.

## Examples

```
x <- -6:16
op <- par(mfrow = c(2, 2))
contour(outer(x, x), method = "edge",
        vfont = c("sans serif", "plain"))
z <- outer(x, sqrt(abs(x)), FUN = "/")
## Should not be necessary:
z[!is.finite(z)] <- NA
image(x, x, z)
contour(x, x, z, col = "pink", add = TRUE, method = "edge",
        vfont = c("sans serif", "plain"))
contour(x, x, z, ylim = c(1, 6), method = "simple",
        labcex = 1)
contour(x, x, z, ylim = c(-6, 6), nlev = 20, lty = 2,
        method = "simple")
par(op)

## Persian Rug Art:
x <- y <- seq(-4*pi, 4*pi, len = 27)
r <- sqrt(outer(x^2, y^2, "+"))
opar <- par(mfrow = c(2, 2), mar = rep(0, 4))
for(f in pi^(0:3))
  contour(cos(r^2)*exp(-r/f),
          drawlabels = FALSE, axes = FALSE, frame = TRUE)

data("volcano")
rx <- range(x <- 10*1:nrow(volcano))
ry <- range(y <- 10*1:ncol(volcano))
ry <- ry + c(-1,1) * (diff(rx) - diff(ry))/2
tcol <- terrain.colors(12)
par(opar); opar <- par(pty = "s", bg = "lightcyan")
plot(x = 0, y = 0, type = "n", xlim = rx, ylim = ry,
     xlab = "", ylab = "")
u <- par("usr")
rect(u[1], u[3], u[2], u[4], col = tcol[8], border = "red")
contour(x, y, volcano, col = tcol[2], lty = "solid",
        add = TRUE, vfont = c("sans serif", "plain"))
title("A Topographic Map of Maunga Whau", font = 4)
```

```
abline(h = 200*0:4, v = 200*0:4, col = "lightgray",  
       lty = 2, lwd = 0.1)  
par(opar)
```

---

**coplot**     *Conditioning Plots*


---

**Description**

This function produces two variants of the **conditioning** plots discussed in the reference below.

**Usage**

```
coplot(formula, data, given.values, panel = points,
       rows, columns,
       show.given = TRUE, col = par("fg"),
       pch = par("pch"),
       bar.bg = c(num = gray(0.8), fac = gray(0.95)),
       xlab = c(x.name, paste("Given :", a.name)),
       ylab = c(y.name, paste("Given :", b.name)),
       subscripts = FALSE,
       axlabels = function(f) abbreviate(levels(f)),
       number = 6, overlap = 0.5, xlim, ylim, ...)
co.intervals(x, number = 6, overlap = 0.5)
```

**Arguments**

- |                     |  |
|---------------------|--|
| <b>formula</b>      | <p>a formula describing the form of conditioning plot. A formula of the form <math>y \sim x \mid a</math> indicates that plots of <math>y</math> versus <math>x</math> should be produced conditional on the variable <math>a</math>. A formula of the form <math>y \sim x \mid a * b</math> indicates that plots of <math>y</math> versus <math>x</math> should be produced conditional on the two variables <math>a</math> and <math>b</math>.</p> <p>All three or four variables may be either numeric or factors. When <math>x</math> or <math>y</math> are factors, the result is almost as if <code>as.numeric()</code> was applied, whereas for factor <math>a</math> or <math>b</math>, the conditioning is adapted (and its graphics if <code>show.given</code> is true).</p> |
| <b>data</b>         | <p>a data frame containing values for any variables in the formula. By default the environment where <code>coplot</code> was called from is used.</p>  |
| <b>given.values</b> | <p>a value or list of two values which determine how the conditioning on <math>a</math> and <math>b</math> is to take place.</p> <p>When there is no <math>b</math> (i.e., conditioning only on <math>a</math>), usually this is a matrix with two columns each row of</p>   |

which gives an interval, to be conditioned on, but is can also be a single vector of numbers or a set of factor levels (if the variable being conditioned on is a factor). In this case (no **b**), the result of `co.intervals` can be used directly as `given.values` argument.

<code>panel</code>	a <code>function(x, y, col, pch, ...)</code> which gives the action to be carried out in each panel of the display. The default is <code>points</code> .
<code>rows</code>	the panels of the plot are laid out in a <code>rows</code> by <code>columns</code> array. <code>rows</code> gives the number of rows in the array.
<code>columns</code>	the number of columns in the panel layout array.
<code>show.given</code>	logical (possibly of length 2 for 2 conditioning variables): should conditioning plots be shown for the corresponding conditioning variables (default <code>TRUE</code> )
<code>col</code>	a vector of colors to be used to plot the points. If too short, the values are recycled.
<code>pch</code>	a vector of plotting symbols or characters. If too short, the values are recycled.
<code>bar.bg</code>	a named vector with components <code>"num"</code> and <code>"fac"</code> giving the background colors for the (shingle) bars, for <b>numeric</b> and <b>factor</b> conditioning variables respectively.
<code>xlab</code>	character; labels to use for the x axis and the first conditioning variable. If only one label is given, it is used for the x axis and the default label is used for the conditioning variable.
<code>ylab</code>	character; labels to use for the y axis and any second conditioning variable.
<code>subscripts</code>	logical: if true the panel function is given an additional (third) argument <code>subscripts</code> giving the subscripts of the data passed to that panel.
<code>axlabels</code>	function for creating axis (tick) labels when x or y are factors.
<code>number</code>	integer; the number of conditioning intervals, for a and b, possibly of length 2. It is only used if the corresponding conditioning variable is not a <b>factor</b> .
<code>overlap</code>	numeric $< 1$ ; the fraction of overlap of the conditioning variables, possibly of length 2 for x and y direction. When <code>overlap &lt; 0</code> there will be <i>gaps</i> between the data slices.

<code>xlim</code>	the range for the x axis.
<code>ylim</code>	the range for the y axis.
<code>...</code>	additional arguments to the panel function.
<code>x</code>	a numeric vector.

## Details

In the case of a single conditioning variable `a`, when both `rows` and `columns` are unspecified, a “close to square” layout is chosen with `columns >= rows`.

In the case of multiple `rows`, the *order* of the panel plots is from the bottom and from the left (corresponding to increasing `a`, typically).

A panel function should not attempt to start a new plot, but just plot within a given coordinate system: thus `plot` and `boxplot` are not panel functions.

## Value

`co.intervals(., number, .)` returns a  $(\text{number} \times 2)$  matrix, say `ci`, where `ci[k,]` is the range of `x` values for the `k`-th interval.

## References

Chambers, J. M. (1992) *Data for models*. Chapter 3 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

Cleveland, W. S. (1993) *Visualizing Data*. New Jersey: Summit Press.

## See Also

`pairs`, `panel.smooth`, `points`.

## Examples

```
## Tonga Trench Earthquakes
data(quakes)
coplot(lat ~ long | depth, data = quakes)
given.depth <-
  co.intervals(quakes$depth, number = 4, overlap = .1)
coplot(lat ~ long | depth, data = quakes,
        given.v = given.depth, rows = 1)

## Conditioning on 2 variables:
ll.dm <- lat ~ long | depth * mag
```

```

coplot(l1.dm, data = quakes)
coplot(l1.dm, data = quakes, number=c(4,7),
       show.given=c(TRUE,FALSE))
coplot(l1.dm, data = quakes, number=c(3,7),
       overlap=c(-.5,.1)) # negative overlap DROPS values

data(warpbreaks)
## given two factors
# to get nicer default labels
Index <- seq(length=nrow(warpbreaks))
coplot(breaks ~ Index | wool * tension, data = warpbreaks,
       show.given = 0:1)
coplot(breaks ~ Index | wool * tension, data = warpbreaks,
       col = "red", bg = "pink", pch = 21,
       bar.bg = c(fac = "light blue"))

## Example with empty panels:
data(state)
attach(data.frame(state.x77)) # don't need 'data' arg below
coplot(Life.Exp ~ Income | Illiteracy * state.region,
       number = 3, panel = function(x, y, ...)
       panel.smooth(x, y, span = .8, ...))
## y ~ factor -- not really sensical, but 'show off':
coplot(Life.Exp ~ state.region | Income * state.division,
       panel = panel.smooth)
detach() # data.frame(state.x77)

```

---

**curve**     *Draw Function Plots*


---

**Description**

Draws a curve corresponding to the given function or expression (in **x**) over the interval [**from**,**to**].

**Usage**

```
curve(expr, from, to, n = 101, add = FALSE, type = "l",
      ylab = NULL, log = NULL, xlim = NULL, ...)
```

```
## S3 method for class 'function':
plot(x, from = 0, to = 1, xlim = NULL, ...)
```

**Arguments**

<b>expr</b>	an expression written as a function of <b>x</b> , or alternatively the name of a function which will be plotted.
<b>x</b>	a ‘vectorizing’ numeric R function.
<b>from,to</b>	the range over which the function will be plotted.
<b>n</b>	integer; the number of <b>x</b> values at which to evaluate.
<b>add</b>	logical; if <b>TRUE</b> add to already existing plot.
<b>xlim</b>	numeric of length 2; if specified, it serves as default for <b>c(from, to)</b> .
<b>type, ylab, log, ...</b>	graphical parameters can also be specified as arguments. <b>plot.function</b> passes all these to <b>curve</b> .

**Details**

The evaluation of **expr** is at **n** points equally spaced over the range [**from**, **to**], possibly adapted to log scale. The points determined in this way are then joined with straight lines. **x(t)** or **expr** (with **x** inside) must return a numeric of the same length as the argument **t** or **x**.

If **add = TRUE**, **c(from,to)** default to **xlim** which defaults to the current x-limits. Further, **log** is taken from the current plot when **add** is true.

This used to be a quick hack which now seems to serve a useful purpose, but can give bad results for functions which are not smooth.



For “expensive” **expressions**, you should use smarter tools.

### See Also

`splinefun` for spline interpolation, `lines`.

### Examples

```
op <- par(mfrow=c(2,2))
curve(x^3-3*x, -2, 2)
curve(x^2-2, add = TRUE, col = "violet")

plot(cos, xlim = c(-pi,3*pi), n = 1001, col = "blue")

chippy <- function(x) sin(cos(x)*exp(-x/2))
curve(chippy, -8, 7, n=2001)
curve(chippy, -8, -5)

for(ll in c("", "x", "y", "xy"))
  curve(log(1+x), 1,100, log=ll,
        sub=paste("log= ",ll,"'",sep=""))
par(op)
```

---

**dev.xxx**      *Control Multiple Devices*

---

**Description**

These functions provide control over multiple graphics devices.

Only one device is the *active* device. This is the device in which all graphics operations occur.

Devices are associated with a name (e.g., "X11" or "postscript") and a number; the "null device" is always device 1.

`dev.off` shuts down the specified (by default the current) device. `graphics.off()` shuts down all open graphics devices.

`dev.set` makes the specified device the active device.

A list of the names of the open devices is stored in `.Devices`. The name of the active device is stored in `.Device`.

**Usage**

```
dev.cur()
dev.list()
dev.next(which = dev.cur())
dev.prev(which = dev.cur())
dev.off(which = dev.cur())
dev.set(which = dev.next())
graphics.off()
```

**Arguments**

**which**                      An integer specifying a device number

**Value**

`dev.cur` returns the number and name of the active device, or 1, the null device, if none is active.

`dev.list` returns the numbers of all open devices, except device 1, the null device. This is a numeric vector with a `names` attribute giving the names, or `NULL` if there is no open device.

`dev.next` and `dev.prev` return the number and name of the next / previous device in the list of devices. The list is regarded as a circular list, and "null device" will be included only if there are no open devices.

`dev.off` returns the name and number of the new active device (after the specified device has been shut down).

`dev.set` returns the name and number of the new active device.

## See Also

Devices, such as `postscript`, etc; `layout` and its links for setting up plotting regions on the current device.

## Examples

```
## Unix-specific example
x11()
plot(1:10)
x11()
plot(rnorm(10))
dev.set(dev.prev())
abline(0,1) # through the 1:10 points
dev.set(dev.next())
abline(h=0, col="gray") # for the residual plot
dev.set(dev.prev())
dev.off(); dev.off() # close the two X devices
```

---

**dev2**     *Copy Graphics Between Multiple Devices*


---

## Description

`dev.copy` copies the graphics contents of the current device to the device specified by **which** or to a new device which has been created by the function specified by **device** (it is an error to specify both **which** and **device**). (If recording is off on the current device, there are no contents to copy: this will result in no plot or an empty plot.) The device copied to becomes the current device.

`dev.print` copies the graphics contents of the current device to a new device which has been created by the function specified by **device** and then shuts the new device.

`dev.copy2eps` is similar to `dev.print` but produces an EPSF output file, in portrait orientation (`horizontal = FALSE`)

`dev.control` allows the user to control the recording of graphics operations in a device. If `displaylist` is "inhibit" ("enable") then recording is turned off (on). It is only safe to change this at the beginning of a plot (just before or just after a new page). Initially recording is on for screen devices, and off for print devices.

## Usage

```
dev.copy(device, ..., which = dev.next())
dev.print(device = postscript, ...)
dev.copy2eps(...)
dev.control(displaylist = c("inhibit", "enable"))
```

## Arguments

<b>device</b>	A device function (e.g., <code>x11</code> , <code>postscript</code> , ...)
<b>...</b>	Arguments to the <b>device</b> function above. For <code>dev.print</code> , this includes <b>which</b> and by default any <code>postscript</code> arguments.
<b>which</b>	A device number specifying the device to copy to
<b>displaylist</b>	A character string: the only valid values are "inhibit" and "enable".

## Details

For `dev.copy2eps`, `width` and `height` are taken from the current device unless otherwise specified. If just one of `width` and `height` is specified, the other is adjusted to preserve the aspect ratio of the device being copied. The default file name is `Rplot.eps`.

The default for `dev.print` is to produce and print a postscript copy, if `options("printcmd")` is set suitably.

`dev.print` is most useful for producing a postscript print (its default) when the following applies. Unless `file` is specified, the plot will be printed. Unless `width`, `height` and `pointsize` are specified the plot dimensions will be taken from the current device, shrunk if necessary to fit on the paper. (`pointsize` is rescaled if the plot is shrunk.) If `horizontal` is not specified and the plot can be printed at full size by switching its value this is done instead of shrinking the plot region.

If `dev.print` is used with a specified device (even `postscript`) it sets the width and height in the same way as `dev.copy2eps`.

## Value

`dev.copy` returns the name and number of the device which has been copied to.

`dev.print` and `dev.copy2eps` return the name and number of the device which has been copied from.

## Note

Most devices (including all screen devices) have a display list which records all of the graphics operations that occur in the device. `dev.copy` copies graphics contents by copying the display list from one device to another device. Also, automatic redrawing of graphics contents following the resizing of a device depends on the contents of the display list.

After the command `dev.control("inhibit")`, graphics operations are not recorded in the display list so that `dev.copy` and `dev.print` will not copy anything and the contents of a device will not be redrawn automatically if the device is resized.

The recording of graphics operations is relatively expensive in terms of memory so the command `dev.control("inhibit")` can be useful if memory usage is an issue.

**See Also**

`dev.cur` and other `dev.xxx` functions

**Examples**

```
x11()
plot(rnorm(10), main="Plot 1")
dev.copy(device=x11)
mtext("Copy 1", 3)
dev.print(width=6, height=6, horizontal=FALSE) # prints it
dev.off(dev.prev())
dev.off()
```

---

dev2bitmap      *Graphics Device for Bitmap Files via GhostScript*


---

## Description

bitmap generates a graphics file. dev2bitmap copies the current graphics device to a file in a graphics format.

## Usage

```
bitmap(file, type = "png256", height = 6, width = 6,
       res = 72, pointsize, ...)
dev2bitmap(file, type = "png256", height = 6, width = 6,
          res = 72, pointsize, ...)
```

## Arguments

file	The output file name, with an appropriate extension.
type	The type of bitmap. the default is "png256".
height	The plot height, in inches.
width	The plot width, in inches.
res	Resolution, in dots per inch.
pointsize	The pointsize to be used for text: defaults to something reasonable given the width and height
...	Other parameters passed to <code>postscript</code> .

## Details

dev2bitmap works by copying the current device to a `postscript` device, and post-processing the output file using `ghostscript`. `bitmap` works in the same way using a `postscript` device and postprocessing the output as “printing”.

You will need a version of `ghostscript` (5.10 and later have been tested): the full path to the executable can be set by the environment variable `R_GSCMD`.

The types available will depend on the version of `ghostscript`, but are likely to include "pcxmono", "pcxgray", "pcx16", "pcx256", "pcx24b", "pcxcmk", "pbm", "pbmraw", "pgm", "pgmraw", "pgnm", "pgnmraw", "pnm", "pnmraw", "ppm", "ppmraw", "pkm", "pkmraw", "tiffcrle", "tiffg3", "tiffg32d", "tiffg4", "tiffllzw",

"tiffpack", "tiff12nc", "tiff24nc", "psmono", "psgray", "psrgb", "bit", "bitrgb", "bitcmyk", "pngmono", "pnggray", "png16", "png256", "png16m", "jpeg", "jpeggray", "pdfwrite".

Note: despite the name of the functions they can produce PDF *via* `type = "pdfwrite"`, and the PDF produced is not bitmapped.

For formats which contain a single image, a file specification like `Rplots%03d.png` can be used: this is interpreted by GhostScript.

For `dev2bitmap` if just one of `width` and `height` is specified, the other is chosen to preserve aspect ratio of the device being copied.

## Value

None.

## See Also

`postscript`, `png` and `jpeg` and on Windows `bmp`.

`pdf` generate PDF directly.

To display an array of data, see `image`.



---

**Devices**     *List of Graphical Devices*

---

**Description**

The following graphics devices are currently available:

- **postscript** Writes PostScript graphics commands to a file
- **pdf** Write PDF graphics commands to a file
- **pictex** Writes LaTeX/PicTeX graphics commands to a file
- **xfig** Device for XFIG graphics file format
- **bitmap** bitmap pseudo-device via **GhostScript** (if available).

The following devices will be available if R was compiled to use them and started with the appropriate ‘`--gui`’ argument:

- **X11** The graphics driver for the X11 Window system
- **png** PNG bitmap device
- **jpeg** JPEG bitmap device
- **GTK**, **GNOME** Graphics drivers for the GNOME GUI.

None of these are available under R CMD BATCH.

**Usage**

```
X11(...)  
postscript(...)  
pdf(...)  
pictex(...)  
png(...)  
jpeg(...)  
GTK(...)  
GNOME(...)  
xfig(...)  
bitmap(...)  
  
dev.interactive()
```

## Details

If no device is open, using a high-level graphics function will cause a device to be opened. Which device is given by `options("device")` which is initially set as the most appropriate for each platform: a screen device in interactive use and `postscript` otherwise.

## Value

`dev.interactive()` returns a logical, `TRUE` iff an interactive (screen) device is in use.

## See Also

The individual help files for further information on any of the devices listed here;

`dev.cur`, `dev.print`, `graphics.off`, `image`, `dev2bitmap`.  
`capabilities` to see if `X11`, `jpeg` and `png` are available.

## Examples

```
## open the default screen device on this platform if no
## device is open
if(dev.cur() == 1) get(getOption("device"))()
```

---

**dotchart**      *Cleveland Dot Plots*

---

**Description**

Draw a Cleveland dot plot.

**Usage**

```
dotchart(x, labels = NULL, groups = NULL, gdata = NULL,
  cex = par("cex"), pch = 21, gpch = 21, bg = par("bg"),
  color = par("fg"), gcolor = par("fg"), lcolor = "gray",
  xlim = range(x[is.finite(x)]),
  main = NULL, xlab = NULL, ylab = NULL, ...)
```

**Arguments**

<b>x</b>	either a vector or matrix of numeric values (NAs are allowed). If <b>x</b> is a matrix the overall plot consists of juxtaposed dotplots for each row.
<b>labels</b>	a vector of labels for each point. For vectors the default is to use <code>names(x)</code> and for matrices the row labels <code>dimnames(x)[[1]]</code> .
<b>groups</b>	an optional factor indicating how the elements of <b>x</b> are grouped. If <b>x</b> is a matrix, <b>groups</b> will default to the columns of <b>x</b> .
<b>gdata</b>	data values for the groups. This is typically a summary such as the median or mean of each group.
<b>cex</b>	the character size to be used. Setting <b>cex</b> to a value smaller than one can be a useful way of avoiding label overlap.
<b>pch</b>	the plotting character or symbol to be used.
<b>gpch</b>	the plotting character or symbol to be used for group values.
<b>bg</b>	the background color of plotting characters or symbols to be used; use <code>par(bg= *)</code> to set the background color of the whole plot.
<b>color</b>	the color(s) to be used for points and labels.
<b>gcolor</b>	the single color to be used for group labels and values.

<code>lcolor</code>	the color(s) to be used for the horizontal lines.
<code>xlim</code>	horizontal range for the plot, see <code>plot.window</code> , e.g.
<code>main</code>	overall title for the plot, see <code>title</code> .
<code>xlab, ylab</code>	axis annotations as in <code>title</code> .
<code>...</code>	graphical parameters can also be specified as arguments.

## Value

This function is invoked for its side effect, which is to produce two variants of dotplots as described in Cleveland (1985).

Dot plots are a reasonable substitute for bar plots.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Cleveland, W. S. (1985) *The Elements of Graphing Data*. Monterey, CA: Wadsworth.

## Examples

```
data(VADeaths)
dotchart(VADeaths, main = "Death Rates in Virginia - 1940")
op <- par(xaxs="i") # 0 -- 100%
dotchart(t(VADeaths), xlim = c(0,100),
         main = "Death Rates in Virginia - 1940")
par(op)
```

---

**filled.contour**      *Level (Contour) Plots*


---

**Description**

This function produces a contour plot with the areas between the contours filled in solid color (Cleveland calls this a level plot). A key showing how the colors map to *z* values is shown to the right of the plot.

**Usage**

```
filled.contour(x = seq(0, 1, len = nrow(z)),
  y = seq(0, 1, len = ncol(z)),
  z,
  xlim = range(x, finite=TRUE),
  ylim = range(y, finite=TRUE),
  zlim = range(z, finite=TRUE),
  levels = pretty(zlim, nlevels), nlevels = 20,
  color.palette = cm.colors,
  col = color.palette(length(levels) - 1),
  plot.title, plot.axes, key.title, key.axes,
  asp = NA, xaxs = "i", yaxs = "i", las = 1,
  axes = TRUE, frame.plot = axes, ...)
```

**Arguments**

- |               |   |
|---------------|---|
| <b>x,y</b>    | locations of grid lines at which the values in <b>z</b> are measured. These must be in ascending order. By default, equally spaced values from 0 to 1 are used. If <b>x</b> is a <b>list</b> , its components <b>x\$x</b> and <b>x\$y</b> are used for <b>x</b> and <b>y</b> , respectively. If the list has component <b>z</b> this is used for <b>z</b> . |
| <b>z</b>      | a matrix containing the values to be plotted (NAs are allowed). Note that <b>x</b> can be used instead of <b>z</b> for convenience.   |
| <b>xlim</b>   | <b>x</b> limits for the plot.   |
| <b>ylim</b>   | <b>y</b> limits for the plot.   |
| <b>zlim</b>   | <b>z</b> limits for the plot.   |
| <b>levels</b> | a set of levels which are used to partition the range of <b>z</b> . Must be <b>strictly</b> increasing (and finite). Areas  |

	with <b>z</b> values between consecutive levels are painted with the same color.
<b>nlevels</b>	if <b>levels</b> is not specified, the range of <b>z</b> , values is divided into approximately this many levels.
<b>color.palette</b>	a color palette function to be used to assign colors in the plot.
<b>col</b>	an explicit set of colors to be used in the plot. This argument overrides any palette function specification.
<b>plot.title</b>	statements which add titles to the main plot.
<b>plot.axes</b>	statements which draw axes (and a <b>box</b> ) on the main plot. This overrides the default axes.
<b>key.title</b>	statements which add titles for the plot key.
<b>key.axes</b>	statements which draw axes on the plot key. This overrides the default axis.
<b>asp</b>	the $y/x$ aspect ratio, see <b>plot.window</b> .
<b>xaxis</b>	the x axis style. The default is to use internal labeling.
<b>yaxis</b>	the y axis style. The default is to use internal labeling.
<b>las</b>	the style of labeling to be used. The default is to use horizontal labeling.
<b>axes, frame.plot</b>	logicals indicating if axes and a box should be drawn, as in <b>plot.default</b> .
<b>...</b>	additional graphical parameters, currently only passed to <b>title()</b> .

## Note

This function currently uses the **layout** function and so is restricted to a full page display. As an alternative consider the **levelplot** function from the **lattice** package which works in multipanel displays.

The output produced by **filled.contour** is actually a combination of two plots; one is the filled contour and one is the legend. Two separate coordinate systems are set up for these two plots, but they are only used internally - once the function has returned these coordinate systems are lost. If you want to annotate the main contour plot, for example to add points, you can specify graphics commands in the **plot.axes** argument. An example is given below.

## Author(s)

Ross Ihaka.

## References

Cleveland, W. S. (1993) *Visualizing Data*. Summit, New Jersey: Hobart.

## See Also

`contour`, `image`, `palette`; `levelplot` from package **lattice**.

## Examples

```
data(volcano)
# simple
filled.contour(volcano, color = terrain.colors, asp = 1)

x <- 10*1:nrow(volcano)
y <- 10*1:ncol(volcano)
filled.contour(x, y, volcano, color = terrain.colors,
  plot.title = title(main = "Topography of Maunga Whau",
    xlab = "Meters North", ylab = "Meters West"),
  plot.axes = { axis(1, seq(100, 800, by = 100))
    axis(2, seq(100, 600, by = 100)) },
  key.title = title(main="Height\n(meters)"),
  key.axes = axis(4, seq(90, 190, by = 10)))

mtext(paste("filled.contour(.) from", R.version.string),
  side = 1, line = 4, adj = 1, cex = .66)

# Annotating a filled contour plot
a <- expand.grid(1:20, 1:20)
b <- matrix(a[,1] + a[,2], 20)
filled.contour(x = 1:20, y = 1:20, z = b,
  plot.axes={ axis(1); axis(2); points(10,10) })

## Persian Rug Art:
x <- y <- seq(-4*pi, 4*pi, len = 27)
r <- sqrt(outer(x^2, y^2, "+"))
filled.contour(cos(r^2)*exp(-r/(2*pi)), axes = FALSE)
## rather, the key should be labeled:
filled.contour(cos(r^2)*exp(-r/(2*pi)), frame.plot = FALSE,
  plot.axes = {})
```

---

**fourfoldplot**      *Fourfold Plots*


---

**Description**

Creates a fourfold display of a 2 by 2 by  $k$  contingency table on the current graphics device, allowing for the visual inspection of the association between two dichotomous variables in one or several populations (strata).

**Usage**

```
fourfoldplot(x, color = c("#99CCFF", "#6699CC"),
             conf.level = 0.95,
             std = c("margins", "ind.max", "all.max"),
             margin = c(1, 2), space = 0.2, main = NULL,
             mfrow = NULL, mfcol = NULL)
```

**Arguments**

- |                   |  |
|-------------------|--|
| <b>x</b>          | a 2 by 2 by $k$ contingency table in array form, or as a 2 by 2 matrix if $k$ is 1.  |
| <b>color</b>      | a vector of length 2 specifying the colors to use for the smaller and larger diagonals of each 2 by 2 table.   |
| <b>conf.level</b> | confidence level used for the confidence rings on the odds ratios. Must be a single nonnegative number less than 1; if set to 0, confidence rings are suppressed.  |
| <b>std</b>        | a character string specifying how to standardize the table. Must be one of "margins", "ind.max", or "all.max", and can be abbreviated by the initial letter. If set to "margins", each 2 by 2 table is standardized to equate the margins specified by <b>margin</b> while preserving the odds ratio. If "ind.max" or "all.max", the tables are either individually or simultaneously standardized to a maximal cell frequency of 1. |
| <b>margin</b>     | a numeric vector with the margins to equate. Must be one of 1, 2, or c(1, 2) (the default), which corresponds to standardizing the row, column, or both margins in each 2 by 2 table. Only used if <b>std</b> equals "margins".  |



<b>space</b>	the amount of space (as a fraction of the maximal radius of the quarter circles) used for the row and column labels.
<b>main</b>	character string for the fourfold title.
<b>mfrow</b>	a numeric vector of the form <code>c(nr, nc)</code> , indicating that the displays for the 2 by 2 tables should be arranged in an <code>nr</code> by <code>nc</code> layout, filled by rows.
<b>mfcol</b>	a numeric vector of the form <code>c(nr, nc)</code> , indicating that the displays for the 2 by 2 tables should be arranged in an <code>nr</code> by <code>nc</code> layout, filled by columns.

## Details

The fourfold display is designed for the display of 2 by 2 by  $k$  tables.

Following suitable standardization, the cell frequencies  $f_{ij}$  of each 2 by 2 table are shown as a quarter circle whose radius is proportional to  $\sqrt{f_{ij}}$  so that its area is proportional to the cell frequency. An association (odds ratio different from 1) between the binary row and column variables is indicated by the tendency of diagonally opposite cells in one direction to differ in size from those in the other direction; color is used to show this direction. Confidence rings for the odds ratio allow a visual test of the null of no association; the rings for adjacent quadrants overlap iff the observed counts are consistent with the null hypothesis.

Typically, the number  $k$  corresponds to the number of levels of a stratifying variable, and it is of interest to see whether the association is homogeneous across strata. The fourfold display visualizes the pattern of association. Note that the confidence rings for the individual odds ratios are not adjusted for multiple testing.

## References

Friendly, M. (1994). A fourfold display for 2 by 2 by  $k$  tables. Technical Report 217, York University, Psychology Department. <http://www.math.yorku.ca/SCS/Papers/4fold/4fold.ps.gz>

## See Also

`mosaicplot`

## Examples

```
data(UCBAdmissions)
## Use the Berkeley admission data as in Friendly (1995).
```

```
x <- aperm(UCBAdmissions, c(2, 1, 3))
dimnames(x)[[2]] <- c("Yes", "No")
names(dimnames(x)) <- c("Sex", "Admit?", "Department")
ftable(x)

## Fourfold display of data aggregated over departments,
## with frequencies standardized to equate the margins for
## admission and sex. Figure 1 in Friendly (1994).
fourfoldplot(margin.table(x, c(1, 2)))

## Fourfold display of x, with frequencies in each table
## standardized to equate the margins for admission and
## sex. Figure 2 in Friendly (1994).
fourfoldplot(x)

## Fourfold display of x, with frequencies in each table
## standardized to equate the margins for admission. but
## not for sex. Figure 3 in Friendly (1994).
fourfoldplot(x, margin = 2)
```

---

**frame**     *Create / Start a New Plot Frame*

---

## Description

This function (**frame** is an alias for **plot.new**) causes the completion of plotting in the current plot (if there is one) and an advance to a new graphics frame. This is used in all high-level plotting functions and also useful for skipping plots when a multi-figure region is in use.

## Usage

```
plot.new()  
frame()
```

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (**frame**.)

## See Also

`plot.window`, `plot.default`.

**Gnome**      *GNOME Desktop Graphics Device*

---

## Description

`gnome` starts a GNOME compatible device driver. GNOME is an acronym for **G**NU **N**etwork **O**bject **M**odel **E**nvironment.

## Usage

```
gnome(display="", width=7, height=7, pointsize=12)
GNOME(display="", width=7, height=7, pointsize=12)
```

## Arguments

<code>display</code>	the display on which the graphics window will appear. The default is to use the value in the user's environment variable <code>DISPLAY</code> .
<code>width</code>	the width of the plotting window in inches.
<code>height</code>	the height of the plotting window in inches.
<code>pointsize</code>	the default pointsize to be used.

## Note

This is still in development state.

The GNOME device is only available when explicitly desired at configure/compile time, see the toplevel 'INSTALL' file.

## Author(s)

Lyndon Drake

## References

<http://www.gnome.org> and <http://www.gtk.org> for the associated GTK+ libraries.

## See Also

`x11`, `Devices`.

## Examples

```
gnome(width=9)
```

---

**gray**     *Gray Level Specification*

---

**Description**

Create a vector of colors from a vector of gray levels.

**Usage**

```
gray(level)
grey(level)
```

**Arguments**

**level**            a vector of desired gray levels between 0 and 1; zero indicates "black" and one indicates "white".

**Details**

The values returned by **gray** can be used with a **col=** specification in graphics functions or in **par**.

**grey** is an alias for **gray**.

**Value**

A vector of “colors” of the same length as **level**.

**See Also**

**rainbow**, **hsv**, **rgb**.

**Examples**

```
gray(0:8 / 8)
```

---

**grid**     *Add Grid to a Plot*

---

**Description**

`grid` adds an `nx` by `ny` rectangular grid to an existing plot.

**Usage**

```
grid(nx = NULL, ny = nx, col = "lightgray", lty = "dotted",  
     lwd = NULL, equilog = TRUE)
```

**Arguments**

<code>nx,ny</code>	number of cells of the grid in x and y direction. When <code>NULL</code> , as per default, the grid aligns with the tick marks on the corresponding <i>default</i> axis (i.e., tick-marks as computed by <code>axTicks</code> ). When <code>NA</code> , no grid lines are drawn in the corresponding direction.
<code>col</code>	character or (integer) numeric; color of the grid lines.
<code>lty</code>	character or (integer) numeric; line type of the grid lines.
<code>lwd</code>	non-negative numeric giving line width of the grid lines; defaults to <code>par("lwd")</code> .
<code>equilog</code>	logical, only used when <i>log</i> coordinates and alignment with the axis tick marks are active. Setting <code>equilog = FALSE</code> in that case gives <i>non equidistant</i> tick aligned grid lines.

**Note**

If more fine tuning is required, use `abline(h = ., v = .)` directly.

**See Also**

`plot`, `abline`, `lines`, `points`.

## Examples

```
plot(1:3)
grid(NA, 5, lwd = 2) # grid only in y-direction

data(iris)
## maybe change the desired number of tick marks:
## par(lab=c(mx,my,7))
op <- par(mfcol = 1:2)
with(iris,
{
  plot(Sepal.Length, Sepal.Width,
       col = as.integer(Species),
       xlim = c(4, 8), ylim = c(2, 4.5), panel.first = grid(),
       main = "with(iris, plot(.., panel.first=grid(), ..))")
  plot(Sepal.Length, Sepal.Width,
       col = as.integer(Species),
       panel.first = grid(3, lty=1,lwd=2),
       main = "... panel.first = grid(3, lty=1,lwd=2), ..")
})
)
par(op)
```



---

**gtk**     *GTK+ Graphics Device*

---

**Description**

This is a graphics device similar to the X11 device but using the Gtk widgets. This is now available via a separate package - `gtkDevice` - and can be used independently of the GNOME GUI for R. This package also allows a device to be embedded within a Gtk-based GUI developed using the `RGtk` package. The `gtkDevice` package is available from CRAN.

**Usage**

```
gtk(display="", width=7, height=7, pointsize=12)
GTK(display="", width=7, height=7, pointsize=12)
```

**Arguments**

<code>display</code>	the display on which the graphics window will appear. The default is to use the value in the user's environment variable <code>DISPLAY</code> .
<code>width</code>	the width of the plotting window in inches.
<code>height</code>	the height of the plotting window in inches.
<code>pointsize</code>	the default pointsize to be used.

**Author(s)**

Original version by Lyndon Drake. Reorganization by Martyn Plummer and Duncan Temple Lang.

**References**

<http://www.gtk.org> for the GTK+ libraries.

**See Also**

`x11`, `Devices`.

---

**Hershey**     *Hershey Vector Fonts in R*

---

**Description**

If the `vfont` argument to one of the text-drawing functions (`text`, `mtext`, `title`, `axis`, and `contour`) is a character vector of length 2, Hershey vector fonts are used to render the text.

These fonts have two advantages:

1. vector fonts describe each character in terms of a set of points; R renders the character by joining up the points with straight lines. This intimate knowledge of the outline of each character means that R can arbitrarily transform the characters, which can mean that the vector fonts look better for rotated and 3d text.
2. this implementation was adapted from the GNU libplot library which provides support for non-ASCII and non-English fonts. This means that it is possible, for example, to produce weird plotting symbols and Japanese characters.

Drawback:

You cannot use mathematical expressions (`plotmath`) with Hershey fonts.

**Usage**

**Hershey**

**Details**

The Hershey characters are organised into a set of fonts, which are specified by a typeface (e.g., `serif` or `sans serif`) and a fontindex or “style” (e.g., `plain` or `italic`). The first element of `vfont` specifies the typeface and the second element specifies the fontindex. The first table produced by `demo(Hershey)` shows the character **a** produced by each of the different fonts.

The available `typeface` and `fontindex` values are available as list components of the variable `Hershey`. The allowed pairs for `(typeface, fontindex)` are:

serif

plain

serif	italic
serif	bold
serif	bold italic
serif	cyrillic
serif	oblique cyrillic
serif	EUC
sans serif	plain
sans serif	italic
sans serif	bold
sans serif	bold italic
script	plain
script	italic
script	bold
gothic english	plain
gothic german	plain
gothic italian	plain
serif symbol	plain
serif symbol	italic
serif symbol	bold
serif symbol	bold italic
sans serif symbol	plain
sans serif symbol	italic

and the indices of these are available as `Hershey$allowed`.

**Escape sequences:** The string to be drawn can include escape sequences, which all begin with a `\`. When R encounters a `\`, rather than drawing the `\`, it treats the subsequent character(s) as a coded description of what to draw.

One useful escape sequence (in the current context) is of the form: `\123`. The three digits following the `\` specify an octal code for a character. For example, the octal code for `p` is 160 so the strings `"p"` and `"\160"` are equivalent. This is useful for producing characters when there is not an appropriate key on your keyboard.

The other useful escape sequences all begin with `\\`. These are described below. Remember that backslashes have to be doubled in R character strings, so they need to be entered with *four* backslashes.

**Symbols:** an entire string of Greek symbols can be produced by selecting the Serif Symbol or Sans Serif Symbol typeface. To allow Greek symbols to be embedded in a string which uses a non-symbol typeface, there are a set of symbol escape sequences of the form `\\ab`. For example, the escape sequence `\\*a` produces a Greek al-

pha. The second table in `demo(Hershey)` shows all of the symbol escape sequences and the symbols that they produce.

**ISO Latin-1:** further escape sequences of the form `\\ab` are provided for producing ISO Latin-1 characters (for example, if you only have a US keyboard). Another option is to use the appropriate octal code. The (non-ASCII) ISO Latin-1 characters are in the range 241...377. For example, `\\366` produces the character `o` with an umlaut. The third table in `demo(Hershey)` shows all of the ISO Latin-1 escape sequences.

**Special Characters:** a set of characters are provided which do not fall into any standard font. These can only be accessed by escape sequence. For example, `\\LI` produces the zodiac sign for Libra, and `\\JU` produces the astronomical sign for Jupiter. The fourth table in `demo(Hershey)` shows all of the special character escape sequences.

**Cyrillic Characters:** cyrillic characters are implemented according to the K018-R encoding. On a US keyboard, these can be produced using the Serif typeface and Cyrillic (or Oblique Cyrillic) fontindex and specifying an octal code in the range 300 to 337 for lower case characters or 340 to 377 for upper case characters. The fifth table in `demo(Hershey)` shows the octal codes for the available cyrillic characters.

**Japanese Characters:** 83 Hiragana, 86 Katakana, and 603 Kanji characters are implemented according to the EUC (Extended Unix Code) encoding. Each character is identified by a unique hexadecimal code. The Hiragana characters are in the range 0x2421 to 0x2473, Katakana are in the range 0x2521 to 0x2576, and Kanji are (scattered about) in the range 0x3021 to 0x6d55.

When using the Serif typeface and EUC fontindex, these characters can be produced by a *pair* of octal codes. Given the hexadecimal code (e.g., 0x2421), take the first two digits and add 0x80 and do the same to the second two digits (e.g., 0x21 and 0x24 become 0xa4 and 0xa1), then convert both to octal (e.g., 0xa4 and 0xa1 become 244 and 241). For example, the first Hiragana character is produced by `\\244\\241`.

It is also possible to use the hexadecimal code directly. This works for all non-EUC fonts by specifying an escape sequence of the form `\\#J1234`. For example, the first Hiragana character is produced by `\\#J2421`.

The Kanji characters may be specified in a third way, using the so-called "Nelson Index", by specifying an escape sequence of the

form `\\#N1234`. For example, the Kanji for “one” is produced by `\\#N0001`.

`demo(Japanese)` shows the available Japanese characters.

**Raw Hershey Glyphs:** all of the characters in the Hershey fonts are stored in a large array. Some characters are not accessible in any of the Hershey fonts. These characters can only be accessed via an escape sequence of the form `\\#H1234`. For example, the fleur-de-lys is produced by `\\#H0746`. The sixth and seventh tables of `demo(Hershey)` shows all of the available raw glyphs.

## References

<http://www.gnu.org/software/plotutils/>

## See Also

`demo(Hershey)`, `text`, `contour`.

Japanese for the Japanese characters in the Hershey fonts.

## Examples

```
str(Hershey)
```

```
## for tables of examples, see demo(Hershey)
```

---

**hist**     *Histograms*


---

**Description**

The generic function **hist** computes a histogram of the given data values. If **plot=TRUE**, the resulting object of class "histogram" is plotted by **plot.histogram**, before it is returned.

**Usage**

```
hist(x, ...)

## Default S3 method:
hist(x, breaks = "Sturges", freq = NULL,
     probability = !freq,
     include.lowest = TRUE, right = TRUE,
     density = NULL, angle = 45, col = NULL, border = NULL,
     main = paste("Histogram of" , xname),
     xlim = range(breaks), ylim = NULL,
     xlab = xname, ylab,
     axes = TRUE, plot = TRUE, labels = FALSE,
     nclass = NULL, ...)
```

**Arguments**

<b>x</b>	a vector of values for which the histogram is desired.
<b>breaks</b>	<p>one of:</p> <ul style="list-style-type: none"> <li>• a vector giving the breakpoints between histogram cells,</li> <li>• a single number giving the number of cells for the histogram,</li> <li>• a character string naming an algorithm to compute the number of cells (see Details),</li> <li>• a function to compute the number of cells.</li> </ul> <p>In the last three cases the number is a suggestion only.</p>
<b>freq</b>	<p>logical; if <b>TRUE</b>, the histogram graphic is a representation of frequencies, the <b>counts</b> component of the result; if <b>FALSE</b>, <i>relative</i> frequencies ("probabilities"), component <b>density</b>, are plotted. Defaults to <b>TRUE</b></p>

	<i>iff</i> <b>breaks</b> are equidistant (and <b>probability</b> is not specified).
<b>probability</b>	an <i>alias</i> for <b>!freq</b> , for S compatibility.
<b>include.lowest</b>	logical; if <b>TRUE</b> , an <b>x[i]</b> equal to the <b>breaks</b> value will be included in the first (or last, for <b>right</b> = <b>FALSE</b> ) bar. This will be ignored (with a warning) unless <b>breaks</b> is a vector.
<b>right</b>	logical; if <b>TRUE</b> , the histograms cells are right-closed (left open) intervals.
<b>density</b>	the density of shading lines, in lines per inch. The default value of <b>NULL</b> means that no shading lines are drawn. Non-positive values of <b>density</b> also inhibit the drawing of shading lines.
<b>angle</b>	the slope of shading lines, given as an angle in degrees (counter-clockwise).
<b>col</b>	a colour to be used to fill the bars. The default of <b>NULL</b> yields unfilled bars.
<b>border</b>	the color of the border around the bars. The default is to use the standard foreground color.
<b>main, xlab, ylab</b>	these arguments to <b>title</b> have useful defaults here.
<b>xlim, ylim</b>	the range of x and y values with sensible defaults. Note that <b>xlim</b> is <i>not</i> used to define the histogram ( <b>breaks</b> ), but only for plotting (when <b>plot</b> = <b>TRUE</b> ).
<b>axes</b>	logical. If <b>TRUE</b> (default), axes are draw if the plot is drawn.
<b>plot</b>	logical. If <b>TRUE</b> (default), a histogram is plotted, otherwise a list of breaks and counts is returned.
<b>labels</b>	logical or character. Additionally draw labels on top of bars, if not <b>FALSE</b> ; see <b>plot.histogram</b> .
<b>nclass</b>	numeric (integer). For S(-PLUS) compatibility only, <b>nclass</b> is equivalent to <b>breaks</b> for a scalar or character argument.
<b>...</b>	further graphical parameters to <b>title</b> and <b>axis</b> .

## Details

The definition of “histogram” differs by source (with country-specific biases). R’s default with equi-spaced breaks (also the default) is to plot

the counts in the cells defined by **breaks**. Thus the height of a rectangle is proportional to the number of points falling into the cell, as is the area *provided* the breaks are equally-spaced.

The default with non-equi-spaced breaks is to give a plot of area one, in which the *area* of the rectangles is the fraction of the data points falling in the cells.

If **right** = **TRUE** (default), the histogram cells are intervals of the form **(a, b]**, i.e., they include their right-hand endpoint, but not their left one, with the exception of the first cell when **include.lowest** is **TRUE**.

For **right** = **FALSE**, the intervals are of the form **[a, b)**, and **include.lowest** really has the meaning of “*include highest*”.

A numerical tolerance of  $10^{-7}$  times the range of the breaks is applied when counting entries on the edges of bins.

The default for **breaks** is “**Sturges**”: see **nclass.Sturges**. Other names for which algorithms are supplied are “**Scott**” and “**FD**” / “**Friedman-Diaconis**” (with corresponding functions **nclass.scott** and **nclass.FD**). Case is ignored and partial matching is used. Alternatively, a function can be supplied which will compute the intended number of breaks as a function of **x**.

## Value

an object of class “**histogram**” which is a list with components:

<b>breaks</b>	the $n + 1$ cell boundaries (= <b>breaks</b> if that was a vector).
<b>counts</b>	$n$ integers; for each cell, the number of <b>x</b> [] inside.
<b>density</b>	values $\hat{f}(x_i)$ , as estimated density values. If <b>all(diff(breaks) == 1)</b> , they are the relative frequencies <b>counts</b> / <b>n</b> and in general satisfy $\sum_i \hat{f}(x_i)(b_{i+1} - b_i) = 1$ , where $b_i = \mathbf{breaks}[i]$ .
<b>intensities</b>	same as <b>density</b> . Deprecated, but retained for compatibility.
<b>mids</b>	the $n$ cell midpoints.
<b>xname</b>	a character string with the actual <b>x</b> argument name.
<b>equidist</b>	logical, indicating if the distances between <b>breaks</b> are all the same.



## Note

The resulting value does *not* depend on the values of the arguments `freq` (or `probability`) or `plot`. This is intentionally different from S. Prior to R 1.7.0, the element `breaks` of the result was adjusted for numerical tolerances. The nominal values are now returned even though tolerances are still used when counting.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Springer.

## See Also

`nclass.Sturges`, `stem`, `density`, `truehist`.

## Examples

```
data(islands)
op <- par(mfrow=c(2, 2))
hist(islands)
str(hist(islands, col="gray", labels = TRUE))
hist(sqrt(islands), br = 12, col="lightblue", border="pink")
# For non-equidistant breaks, counts should NOT be
# graphed unscaled:
r <- hist(sqrt(islands),
          br = c(4*0:5, 10*3:5, 70, 100, 140),
          col='blue1')
text(r$mids, r$density, r$counts, adj=c(.5, -.5),
     col='blue3')
sapply(r[2:3], sum)
sum(r$density * diff(r$breaks)) # == 1
lines(r, lty = 3, border = "purple") # lines.histogram(*)
par(op)

str(hist(islands, plot= FALSE))          # 5 breaks
str(hist(islands, br=12, plot= FALSE)) # 10 (~= 12) breaks
str(hist(islands, br=c(12,20,36,80,200,1000,17000),
          plot = FALSE))
hist(islands, br=c(12,20,36,80,200,1000,17000),
     freq = TRUE, main = "WRONG histogram") # and warning
```

---

**hsv**     *HSV Color Specification*

---

**Description**

Create a vector of colors from vectors specifying hue, saturation and value.

**Usage**

```
hsv(h=1, s=1, v=1, gamma=1)
```

**Arguments**

<b>h,s,v</b>	numeric vectors of values in the range $[0,1]$ for “hue”, “saturation” and “value” to be combined to form a vector of colors. Values in shorter arguments are recycled.
<b>gamma</b>	a “gamma correction”, $\gamma$

**Value**

This function creates a vector of “colors” corresponding to the given values in HSV space. The values returned by **hsv** can be used with a **col=** specification in graphics functions or in **par**.

**Gamma correction**

For each color,  $(r, g, b)$  in RGB space (with all values in  $[0, 1]$ ), the final color corresponds to  $(r^\gamma, g^\gamma, b^\gamma)$ .

**See Also**

**rainbow**, **rgb**, **gray**.

**Examples**

```
hsv(.5,.5,.5)

## Look at gamma effect:
n <- 20; y <- -sin(3*pi*((1:n)-1/2)/n)
op <- par(mfrow=c(3,2),mar=rep(1.5,4))
for(gamma in c(.4, .6, .8, 1, 1.2, 1.5))
```

```
plot(y, axes = FALSE, frame.plot = TRUE,
      xlab = "", ylab = "", pch = 21, cex = 30,
      bg = rainbow(n, start=.85, end=.1, gamma = gamma),
      main = paste("Red tones; gamma=",format(gamma)))
par(op)
```

---

**identify**      *Identify Points in a Scatter Plot*


---

**Description**

**identify** reads the position of the graphics pointer when the (first) mouse button is pressed. It then searches the coordinates given in **x** and **y** for the point closest to the pointer. If this point is close to the pointer, its index will be returned as part of the value of the call.

**Usage**

```
identify(x, ...)
```

```
## Default S3 method:
```

```
identify(x, y = NULL, labels = seq(along = x), pos = FALSE,
        n = length(x), plot = TRUE, offset = 0.5, ...)
```

**Arguments**

<b>x,y</b>	coordinates of points in a scatter plot. Alternatively, any object which defines coordinates (a plotting structure, time series etc.) can be given as <b>x</b> and <b>y</b> left undefined.
<b>labels</b>	an optional vector, the same length as <b>x</b> and <b>y</b> , giving labels for the points.
<b>pos</b>	if <b>pos</b> is <b>TRUE</b> , a component is added to the return value which indicates where text was plotted relative to each identified point (1=below, 2=left, 3=above and 4=right).
<b>n</b>	the maximum number of points to be identified.
<b>plot</b>	if <b>plot</b> is <b>TRUE</b> , the labels are printed at the points and if <b>FALSE</b> they are omitted.
<b>offset</b>	the distance (in character widths) which separates the label from identified points.
<b>...</b>	further arguments to <b>par()</b> .

## Details

If in addition, `plot` is `TRUE`, the point is labelled with the corresponding element of `text`.

The labels are placed either below, to the left, above or to the right of the identified point, depending on where the cursor was.

The identification process is terminated by pressing any mouse button other than the first.

On most devices which support `locator`, successful selection of a point is indicated by a bell sound unless `options(locatorBell=FALSE)`

## Value

If `pos` is `FALSE`, an integer vector containing the indexes of the identified points.

If `pos` is `TRUE`, a list containing a component `ind`, indicating which points were identified and a component `pos`, indicating where the labels were placed relative to the identified points.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`locator`

---

**image**     *Display a Color Image*


---

## Description

Creates a grid of colored or gray-scale rectangles with colors corresponding to the values in **z**. This can be used to display three-dimensional or spatial data aka “images”. This is a generic function.

The functions `heat.colors`, `terrain.colors` and `topo.colors` create heat-spectrum (red to white) and topographical color schemes suitable for displaying ordered data, with **n** giving the number of colors desired.

## Usage

```
image(x, ...)
```

```
## Default S3 method:
```

```
image(x, y, z, zlim, xlim, ylim, col = heat.colors(12),
      add = FALSE, xaxs = "i", yaxs = "i", xlab, ylab,
      breaks, oldstyle = FALSE, ...)
```

## Arguments

- |                   |  |
|-------------------|--|
| <b>x,y</b>        | locations of grid lines at which the values in <b>z</b> are measured. These must be in (strictly) ascending order. By default, equally spaced values from 0 to 1 are used. If <b>x</b> is a <b>list</b> , its components <b>x\$x</b> and <b>x\$y</b> are used for <b>x</b> and <b>y</b> , respectively. If the list has component <b>z</b> this is used for <b>z</b> . |
| <b>z</b>          | a matrix containing the values to be plotted (NAs are allowed). Note that <b>x</b> can be used instead of <b>z</b> for convenience.  |
| <b>zlim</b>       | the minimum and maximum <b>z</b> values for which colors should be plotted. Each of the given colors will be used to color an equispaced interval of this range. The <i>midpoints</i> of the intervals cover the range, so that values just outside the range will be plotted.   |
| <b>xlim, ylim</b> | ranges for the plotted <b>x</b> and <b>y</b> values, defaulting to the range of the finite values of <b>x</b> and <b>y</b> .   |

<code>col</code>	a list of colors such as that generated by <code>rainbow</code> , <code>heat.colors</code> , <code>topo.colors</code> , <code>terrain.colors</code> or similar functions.
<code>add</code>	logical; if <code>TRUE</code> , add to current plot (and disregard the following arguments). This is rarely useful because <code>image</code> “paints” over existing graphics.
<code>xaxs, yaxs</code>	style of x and y axis. The default “i” is appropriate for images. See <code>par</code> .
<code>xlab, ylab</code>	each a character string giving the labels for the x and y axis. Default to the ‘call names’ of x or y, or to “” if these were unspecified.
<code>breaks</code>	a set of breakpoints for the colours: must give one more breakpoint than colour.
<code>oldstyle</code>	logical. If true the midpoints of the colour intervals are equally spaced, and <code>zlim[1]</code> and <code>zlim[2]</code> were taken to be midpoints. (This was the default prior to R 1.1.0.) The current default is to have colour intervals of equal lengths between the limits.
<code>...</code>	graphical parameters for <code>plot</code> may also be passed as arguments to this function.

## Details

The length of `x` should be equal to the `nrow(z)+1` or `nrow(z)`. In the first case `x` specifies the boundaries between the cells: in the second case `x` specifies the midpoints of the cells. Similar reasoning applies to `y`. It probably only makes sense to specify the midpoints of an equally-spaced grid. If you specify just one row or column and a length-one `x` or `y`, the whole user area in the corresponding direction is filled.

If `breaks` is specified then `zlim` is unused and the algorithm used follows `cut`, so intervals are closed on the right and open on the left except for the lowest interval.

## Note

Based on a function by Thomas Lumley.

## See Also

`filled.contour` or `heatmap` which can look nicer (but are less modular), `contour`;

`heat.colors`, `topo.colors`, `terrain.colors`, `rainbow`, `hsv`, `par`.

## Examples

```
x <- y <- seq(-4*pi, 4*pi, len=27)
r <- sqrt(outer(x^2, y^2, "+"))
image(z = z <- cos(r^2)*exp(-r/6), col=gray((0:32)/32))
image(z, axes = FALSE, main = "Math can be beautiful ...",
      xlab = expression(cos(r^2) * e^{-r/6}))
contour(z, add = TRUE, drawlabels = FALSE)

data(volcano)
x <- 10*(1:nrow(volcano))
y <- 10*(1:ncol(volcano))
image(x, y, volcano, col = terrain.colors(100),
      axes = FALSE)
contour(x, y, volcano, levels = seq(90, 200, by=5),
      add = TRUE, col = "peru")
axis(1, at = seq(100, 800, by = 100))
axis(2, at = seq(100, 600, by = 100))
box()
title(main = "Maunga Whau Volcano", font.main = 4)
```



---

**interaction.plot**      *Two-way Interaction Plot*


---

**Description**

Plots the mean (or other summary) of the response for two-way combinations of factors, thereby illustrating possible interactions.

**Usage**

```
interaction.plot(x.factor, trace.factor, response,
  fun = mean, type = c("l", "p"), legend = TRUE,
  trace.label=deparse(substitute(trace.factor)),
  fixed=FALSE, xlab = deparse(substitute(x.factor)),
  ylab = ylabel, ylim = range(cells, na.rm=TRUE),
  lty = nc:1, col = 1, pch = c(1:9, 0, letters),
  xpd = NULL, leg.bg = par("bg"), leg.bty = "n",
  xtick = FALSE, xaxt = par("xaxt"), axes = TRUE, ...)
```

**Arguments**

<b>x.factor</b>	a factor whose levels will form the x axis.
<b>trace.factor</b>	another factor whose levels will form the traces.
<b>response</b>	a numeric variable giving the response
<b>fun</b>	the function to compute the summary. Should return a single real value.
<b>type</b>	the type of plot: lines or points.
<b>legend</b>	logical. Should a legend be included?
<b>trace.label</b>	overall label for the legend.
<b>fixed</b>	logical. Should the legend be in the order of the levels of <b>trace.factor</b> or in the order of the traces at their right-hand ends?
<b>xlab,ylab</b>	the x and y label of the plot each with a sensible default.
<b>ylim</b>	numeric of length 2 giving the y limits for the plot.
<b>lty</b>	line type for the lines drawn, with sensible default.
<b>col</b>	the color to be used for plotting.
<b>pch</b>	a vector of plotting symbols or characters, with sensible default.

**xpd** determines clipping behaviour for the **legend** used, see **par(xpd)**. Per default, the legend is *not* clipped at the figure border.

**leg.bg, leg.bty** arguments passed to **legend()**.

**xtick** logical. Should tick marks be used on the x axis?

**xaxt, axes, ...** graphics parameters to be passed to the plotting routines.

## Details

By default the levels of **x.factor** are plotted on the x axis in their given order, with extra space left at the right for the legend (if specified). If **x.factor** is an ordered factor and the levels are numeric, these numeric values are used for the x axis.

The response and hence its summary can contain missing values. If so, the missing values and the line segments joining them are omitted from the plot (and this can be somewhat disconcerting).

The graphics parameters **xlab**, **ylab**, **ylim**, **lty**, **col** and **pch** are given suitable defaults (and **xlim** and **xaxs** are set and cannot be overridden). The defaults are to cycle through the line types, use the foreground colour, and to use the symbols 1:9, 0, and the capital letters to plot the traces.

## Note

Some of the argument names and the precise behaviour are chosen for S-compatibility.

## References

Chambers, J. M., Freeny, A and Heiberger, R. M. (1992) *Analysis of variance; designed experiments*. Chapter 5 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## Examples

```
data(ToothGrowth)
attach(ToothGrowth)
interaction.plot(dose, supp, len, fixed=TRUE)
dose <- ordered(dose)
interaction.plot(dose, supp, len, fixed=TRUE, col = 2:3,
```

```
leg.bty = "o")
detach()

data(OrchardSprays)
with(OrchardSprays, {
  interaction.plot(treatment, rowpos, decrease)
  interaction.plot(rowpos, treatment, decrease,
                  cex.axis=0.8)
  ## order the rows by their mean effect
  rowpos <- factor(rowpos, levels = sort.list(tapply(
    decrease, rowpos, mean)))
  interaction.plot(rowpos, treatment, decrease, col = 2:9,
                  lty = 1)
})

data(esoph)
with(esoph, {
  interaction.plot(agegp, alcgp, ncases/ncontrols)
  interaction.plot(agegp, tobgp, ncases/ncontrols,
                  trace.label="tobacco",
                  fixed=TRUE, xaxt = "n")
})
```

---

## Japanese *Japanese characters in R*

---

### Description

The implementation of Hershey vector fonts provides a large number of Japanese characters (Hiragana, Katakana, and Kanji).

### Details

Without keyboard support for typing Japanese characters, the only way to produce these characters is to use special escape sequences: see *Hershey*.

For example, the Hiragana character for the sound "ka" is produced by `\\#J242b` and the Katakana character for this sound is produced by `\\#J252b`. The Kanji ideograph for "one" is produced by `\\#J306c` or `\\#N0001`.

The output from `demo(Japanese)` shows tables of the escape sequences for the available Japanese characters.

### References

<http://www.gnu.org/software/plotutils/>

### See Also

`demo(Japanese)`, `Hershey`, `text`, `contour`

### Examples

```
plot(1:9, type="n", axes=FALSE, frame=TRUE, ylab="",
     main="example(Japanese)", xlab="using Hershey fonts")
par(cex=3)
Vf <- c("serif", "plain")
text(4, 2, "\\#J2438\\#J2421\\#J2451\\#J2473", vfont = Vf)
text(4, 4, "\\#J2538\\#J2521\\#J2551\\#J2573", vfont = Vf)
text(4, 6, "\\#J467c\\#J4b5c", vfont = Vf)
text(4, 8, "Japan", vfont = Vf)
par(cex=1)
text(8, 2, "Hiragana")
text(8, 4, "Katakana")
text(8, 6, "Kanji")
text(8, 8, "English")
```

---

**jitter**     *Add ‘Jitter’ (Noise) to Numbers*

---

## Description

Add a small amount of noise to a numeric vector.

## Usage

```
jitter(x, factor=1, amount = NULL)
```

## Arguments

<b>x</b>	numeric to which <i>jitter</i> should be added.
<b>factor</b>	numeric
<b>amount</b>	numeric; if positive, used as <i>amount</i> (see below), otherwise, if = 0 the default is <code>factor * z/50</code> . Default (NULL): <code>factor * d/5</code> where <i>d</i> is about the smallest difference between <i>x</i> values.

## Details

The result, say *r*, is `r <- x + runif(n, -a, a)` where `n <- length(x)` and *a* is the `amount` argument (if specified).

Let `z <- max(x) - min(x)` (assuming the usual case). The amount *a* to be added is either provided as *positive* argument `amount` or otherwise computed from *z*, as follows:

If `amount == 0`, we set `a <- factor * z/50` (same as *S*).

If `amount` is NULL (*default*), we set `a <- factor * d/5` where *d* is the smallest difference between adjacent unique (apart from fuzz) *x* values.

## Value

`jitter(x, ...)` returns a numeric of the same length as *x*, but with an `amount` of noise added in order to break ties.

## Author(s)

Werner Stahel and Martin Maechler, ETH Zurich

## References

- Chambers, J. M., Cleveland, W. S., Kleiner, B. and Tukey, P.A. (1983) *Graphical Methods for Data Analysis*. Wadsworth; figures 2.8, 4.22, 5.4.
- Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

## See Also

`rug` which you may want to combine with `jitter`.

## Examples

```
round(jitter(c(rep(1,3), rep(1.2, 4), rep(3,3))), 3)
## These two 'fail' with S-plus 3.x:
jitter(rep(0, 7))
jitter(rep(10000,5))
```

---

**layout**      *Specifying Complex Plot Arrangements*

---

**Description**

`layout` divides the device up into as many rows and columns as there are in matrix `mat`, with the column-widths and the row-heights specified in the respective arguments.

**Usage**

```
layout(mat,  
       widths = rep(1, dim(mat)[2]),  
       heights= rep(1, dim(mat)[1]),  
       respect= FALSE)
```

```
layout.show(n = 1)  
lcm(x)
```

**Arguments**

- |                      |   |
|----------------------|---|
| <code>mat</code>     | a matrix object specifying the location of the next $N$ figures on the output device. Each value in the matrix must be 0 or a positive integer. If $N$ is the largest positive integer in the matrix, then the integers $\{1, \dots, N - 1\}$ must also appear at least once in the matrix. |
| <code>widths</code>  | a vector of values for the widths of columns on the device. Relative widths are specified with numeric values. Absolute widths (in centimetres) are specified with the <code>lcm()</code> function (see examples).  |
| <code>heights</code> | a vector of values for the heights of rows on the device. Relative and absolute heights can be specified, see <code>widths</code> above.  |
| <code>respect</code> | either a logical value or a matrix object. If the latter, then it must have the same dimensions as <code>mat</code> and each value in the matrix must be either 0 or 1.   |
| <code>n</code>       | number of figures to plot.  |
| <code>x</code>       | a dimension to be interpreted as a number of centimetres.   |

## Details

Figure  $i$  is allocated a region composed from a subset of these rows and columns, based on the rows and columns in which  $i$  occurs in `mat`.

The `respect` argument controls whether a unit column-width is the same physical measurement on the device as a unit row-height.

`layout.show(n)` plots (part of) the current layout, namely the outlines of the next `n` figures.

`1cm` is a trivial function, to be used as *the* interface for specifying absolute dimensions for the `widths` and `heights` arguments of `layout()`.

## Value

`layout` returns the number of figures,  $N$ , see above.

## Author(s)

Paul R. Murrell

## References

Murrell, P. R. (1999) Layouts: A mechanism for arranging plots on a page. *Journal of Computational and Graphical Statistics*, **8**, 121-134. Chapter 5 of Paul Murrell's Ph.D. thesis.

## See Also

`par` with arguments `mfrow`, `mfcol`, or `mfg`.

## Examples

```
# save default, for resetting...
def.par <- par(no.readonly = TRUE)

## divide the device into two rows and two columns allocate
## figure 1 all of row 1 allocate figure 2 the intersection
## of column 2 and row 2
layout(matrix(c(1,1,0,2), 2, 2, byrow = TRUE))
## show the regions that have been allocated to each plot
layout.show(2)

## divide device into two rows and two columns allocate
## figure 1 and figure 2 as above respect relations between
## widths and heights
```



```
nf <- layout(matrix(c(1,1,0,2), 2, 2, byrow=TRUE),
                 respect=TRUE)
layout.show(nf)

## create single figure which is 5cm square
nf <- layout(matrix(1), widths=lcm(5), heights=lcm(5))
layout.show(nf)

## Create a scatterplot with marginal histograms
x <- pmin(3, pmax(-3, rnorm(50)))
y <- pmin(3, pmax(-3, rnorm(50)))
xhist <- hist(x, breaks=seq(-3,3,0.5), plot=FALSE)
yhist <- hist(y, breaks=seq(-3,3,0.5), plot=FALSE)
top <- max(c(xhist$counts, yhist$counts))
xrange <- c(-3,3)
yrange <- c(-3,3)
nf <- layout(matrix(c(2,0,1,3),2,2,byrow=TRUE), c(3,1),
                  c(1,3), TRUE)
layout.show(nf)

par(mar=c(3,3,1,1))
plot(x, y, xlim=xrange, ylim=yrange, xlab="", ylab="")
par(mar=c(0,3,1,1))
barplot(xhist$counts, axes=FALSE, ylim=c(0, top), space=0)
par(mar=c(3,0,1,1))
barplot(yhist$counts, axes=FALSE, xlim=c(0, top), space=0,
        horiz=TRUE)

par(def.par) # reset to default
```

---

**legend**     *Add Legends to Plots*


---

**Description**

This function can be used to add legends to plots. Note that a call to the function `locator` can be used in place of the `x` and `y` arguments.

**Usage**

```
legend(x, y = NULL, legend, fill = NULL, col = "black",
      lty, lwd, pch, angle = NULL, density = NULL, bty = "o",
      bg = par("bg"), pt.bg = NA, cex = 1, xjust = 0, yjust = 1,
      x.intersp = 1, y.intersp = 1, adj = c(0, 0.5),
      text.width = NULL, merge = do.lines && has.pch,
      trace = FALSE, plot = TRUE, ncol = 1, horiz = FALSE)
```

**Arguments**

<code>x, y</code>	the <code>x</code> and <code>y</code> co-ordinates to be used to position the legend. They can be specified in any way which is accepted by <code>xy.coords</code> : See Details.
<code>legend</code>	a vector of text values or an <b>expression</b> of length $\geq 1$ , or a <b>call</b> (as resulting from <code>substitute</code> ) to appear in the legend.
<code>fill</code>	if specified, this argument will cause boxes filled with the specified colors (or shaded in the specified colors) to appear beside the legend text.
<code>col</code>	the color of points or lines appearing in the legend.
<code>lty, lwd</code>	the line types and widths for lines appearing in the legend. One of these two <i>must</i> be specified for line drawing.
<code>pch</code>	the plotting symbols appearing in the legend, either as vector of 1-character strings, or one (multi character) string. <i>Must</i> be specified for symbol drawing.
<code>angle</code>	angle of shading lines.
<code>density</code>	the density of shading lines, if numeric and positive. If <code>NULL</code> or negative or <code>NA</code> color filling is assumed.
<code>bty</code>	the type of box to be drawn around the legend. The allowed values are <code>"o"</code> (the default) and <code>"n"</code> .

<code>bg</code>	the background color for the legend box. (Note that this is only used if <code>bty = "n"</code> .)
<code>pt.bg</code>	the background color for the points.
<code>cex</code>	character expansion factor <b>relative</b> to current <code>par("cex")</code> .
<code>xjust</code>	how the legend is to be justified relative to the legend x location. A value of 0 means left justified, 0.5 means centered and 1 means right justified.
<code>yjust</code>	the same as <code>xjust</code> for the legend y location.
<code>x.intersp</code>	character interspacing factor for horizontal (x) spacing.
<code>y.intersp</code>	the same for vertical (y) line distances.
<code>adj</code>	numeric of length 1 or 2; the string adjustment for legend text. Useful for y-adjustment when <code>labels</code> are plotmath expressions.
<code>text.width</code>	the width of the legend text in x ("user") coordinates. Defaults to the proper value computed by <code>strwidth(legend)</code> .
<code>merge</code>	logical; if <code>TRUE</code> , "merge" points and lines but not filled boxes. Defaults to <code>TRUE</code> if there are points and lines.
<code>trace</code>	logical; if <code>TRUE</code> , shows how <code>legend</code> does all its magical computations.
<code>plot</code>	logical. If <code>FALSE</code> , nothing is plotted but the sizes are returned.
<code>ncol</code>	the number of columns in which to set the legend items (default is 1, a vertical legend).
<code>horiz</code>	logical; if <code>TRUE</code> , set the legend horizontally rather than vertically (specifying <code>horiz</code> overrides the <code>ncol</code> specification).

## Details

Arguments `x`, `y`, `legend` are interpreted in a non-standard way to allow the coordinates to be specified *via* one or two arguments. If `legend` is missing and `y` is not numeric, it is assumed that the second argument is intended to be `legend` and that the first argument specifies the coordinates.

The coordinates can be specified in any way which is accepted by `xy.coords`. If this gives the coordinates of one point, it is used as the

top-left coordinate of the rectangle containing the legend. If it gives the coordinates of two points, these specify opposite corners of the rectangle (either pair of corners, in any order).

“Attribute” arguments such as `col`, `pch`, `lty`, etc, are recycled if necessary. `merge` is not.

Points are drawn *after* lines in order that they can cover the line with their background color `pt.bg`, if applicable.

See the examples for how to right-justify labels.

## Value

A list with list components

<code>rect</code>	a list with components
	<code>w</code> , <code>h</code> positive numbers giving <b>w</b> idth and <b>h</b> eight of the legend's box.
	<code>left</code> , <code>top</code> x and y coordinates of upper left corner of the box.
<code>text</code>	a list with components
	<code>x</code> , <code>y</code> numeric vectors of length <code>length(legend)</code> , giving the x and y coordinates of the legend's text(s).

returned invisibly.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`plot`, `barplot` which uses `legend()`, and `text` for more examples of math expressions.

## Examples

```
## Run the example in '?matplot' or the following:
leg.txt <- c("Setosa      Petals", "Setosa      Sepals",
            "Versicolor Petals", "Versicolor Sepals")
y.leg <- c(4.5, 3, 2.1, 1.4, .7)
cexv  <- c(1.2, 1, 4/5, 2/3, 1/2)
matplot(c(1,8), c(0,4.5), type = "n",
        xlab = "Length", ylab = "Width",
```

```

    main = "Petal and Sepal Dimensions in Iris Blossoms")
  for (i in seq(cexv)) {
    text(1, y.leg[i]-.1, paste("cex=",formatC(cexv[i])),
        cex=.8, adj = 0)
    legend(3, y.leg[i], leg.txt, pch = "sSvV", col = c(1, 3),
        cex = cexv[i])
  }

## 'merge = TRUE' for merging lines & points:
x <- seq(-pi, pi, len = 65)
plot(x, sin(x), type = "l", ylim = c(-1.2, 1.8), col = 3,
     lty = 2)
points(x, cos(x), pch = 3, col = 4)
lines(x, tan(x), type = "b", lty = 1, pch = 4, col = 6)
title("legend(...,lty=c(2,-1,1),pch=c(-1,3,4),merge=TRUE)",
     cex.main = 1.1)
legend(-1, 1.9, c("sin", "cos", "tan"), col = c(3,4,6),
     lty = c(2, -1, 1), pch = c(-1, 3, 4), merge = TRUE,
     bg='gray90')

## right-justifying a set of labels: thanks to Uwe Ligges
x <- 1:5; y1 <- 1/x; y2 <- 2/x
plot(rep(x, 2), c(y1, y2), type="n", xlab="x", ylab="y")
lines(x, y1); lines(x, y2, lty=2)
temp <- legend(5, 2, legend = c(" ", " "),
     text.width = strwidth("1,000,000"),
     lty = 1:2, xjust = 1, yjust = 1)
text(temp$rect$left + temp$rect$w, temp$text$y,
     c("1,000", "1,000,000"), pos=2)

## log scaled Examples
leg.txt <- c("a one", "a two")

par(mfrow = c(2,2))
for(ll in c("", "x", "y", "xy")) {
  plot(2:10, log=ll, main=paste("log = '",ll,"'", sep=""))
  abline(1,1)
  lines(2:3,3:4, col=2) #
  points(2,2, col=3)   #
  rect(2,3,3,2, col=4)
  text(c(3,3),2:3, c("rect(2,3,3,2, col=4)",
    "text(c(3,3),2:3,\"c(rect(...)\")\"", adj = c(0,.3))
  legend(list(x=2,y=8), legend = leg.txt, col=2:3, pch=1:2,

```

```

        lty=1, merge=TRUE) #, trace=TRUE)
}
par(mfrow=c(1,1))

## Math expressions:
x <- seq(-pi, pi, len = 65)
plot(x, sin(x), type="l", col = 2,
      xlab = expression(phi), ylab = expression(f(phi)))
abline(h=-1:1, v=pi/2*(-6:6), col="gray90")
lines(x, cos(x), col = 3, lty = 2)
# 2 ways
ex.cs1 <- expression(plain(sin) * phi, paste("cos", phi))
# adj y !
str(legend(-3,.9,ex.cs1,lty=1:2,plot=FALSE,adj=c(0,.6)))
legend(-3, .9, ex.cs1, lty=1:2, col=2:3, adj = c(0, .6))

x <- rexp(100, rate = .5)
hist(x, main = "Mean and Median of a Skewed Distribution")
abline(v = mean(x), col=2, lty=2, lwd=2)
abline(v = median(x), col=3, lty=3, lwd=2)
ex12 <- expression(bar(x) == sum(over(x[i], n), i==1, n),
                    hat(x) == median(x[i], i==1,n))
str(legend(4.1, 30, ex12, col = 2:3, lty=2:3, lwd=2))

## 'Filled' boxes -- for more, see example(plotfactor)
op <- par(bg="white") # to get an opaque box for the legend
data(PlantGrowth)
plot(cut(weight, 3) ~ group, data = PlantGrowth,
      col = NULL, density = 16*(1:3))
par(op)

## Using 'ncol' :
x <- 0:64/64
matplot(x, outer(x, 1:7, function(x, k) sin(k * pi * x)),
        type = "o", col = 1:7, ylim = c(-1, 1.5), pch = "*")
op <- par(bg="antiquewhite1")
legend(0, 1.5, paste("sin(",1:7,"pi * x)"), col=1:7,
      lty=1:7, pch = "*", ncol = 4, cex=.8)
legend(.8,1.2, paste("sin(",1:7,"pi * x)"), col=1:7,
      lty=1:7, pch = "*",cex=.8)
legend(0, -.1, paste("sin(",1:4,"pi * x)"), col=1:4,
      lty=1:4, ncol=2, cex=.8)
legend(0, -.4, paste("sin(",5:7,"pi * x)"), col=5:7,

```

```
      pch=24, ncol=2, cex=1.5, pt.bg="pink")
par(op)

## point covering line :
y <- sin(3*pi*x)
plot(x, y, type="l", col="blue",
      main = "points with bg & legend(*, pt.bg)")
points(x, y, pch=21, bg="white")
legend(.4,1, "sin(c x)", pch=21, pt.bg="white", lty=1,
       col = "blue")
```

---

<b>lines</b>	<i>Add Connected Line Segments to a Plot</i>
--------------	--

---

## Description

A generic function taking coordinates given in various ways and joining the corresponding points with line segments.

## Usage

```
lines(x, ...)  
  
## Default S3 method:  
lines(x, y = NULL, type = "l", col = par("col"),  
      lty = par("lty"), ...)
```

## Arguments

<b>x, y</b>	coordinate vectors of points to join.
<b>type</b>	character indicating the type of plotting; actually any of the <b>types</b> as in <b>plot</b> .
<b>col</b>	color to use. This can be vector of length greater than one, but only the first value will be used.
<b>lty</b>	line type to use.
<b>...</b>	Further graphical parameters (see <b>par</b> ) may also be supplied as arguments, particularly, line type, <b>lty</b> and line width, <b>lwd</b> .

## Details

The coordinates can be passed to **lines** in a plotting structure (a list with **x** and **y** components), a time series, etc. See **xy.coords**.

The coordinates can contain **NA** values. If a point contains **NA** in either its **x** or **y** value, it is omitted from the plot, and lines are not drawn to or from such points. Thus missing values can be used to achieve breaks in lines.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.



**See Also**

`points`, `plot`, and the underlying “primitive” `plot.xy`.  
`par` for how to specify colors.

**Examples**

```
data(cars)
# draw a smooth line through a scatter plot
plot(cars, main="Stopping Distance versus Speed")
lines(lowess(cars))
```

---

**locator**      *Graphical Input*


---

## Description

Reads the position of the graphics cursor when the (first) mouse button is pressed.

## Usage

```
locator(n = 512, type = "n", ...)
```

## Arguments

<b>n</b>	the maximum number of points to locate.
<b>type</b>	One of "n", "p", "l" or "o". If "p" or "o" the points are plotted; if "l" or "o" they are joined by lines.
<b>...</b>	additional graphics parameters used if <b>type</b> != "n" for plotting the locations.

## Details

Unless the process is terminated prematurely by the user (see below) at most **n** positions are determined.

The identification process can be terminated by pressing any mouse button other than the first.

The current graphics parameters apply just as if `plot.default` has been called with the same value of **type**. The plotting of the points and lines is subject to clipping, but locations outside the current clipping rectangle will be returned.

On most devices which support **locator**, successful selection of a point is indicated by a bell sound unless `options(locatorBell=FALSE)` has been set.

If the window is resized or hidden and then exposed before the input process has terminated, any lines or points drawn by **locator** will disappear. These will reappear once the input process has terminated and the window is resized or hidden and exposed again. This is because the points and lines drawn by **locator** are not recorded in the device's display list until the input process has terminated.

**Value**

A list containing **x** and **y** components which are the coordinates of the identified points in the user coordinate system, i.e., the one specified by `par("usr")`.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`identify`

---

**matplot**     *Plot Columns of Matrices*

---

**Description**

Plot the columns of one matrix against the columns of another.

**Usage**

```
matplot(x, y, type = "p", lty = 1:5, lwd = 1, pch = NULL,
        col = 1:6, cex = NULL, xlab = NULL, ylab = NULL,
        xlim = NULL, ylim = NULL, ...,
        add = FALSE, verbose = getOption("verbose"))
matpoints(x, y, type = "p", lty = 1:5, lwd = 1, pch = NULL,
          col = 1:6, ...)
matlines (x, y, type = "l", lty = 1:5, lwd = 1, pch = NULL,
          col = 1:6, ...)
```

**Arguments**

- |                |  |
|----------------|--|
| <b>x,y</b>     | vectors or matrices of data for plotting. The number of rows should match. If one of them are missing, the other is taken as <b>y</b> and an <b>x</b> vector of <b>1:n</b> is used. Missing values (NAs) are allowed.  |
| <b>type</b>    | character string (length 1 vector) or vector of 1-character strings indicating the type of plot for each column of <b>y</b> , see <b>plot</b> for all possible <b>types</b> . The first character of <b>type</b> defines the first plot, the second character the second, etc. Characters in <b>type</b> are cycled through; e.g., " <b>p1</b> " alternately plots points and lines. |
| <b>lty,lwd</b> | vector of line types and widths. The first element is for the first column, the second element for the second column, etc., even if lines are not plotted for all columns. Line types will be used cyclically until all plots are drawn.   |
| <b>pch</b>     | character string or vector of 1-characters or integers for plotting characters, see <b>points</b> . The first character is the plotting-character for the first plot, the second for the second, etc. The default is the digits (1 through 9, 0) then the letters.   |

<code>col</code>	vector of colors. Colors are used cyclically.
<code>cex</code>	vector of character expansion sizes, used cyclically.
<code>xlab, ylab</code>	titles for x and y axes, as in <code>plot</code> .
<code>xlim, ylim</code>	ranges of x and y axes, as in <code>plot</code> .
<code>...</code>	Graphical parameters (see <code>par</code> ) and any further arguments of <code>plot</code> , typically <code>plot.default</code> , may also be supplied as arguments to this function. Hence, the high-level graphics control arguments described under <code>par</code> and the arguments to <code>title</code> may be supplied to this function.
<code>add</code>	logical. If <code>TRUE</code> , plots are added to current one, using <code>points</code> and <code>lines</code> .
<code>verbose</code>	logical. If <code>TRUE</code> , write one line of what is done.

## Details

Points involving missing values are not plotted.

The first column of `x` is plotted against the first column of `y`, the second column of `x` against the second column of `y`, etc. If one matrix has fewer columns, plotting will cycle back through the columns again. (In particular, either `x` or `y` may be a vector, against which all columns of the other argument will be plotted.)

The first element of `col`, `cex`, `lty`, `lwd` is used to plot the axes as well as the first line.

Because plotting symbols are drawn with lines and because these functions may be changing the line style, you should probably specify `lty=1` when using plotting symbols.

## Side Effects

Function `matplot` generates a new plot; `matpoints` and `matlines` add to the current one.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`plot`, `points`, `lines`, `matrix`, `par`.

**Examples**

```

# almost identical to plot(*)
matplot((-4:5)^2, main = "Quadratic")
sines <- outer(1:20, 1:4, function(x,y) sin(x/20 * pi * y))
matplot(sines, pch = 1:4, type = "o",
        col = rainbow(ncol(sines)))

x <- 0:50/50
matplot(x, outer(x, 1:8, function(x, k) sin(k*pi * x)),
        ylim = c(-2,2), type = "plobcsSh",
        main= "matplot(,type = \"plobcsSh\" )")
## pch & type = vector of 1-chars :
matplot(x, outer(x, 1:4, function(x, k) sin(k*pi * x)),
        pch = letters[1:4], type = c("b","p","o"))

data(iris) # is data.frame with 'Species' factor
table(iris$Species)
iS <- iris$Species == "setosa"
iV <- iris$Species == "versicolor"
op <- par(bg = "bisque")
matplot(c(1, 8), c(0, 4.5), type= "n",
        xlab = "Length", ylab = "Width",
        main = "Petal and Sepal Dimension in Iris Blossoms")
matpoints(iris[iS,c(1,3)], iris[iS,c(2,4)], pch = "sS",
        col = c(2,4))
matpoints(iris[iV,c(1,3)], iris[iV,c(2,4)], pch = "vV",
        col = c(2,4))
legend(1, 4, c("    Setosa Petals", "    Setosa Sepals",
               "Versicolor Petals", "Versicolor Sepals"),
        pch = "sSvV", col = rep(c(2,4), 2))

nam.var <- colnames(iris)[-5]
nam.spec <- as.character(iris[1+50*0:2, "Species"])
iris.S <- array(NA, dim = c(50,4,3),
               dimnames = list(NULL, nam.var, nam.spec))
for(i in 1:3)
  iris.S[,i] <- data.matrix(iris[1:50+50*(i-1), -5])

matplot(iris.S[, "Petal.Length", ], iris.S[, "Petal.Width", ],
        pch="SCV",
        col = rainbow(3, start = .8, end = .1),
        sub = paste(c("S", "C", "V"), dimnames(iris.S)[[3]],
                    sep = "=", collapse= " ", ),

```

```
main = "Fisher's Iris Data")
```

---

**mosaicplot**      *Mosaic Plots*


---

**Description**

Plots a mosaic on the current graphics device.

**Usage**

```
mosaicplot(x, ...)

## Default S3 method:
mosaicplot(x, main = deparse(substitute(x)),
  sub = NULL, xlab = NULL, ylab = NULL,
  sort = NULL, off = NULL, dir = NULL,
  color = FALSE, shade = FALSE, margin = NULL,
  cex.axis = 0.66, las = par("las"),
  type = c("pearson", "deviance", "FT"), ...)

## S3 method for class 'formula':
mosaicplot(formula, data = NULL, ...,
  main = deparse(substitute(data)), subset)
```

**Arguments**

<b>x</b>	a contingency table in array form, with optional category labels specified in the <code>dimnames(x)</code> attribute. The table is best created by the <code>table()</code> command.
<b>main</b>	character string for the mosaic title.
<b>sub</b>	character string for the mosaic sub-title (at bottom).
<b>xlab,ylab</b>	x- and y-axis labels used for the plot; by default, the first and second element of <code>names(dimnames(X))</code> (i.e., the name of the first and second variable in <code>X</code> ).
<b>sort</b>	vector ordering of the variables, containing a permutation of the integers <code>1:length(dim(x))</code> (the default).
<b>off</b>	vector of offsets to determine percentage spacing at each level of the mosaic (appropriate values are between 0 and 20, and the default is 10 at each level). There should be one offset for each dimension of the contingency table.



<b>dir</b>	vector of split directions (" <b>v</b> " for vertical and " <b>h</b> " for horizontal) for each level of the mosaic, one direction for each dimension of the contingency table. The default consists of alternating directions, beginning with a vertical split.
<b>color</b>	logical or (recycling) vector of colors for color shading, used only when <b>shade</b> is <b>FALSE</b> . The default <b>color=FALSE</b> gives empty boxes with no shading.
<b>shade</b>	a logical indicating whether to produce extended mosaic plots, or a numeric vector of at most 5 distinct positive numbers giving the absolute values of the cut points for the residuals. By default, <b>shade</b> is <b>FALSE</b> , and simple mosaics are created. Using <b>shade = TRUE</b> cuts absolute values at 2 and 4.
<b>margin</b>	a list of vectors with the marginal totals to be fit in the log-linear model. By default, an independence model is fitted. See <b>loglin</b> for further information.
<b>cex.axis</b>	The magnification to be used for axis annotation, as a multiple of <b>par("cex")</b> .
<b>las</b>	numeric; the style of axis labels, see <b>par</b> .
<b>type</b>	a character string indicating the type of residual to be represented. Must be one of " <b>pearson</b> " (giving components of Pearson's $\chi^2$ ), " <b>deviance</b> " (giving components of the likelihood ratio $\chi^2$ ), or " <b>FT</b> " for the Freeman-Tukey residuals. The value of this argument can be abbreviated.
<b>formula</b>	a formula, such as <b>y ~ x</b> .
<b>data</b>	a data frame (or list), or a contingency table from which the variables in <b>formula</b> should be taken.
<b>...</b>	further arguments to be passed to or from methods.
<b>subset</b>	an optional vector specifying a subset of observations in the data frame to be used for plotting.

## Details

This is a generic function. It currently has a default method (**mosaicplot.default**) and a formula interface (**mosaicplot.formula**).

Extended mosaic displays show the standardized residuals of a log-linear model of the counts from by the color and outline of the mosaic's

tiles. (Standardized residuals are often referred to a standard normal distribution.) Negative residuals are drawn in shaded of red and with broken outlines; positive ones are drawn in blue with solid outlines.

For the formula method, if **data** is an object inheriting from classes "**table**" or "**fable**", or an array with more than 2 dimensions, it is taken as a contingency table, and hence all entries should be nonnegative. In this case, the left-hand side of **formula** should be empty, and the variables on the right-hand side should be taken from the names of the **dimnames** attribute of the contingency table. A marginal table of these variables is computed, and a mosaic of this table is produced.

Otherwise, **data** should be a data frame or matrix, list or environment containing the variables to be cross-tabulated. In this case, after possibly selecting a subset of the data as specified by the **subset** argument, a contingency table is computed from the variables given in **formula**, and a mosaic is produced from this.

See Emerson (1998) for more information and a case study with television viewer data from Nielsen Media Research.

## Author(s)

S-PLUS original by John Emerson. Originally modified and enhanced for R by KH.

## References

Hartigan, J.A., and Kleiner, B. (1984) A mosaic of television ratings. *The American Statistician*, **38**, 32–35.

Emerson, J. W. (1998) Mosaic displays in S-PLUS: a general implementation and a case study. *Statistical Computing and Graphics Newsletter (ASA)*, **9**, 1, 17–23.

Friendly, M. (1994) Mosaic displays for multi-way contingency tables. *Journal of the American Statistical Association*, **89**, 190–200.

The home page of Michael Friendly (<http://www.math.yorku.ca/SCS/friendly.html>) provides information on various aspects of graphical methods for analyzing categorical data, including mosaic plots.

## See Also

`assocplot`, `loglin`.

## Examples

```
data(Titanic)
mosaicplot(Titanic, main = "Survival on the Titanic",
           color = TRUE)
## Formula interface for tabulated data:
mosaicplot(~ Sex + Age + Survived, data = Titanic,
           color = TRUE)

data(HairEyeColor)
mosaicplot(HairEyeColor, shade = TRUE)
## Independence model of hair and eye color and sex.
## Indicates that there are significantly more blue eyed
## blonde females than expected in the case of independence
## (and too few brown eyed blonde females).

mosaicplot(HairEyeColor, shade = TRUE,
           margin = list(c(1,2), 3))
## Model of joint independence of sex from hair and eye
## color. Males are underrepresented among people with
## brown hair and eyes, and are overrepresented among
## people with brown hair and blue eyes, but not
## "significantly".

## Formula interface for raw data: visualize
## crosstabulation of numbers of gears and carburettors in
## Motor Trend car data.
data(mtcars)
mosaicplot(~ gear + carb, data = mtcars, color = TRUE,
           las = 1)
# color recycling
mosaicplot(~ gear + carb, data = mtcars, color = 2:3,
           las = 1)
```

---

**mtext**      *Write Text into the Margins of a Plot*

---

## Description

Text is written in one of the four margins of the current figure region or one of the outer margins of the device region.

## Usage

```
mtext(text, side = 3, line = 0, outer = FALSE, at = NA,  
      adj = NA, cex = NA, col = NA, font = NA, vfont = NULL,  
      ...)
```

## Arguments

<b>text</b>	one or more character strings or expressions.
<b>side</b>	on which side of the plot (1=bottom, 2=left, 3=top, 4=right).
<b>line</b>	on which MARGin line, starting at 0 counting outwards.
<b>outer</b>	use outer margins if available.
<b>at</b>	give location in user-coordinates. If <code>length(at)==0</code> (the default), the location will be determined by <code>adj</code> .
<b>adj</b>	adjustment for each string. For strings parallel to the axes, <code>adj=0</code> means left or bottom alignment, and <code>adj=1</code> means right or top alignment. If <code>adj</code> is not a finite value (the default), the value <code>par("las")</code> determines the adjustment. For strings plotted parallel to the axis the default is to centre the string.
<b>...</b>	Further graphical parameters (see <code>text</code> and <code>par</code> ) ; currently supported are:
<b>cex</b>	character expansion factor (default = 1).
<b>col</b>	color to use.
<b>font</b>	font for text.
<b>vfont</b>	vector font for text.

## Details

The “user coordinates” in the outer margins always range from zero to one, and are not affected by the user coordinates in the figure region(s) — R is differing here from other implementations of S.

The arguments **side**, **line**, **at**, **at**, **adj**, the further graphical parameters and even **outer** can be vectors, and recycling will take place to plot as many strings as the longest of the vector arguments. Note that a vector **adj** has a different meaning from **text**.

**adj** = 0.5 will centre the string, but for **outer**=TRUE on the device region rather than the plot region.

Parameter **las** will determine the orientation of the string(s). For strings plotted perpendicular to the axis the default justification is to place the end of the string nearest the axis on the specified line.

Note that if the text is to be plotted perpendicular to the axis, **adj** determines the justification of the string *and* the position along the axis unless **at** is specified.

## Side Effects

The given text is written onto the current plot.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**title**, **text**, **plot**, **par**; **plotmath** for details on mathematical annotation.

## Examples

```
plot(1:10, (-4:5)^2, main="Parabola Points", xlab="xlab")
mtext("10 of them")
for(s in 1:4)
  mtext(paste("mtext(..., line= -1, {side, col, font} = ",
             s, ", cex = ", (1+s)/2, ")"), line = -1,
        side=s, col=s, font=s, cex= (1+s)/2)
mtext("mtext(..., line= -2)", line = -2)
mtext("mtext(..., line= -2, adj = 0)", line = -2, adj = 0)
## log axis :
plot(1:10, exp(1:10), log='y', main="log='y'", xlab="xlab")
```

```
for(s in 1:4) mtext(paste("mtext(...,side=",s,""), side=s)
```

---

**n2mfrow**      *Compute Default mfrow From Number of Plots*

---

**Description**

Easy setup for plotting multiple figures (in a rectangular layout) on one page. This computes a sensible default for `par(mfrow)`.

**Usage**

```
n2mfrow(nr.plots)
```

**Arguments**

`nr.plots`            integer; the number of plot figures you'll want to draw.

**Value**

A length two integer vector `nr`, `nc` giving the number of rows and columns, fulfilling `nr >= nc >= 1` and `nr * nc >= nr.plots`.

**Author(s)**

Martin Maechler

**See Also**

`par`, `layout`.

**Examples**

```
n2mfrow(8) # 3 x 3

n <- 5 ; x <- seq(-2,2, len=51)
## suppose now that 'n' is not known {inside function}
op <- par(mfrow = n2mfrow(n))
for (j in 1:n)
  plot(x, x^j, main = substitute(x^ exp, list(exp = j)),
       type='l', col="blue")

sapply(1:10, n2mfrow)
```

---

**pairs**     *Scatterplot Matrices*


---

**Description**

A matrix of scatterplots is produced.

**Usage**

```
pairs(x, ...)
```

```
## S3 method for class 'formula':
pairs(formula, data = NULL, ..., subset)
```

```
## Default S3 method:
pairs(x, labels, panel = points, ...,
      lower.panel = panel, upper.panel = panel,
      diag.panel = NULL, text.panel = textPanel,
      label.pos = 0.5 + has.diag/3,
      cex.labels = NULL, font.labels = 1,
      rowlattop = TRUE, gap = 1)
```

**Arguments**

<code>x</code>	the coordinates of points given as columns of a matrix.
<code>formula</code>	a formula, such as <code>y ~ x</code> .
<code>data</code>	a data.frame (or list) from which the variables in <code>formula</code> should be taken.
<code>subset</code>	an optional vector specifying a subset of observations to be used for plotting.
<code>labels</code>	the names of the variables.
<code>panel</code>	<code>function(x,y,...)</code> which is used to plot the contents of each panel of the display.
<code>...</code>	graphical parameters can be given as arguments to <code>plot</code> .
<code>lower.panel</code> , <code>upper.panel</code>	separate panel functions to be used below and above the diagonal respectively.
<code>diag.panel</code>	optional <code>function(x, ...)</code> to be applied on the diagonals.



<code>text.panel</code>	optional function( <code>x</code> , <code>y</code> , <code>labels</code> , <code>cex</code> , <code>font</code> , ... ) to be applied on the diagonals.
<code>label.pos</code>	y position of labels in the text panel.
<code>cex.labels</code> , <code>font.labels</code>	graphics parameters for the text panel.
<code>row1attop</code>	logical. Should the layout be matrix-like with row 1 at the top, or graph-like with row 1 at the bottom?
<code>gap</code>	Distance between subplots, in margin lines.

## Details

The  $ij$ th scatterplot contains `x[,i]` plotted against `x[,j]`. The “scatterplot” can be customised by setting panel functions to appear as something completely different. The off-diagonal panel functions are passed the appropriate columns of `x` as `x` and `y`: the diagonal panel function (if any) is passed a single column, and the `text.panel` function is passed a single (`x`, `y`) location and the column name.

The graphical parameters `pch` and `col` can be used to specify a vector of plotting symbols and colors to be used in the plots.

The graphical parameter `oma` will be set by `pairs.default` unless supplied as an argument.

A panel function should not attempt to start a new plot, but just plot within a given coordinate system: thus `plot` and `boxplot` are not panel functions.

## Author(s)

Enhancements for R 1.0.0 contributed by Dr. Jens Oehlschlaegel-Akiyoshi and R-core members.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## Examples

```
data(iris)
pairs(iris[1:4], main = "Anderson's Iris Data, 3 species",
      pch = 21,
      bg = c("red", "green3", "blue")[unclass(iris$Species)])

## formula method
```

```

data(swiss)
pairs(~ Fertility + Education + Catholic, data = swiss,
      subset = Education < 20, main = "Swiss data,
      Education < 20")

data(USJudgeRatings)
pairs(USJudgeRatings)

## put histograms on the diagonal
panel.hist <- function(x, ...)
{
  usr <- par("usr"); on.exit(par(usr))
  par(usr = c(usr[1:2], 0, 1.5) )
  h <- hist(x, plot = FALSE)
  breaks <- h$breaks; nB <- length(breaks)
  y <- h$counts; y <- y/max(y)
  rect(breaks[-nB], 0, breaks[-1], y, col="cyan", ...)
}
pairs(USJudgeRatings[1:5], panel=panel.smooth,
      cex = 1.5, pch = 24, bg="light blue",
      diag.panel=panel.hist, cex.labels = 2, font.labels=2)

## put (absolute) correlations on the upper panels, with
## size proportional to the correlations.
panel.cor <- function(x, y, digits=2, prefix="", cex.cor)
{
  usr <- par("usr"); on.exit(par(usr))
  par(usr = c(0, 1, 0, 1))
  r <- abs(cor(x, y))
  txt <- format(c(r, 0.123456789), digits=digits)[1]
  txt <- paste(prefix, txt, sep="")
  if(missing(cex.cor)) cex <- 0.8/strwidth(txt)
  text(0.5, 0.5, txt, cex = cex * r)
}
pairs(USJudgeRatings,
      lower.panel=panel.smooth, upper.panel=panel.cor)

```

---

**palette**     *Set or View the Graphics Palette*

---

**Description**

View or manipulate the color palette which is used when a `col=` has a numeric index.

**Usage**

```
palette(value)
```

**Arguments**

**value**                    an optional character vector.

**Details**

If **value** has length 1, it is taken to be the name of a built in color palette. If **value** has length greater than 1 it is assumed to contain a description of the colors which are to make up the new palette (either by name or by RGB levels).

If **value** is omitted or has length 0, no change is made the current palette.

Currently, the only built-in palette is "default".

**Value**

The palette which *was* in effect. This is `invisible` unless the argument is omitted.

**See Also**

`colors` for the vector of built-in “named” colors; `hsv`, `gray`, `rainbow`, `terrain.colors`,...to construct colors;

`col2rgb` for translating colors to RGB 3-vectors.

**Examples**

```
palette()           # obtain the current palette
palette(rainbow(6)) # six color rainbow

# gray scales; print old palette
(palette(gray(seq(0,.9,len=25))))
matplot(outer(1:100,1:30), type='l', lty=1,lwd=2, col=1:30,
        main = "Gray Scales Palette",
        sub = "palette(gray(seq(0,.9,len=25)))")
palette("default")  # reset back to the default
```

---

**Palettes**      *Color Palettes*

---

**Description**

Create a vector of `n` “contiguous” colors.

**Usage**

```
rainbow(n, s = 1, v = 1, start = 0, end = max(1, n - 1)/n,  
        gamma = 1)  
heat.colors(n)  
terrain.colors(n)  
topo.colors(n)  
cm.colors(n)
```

**Arguments**

<code>n</code>	the number of colors ( $\geq 1$ ) to be in the palette.
<code>s, v</code>	the “saturation” and “value” to be used to complete the HSV color descriptions.
<code>start</code>	the (corrected) hue in $[0,1]$ at which the rainbow begins.
<code>end</code>	the (corrected) hue in $[0,1]$ at which the rainbow ends.
<code>gamma</code>	the gamma correction, see argument <code>gamma</code> in <code>hsv</code> .

**Details**

Conceptually, all of these functions actually use (parts of) a line cut out of the 3-dimensional color space, parametrized by `hsv(h, s, v, gamma)`, where `gamma`= 1 for the `foo.colors` function, and hence, equispaced hues in RGB space tend to cluster at the red, green and blue primaries.

Some applications such as contouring require a palette of colors which do not “wrap around” to give a final color close to the starting one.

With `rainbow`, the parameters `start` and `end` can be used to specify particular subranges of hues. The following values can be used when generating such a subrange: red=0, yellow= $\frac{1}{6}$ , green= $\frac{2}{6}$ , cyan= $\frac{3}{6}$ , blue= $\frac{4}{6}$  and magenta= $\frac{5}{6}$ .

## Value

A character vector, `cv`, of color names. This can be used either to create a user-defined color palette for subsequent graphics by `palette(cv)`, a `col=` specification in graphics functions or in `par`.

## See Also

`colors`, `palette`, `hsv`, `rgb`, `gray` and `col2rgb` for translating to RGB numbers.

## Examples

```
# A Color Wheel
pie(rep(1,12), col=rainbow(12))

# Some palettes
demo.pal <-
  function(n, border = if (n<32) "light gray" else NA,
           main = paste("color palettes; n=", n),
           ch.col = c("rainbow(n, start=.7, end=.1)",
                      "heat.colors(n)", "terrain.colors(n)",
                      "topo.colors(n)", "cm.colors(n)"))
  {
    nt <- length(ch.col)
    i <- 1:n; j <- n / nt; d <- j/6; dy <- 2*d
    plot(i,i+d, type="n", yaxt="n", ylab="", main=main)
    for (k in 1:nt) {
      rect(i-.5, (k-1)*j+ dy, i+.4, k*j,
           col = eval(parse(text=ch.col[k])), border = border)
      text(2*j, k * j +dy/4, ch.col[k])
    }
  }
n <- if(.Device == "postscript") 64 else 16
# For screen, larger n may give color allocation problem
demo.pal(n)
```

---

**panel.smooth**     *Simple Panel Plot*

---

**Description**

An example of a simple useful **panel** function to be used as argument in e.g., **coplot** or **pairs**.

**Usage**

```
panel.smooth(x, y, col = par("col"), bg = NA,  
             pch = par("pch"), cex = 1, col.smooth = "red",  
             span = 2/3, iter=3, ...)
```

**Arguments**

<b>x,y</b>	numeric vectors of the same length
<b>col,bg,pch,cex</b>	numeric or character codes for the color(s), point type and size of <b>points</b> ; see also <b>par</b> .
<b>col.smooth</b>	color to be used by <b>lines</b> for drawing the smooths.
<b>span</b>	smoothing parameter <b>f</b> for <b>lowess</b> , see there.
<b>iter</b>	number of robustness iterations for <b>lowess</b> .
<b>...</b>	further arguments to <b>lines</b> .

**See Also**

**coplot** and **pairs** where **panel.smooth** is typically used; **lowess**.

**Examples**

```
data(swiss)  
# emphasize the smooths  
pairs(swiss, panel = panel.smooth, pch = ".")  
pairs(swiss, panel = panel.smooth, lwd = 2, cex= 1.5,  
      col="blue")
```

---

## par      *Set or Query Graphical Parameters*

---

### Description

`par` can be used to set or query graphical parameters. Parameters can be set by specifying them as arguments to `par` in `tag = value` form, or by passing them as a list of tagged values.

### Usage

```
par(..., no.readonly = FALSE)

<highlevel plot> (... , <tag> = <value>)
```

### Arguments

<code>...</code>	arguments in <code>tag = value</code> form, or a list of tagged values. The tags must come from the graphical parameters described below.
<code>no.readonly</code>	logical; if <code>TRUE</code> and there are no other arguments, only parameters are returned which can be set by a subsequent <code>par()</code> call.

### Details

Parameters are queried by giving one or more character vectors to `par`. `par()` (no arguments) or `par(no.readonly=TRUE)` is used to get *all* the graphical parameters (as a named list). Their names are currently taken from the variable `.Pars`. `.Pars.readonly` contains the names of the `par` arguments which are *readonly*.

***R.O.*** indicates *read-only arguments*: These may only be used in queries, i.e., they do *not* set anything.

All but these ***R.O.*** and the following *low-level arguments* can be set as well in high-level and mid-level plot functions, such as `plot`, `points`, `lines`, `axis`, `title`, `text`, `mtext`:

- "ask"
- "fig", "fin"
- "mai", "mar", "mex"
- "mfrow", "mfcoll", "mfg"



- "new"
- "oma", "omd", "omi"
- "pin", "plt", "ps", "pty"
- "usr"
- "xlog", "ylog"

## Value

When parameters are set, their former values are returned in an invisible named list. Such a list can be passed as an argument to **par** to restore the parameter values. Use **par(no.readonly = TRUE)** for the full list of parameters that can be restored.

When just one parameter is queried, the value is a character string. When two or more parameters are queried, the result is a list of character strings, with the list names giving the parameters.

Note the inconsistency: setting one parameter returns a list, but querying one parameter returns a vector.

## Graphical Parameters

- adj** The value of **adj** determines the way in which text strings are justified. A value of 0 produces left-justified text, 0.5 centered text and 1 right-justified text. (Any value in  $[0, 1]$  is allowed, and on most devices values outside that interval will also work.) Note that the **adj** argument of **text** also allows **adj = c(x, y)** for different adjustment in x- and y- direction.
- ann** If set to **FALSE**, high-level plotting functions do not annotate the plots they produce with axis and overall titles. The default is to do annotation.
- ask** logical. If **TRUE**, the user is asked for input, before a new figure is drawn.
- bg** The color to be used for the background of plots. A description of how colors are specified is given below.
- bty** A character string which determined the type of box which is drawn about plots. If **bty** is one of "o", "l", "7", "c", "u", or "]" the resulting box resembles the corresponding upper case letter. A value of "n" suppresses the box.
- cex** A numerical value giving the amount by which plotting text and symbols should be scaled relative to the default.
- cex.axis** The magnification to be used for axis annotation relative to the current.

- cex.lab** The magnification to be used for x and y labels relative to the current.
- cex.main** The magnification to be used for main titles relative to the current.
- cex.sub** The magnification to be used for sub-titles relative to the current.
- cin** *R.O.*; character size (**width,height**) in inches.
- col** A specification for the default plotting color. A description of how colors are specified is given below.
- col.axis** The color to be used for axis annotation.
- col.lab** The color to be used for x and y labels.
- col.main** The color to be used for plot main titles.
- col.sub** The color to be used for plot sub-titles.
- cra** *R.O.*; size of default character (**width,height**) in “rasters” (pixels).
- crt** A numerical value specifying (in degrees) how single characters should be rotated. It is unwise to expect values other than multiples of 90 to work. Compare with **srt** which does string rotation.
- csi** *R.O.*; height of (default sized) characters in inches.
- cxy** *R.O.*; size of default character (**width,height**) in user coordinate units. **par("cxy")** is **par("cin")/par("pin")** scaled to user coordinates. Note that **c(strwidth(ch), strwidth(ch))** for a given string **ch** is usually much more precise.
- din** *R.O.*; the device dimensions in inches.
- err** (*Unimplemented*; R is silent when points outside the plot region are *not* plotted.) The degree of error reporting desired.
- fg** The color to be used for the foreground of plots. This is the default color is used for things like axes and boxes around plots. A description of how colors are specified is given below.
- fig** A numerical vector of the form **c(x1, x2, y1, y2)** which gives the (NDC) coordinates of the figure region in the display region of the device.
- fin** A numerical vector of the form **c(x, y)** which gives the size of the figure region in inches.
- font** An integer which specifies which font to use for text. If possible, device drivers arrange so that 1 corresponds to plain text, 2 to bold face, 3 to italic and 4 to bold italic.
- font.axis** The font to be used for axis annotation.

**font.lab** The font to be used for x and y labels.

**font.main** The font to be used for plot main titles.

**font.sub** The font to be used for plot sub-titles.

**gamma** the gamma correction, see argument **gamma** to **hsv**.

**lab** A numerical vector of the form `c(x, y, len)` which modifies the way that axes are annotated. The values of `x` and `y` give the (approximate) number of tickmarks on the x and y axes and `len` specifies the label size. The default is `c(5, 5, 7)`. *Currently, len is unimplemented.*

**las** numeric in `{0,1,2,3}`; the style of axis labels.

**0:** always parallel to the axis [*default*],

**1:** always horizontal,

**2:** always perpendicular to the axis,

**3:** always vertical.

Note that other string/character rotation (via argument **srt** to **par**) does *not* affect the axis labels.

**lty** The line type. Line types can either be specified as an integer (0=blank, 1=solid, 2=dashed, 3=dotted, 4=dotdash, 5=longdash, 6=twodash) or as one of the character strings "blank", "solid", "dashed", "dotted", "dotdash", "longdash", or "twodash", where "blank" uses 'invisible lines' (i.e., doesn't draw them).

Alternatively, a string of up to 8 characters (from `c(1:9, "A":"F")`) may be given, giving the length of line segments which are alternatively drawn and skipped. See section 'Line Type Specification' below.

**lwd** The line width, a *positive* number, defaulting to 1.

**mai** A numerical vector of the form `c(bottom, left, top, right)` which gives the margin size specified in inches.

**mar** A numerical vector of the form `c(bottom, left, top, right)` which gives the lines of margin to be specified on the four sides of the plot. The default is `c(5, 4, 4, 2) + 0.1`.

**mex** **mex** is a character size expansion factor which is used to describe coordinates in the margins of plots.

**mfcol**, **mfrow** A vector of the form `c(nr, nc)`. Subsequent figures will be drawn in an `nr`-by-`nc` array on the device by *columns* (**mfcol**), or *rows* (**mfrow**), respectively.

Consider the alternatives, **layout** and **split.screen**.

- mfg** A numerical vector of the form `c(i, j)` where `i` and `j` indicate which figure in an array of figures is to be drawn next (if setting) or is being drawn (if enquiring). The array must already have been set by `mfc` or `mfrow`.
- For compatibility with S, the form `c(i, j, nr, nc)` is also accepted, when `nr` and `nc` should be the current number of rows and number of columns. Mismatches will be ignored, with a warning.
- mgp** The margin line (in `mex` units) for the axis title, axis labels and axis line. The default is `c(3, 1, 0)`.
- mkh** The height in inches of symbols to be drawn when the value of `pch` is an integer. *Completely ignored currently.*
- new** logical, defaulting to `FALSE`. If set to `TRUE`, the next high-level plotting command (actually `plot.new`) should *not clean* the frame before drawing “as if it was on a *new* device”.
- oma** A vector of the form `c(bottom, left, top, right)` giving the size of the outer margins in lines of text.
- omd** A vector of the form `c(x1, x2, y1, y2)` giving the outer margin region in NDC (= normalized device coordinates), i.e., as fraction (in `[0, 1]`) of the device region.
- omi** A vector of the form `c(bottom, left, top, right)` giving the size of the outer margins in inches.
- pch** Either an integer specifying a symbol or a single character to be used as the default in plotting points.
- pin** The width and height of the current plot in inches.
- plt** A vector of the form `c(x1, x2, y1, y2)` giving the coordinates of the plot region as fractions of the current figure region.
- ps** integer; the pointsize of text and symbols.
- pty** A character specifying the type of plot region to be used; “s” generates a square plotting region and “m” generates the maximal plotting region.
- smo** (*Unimplemented*) a value which indicates how smooth circles and circular arcs should be.
- srt** The string rotation in degrees. See the comment about `crt`.
- tck** The length of tick marks as a fraction of the smaller of the width or height of the plotting region. If `tck >= 0.5` it is interpreted as a fraction of the relevant side, so if `tck=1` grid lines are drawn. The default setting (`tck = NA`) is to use `tc1 = -0.5` (see below).
- tc1** The length of tick marks as a fraction of the height of a line of text. The default value is `-0.5`; setting `tc1 = NA` sets `tck = -0.01` which is S’ default.

- tmag** A number specifying the enlargement of text of the main title relative to the other annotating text of the plot.
- type** character; the default plot type desired, see `plot.default(type=...)`, defaulting to "p".
- usr** A vector of the form `c(x1, x2, y1, y2)` giving the extremes of the user coordinates of the plotting region. When a logarithmic scale is in use (i.e., `par("xlog")` is true, see below), then the x-limits will be  $10^{\sim \text{par}(\text{"usr"})[1:2]}$ . Similarly for the y-axis.
- xaxp** A vector of the form `c(x1, x2, n)` giving the coordinates of the extreme tick marks and the number of intervals between tick-marks when `par("xlog")` is false. Otherwise, when *log* coordinates are active, the three values have a different meaning: For a small range, *n* is *negative*, and the ticks are as in the linear case, otherwise, *n* is in `1:3`, specifying a case number, and *x1* and *x2* are the lowest and highest power of 10 inside the user coordinates, `par("usr")[1:2]`. See `axTicks()` for more details.
- xaxis** The style of axis interval calculation to be used for the x-axis. Possible values are "r", "i", "e", "s", "d". The styles are generally controlled by the range of data or `xlim`, if given. Style "r" (regular) first extends the data range by 4 percent and then finds an axis with pretty labels that fits within the range. Style "i" (internal) just finds an axis with pretty labels that fits within the original data range. Style "s" (standard) finds an axis with pretty labels within which the original data range fits. Style "e" (extended) is like style "s", except that it is also ensured that there is room for plotting symbols within the bounding box. Style "d" (direct) specifies that the current axis should be used on subsequent plots. (*Only "r" and "i" styles are currently implemented*)
- xaxt** A character which specifies the axis type. Specifying "n" causes an axis to be set up, but not plotted. The standard value is "s": for compatibility with S values "l" and "e" are accepted but are equivalent to "s".
- xlog** logical value (see `log` in `plot.default`). If TRUE, a logarithmic scale is in use (e.g., after `plot(*, log = "x")`). For a new device, it defaults to FALSE, i.e., linear scale.
- xpd** A logical value or NA. If FALSE, all plotting is clipped to the plot region, if TRUE, all plotting is clipped to the figure region, and if NA, all plotting is clipped to the device region.
- yaxp** A vector of the form `c(y1, y2, n)` giving the coordinates of the extreme tick marks and the number of intervals between tick-marks unless for log coordinates, see **xaxp** above.

- yaxs** The style of axis interval calculation to be used for the y-axis. See **xaxs** above.
- yaxt** A character which specifies the axis type. Specifying "n" causes an axis to be set up, but not plotted.
- ylog** a logical value; see **xlog** above.

## Color Specification

Colors can be specified in several different ways. The simplest way is with a character string giving the color name (e.g., "red"). A list of the possible colors can be obtained with the function **colors**. Alternatively, colors can be specified directly in terms of their RGB components with a string of the form "#RRGGBB" where each of the pairs RR, GG, BB consist of two hexadecimal digits giving a value in the range 00 to FF. Colors can also be specified by giving an index into a small table of colors, the **palette**. This provides compatibility with S. Index 0 corresponds to the background color.

Additionally, "transparent" or (integer) NA is *transparent*, useful for filled areas (such as the background!), and just invisible for things like lines or text.

The functions **rgb**, **hsv**, **gray** and **rainbow** provide additional ways of generating colors.

## Line Type Specification

Line types can either be specified by giving an index into a small built in table of line types (1 = solid, 2 = dashed, etc, see **lty** above) or directly as the lengths of on/off stretches of line. This is done with a string of an even number (up to eight) of characters, namely non-zero (hexadecimal) digits which give the lengths in consecutive positions in the string. For example, the string "33" specifies three units on followed by three off and "3313" specifies three units on followed by three off followed by one on and finally three off. The 'units' here are (on most devices) proportional to **lwd**, and with **lwd** = 1 are in pixels or points.

The five standard dash-dot line types (**lty** = 2:6) correspond to **c("44", "13", "1343", "73", "2262")**.

Note that NA is not a valid value for **lty**.

## Note

The effect of restoring all the (settable) graphics parameters as in the examples is hard to predict if the device has been resized. Several of

them are attempting to set the same things in different ways, and those last in the alphabet will win. In particular, the settings of `mai`, `mar`, `pin`, `plt` and `pty` interact, as do the outer margin settings, the figure layout and figure region size.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`plot.default` for some high-level plotting parameters; `colors`, `gray`, `rainbow`, `rgb`; `options` for other setup parameters; graphic devices `x11`, `postscript` and setting up device regions by `layout` and `split.screen`.

## Examples

```
op <- par(mfrow = c(2, 2), # 2 x 2 pictures on one plot
          pty = "s")      # square plotting region,
                          # independent of device size

## At end of plotting, reset to previous settings:
par(op)

## Alternatively,
# the whole list of settable par's.
op <- par(no.readonly = TRUE)
## do lots of plotting and par(.) calls, then reset:
par(op)

par("ylog") # FALSE
plot(1 : 12, log = "y")
par("ylog") # TRUE

plot(1:2, xaxs = "i") # 'inner axis' w/o extra space
stopifnot(par("xaxp")[1:2] == 1:2 &&
           par("usr") [1:2] == 1:2)

( nr.prof <-
  c(prof.pilots=16,lawyers=11,farmers=10,salesmen=9,
    physicians=9,mechanics=6,policemen=6,managers=6,
    engineers=5,teachers=4,housewives=3,students=3,
```

```
    armed.forces=1))
par(las = 3)
barplot(rbind(nr.prof)) # R 0.63.2: shows alignment problem
par(las = 0) # reset to default

## 'fg' use:
plot(1:12, type = "b",
     main="'fg' : axes, ticks and box in gray",
     fg = gray(0.7), bty="7" , sub=R.version.string)

ex <- function() {
  # all par settings which could be changed.
  old.par <- par(no.readonly = TRUE)

  on.exit(par(old.par))
  ## ...
  ## ... do lots of par() settings and plots
  ## ...
  invisible() # now, par(old.par) will be executed
}
ex()
```



---

**pdf**     *PDF Graphics Device*

---

**Description**

pdf starts the graphics device driver for producing PDF graphics.

**Usage**

```
pdf(file = ifelse(onefile, "Rplots.pdf", "Rplot%03d.pdf"),
     width=6, height=6, onefile = TRUE, family = "Helvetica",
     title = "R Graphics Output", encoding, bg, fg, pointsize)
```

**Arguments**

<b>file</b>	a character string giving the name of the file.
<b>width, height</b>	the width and height of the graphics region in inches.
<b>onefile</b>	logical: if true (the default) allow multiple figures in one file. If false, generate a file name containing the page number.
<b>family</b>	the font family to be used, one of "AvantGarde", "Bookman", "Courier", "Helvetica", "Helvetica-Narrow", "NewCenturySchoolbook", "Palatino" or "Times".
<b>title</b>	title string to embed in the file.
<b>encoding</b>	the name of an encoding file. Defaults to "ISOLatin1.enc" in the 'R_HOME/afm' directory, which is used if the path does not contain a path separator. An extension ".enc" can be omitted.
<b>pointsize</b>	the default point size to be used.
<b>bg</b>	the default background color to be used.
<b>fg</b>	the default foreground color to be used.

**Details**

pdf() opens the file **file** and the PDF commands needed to plot any graphics requested are sent to that file.

See **postscript** for details of encodings, as the internal code is shared between the drivers. The native PDF encoding is given in file 'PDFDoc.enc'.

`pdf` writes uncompressed PDF. It is primarily intended for producing PDF graphics for inclusion in other documents, and PDF-includers such as `pdftex` are usually able to handle compression.

At present the PDF is fairly simple, with each page being represented as a single stream. The R graphics model does not distinguish graphics objects at the level of the driver interface.

## Note

Acrobat Reader does not use the fonts specified but rather emulates them from multiple-master fonts. This can be seen in imprecise centering of characters, for example the multiply and divide signs in Helvetica.

## See Also

Devices, `postscript`

## Examples

```
## Test function for encodings
TestChars <- function(encoding="ISOLatin1")
{
  pdf(encoding=encoding)
  par(pty="s")
  plot(c(0,15), c(0,15), type="n", xlab="", ylab="")
  title(paste("Centred chars in encoding", encoding))
  grid(15, 15, lty=1)
  for(i in c(32:255)) {
    x <- i
    y <- i
    points(x, y, pch=i)
  }
  dev.off()
}
## there will be many warnings.
TestChars("ISOLatin2")
## doesn't view properly in US-spec Acrobat 5.05, but
## gs7.04 works. Lots of characters are not centred.
```

---

**persp**     *Perspective Plots*


---

**Description**

This function draws perspective plots of surfaces over the x-y plane. `persp` is a generic function.

**Usage**

```
persp(x, ...)

## Default S3 method:
persp(x = seq(0, 1, len = nrow(z)),
      y = seq(0, 1, len = ncol(z)), z, xlim = range(x),
      ylim = range(y), zlim = range(z, na.rm = TRUE),
      xlab = NULL, ylab = NULL, zlab = NULL, main = NULL,
      sub = NULL, theta = 0, phi = 15, r = sqrt(3), d = 1,
      scale = TRUE, expand = 1, col = "white", border = NULL,
      ltheta = -135, lphi = 0, shade = NA, box = TRUE,
      axes = TRUE, nticks = 5, ticktype = "simple", ...)
```

**Arguments**

<b>x, y</b>	locations of grid lines at which the values in <b>z</b> are measured. These must be in ascending order. By default, equally spaced values from 0 to 1 are used. If <b>x</b> is a <b>list</b> , its components <b>x\$x</b> and <b>x\$y</b> are used for <b>x</b> and <b>y</b> , respectively.
<b>z</b>	a matrix containing the values to be plotted (NAs are allowed). Note that <b>x</b> can be used instead of <b>z</b> for convenience.
<b>xlim, ylim, zlim</b>	x-, y- and z-limits. The plot is produced so that the rectangular volume defined by these limits is visible.
<b>xlab, ylab, zlab</b>	titles for the axes. N.B. These must be character strings; expressions are not accepted. Numbers will be coerced to character strings.
<b>main, sub</b>	main and sub title, as for <b>title</b> .
<b>theta, phi</b>	angles defining the viewing direction. <b>theta</b> gives the azimuthal direction and <b>phi</b> the colatitude.

<b>r</b>	the distance of the eyepoint from the centre of the plotting box.
<b>d</b>	a value which can be used to vary the strength of the perspective transformation. Values of <b>d</b> greater than 1 will lessen the perspective effect and values less and 1 will exaggerate it.
<b>scale</b>	before viewing the x, y and z coordinates of the points defining the surface are transformed to the interval [0,1]. If <b>scale</b> is TRUE the x, y and z coordinates are transformed separately. If <b>scale</b> is FALSE the coordinates are scaled so that aspect ratios are retained. This is useful for rendering things like DEM information.
<b>expand</b>	an expansion factor applied to the z coordinates. Often used with $0 < \text{expand} < 1$ to shrink the plotting box in the z direction.
<b>col</b>	the color(s) of the surface facets. Transparent colours are ignored. This is recycled to the $(nx - 1)(ny - 1)$ facets.
<b>border</b>	the color of the line drawn around the surface facets. A value of NA will disable the drawing of borders. This is sometimes useful when the surface is shaded.
<b>ltheta, lphi</b>	if finite values are specified for <b>ltheta</b> and <b>lphi</b> , the surface is shaded as though it was being illuminated from the direction specified by azimuth <b>ltheta</b> and colatitude <b>lphi</b> .
<b>shade</b>	the shade at a surface facet is computed as $((1+d)/2)^{\text{shade}}$ , where <b>d</b> is the dot product of a unit vector normal to the facet and a unit vector in the direction of a light source. Values of <b>shade</b> close to one yield shading similar to a point light source model and values close to zero produce no shading. Values in the range 0.5 to 0.75 provide an approximation to daylight illumination.
<b>box</b>	should the bounding box for the surface be displayed. The default is TRUE.
<b>axes</b>	should ticks and labels be added to the box. The default is TRUE. If <b>box</b> is FALSE then no ticks or labels are drawn.
<b>ticktype</b>	character: "simple" draws just an arrow parallel to the axis to indicate direction of increase; "detailed" draws normal ticks as per 2D plots.

**nticks**                the (approximate) number of tick marks to draw on the axes. Has no effect if **ticktype** is "simple".

**...**                additional graphical parameters (see **par**).

## Details

The plots are produced by first transforming the coordinates to the interval  $[0,1]$ . The surface is then viewed by looking at the origin from a direction defined by **theta** and **phi**. If **theta** and **phi** are both zero the viewing direction is directly down the negative y axis. Changing **theta** will vary the azimuth and changing **phi** the colatitude.

## Value

The *viewing transformation matrix*, say **VT**, a  $4 \times 4$  matrix suitable for projecting 3D coordinates  $(x, y, z)$  into the 2D plane using homogenous 4D coordinates  $(x, y, z, t)$ . It can be used to superimpose additional graphical elements on the 3D plot, by **lines()** or **points()**, e.g. using the function **trans3d** given in the last examples section below.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**contour** and **image**.

## Examples

```
## More examples in demo(persp)
# (1) The Obligatory Mathematical surface.
#     Rotated sinc function.
x <- seq(-10, 10, length= 30)
y <- x
f <- function(x,y) { r <- sqrt(x^2+y^2); 10 * sin(r)/r }
z <- outer(x, y, f)
z[is.na(z)] <- 1
op <- par(bg = "white")
persp(x, y, z, theta = 30, phi = 30, expand = 0.5,
      col = "lightblue")
persp(x, y, z, theta = 30, phi = 30, expand = 0.5,
      col = "lightblue", ltheta = 120, shade = 0.75,
      ticktype = "detailed", xlab = "X", ylab = "Y",
```

```

      zlab = "Sinc( r )") -> res
round(res, 3)

# (2) Add to existing persp plot :

trans3d <- function(x,y,z, pmat) {
  tr <- cbind(x,y,z,1) %*% pmat
  list(x = tr[,1]/tr[,4], y= tr[,2]/tr[,4])
}
xE <- c(-10,10); xy <- expand.grid(xE, xE)
points(trans3d(xy[,1], xy[,2], 6, pm=res), col=2, pch=16)
lines (trans3d(x, y=10, z= 6 + sin(x), pm = res), col = 3)

phi <- seq(0, 2*pi, len = 201)
r1 <- 7.725 # radius of 2nd maximum
xr <- r1 * cos(phi)
yr <- r1 * sin(phi)
## (no hidden lines)
lines(trans3d(xr,yr, f(xr,yr), res), col = "pink", lwd=2)

# (3) Visualizing a simple DEM model

data(volcano)
z <- 2 * volcano      # Exaggerate the relief
x <- 10 * (1:nrow(z)) # 10 meter spacing (S to N)
y <- 10 * (1:ncol(z)) # 10 meter spacing (E to W)
## Don't draw the grid lines : border = NA
par(bg = "slategray")
persp(x, y, z, theta = 135, phi = 30, col = "green3",
      scale = FALSE, ltheta = -120, shade = 0.75, border = NA,
      box = FALSE)
par(op)

```

---

**pictex**     *A PicTeX Graphics Driver*

---

**Description**

This function produces graphics suitable for inclusion in TeX and LaTeX documents.

**Usage**

```
pictex(file = "Rplots.tex", width = 5, height = 4,  
       debug = FALSE, bg = "white", fg = "black")
```

**Arguments**

<b>file</b>	the file where output will appear.
<b>width</b>	The width of the plot in inches.
<b>height</b>	the height of the plot in inches.
<b>debug</b>	should debugging information be printed.
<b>bg</b>	the background color for the plot.
<b>fg</b>	the foreground color for the plot.

**Details**

This driver does not have any font metric information, so the use of `plotmath` is not supported.

Multiple plots will be placed as separate environments in the output file.

**Author(s)**

This driver was provided by Valerio Aimale of the Department of Internal Medicine, University of Genoa, Italy.

**References**

Knuth, D. E. (1984) *The TeXbook*. Reading, MA: Addison-Wesley.  
Lamport, L. (1994) *LATEX: A Document Preparation System*. Reading, MA: Addison-Wesley.  
Goossens, M., Mittelbach, F. and Samarin, A. (1994) *The LATEX Companion*. Reading, MA: Addison-Wesley.

## See Also

postscript, Devices.

## Examples

```
pictex()
plot(1:11, (-5:5)^2, type='b', main="Simple Example Plot")
dev.off()
##

%% LaTeX Example
\documentclass{article}
\usepackage{pictex}
\begin{document}
%...
\begin{figure}[h]
  \centerline{\input{Rplots.tex}}
  \caption{}
\end{figure}
%...
\end{document}

%%-- TeX Example --
\input pictex
$$ \input Rplots.tex $$

##
unlink("Rplots.tex")
```



---

**pie**     *Pie Charts*

---

**Description**

Draw a pie chart.

**Usage**

```
pie(x, labels = names(x), edges = 200, radius = 0.8,  
    density = NULL, angle = 45, col = NULL, border = NULL,  
    lty = NULL, main = NULL, ...)
```

**Arguments**

<b>x</b>	a vector of positive quantities. The values in <b>x</b> are displayed as the areas of pie slices.
<b>labels</b>	a vector of character strings giving names for the slices. For empty or NA labels, no pointing line is drawn either.
<b>edges</b>	the circular outline of the pie is approximated by a polygon with this many edges.
<b>radius</b>	the pie is drawn centered in a square box whose sides range from $-1$ to $1$ . If the character strings labeling the slices are long it may be necessary to use a smaller radius.
<b>density</b>	the density of shading lines, in lines per inch. The default value of <b>NULL</b> means that no shading lines are drawn. Non-positive values of <b>density</b> also inhibit the drawing of shading lines.
<b>angle</b>	the slope of shading lines, given as an angle in degrees (counter-clockwise).
<b>col</b>	a vector of colors to be used in filling or shading the slices. If missing a set of 6 pastel colours is used, unless <b>density</b> is specified when <b>par("fg")</b> is used.
<b>border, lty</b>	(possibly vectors) arguments passed to <b>polygon</b> which draws each slice.
<b>main</b>	an overall title for the plot.
<b>...</b>	graphical parameters can be given as arguments to <b>pie</b> . They will affect the main title and labels only.

## Note

Pie charts are a very bad way of displaying information. The eye is good at judging linear measures and bad at judging relative areas. A bar chart or dot chart is a preferable way of displaying this type of data.

Cleveland (1985), page 264: “Data that can be shown by pie charts always can be shown by a dot chart. This means that judgements of position along a common scale can be made instead of the less accurate angle judgements.” This statement is based on the empirical investigations of Cleveland and McGill as well as investigations by perceptual psychologists.

Prior to R 1.5.0 this was known as `piechart`, which is the name of a Trellis function, so the name was changed to be compatible with S.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Cleveland, W. S. (1985) *The elements of graphing data*. Wadsworth: Monterey, CA, USA.

## See Also

`dotchart`.

## Examples

```
pie(rep(1, 24), col = rainbow(24), radius = 0.9)

pie.sales <- c(0.12, 0.3, 0.26, 0.16, 0.04, 0.12)
names(pie.sales) <- c("Blueberry", "Cherry",
  "Apple", "Boston Cream", "Other", "Vanilla Cream")
pie(pie.sales) # default colours
pie(pie.sales,
  col = c("purple", "violetred1", "green3", "cornsilk",
    "cyan", "white"))
pie(pie.sales, col = gray(seq(0.4,1.0,length=6)))
pie(pie.sales, density = 10, angle = 15 + 10 * 1:6)

n <- 200
pie(rep(1,n), labels="", col=rainbow(n), border=NA,
  main = "pie(*, labels=\"\\\", col=rainbow(n), border=NA,...")
```

---

**plot**      *Generic X-Y Plotting*


---

**Description**

Generic function for plotting of R objects. For more details about the graphical parameter arguments, see **par**.

**Usage**

```
plot(x, y, ...)
```

**Arguments**

- |             |   |
|-------------|---|
| <b>x</b>    | the coordinates of points in the plot. Alternatively, a single plotting structure, function or <i>any R object with a <b>plot</b> method</i> can be provided.   |
| <b>y</b>    | the y coordinates of points in the plot, <i>optional</i> if <b>x</b> is an appropriate structure.   |
| <b>...</b>  | graphical parameters can be given as arguments to <b>plot</b> . Many methods will also accept the following arguments:  |
| <b>type</b> | <p>what type of plot should be drawn. Possible types are</p> <ul style="list-style-type: none"> <li>• <b>"p"</b> for <b>p</b>oints,</li> <li>• <b>"l"</b> for <b>l</b>ines,</li> <li>• <b>"b"</b> for <b>b</b>oth,</li> <li>• <b>"c"</b> for the lines part alone of <b>"b"</b>,</li> <li>• <b>"o"</b> for both <b>"o</b>verplotted",</li> <li>• <b>"h"</b> for <b>"h</b>istogram" like (or <b>"h</b>igh-density") vertical lines,</li> <li>• <b>"s"</b> for stair steps,</li> <li>• <b>"S"</b> for other steps, see <i>Details</i> below,</li> <li>• <b>"n"</b> for no plotting.</li> </ul> <p>All other <b>types</b> give a warning or an error; using, e.g., <b>type = "punkte"</b> being equivalent to <b>type = "p"</b> for S compatibility.</p> |
| <b>main</b> | an overall title for the plot: see <b>title</b> .   |
| <b>sub</b>  | a sub title for the plot: see <b>title</b> .  |
| <b>xlab</b> | a title for the x axis: see <b>title</b> .  |
| <b>ylab</b> | a title for the y axis: see <b>title</b> .  |

## Details

For simple scatter plots, `plot.default` will be used. However, there are `plot` methods for many R objects, including `functions`, `data.frames`, `density` objects, etc. Use `methods(plot)` and the documentation for these.

The two step types differ in their x-y preference: Going from  $(x_1, y_1)$  to  $(x_2, y_2)$  with  $x_1 < x_2$ , `type = "s"` moves first horizontal, then vertical, whereas `type = "S"` moves the other way around.

## See Also

`plot.default`, `plot.formula` and other methods; `points`, `lines`, `par`.

## Examples

```
data(cars)
plot(cars)
lines(lowess(cars))
```

```
plot(sin, -pi, 2*pi)
```

```
## Discrete Distribution Plot:
```

```
plot(table(rpois(100,5)), type = "h", col = "red", lwd=10,
      main="rpois(100,lambda=5)")
```

```
## Simple quantiles/ECDF, see ecdf() {library(stepfun)} for
## a better one:
```

```
plot(x <- sort(rnorm(47)), type = "s", main =
     "plot(x, type = \"s\")")
points(x, cex = .5, col = "dark red")
```

---

**plot.data.frame**      *Plot Method for Data Frames*

---

**Description**

`plot.data.frame`, a method of the `plot` generic, uses `stripchart` for *one* variable, `plot.default` (scatterplot) for *two* variables, and `pairs` (scatterplot matrix) otherwise.

**Usage**

```
## S3 method for class 'data.frame':  
plot(x, ...)
```

**Arguments**

<code>x</code>	object of class <code>data.frame</code> .
<code>...</code>	further arguments to <code>stripchart</code> , <code>plot.default</code> or <code>pairs</code> .

**See Also**

`data.frame`

**Examples**

```
data(OrchardSprays)  
plot(OrchardSprays[1], method="jitter")  
plot(OrchardSprays[c(4,1)])  
plot(OrchardSprays)
```

---

**plot.default**      *The Default Scatterplot Function*


---

**Description**

Draw a scatter plot with “decorations” such as axes and titles in the active graphics window.

**Usage**

```
## Default S3 method:
plot(x, y = NULL, type = "p", xlim = NULL, ylim = NULL,
     log="", main = NULL, sub = NULL, xlab = NULL, ylab = NULL,
     ann = par("ann"), axes = TRUE, frame.plot = axes,
     panel.first = NULL, panel.last = NULL,
     col = par("col"), bg = NA, pch = par("pch"),
     cex = 1, lty = par("lty"), lab = par("lab"),
     lwd = par("lwd"), asp = NA, ...)
```

**Arguments**

<b>x,y</b>	the x and y arguments provide the x and y coordinates for the plot. Any reasonable way of defining the coordinates is acceptable. See the function <code>xy.coords</code> for details.
<b>type</b>	1-character string giving the type of plot desired. The following values are possible, for details, see <code>plot</code> : "p" for points, "l" for lines, "o" for overplotted points and lines, "b", "c" for (empty if "c") points joined by lines, "s" and "S" for stair steps and "h" for histogram-like vertical lines. Finally, "n" does not produce any points or lines.
<b>xlim</b>	the x limits (min,max) of the plot.
<b>ylim</b>	the y limits of the plot.
<b>log</b>	a character string which contains "x" if the x axis is to be logarithmic, "y" if the y axis is to be logarithmic and "xy" or "yx" if both axes are to be logarithmic.
<b>main</b>	a main title for the plot.
<b>sub</b>	a sub title for the plot.
<b>xlab</b>	a label for the x axis.

<code>ylab</code>	a label for the y axis.
<code>ann</code>	a logical value indicating whether the default annotation (title and x and y axis labels) should appear on the plot.
<code>axes</code>	a logical value indicating whether axes should be drawn on the plot.
<code>frame.plot</code>	a logical indicating whether a box should be drawn around the plot.
<code>panel.first</code>	an expression to be evaluated after the plot axes are set up but before any plotting takes place. This can be useful for drawing background grids or scatterplot smooths.
<code>panel.last</code>	an expression to be evaluated after plotting has taken place.
<code>col</code>	The colors for lines and points. Multiple colors can be specified so that each point can be given its own color. If there are fewer colors than points they are recycled in the standard fashion. Lines will all be plotted in the first colour specified.
<code>bg</code>	background color for open plot symbols, see <code>points</code> .
<code>pch</code>	a vector of plotting characters or symbols: see <code>points</code> .
<code>cex</code>	a numerical vector giving the amount by which plotting text and symbols should be scaled relative to the default.
<code>lty</code>	the line type, see <code>par</code> .
<code>lab</code>	the specification for the (approximate) numbers of tick marks on the x and y axes.
<code>lwd</code>	the line width <b>not yet supported for postscript</b> .
<code>asp</code>	the $y/x$ aspect ratio, see <code>plot.window</code> .
<code>...</code>	graphical parameters as in <code>par</code> may also be passed as arguments.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Cleveland, W. S. (1985) *The Elements of Graphing Data*. Monterey, CA: Wadsworth.

**See Also**

plot, plot.window, xy.coords.

**Examples**

```
data(cars)
Speed <- cars$speed
Distance <- cars$dist
plot(Speed, Distance, panel.first = grid(8,8),
      pch = 0, cex = 1.2, col = "blue")
plot(Speed, Distance, panel.first =
      lines(lowess(Speed, Distance), lty = "dashed"),
      pch = 0, cex = 1.2, col = "blue")

## Show the different plot types
x <- 0:12
y <- sin(pi/5 * x)
op <- par(mfrow = c(3,3), mar = .1+ c(2,2,3,1))
for (tp in c("p","l","b", "c","o","h", "s","S","n")) {
  plot(y ~ x, type = tp,
        main = paste("plot(*, type = \"",tp,"\"",sep=""))
  if(tp == "S") {
    lines(x,y, type = "s", col = "red", lty = 2)
    mtext("lines(*, type = \"s\\", ...)", col="red", cex=.8)
  }
}
par(op)

## Log-Log Plot with custom axes
lx <- seq(1,5, length=41)
yl <- expression(e^{-frac(1,2) * {log[10](x)}^2})
y <- exp(-.5*lx^2)
op <- par(mfrow=c(2,1), mar=par("mar")+c(0,1,0,0))
plot(10^lx, y, log="xy", type="l", col="purple",
      main="Log-Log plot", ylab=yl, xlab="x")
plot(10^lx, y, log="xy", type="o", pch='.',
      col = "forestgreen",
      main = "Log-Log plot with custom axes",
      ylab=yl, xlab="x",
      axes = FALSE, frame.plot = TRUE)
axis(1, at = my.at <- 10^(1:5),
      labels = formatC(my.at, format="fg"))
at.y <- 10^(-5:-1)
```



```
axis(2, at = at.y,  
     labels = formatC(at.y, format="fg"), col.axis="red")  
par(op)
```

---

**plot.density**     *Plot Method for Kernel Density Estimation*

---

**Description**

The `plot` method for density objects.

**Usage**

```
## S3 method for class 'density':  
plot(x, main = NULL, xlab = NULL, ylab = "Density",  
     type = "l", zero.line = TRUE, ...)
```

**Arguments**

`x`                    a “density” object.  
`main`, `xlab`, `ylab`, `type`           plotting parameters with useful defaults.  
`...`                further plotting parameters.  
`zero.line`           logical; if `TRUE`, add a base line at  $y = 0$

**Value**

None.

**References****See Also**

`density`.

---

**plot.design**      *Plot Univariate Effects of a ‘Design’ or Model*

---

**Description**

Plot univariate effects of one or more **factors**, typically for a designed experiment as analyzed by `aov()`. Further, in S this is a method of the `plot` generic function for `design` objects.

**Usage**

```
plot.design(x, y = NULL, fun = mean, data = NULL, ...,
  ylim = NULL, xlab = "Factors", ylab = NULL, main = NULL,
  ask = NULL, xaxt = par("xaxt"), axes = TRUE,
  xtick = FALSE)
```

**Arguments**

<b>x</b>	either a data frame containing the design factors and optionally the response, or a <b>formula</b> or <b>terms</b> object.
<b>y</b>	the response, if not given in <b>x</b> .
<b>fun</b>	a function (or name of one) to be applied to each subset. It must return one number for a numeric (vector) input.
<b>data</b>	data frame containing the variables referenced by <b>x</b> when that is formula like.
<b>...</b>	graphical arguments such as <b>col</b> , see <b>par</b> .
<b>ylim</b>	range of y values, as in <b>plot.default</b> .
<b>xlab</b>	x axis label, see <b>title</b> .
<b>ylab</b>	y axis label with a “smart” default.
<b>main</b>	main title, see <b>title</b> .
<b>ask</b>	logical indicating if the user should be asked before a new page is started – in the case of multiple y’s.
<b>xaxt</b>	character giving the type of x axis.
<b>axes</b>	logical indicating if axes should be drawn.
<b>xtick</b>	logical indicating if “ticks” (one per factor) should be drawn on the x axis.

## Details

The supplied function will be called once for each level of each factor in the design and the plot will show these summary values. The levels of a particular factor are shown along a vertical line, and the overall value of `fun()` for the response is drawn as a horizontal line.

This is a new R implementation which will not be completely compatible to the earlier S implementations. This is not a bug but might still change.

## Note

A big effort was taken to make this closely compatible to the S version. However, `col` (and `fg`) specification has different effects.

## Author(s)

Roberto Frisullo and Martin Maechler

## References

Chambers, J. M. and Hastie, T. J. eds (1992) *Statistical Models in S*. Chapman & Hall, London, **the white book**, pp. 546–7 (and 163–4).

Freeny, A. E. and Landwehr, J. M. (1990) Displays for data from large designed experiments; Computer Science and Statistics: Proc. 22nd SympInterface, 117–126, Springer Verlag.

## See Also

`interaction.plot` for a “standard graphic” of designed experiments.

## Examples

```
data(warpbreaks)
# automatic for data frame with one numeric var.
plot.design(warpbreaks)

Form <- breaks ~ wool + tension
summary(fm1 <- aov(Form, data = warpbreaks))
# same as above
plot.design(      Form, data = warpbreaks, col = 2)

## More than one y :
data(esoph)
str(esoph)
```

```
## two plots; if interactive you are "ask"ed
plot.design(esoph)

## or rather, compare mean and median:
op <- par(mfcol = 1:2)
plot.design(ncases/ncontrols ~ ., data = esoph,
            ylim = c(0,0.8))
plot.design(ncases/ncontrols ~ ., data = esoph,
            ylim = c(0,0.8), fun = median)
par(op)
```

---

**plot.factor**      *Plotting Factor Variables*

---

**Description**

This function implements a “scatterplot” method for **factor** arguments of the *generic* **plot** function. Actually, **boxplot** or **barplot** are used when appropriate.

**Usage**

```
## S3 method for class 'factor':  
plot(x, y, legend.text = levels(y), ...)
```

**Arguments**

<b>x,y</b>	numeric or factor. <b>y</b> may be missing.
<b>legend.text</b>	a vector of text used to construct a legend for the plot. Only used if <b>y</b> is present and a factor.
<b>...</b>	Further arguments to <b>plot</b> , see also <b>par</b> .

**See Also**

**plot.default**, **plot.formula**, **barplot**, **boxplot**.

**Examples**

```
data(PlantGrowth)  
# plot.data.frame  
plot(PlantGrowth)  
# numeric vector ~ factor  
plot(weight ~ group, data = PlantGrowth)  
# factor ~ factor  
plot(cut(weight, 2) ~ group, data = PlantGrowth)  
## passing "... " to barplot() eventually:  
plot(cut(weight, 3) ~ group, data = PlantGrowth,  
      density = 16*(1:3), col=NULL)  
  
# extremely silly  
plot(PlantGrowth$group, axes=FALSE, main="no axes")
```

---

**plot.formula**      *Formula Notation for Scatterplots*


---

**Description**

Specify a scatterplot or add points or lines via a formula.

**Usage**

```
## S3 method for class 'formula':
plot(formula, data = parent.frame(), ..., subset,
      ylab = varnames[response], ask = TRUE)

## S3 method for class 'formula':
points(formula, data = parent.frame(), ..., subset)

## S3 method for class 'formula':
lines(formula, data = parent.frame(), ..., subset)
```

**Arguments**

<b>formula</b>	a formula, such as $y \sim x$ .
<b>data</b>	a data.frame (or list) from which the variables in <b>formula</b> should be taken.
<b>...</b>	Further graphical parameters may also be passed as arguments, see <b>par</b> . <b>horizontal = TRUE</b> is also accepted.
<b>subset</b>	an optional vector specifying a subset of observations to be used in the fitting process.
<b>ylab</b>	the y label of the plot(s).
<b>ask</b>	logical, see <b>par</b> .

**Details**

Both the terms in the formula and the **...** arguments are evaluated in **data** enclosed in **parent.frame()** if **data** is a list or a data frame. The terms of the formula and those arguments in **...** that are of the same length as **data** are subjected to the subsetting specified in **subset**. If the formula in **plot.formula** contains more than one non-response term, a series of plots of y against each term is given. A plot against the running index can be specified as **plot(y~1)**.

If `y` is an object (ie. has a `class` attribute) then `plot.formula` looks for a plot method for that class first. Otherwise, the class of `x` will determine the type of the plot. For factors this will be a parallel boxplot, and argument `horizontal = TRUE` can be used (see `boxplot`).

## Value

These functions are invoked for their side effect of drawing in the active graphics device.

## See Also

`plot.default`, `plot.factor`.

## Examples

```
data(airquality)
op <- par(mfrow=c(2,1))
plot(Ozone ~ Wind, data = airquality,
     pch=as.character(Month))
plot(Ozone ~ Wind, data = airquality,
     pch=as.character(Month),
     subset = Month != 7)
par(op)
```



---

**plot.histogram**      *Plot Histograms*


---

**Description**

These are methods for objects of class "**histogram**", typically produced by **hist**.

**Usage**

```
## S3 method for class 'histogram':
plot(x, freq = equidist, density = NULL, angle = 45,
     col = NULL, border = par("fg"), lty = NULL,
     main = paste("Histogram of", x$xname), sub = NULL,
     xlab = x$xname, ylab, xlim = range(x$breaks), ylim = NULL,
     axes = TRUE, labels = FALSE, add = FALSE, ...)

## S3 method for class 'histogram':
lines(x, ...)
```

**Arguments**

<b>x</b>	a <b>histogram</b> object, or a list with components <b>density</b> , <b>mid</b> , etc, see <b>hist</b> for information about the components of <b>x</b> .
<b>freq</b>	logical; if <b>TRUE</b> , the histogram graphic is to present a representation of frequencies, i.e. <b>x\$count</b> s; if <b>FALSE</b> , <i>relative</i> frequencies ("probabilities"), i.e., <b>x\$density</b> , are plotted. The default is true for equidistant <b>breaks</b> and false otherwise.
<b>col</b>	a colour to be used to fill the bars. The default of <b>NULL</b> yields unfilled bars.
<b>border</b>	the color of the border around the bars.
<b>angle, density</b>	select shading of bars by lines: see <b>rect</b> .
<b>lty</b>	the line type used for the bars, see also <b>lines</b> .
<b>main, sub, xlab, ylab</b>	these arguments to <b>title</b> have useful defaults here.
<b>xlim, ylim</b>	the range of x and y values with sensible defaults.
<b>axes</b>	logical, indicating if axes should be drawn.

<code>labels</code>	logical or character. Additionally draw labels on top of bars, if not <code>FALSE</code> ; if <code>TRUE</code> , draw the counts or rounded densities; if <code>labels</code> is a character, draw itself.
<code>add</code>	logical. If <code>TRUE</code> , only the bars are added to the current plot. This is what <code>lines.histogram(*)</code> does.
<code>...</code>	further graphical parameters to <code>title</code> and <code>axis</code> .

## Details

`lines.histogram(*)` is the same as `plot.histogram(*, add = TRUE)`.

## See Also

`hist`, `stem`, `density`.

## Examples

```
data(women)
str(wwt <- hist(women$weight, nc= 7, plot = FALSE))
# default main & xlab using wwt$xname
plot(wwt, labels = TRUE)
plot(wwt, border = "dark blue", col = "light blue",
      main = "Histogram of 15 women's weights",
      xlab = "weight [pounds]")

## Fake "lines" example, using non-default labels:
w2 <- wwt; w2$counts <- w2$counts - 1
lines(w2, col = "Midnight Blue",
      labels = ifelse(w2$counts, "> 1", "1"))
```

---

**plot.lm**     *Plot Diagnostics for an lm Object*


---

**Description**

Four plots (selectable by **which**) are currently provided: a plot of residuals against fitted values, a Scale-Location plot of  $\sqrt{|residuals|}$  against fitted values, a Normal Q-Q plot, and a plot of Cook's distances versus row labels.

**Usage**

```
## S3 method for class 'lm':
plot(x, which = 1:4,
     caption = c("Residuals vs Fitted",
                 "Normal Q-Q plot",
                 "Scale-Location plot",
                 "Cook's distance plot"),
     panel = points,
     sub.caption = deparse(x$call), main = "",
     ask =
       prod(par("mfcol")) < length(which) && dev.interactive(),
     ...,
     id.n = 3, labels.id = names(residuals(x)), cex.id = 0.75)
```

**Arguments**

<b>x</b>	lm object, typically result of <b>lm</b> or <b>glm</b> .
<b>which</b>	If a subset of the plots is required, specify a subset of the numbers 1:4.
<b>caption</b>	Captions to appear above the plots
<b>panel</b>	Panel function. A useful alternative to <b>points</b> is <b>panel.smooth</b> .
<b>sub.caption</b>	common title—above figures if there are multiple; used as <b>sub</b> ( <b>s.title</b> ) otherwise.
<b>main</b>	title to each plot—in addition to the above <b>caption</b> .
<b>ask</b>	logical; if <b>TRUE</b> , the user is <i>asked</i> before each plot, see <b>par(ask=.)</b> .
<b>...</b>	other parameters to be passed through to plotting functions.

<code>id.n</code>	number of points to be labelled in each plot, starting with the most extreme.
<code>labels.id</code>	vector of labels, from which the labels for extreme points will be chosen. <code>NULL</code> uses observation numbers.
<code>cex.id</code>	magnification of point labels.

## Details

`sub.caption`—by default the function call—is shown as a subtitle (under the x-axis title) on each plot when plots are on separate pages, or as a subtitle in the outer margin (if any) when there are multiple plots per page.

The “Scale-Location” plot, also called “Spread-Location” or “S-L” plot, takes the square root of the absolute residuals in order to diminish skewness ( $\sqrt{|E|}$  is much less skewed than  $|E|$  for Gaussian zero-mean  $E$ ).

This ‘S-L’ and the Q-Q plot use *standardized* residuals which have identical variance (under the hypothesis). They are given as  $R_i/(s \times \sqrt{1 - h_{ii}})$  where  $h_{ii}$  are the diagonal entries of the hat matrix, `influence()` `$hat`, see also `hat`.

## Author(s)

John Maindonald and Martin Maechler.

## References

- Belsley, D. A., Kuh, E. and Welsch, R. E. (1980) *Regression Diagnostics*. New York: Wiley.
- Cook, R. D. and Weisberg, S. (1982) *Residuals and Influence in Regression*. London: Chapman and Hall.
- Hinkley, D. V. (1975) On power transformations to symmetry. *Biometrika* **62**, 101–111.
- McCullagh, P. and Nelder, J. A. (1989) *Generalized Linear Models*. London: Chapman and Hall.

## See Also

`termplot`, `lm.influence`, `cooks.distance`.

## Examples

```
## Analysis of the life-cycle savings data given in
## Belsley, Kuh and Welsch.
data(LifeCycleSavings)
plot(lm.SR <- lm(sr ~ pop15 + pop75 + dpi + ddpi,
                 data = LifeCycleSavings))

## 4 plots on 1 page; allow room for printing model formula
## in outer margin:
par(mfrow = c(2, 2), oma = c(0, 0, 2, 0))
plot(lm.SR)
plot(lm.SR, id.n = NULL)           # no id's
plot(lm.SR, id.n = 5, labels.id = NULL) # 5 id numbers

## Fit a smooth curve, where applicable:
plot(lm.SR, panel = panel.smooth)
## Gives a smoother curve
plot(lm.SR, panel = function(x,y) panel.smooth(x, y,
span = 1))

par(mfrow=c(2,1)) # same oma as above
plot(lm.SR, which = 1:2,
      sub.caption = "Saving Rates, n=50, p=5")
```

---

**plot.table**      *Plot Methods for ‘table’ Objects*

---

**Description**

This is a method of the generic `plot` function for (contingency) `table` objects. Whereas for two- and more dimensional tables, a `mosaicplot` is drawn, one-dimensional ones are plotted “bar like”.

**Usage**

```
## S3 method for class 'table':  
plot(x, type = "h", ylim = c(0, max(x)), lwd = 2,  
      xlab = NULL, ylab = NULL, frame.plot = is.num, ...)
```

**Arguments**

<code>x</code>	a <code>table</code> (like) object.
<code>type</code>	plotting type.
<code>ylim</code>	range of y-axis.
<code>lwd</code>	line width for bars when <code>type = "h"</code> is used in the 1D case.
<code>xlab, ylab</code>	x- and y-axis labels.
<code>frame.plot</code>	logical indicating if a frame ( <code>box</code> ) should be drawn in the 1D case. Defaults to true when <code>x</code> has <code>dimnames</code> coerceable to numbers.
<code>...</code>	further graphical arguments, see <code>plot.default</code> .

**Details**

The current implementation (R 1.2) is somewhat experimental and will be improved and extended.

**See Also**

`plot.factor`, the `plot` method for factors.

**Examples**

```
## 1-d tables
(Poiss.tab <- table(N = rpois(200, lam= 5)))
plot(Poiss.tab, main = "plot(table(rpois(200, lam=5)))")

data(state)
plot(table(state.division))

## 4-D :
data(Titanic)
plot(Titanic, main = "plot(Titanic, main= *)")
```

---

**plot.ts**     *Plotting Time-Series Objects*


---

**Description**

Plotting method for objects inheriting from class "ts".

**Usage**

```
## S3 method for class 'ts':
plot(x, y = NULL, plot.type = c("multiple", "single"),
      xy.labels, xy.lines, panel = lines, nc, ...)

## S3 method for class 'ts':
lines(x, ...)
```

**Arguments**

<code>x, y</code>	time series objects, usually inheriting from class "ts".
<code>plot.type</code>	for multivariate time series, should the series be plotted separately (with a common time axis) or on a single plot?
<code>xy.labels</code>	logical, indicating if <code>text()</code> labels should be used for an x-y plot, <i>or</i> character, supplying a vector of labels to be used. The default is to label for up to 150 points, and not for more.
<code>xy.lines</code>	logical, indicating if <code>lines</code> should be drawn for an x-y plot. Defaults to the value of <code>xy.labels</code> if that is logical, otherwise to <code>TRUE</code> .
<code>panel</code>	a <code>function(x, col, bg, pch, type, ...)</code> which gives the action to be carried out in each panel of the display for <code>plot.type="multiple"</code> . The default is <code>lines</code> .
<code>nc</code>	the number of columns to use when <code>type="multiple"</code> . Defaults to 1 for up to 4 series, otherwise to 2.
<code>...</code>	additional graphical arguments, see <code>plot</code> , <code>plot.default</code> and <code>par</code> .



## Details

If `y` is missing, this function creates a time series plot, for multivariate series of one of two kinds depending on `plot.type`.

If `y` is present, both `x` and `y` must be univariate, and a “scatter” plot `y ~ x` will be drawn, enhanced by using `text` if `xy.labels` is `TRUE` or `character`, and `lines` if `xy.lines` is `TRUE`.

## See Also

`ts` for basic time series construction and access functionality.

## Examples

```
## Multivariate
z <- ts(matrix(rt(300, df = 3), 100, 3), start=c(1961, 1),
          frequency=12)
plot(z, type = "b")      # multiple
plot(z, plot.type="single", lty=1:3, col=4:2)
```

```
## A phase plot:
data(nhtemp)
plot(nhtemp, c(nhtemp[-1], NA), cex = .8, col="blue",
     main = "Lag plot of New Haven temperatures")
## a clearer way to do this would be
library(ts)
plot(nhtemp, lag(nhtemp, 1), cex = .8, col="blue",
     main = "Lag plot of New Haven temperatures")
```

```
library(ts) # normally loaded
```

```
data(sunspots)
## xy.lines and xy.labels are FALSE for large series:
plot(lag(sunspots, 1), sunspots, pch = ".")
```

```
data(EuStockMarkets)
SMI <- EuStockMarkets[, "SMI"]
plot(lag(SMI, 1), SMI, pch = ".")
plot(lag(SMI, 20), SMI, pch = ".", log = "xy",
     main = "4 weeks lagged SMI stocks -- log scale",
     xy.lines= TRUE)
```

---

<code>plot.window</code>	<i>Set up World Coordinates for Graphics Window</i>
--------------------------	---

---

## Description

This function sets up the world coordinate system for a graphics window. It is called by higher level functions such as `plot.default` (*after* `plot.new`).

## Usage

```
plot.window(xlim, ylim, log = "", asp = NA, ...)
```

## Arguments

<code>xlim, ylim</code>	numeric of length 2, giving the x and y coordinates ranges.
<code>log</code>	character; indicating which axes should be in log scale.
<code>asp</code>	numeric, giving the <b>aspect</b> ratio y/x.
<code>...</code>	further graphical parameters as in <b>par</b> .

## Details

Note that if **asp** is a finite positive value then the window is set up so that one data unit in the x direction is equal in length to **asp** × one data unit in the y direction.

The special case **asp** == 1 produces plots where distances between points are represented accurately on screen. Values with **asp** > 1 can be used to produce more accurate maps when using latitude and longitude.

Usually, one should rather use the higher level functions such as `plot`, `hist`, `image`, ..., instead and refer to their help pages for explanation of the arguments.

## See Also

`xy.coords`, `plot.xy`, `plot.default`.

## Examples

```
## An example for the use of 'asp' :
library(mva) # normally loaded
data(eurodist)
loc <- cmdscale(eurodist)
rx <- range(x <- loc[,1])
ry <- range(y <- -loc[,2])
plot(x, y, type="n", asp=1, xlab="", ylab="")
abline(h = pretty(rx, 10), v = pretty(ry, 10),
       col = "lightgray")
text(x, y, names(eurodist), cex=0.8)
```

---

**plot.xy**     *Basic Internal Plot Function*

---

**Description**

This is *the* internal function that does the basic plotting of points and lines. Usually, one should rather use the higher level functions instead and refer to their help pages for explanation of the arguments.

**Usage**

```
plot.xy(xy, type, pch=1, lty="solid", col=par("fg"), bg=NA,
        cex=1, ...)
```

**Arguments**

<code>xy</code>	A four-element list as results from <code>xy.coords</code> .
<code>type</code>	1 character code.
<code>pch</code>	character or integer code for kind of points/lines, see <code>points.default</code> .
<code>lty</code>	line type code, see <code>lines</code> .
<code>col</code>	color code or name, see <code>colors</code> , <code>palette</code> .
<code>bg</code>	background (“fill”) color for open plot symbols.
<code>cex</code>	character expansion.
<code>...</code>	further graphical parameters.

**See Also**

`plot`, `plot.default`, `points`, `lines`.

**Examples**

```
# to see how it calls "plot.xy(xy.coords(x, y), ...)"
points.default
```

---

**plotmath**     *Mathematical Annotation in R*


---

**Description**

If the `text` argument to one of the text-drawing functions (`text`, `mtext`, `axis`) in R is an expression, the argument is interpreted as a mathematical expression and the output will be formatted according to TeX-like rules. Expressions can also be used for titles, subtitles and x- and y-axis labels (but not for axis labels on `persp` plots).

**Details**

A mathematical expression must obey the normal rules of syntax for any R expression, but it is interpreted according to very different rules than for normal R expressions.

It is possible to produce many different mathematical symbols, generate sub- or superscripts, produce fractions, etc.

The output from `demo(plotmath)` includes several tables which show the available features. In these tables, the columns of grey text show sample R expressions, and the columns of black text show the resulting output.

The available features are also described in the tables below:

**Syntax**

```
x + y
x - y
x*y
x/y
x %+-% y
x %/% y
x %*% y
x[i]
x^2
paste(x, y, z)
sqrt(x)
sqrt(x, y)
x == y
x != y
x < y
x <= y
```

**Meaning**

```
x plus y
x minus y
juxtapose x and y
x forwardslash y
x plus or minus y
x divided by y
x times y
x subscript i
x superscript 2
juxtapose x, y, and z
square root of x
yth root of x
x equals y
x is not equal to y
x is less than y
x is less than or equal to y
```

<code>x &gt; y</code>	$x$ is greater than $y$
<code>x &gt;= y</code>	$x$ is greater than or equal to $y$
<code>x %~=% y</code>	$x$ is approximately equal to $y$
<code>x %=% y</code>	$x$ and $y$ are congruent
<code>x %==% y</code>	$x$ is defined as $y$
<code>x %prop% y</code>	$x$ is proportional to $y$
<code>plain(x)</code>	draw $x$ in normal font
<code>bold(x)</code>	draw $x$ in bold font
<code>italic(x)</code>	draw $x$ in italic font
<code>bolditalic(x)</code>	draw $x$ in bolditalic font
<code>list(x, y, z)</code>	comma-separated list
<code>...</code>	ellipsis (height varies)
<code>cdots</code>	ellipsis (vertically centred)
<code>ldots</code>	ellipsis (at baseline)
<code>x %subset% y</code>	$x$ is a proper subset of $y$
<code>x %subsetq% y</code>	$x$ is a subset of $y$
<code>x %notsubset% y</code>	$x$ is not a subset of $y$
<code>x %supset% y</code>	$x$ is a proper superset of $y$
<code>x %supsetq% y</code>	$x$ is a superset of $y$
<code>x %in% y</code>	$x$ is an element of $y$
<code>x %notin% y</code>	$x$ is not an element of $y$
<code>hat(x)</code>	$x$ with a circumflex
<code>tilde(x)</code>	$x$ with a tilde
<code>dot(x)</code>	$x$ with a dot
<code>ring(x)</code>	$x$ with a ring
<code>bar(xy)</code>	$xy$ with bar
<code>widehat(xy)</code>	$xy$ with a wide circumflex
<code>widetilde(xy)</code>	$xy$ with a wide tilde
<code>x %&lt;-&gt;% y</code>	$x$ double-arrow $y$
<code>x %-&gt;% y</code>	$x$ right-arrow $y$
<code>x %&lt;-% y</code>	$x$ left-arrow $y$
<code>x %up% y</code>	$x$ up-arrow $y$
<code>x %down% y</code>	$x$ down-arrow $y$
<code>x %&lt;=&gt;% y</code>	$x$ is equivalent to $y$
<code>x %=&gt;% y</code>	$x$ implies $y$
<code>x %&lt;=% y</code>	$y$ implies $x$
<code>x %dblup% y</code>	$x$ double-up-arrow $y$
<code>x %dbldown% y</code>	$x$ double-down-arrow $y$
<code>alpha - omega</code>	Greek symbols
<code>Alpha - Omega</code>	uppercase Greek symbols
<code>infinity</code>	infinity symbol
<code>partialdiff</code>	partial differential symbol
<code>32*degree</code>	32 degrees

<code>60*minute</code>	60 minutes of angle
<code>30*second</code>	30 seconds of angle
<code>displaystyle(x)</code>	draw x in normal size (extra spacing)
<code>textstyle(x)</code>	draw x in normal size
<code>scriptstyle(x)</code>	draw x in small size
<code>scriptscriptstyle(x)</code>	draw x in very small size
<code>x ~~ y</code>	put extra space between x and y
<code>x + phantom(0) + y</code>	leave gap for "0", but don't draw it
<code>x + over(1, phantom(0))</code>	leave vertical gap for "0" (don't draw)
<code>frac(x, y)</code>	x over y
<code>over(x, y)</code>	x over y
<code>atop(x, y)</code>	x over y (no horizontal bar)
<code>sum(x[i], i==1, n)</code>	sum x[i] for i equals 1 to n
<code>prod(plain(P)(X==x), x)</code>	product of P(X=x) for all values of x
<code>integral(f(x)*dx, a, b)</code>	definite integral of f(x) wrt x
<code>union(A[i], i==1, n)</code>	union of A[i] for i equals 1 to n
<code>intersect(A[i], i==1, n)</code>	intersection of A[i]
<code>lim(f(x), x %-&gt;% 0)</code>	limit of f(x) as x tends to 0
<code>min(g(x), x &gt; 0)</code>	minimum of g(x) for x greater than 0
<code>inf(S)</code>	infimum of S
<code>sup(S)</code>	supremum of S
<code>x^y + z</code>	normal operator precedence
<code>x^(y + z)</code>	visible grouping of operands
<code>x^{y + z}</code>	invisible grouping of operands
<code>group("(", list(a, b), ")")</code>	specify left and right delimiters
<code>bgroup("(", atop(x,y), ")")</code>	use scalable delimiters
<code>group(lceil, x, rceil)</code>	special delimiters

## References

Murrell, P. and Ihaka, R. (2000) An approach to providing mathematical annotation in plots. *Journal of Computational and Graphical Statistics*, **9**, 582–599.

## See Also

`demo(plotmath)`, `axis`, `mtext`, `text`, `title`, `substitute quote`, `bquote`

## Examples

```
x <- seq(-4, 4, len = 101)
y <- cbind(sin(x), cos(x))
matplot(x, y, type = "l", xaxt = "n",
```

```

main = expression(paste(plain(sin) * phi, " and ",
                          plain(cos) * phi)),
# only 1st is taken
ylab = expression("sin" * phi, "cos" * phi),
xlab = expression(paste("Phase Angle ", phi)),
col.main = "blue")
axis(1, at = c(-pi, -pi/2, 0, pi/2, pi),
     lab = expression(-pi, -pi/2, 0, pi/2, pi))

## How to combine "math" and numeric variables :
plot(1:10, type="n", xlab="", ylab="",
     main = "plot math & numbers")
theta <- 1.23 ; mtext(bquote(hat(theta) == .(theta)))
for(i in 2:9)
  text(i,i+1,
       substitute(list(xi,eta) == group("(",list(x,y),")"),
                  list(x=i, y=i+1)))

plot(1:10, 1:10)
text(4, 9,
     expression(hat(beta) == (X^t * X)^{-1} * X^t * y))
text(4, 8.4,
     "expression(hat(beta) == (X^t * X)^{-1} * X^t * y)",
     cex = .8)
text(4, 7,
     expression(bar(x) == sum(frac(x[i], n), i==1, n)))
text(4, 6.4,
     "expression(bar(x) == sum(frac(x[i], n), i==1, n))",
     cex = .8)
text(8, 5,
     expression(paste(frac(1, sigma*sqrt(2*pi)), " ",
                       plain(e)^{frac(-(x-mu)^2, 2*sigma^2)})),
     cex = 1.2)

```



---

**png**     *JPEG and PNG graphics devices*

---

**Description**

A graphics device for JPEG or PNG format bitmap files.

**Usage**

```
jpeg(filename = "Rplot%03d.jpeg", width = 480, height = 480,  
      pointsize = 12, quality = 75, bg = "white", ...)  
png(filename = "Rplot%03d.png", width = 480, height = 480,  
     pointsize = 12, bg = "white", ...)
```

**Arguments**

<b>filename</b>	the name of the output file. The page number is substituted if an integer format is included in the character string. (The result must be less than <code>PATH_MAX</code> characters long, and may be truncated if not.) Tilde expansion is performed where supported by the platform.
<b>width</b>	the width of the device in pixels.
<b>height</b>	the height of the device in pixels.
<b>pointsize</b>	the default pointsize of plotted text.
<b>quality</b>	the ‘quality’ of the JPEG image, as a percentage. Smaller values will give more compression but also more degradation of the image.
<b>bg</b>	default background colour.
<b>...</b>	additional arguments to the X11 device.

**Details**

Plots in PNG and JPEG format can easily be converted to many other bitmap formats, and both can be displayed in most modern web browsers. The PNG format is lossless and is best for line diagrams and blocks of solid colour. The JPEG format is lossy, but may be useful for image plots, for example.

**png** supports transparent backgrounds: use **bg** = **"transparent"**. Not all PNG viewers render files with transparency correctly. When transparency is in use a very light grey is used as the background and so will

appear as transparent if used in the plot. This allows opaque white to be used, as on the example.

R can be compiled without support for either or both of these devices: this will be reported if you attempt to use them on a system where they are not supported. They will not be available if R has been started with ‘`--gui=none`’ (and will give a different error message), and they may not be usable unless the X11 display is available to the owner of the R process.

## Value

A plot device is opened: nothing is returned to the R interpreter.

## Warning

If you plot more than one page on one of these devices and do not include something like `%d` for the sequence number in `file`, the file will contain the last page plotted.

## Note

These are based on the `X11` device, so the additional arguments to that device work, but are rarely appropriate. The colour handling will be that of the `X11` device in use.

## Author(s)

Guido Masarotto and Brian Ripley

## See Also

Devices, `dev.print`

`capabilities` to see if these devices are supported by this build of R.

`bitmap` provides an alternative way to generate PNG and JPEG plots that does not depend on accessing the X11 display but does depend on having GhostScript installed.

## Examples

```
## these examples will work only if the devices are
## available and the X11 display is available.

## copy current plot to a PNG file
dev.print(png, file="myplot.png", width=480, height=480)
```

```
png(file="myplot.png", bg="transparent")
plot(1:10)
rect(1, 5, 3, 7, col="white")
dev.off()
```

```
jpeg(file="myplot.jpeg")
example(rect)
dev.off()
```

---

**points**     *Add Points to a Plot*

---

## Description

**points** is a generic function to draw a sequence of points at the specified coordinates. The specified character(s) are plotted, centered at the coordinates.

## Usage

```
points(x, ...)
```

```
## Default S3 method:
```

```
points(x, y = NULL, type = "p", pch = par("pch"),  
       col = par("col"), bg = NA, cex = 1, ...)
```

## Arguments

**x, y**            coordinate vectors of points to plot.

**type**           character indicating the type of plotting; actually any of the **types** as in **plot**.

**pch**            plotting “character”, i.e., symbol to use. **pch** can either be a **character** or an integer code for a set of graphics symbols. The full set of S symbols is available with **pch=0:18**, see the last picture from **example(points)**, i.e., the examples below.

In addition, there is a special set of R plotting symbols which can be obtained with **pch=19:25** and **21:25** can be colored and filled with different colors:

- **pch=19**: solid circle,
- **pch=20**: bullet (smaller circle),
- **pch=21**: circle,
- **pch=22**: square,
- **pch=23**: diamond,
- **pch=24**: triangle point-up,
- **pch=25**: triangle point down.

Values **pch=26:32** are currently unused, and **pch=32:255** give the text symbol in the encoding in use (see **postscript**).

<code>col</code>	color code or name, see <code>par</code> .
<code>bg</code>	background (“fill”) color for open plot symbols
<code>cex</code>	character expansion: a numerical vector.
<code>...</code>	Further graphical parameters (see <code>plot.xy</code> and <code>par</code> ) may also be supplied as arguments.

## Details

The coordinates can be passed in a plotting structure (a list with `x` and `y` components), a two-column matrix, a time series, .... See `xy.coords`.

Arguments `pch`, `col`, `bg` and `cex` can be vectors (which will be recycled as needed) giving a value for each point plotted. Points whose `x`, `y`, `pch`, `col` or `cex` value is `NA` are omitted from the plot.

Graphical parameters are permitted as arguments to this function.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`plot`, `lines`, and the underlying “primitive” `plot.xy`.

## Examples

```
plot(-4:4, -4:4, type = "n") # setting up coord. system
points(rnorm(200), rnorm(200), col = "red")
points(rnorm(100)/2, rnorm(100)/2, col = "blue", cex = 1.5)

op <- par(bg = "light blue")
x <- seq(0,2*pi, len=51)
## something "between type='b' and type='o'":
plot(x, sin(x), type="o", pch=21, bg=par("bg"),
     col = "blue", cex=.6,
     main='plot(..., type="o", pch=21, bg=par("bg"))')
par(op)

## Showing all the extra & some char graphics symbols
Pex <- 3 ## good for both .Device=="postscript" and "x11"
ipch <- 1:(np <- 25+11); k <- floor(sqrt(np));
dd <- c(-1,1)/2
rx <- dd + range(ix <- (ipch-1) %/% k)
```

```

ry <- dd + range(iy <- 3 + (k-1)-(ipch-1) %% k)
pch <- as.list(ipch)
pch[25+ 1:11] <-
  as.list(c("o", "O", "+", "-", ":", "|", "%", "#"))
plot(rx, ry, type="n", axes = FALSE, xlab = "", ylab = "",
      main = paste("symbols: points(... pch=*, cex=", Pex, ")"))
abline(v = ix, h = iy, col = "lightgray", lty = "dotted")
for(i in 1:np) {
  pc <- pch[[i]]
  points(ix[i], iy[i], pch = pc, col = "red",
         bg = "yellow", cex = Pex)
  ## red symbols with a yellow interior (where available)
  text(ix[i] - .3, iy[i], pc, col = "brown", cex = 1.2)
}

```

---

**polygon**     *Polygon Drawing*


---

**Description**

`polygon` draws the polygons whose vertices are given in `x` and `y`.

**Usage**

```
polygon(x, y = NULL, density = NULL, angle = 45,
        border = NULL, col = NA, lty = NULL, xpd = NULL, ...)
```

**Arguments**

<code>x,y</code>	vectors containing the coordinates of the vertices of the polygon.
<code>density</code>	the density of shading lines, in lines per inch. The default value of <code>NULL</code> means that no shading lines are drawn. A zero value of <code>density</code> means no shading lines whereas negative values (and <code>NA</code> ) suppress shading (and so allow color filling).
<code>angle</code>	the slope of shading lines, given as an angle in degrees (counter-clockwise).
<code>col</code>	the color for filling the polygon. The default, <code>NA</code> , is to leave polygons unfilled.
<code>border</code>	the color to draw the border. The default, <code>NULL</code> , uses <code>par("fg")</code> . Use <code>border = NA</code> to omit borders. For compatibility with S, <code>border</code> can also be logical, in which case <code>FALSE</code> is equivalent to <code>NA</code> (borders omitted) and <code>TRUE</code> is equivalent to <code>NULL</code> (use the foreground colour),
<code>lty</code>	the line type to be used, as in <code>par</code> .
<code>xpd</code>	(where) should clipping take place? Defaults to <code>par("xpd")</code> .
<code>...</code>	graphical parameters can be given as arguments to <code>polygon</code> .

## Details

The coordinates can be passed in a plotting structure (a list with **x** and **y** components), a two-column matrix, .... See **xy.coords**.

It is assumed that the polygon is closed by joining the last point to the first point.

The coordinates can contain missing values. The behaviour is similar to that of **lines**, except that instead of breaking a line into several lines, NA values break the polygon into several complete polygons (including closing the last point to the first point). See the examples below.

When multiple polygons are produced, the values of **density**, **angle**, **col**, **border**, and **lty** are recycled in the usual manner.

## Bugs

The present shading algorithm can produce incorrect results for self-intersecting polygons.

## Author(s)

The code implementing polygon shading was donated by Kevin Buhr.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**segments** for even more flexibility, **lines**, **rect**, **box**, **abline**.

**par** for how to specify colors.

## Examples

```
x <- c(1:9,8:1)
y <- c(1,2*(5:3),2,-1,17,9,8,2:9)
op <- par(mfcol=c(3,1))
for(xpd in c(FALSE,TRUE,NA)) {
  plot(1:10, main=paste("xpd =", xpd)) ;
  box("figure", col = "pink", lwd=3)
  polygon(x,y, xpd=xpd, col = "orange", lty=2, lwd=2,
          border = "red")
}
par(op)
```



```

n <- 100
xx <- c(0:n, n:0)
yy <- c(c(0,cumsum(rnorm(n))), rev(c(0,cumsum(rnorm(n)))))
plot  (xx, yy, type="n", xlab="Time", ylab="Distance")
polygon(xx, yy, col="gray", border = "red")
title("Distance Between Brownian Motions")

# Multiple polygons from NA values
# and recycling of col, border, and lty
op <- par(mfrow=c(2,1))
plot(c(1,9), 1:2, type="n")
polygon(1:9, c(2,1,2,1,1,2,1,2,1),
       col=c("red", "blue"),
       border=c("green", "yellow"),
       lwd=3, lty=c("dashed", "solid"))
plot(c(1,9), 1:2, type="n")
polygon(1:9, c(2,1,2,1,NA,2,1,2,1),
       col=c("red", "blue"),
       border=c("green", "yellow"),
       lwd=3, lty=c("dashed", "solid"))
par(op)

# Line-shaded polygons
plot(c(1,9), 1:2, type="n")
polygon(1:9, c(2,1,2,1,NA,2,1,2,1),
       density=c(10, 20), angle=c(-45, 45))

```

---

**postscript**     *PostScript Graphics*


---

**Description**

`postscript` starts the graphics device driver for producing PostScript graphics.

The auxiliary function `ps.options` can be used to set and view (if called without arguments) default values for the arguments to `postscript`.

**Usage**

```
postscript(file = ifelse(onefile, "Rplots.ps",
  "Rplot%03d.ps"), onefile = TRUE, paper, family, encoding,
  bg, fg, width, height, horizontal, pointsize, pagecentre,
  print.it, command, title = "R Graphics Output")

ps.options(paper, horizontal, width, height, family,
  encoding, pointsize, bg, fg, onefile = TRUE,
  print.it = FALSE, append = FALSE, reset = FALSE,
  override.check = FALSE)
.PostScript.Options
```

**Arguments**

- |                   |   |
|-------------------|---|
| <b>file</b>       | a character string giving the name of the file. If it is "", the output is piped to the command given by the argument <code>command</code> . If it is " cmd", the output is piped to the command given by 'cmd'.<br>For use with <code>onefile=FALSE</code> give a <code>printf</code> format such as "Rplot%03d.ps" (the default in that case).  |
| <b>paper</b>      | the size of paper in the printer. The choices are "a4", "letter", "legal" and "executive" (and these can be capitalized). Also, "special" can be used, when the <code>width</code> and <code>height</code> specify the paper size. A further choice is "default", which is the default. If this is selected, the papersize is taken from the option "papersize" if that is set and to "a4" if it is unset or empty. |
| <b>horizontal</b> | the orientation of the printed image, a logical. Defaults to true, that is landscape orientation.   |

<code>width, height</code>	the width and height of the graphics region in inches. The default is to use the entire page less a 0.25 inch border on each side.
<code>family</code>	the font family to be used. EITHER a single character string OR a character vector of length four or five. See the section ‘Families’.
<code>encoding</code>	the name of an encoding file. Defaults to "ISO-Latin1.enc" in the 'R_HOME/afm' directory, which is used if the path does not contain a path separator. An extension ".enc" can be omitted.
<code>pointsize</code>	the default point size to be used.
<code>bg</code>	the default background color to be used. If "transparent" (or an equivalent specification), no background is painted.
<code>fg</code>	the default foreground color to be used.
<code>onefile</code>	logical: if true (the default) allow multiple figures in one file. If false, generate a file name containing the page number and use an EPSF header and no DocumentMedia comment.
<code>pagecentre</code>	logical: should the device region be centred on the page: defaults to true.
<code>print.it</code>	logical: should the file be printed when the device is closed? (This only applies if <code>file</code> is a real file name.)
<code>command</code>	the command to be used for “printing”. Defaults to option "printcmd"; this can also be selected as "default".
<code>append</code>	logical; currently <b>disregarded</b> ; just there for compatibility reasons.
<code>reset, override.check</code>	logical arguments passed to <code>check.options</code> . See the Examples.
<code>title</code>	title string to embed in the file.

## Details

`postscript` opens the file `file` and the PostScript commands needed to plot any graphics requested are stored in that file. This file can then be printed on a suitable device to obtain hard copy.

A postscript plot can be printed via `postscript` in two ways.

1. Setting `print.it = TRUE` causes the command given in argument `command` to be called with argument `"file"` when the device is closed. Note that the plot file is not deleted unless `command` arranges to delete it.
2. `file=""` or `file="|cmd"` can be used to print using a pipe on systems that support `'popen'`. Failure to open the command will probably be reported to the terminal but not to `'popen'`, in which case close the device by `dev.off` immediately.

The postscript produced by R is EPS (*Encapsulated PostScript*) compatible, and can be included into other documents, e.g., into LaTeX, using

`includegraphics{<filename>}`. For use in this way you will probably want to set `horizontal = FALSE`, `onefile = FALSE`, `paper = "special"`.

Most of the PostScript prologue used is taken from the R character vector `.ps.prolog`. This is marked in the output, and can be changed by changing that vector. (This is only advisable for PostScript experts.)

`ps.options` needs to be called before calling `postscript`, and the default values it sets can be overridden by supplying arguments to `postscript`.

## Families

The argument `family` specifies the font family to be used. In normal use it is one of `"AvantGarde"`, `"Bookman"`, `"Courier"`, `"Helvetica"`, `"Helvetica-Narrow"`, `"NewCenturySchoolbook"`, `"Palatino"` or `"Times"`, and refers to the standard Adobe PostScript fonts of those names which are included (or cloned) in all common PostScript devices.

Many PostScript emulators (including those based on `ghostscript`) use the URW equivalents of these fonts, which are `"URWGothic"`, `"URWBookman"`, `"NimbusMon"`, `"NimbusSan"`, `"NimbusSanCond"`, `"CenturySch"`, `"URWPalladio"` and `"NimbusRom"` respectively. If your PostScript device is using URW fonts, you will obtain access to more characters and more appropriate metrics by using these names. To make these easier to remember, `"URWHelvetica" == "NimbusSan"` and `"URWTimes" == "NimbusRom"` are also supported.

It is also possible to specify `family="ComputerModern"`. This is intended to use with the Type 1 versions of the TeX CM fonts. It will normally be possible to include such output in TeX or LaTeX provided it is processed with `dvips -Ppfb -j0` or the equivalent on your system. (`-j0` turns off font subsetting.)

If the second form of argument `"family"` is used, it should be a character vector of four or five paths to Adobe Font Metric files for the regular, bold, italic, bold italic and (optionally) symbol fonts to be used. If these paths do not contain the file separator, they are taken to refer to files in the R directory `'R_HOME/afm'`. Thus the default Helvetica family can be specified by `family = c("hv_____.afm", "hvb____.afm", "hvo____.afm", "hvbo____.afm", "sy_____.afm")`. It is the user's responsibility to check that suitable fonts are made available, and that they contain the needed characters when re-encoded. The fontnames used are taken from the `FontName` fields of the afm files. The software including the PostScript plot file should either embed the font outlines (usually from `' .pfb'` or `' .pfa'` files) or use DSC comments to instruct the print spooler to do so.

## Encodings

Encodings describe which glyphs are used to display the character codes (in the range 0–255). By default R uses ISOLatin1 encoding, and the examples for `text` are in that encoding. However, the encoding used on machines running R may well be different, and by using the `encoding` argument the glyphs can be matched to encoding in use.

None of this will matter if only ASCII characters (codes 32–126) are used as all the encodings agree over that range. Some encodings are supersets of ISOLatin1, too. However, if accented and special characters do not come out as you expect, you may need to change the encoding. Three other encodings are supplied with R: `"WinAnsi.enc"` and `"MacRoman.enc"` correspond to the encodings normally used on Windows and MacOS (at least by Adobe), and `"PDFDoc.enc"` is the first 256 characters of the Unicode encoding, the standard for PDF.

If you change the encoding, it is your responsibility to ensure that the PostScript font contains the glyphs used. One issue here is the Euro symbol which is in the WinAnsi and MacRoman encodings but may well not be in the PostScript fonts. (It is in the URW variants; it is not in the supplied Adobe Font Metric files.)

There is one exception. Character 45 ("-") is always set as minus (its value in Adobe ISOLatin1) even though it is hyphen in the other encodings. Hyphen is available as character 173 (octal 0255) in ISOLatin1.

## Author(s)

Support for Computer Modern fonts is based on a contribution by Brian D'Urso.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

Devices, `check.options` which is called from both `ps.options` and `postscript`.

## Examples

```
# open the file "foo.ps" for graphics output
postscript("foo.ps")
# produce the desired graph(s)
dev.off()           # turn off the postscript device
postscript("|lp -dlw")
# produce the desired graph(s)
dev.off()           # plot will appear on printer

# for URW PostScript devices
postscript("foo.ps", family = "NimbusSan")

# for inclusion in Computer Modern TeX documents, perhaps
postscript("cm_test.eps", width = 4.0, height = 3.0,
  horizontal = FALSE, onefile = FALSE, paper = "special",
  family = "ComputerModern")
# The resultant postscript file can be used by dvips -Ppfb
# -j0.

# To test out encodings, you can use
TestChars <-
  function (encoding="ISOLatin1", family="URWHelvetica")
  {
    postscript(encoding=encoding, family=family)
    par(pty="s")
    plot(c(0,15), c(0,15), type="n", xlab="", ylab="")
    title(paste("Centred chars in encoding", encoding))
    grid(15, 15, lty=1)
    for(i in c(32:255)) {
      x <- i
      y <- i
      points(x, y, pch=i)
    }
  }
```

```
    dev.off()
}
## there will be many warnings. We use URW to get a
## complete enough set of font metrics.
TestChars()
TestChars("ISOLatin2")
TestChars("WinAnsi")

stopifnot(
  unlist(ps.options()) == unlist(.PostScript.Options)
)
ps.options(bg = "pink")
str(ps.options(reset = TRUE))

### ---- error checking of arguments: ----
ps.options(width=0:12, onefile=0, bg=pi)
# override the check for 'onefile', but not the others:
str(ps.options(width=0:12, onefile=1, bg=pi,
               override.check = c(FALSE,TRUE,FALSE)))
```

---

**ppoints**      *Ordinates for Probability Plotting*

---

**Description**

Generates the sequence of “probability” points  $(1:m - a)/(m + (1-a)-a)$  where  $m$  is either  $n$ , if `length(n)==1`, or `length(n)`.

**Usage**

```
ppoints(n, a = ifelse(n <= 10, 3/8, 1/2))
```

**Arguments**

<code>n</code>	either the number of points generate or a vector of observations.
<code>a</code>	the offset fraction to be used; typically in $(0, 1)$ .

**Details**

If  $0 < a < 1$ , the resulting values are within  $(0, 1)$  (excluding boundaries). In any case, the resulting sequence is symmetric in  $[0, 1]$ , i.e., `p + rev(p) == 1`.

`ppoints()` is used in `qqplot` and `qqnorm` to generate the set of probabilities at which to evaluate the inverse distribution.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`qqplot`, `qqnorm`.

**Examples**

```
ppoints(4) # the same as ppoints(1:4)
ppoints(10)
ppoints(10, a=1/2)
```



---

**preplot**     *Pre-computations for a Plotting Object*

---

**Description**

Compute an object to be used for plots relating to the given model object.

**Usage**

```
preplot(object, ...)
```

**Arguments**

<code>object</code>	a fitted model object.
<code>...</code>	additional arguments for specific methods.

**Details**

Only the generic function is currently provided in base R, but some add-on packages have methods. Principally here for S compatibility.

**Value**

An object set up to make a plot that describes `object`.

---

**pretty**     *Pretty Breakpoints*


---

**Description**

Compute a sequence of about  $n+1$  equally spaced nice values which cover the range of the values in `x`. The values are chosen so that they are 1, 2 or 5 times a power of 10.

**Usage**

```
pretty(x, n = 5, min.n = n %% 3, shrink.sml = 0.75,
      high.u.bias = 1.5, u5.bias = .5 + 1.5*high.u.bias,
      eps.correct = 0)
```

**Arguments**

<code>x</code>	numeric vector
<code>n</code>	integer giving the <i>desired</i> number of intervals. Non-integer values are rounded down.
<code>min.n</code>	nonnegative integer giving the <i>minimal</i> number of intervals. If <code>min.n == 0</code> , <code>pretty(.)</code> may return a single value.
<code>shrink.sml</code>	positive numeric by which a default scale is shrunk in the case when <code>range(x)</code> is “very small” (usually 0).
<code>high.u.bias</code>	non-negative numeric, typically $> 1$ . The interval unit is determined as $\{1,2,5,10\}$ times <code>b</code> , a power of 10. Larger <code>high.u.bias</code> values favor larger units.
<code>u5.bias</code>	non-negative numeric multiplier favoring factor 5 over 2. Default and “optimal”: <code>u5.bias = .5 + 1.5*high.u.bias</code> .
<code>eps.correct</code>	integer code, one of $\{0,1,2\}$ . If non-0, an “ <i>epsilon correction</i> ” is made at the boundaries such that the result boundaries will be outside <code>range(x)</code> ; in the <i>small</i> case, the correction is only done if <code>eps.correct &gt;= 2</code> .

## Details

Let  $d \leftarrow \max(x) - \min(x) \geq 0$ . If  $d$  is not (very close) to 0, we let  $c \leftarrow d/n$ , otherwise more or less  $c \leftarrow \max(\text{abs}(\text{range}(x))) * \text{shrink} \cdot \text{sm1} / \min.n$ . Then, the 10 base  $b$  is  $10^{\lfloor \log_{10}(c) \rfloor}$  such that  $b \leq c < 10b$ .

Now determine the basic *unit*  $u$  as one of  $\{1, 2, 5, 10\}b$ , depending on  $c/b \in [1, 10)$  and the two “*bias*” coefficients,  $h = \text{high.u.bias}$  and  $f = \text{u5.bias}$ .

.....

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## Examples

```
pretty(1:15)          # 0  2  4  6  8 10 12 14 16
pretty(1:15, h=2)     # 0  5 10 15
pretty(1:15, n=4)     # 0  5 10 15
pretty(1:15 * 2)      # 0  5 10 15 20 25 30
pretty(1:20)          # 0  5 10 15 20
pretty(1:20, n=2)     # 0 10 20
pretty(1:20, n=10)    # 0  2  4 ... 20

for(k in 5:11) {
  cat("k=", k, ": ");
  print(diff(range(pretty(100 + c(0, pi*10^-k)))))
}

## more bizarre, when min(x) == max(x):
pretty(pi)

add.names <- function(v) { names(v) <- paste(v); v }
str(lapply(add.names(-10:20), pretty))
str(lapply(add.names(0:20),    pretty, min = 0))
sapply(  add.names(0:20),    pretty, min = 4)

pretty(1.234e100)
pretty(1001.1001)
pretty(1001.1001, shrink = .2)
for(k in -7:3)
  cat("shrink=", formatC(2^k, wid=9), ":",
      formatC(pretty(1001.1001, shrink = 2^k), wid=6), "\n")
```

---

**qqnorm**      *Quantile-Quantile Plots*


---

**Description**

`qqnorm` is a generic function the default method of which produces a normal QQ plot of the values in `y`. `qqline` adds a line to a normal quantile-quantile plot which passes through the first and third quartiles.

`qqplot` produces a QQ plot of two datasets.

Graphical parameters may be given as arguments to `qqnorm`, `qqplot` and `qqline`.

**Usage**

```
qqnorm(y, ...)
## Default S3 method:
qqnorm(y, ylim, main = "Normal Q-Q Plot",
       xlab = "Theoretical Quantiles",
       ylab = "Sample Quantiles", plot.it = TRUE, datax = FALSE,
       ...)
qqline(y, datax = FALSE, ...)
qqplot(x, y, plot.it = TRUE, xlab = deparse(substitute(x)),
       ylab = deparse(substitute(y)), ...)
```

**Arguments**

<code>x</code>	The first sample for <code>qqplot</code> .
<code>y</code>	The second or only data sample.
<code>xlab</code> , <code>ylab</code> , <code>main</code>	plot labels.
<code>plot.it</code>	logical. Should the result be plotted?
<code>datax</code>	logical. Should data values be on the x-axis?
<code>ylim</code> , ...	graphical parameters.

**Value**

For `qqnorm` and `qqplot`, a list with components

<code>x</code>	The x coordinates of the points that were/would be plotted
<code>y</code>	The original <code>y</code> vector, i.e., the corresponding <code>y</code> coordinates <i>including NAs</i> .

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

ppoints.

## Examples

```
y <- rt(200, df = 5)
qqnorm(y); qqline(y, col = 2)
qqplot(y, rt(300, df = 5))
data(precip)
qqnorm(precip, ylab = "Precipitation [in/yr] - US cities")
```

---

**quartz**     *MacOS X Quartz device*

---

## Description

**quartz** starts a graphics device driver for the MacOS X System. This can only be done on machines that run MacOS X.

## Usage

```
quartz(display = "", width = 5, height = 5, pointsize = 12,  
       family = "Helvetica", antialias=TRUE, autorefresh=TRUE)
```

## Arguments

<b>display</b>	the display on which the graphics window will appear. The default is to use the value in the user's environment variable <b>DISPLAY</b> .
<b>width</b>	the width of the plotting window in inches.
<b>height</b>	the height of the plotting window in inches.
<b>pointsize</b>	the default pointsize to be used.
<b>family</b>	this is the family name of the Postscript font that will be used by the device.
<b>antialias</b>	whether to use antialiasing. It is never the case to set it <b>FALSE</b>
<b>autorefresh</b>	logical specifying if realtime refreshing should be done. If <b>FALSE</b> , the system is charged to refresh the context of the device window.

## Details

Quartz is the graphic engine based on the PDF format. It is used by the graphic interface of MacOS X to render high quality graphics. As PDF it is device independent and can be rescaled without loss of definition.

Calling **quartz()** sets **.Device** to "quartz".

## See Also

**Devices.**

---

**recordPlot**      *Record and Replay Plots*

---

**Description**

Functions to save the current plot in an R variable, and to replay it.

**Usage**

```
recordPlot()  
replayPlot(x)
```

**Arguments**

**x**                      A saved plot.

**Details**

These functions record and replay the displaylist of the current graphics device. The returned object is of class "**recordedplot**", and **replayPlot** acts as a **print** method for that class.

The format of recorded plots was changed in R 1.4.0: plots saved in earlier versions can still be replayed.

**Value**

**recordPlot** returns an object of class "**recordedplot**", a list with components:

**displaylist**      The saved display list, as a pairlist.  
**gpar**              The graphics state, as an integer vector.

**replayPlot** has no return value.

---

**rect**     *Draw a Rectangle*

---

## Description

**rect** draws a rectangle (or sequence of rectangles) with the given coordinates, fill and border colors.

## Usage

```
rect(xleft, ybottom, xright, ytop, density = NULL,
     angle = 45, col = NULL, border = NULL, lty = NULL,
     lwd = par("lwd"), xpd = NULL, ...)
```

## Arguments

<b>xleft</b>	a vector (or scalar) of left x positions.
<b>ybottom</b>	a vector (or scalar) of bottom y positions.
<b>xright</b>	a vector (or scalar) of right x positions.
<b>ytop</b>	a vector (or scalar) of top y positions.
<b>density</b>	the density of shading lines, in lines per inch. The default value of <b>NULL</b> means that no shading lines are drawn. A zero value of <b>density</b> means no shading lines whereas negative values (and <b>NA</b> ) suppress shading (and so allow color filling).
<b>angle</b>	angle (in degrees) of the shading lines.
<b>col</b>	color(s) to fill or shade the rectangle(s) with. The default <b>NULL</b> , or also <b>NA</b> do not fill, i.e., draw transparent rectangles, unless <b>density</b> is specified.
<b>border</b>	color for rectangle border(s).
<b>lty</b>	line type for borders; defaults to <b>"solid"</b> .
<b>lwd</b>	width for borders.
<b>xpd</b>	logical ( <b>"expand"</b> ); defaults to <b>par("xpd")</b> . See <b>par(xpd= )</b> .
<b>...</b>	other graphical parameters can be given as arguments.



## Details

The positions supplied, i.e., `xleft`, ..., are relative to the current plotting region. If the x-axis goes from 100 to 200 then `xleft` must be larger than 100 and `xright` must be less than 200.

It is a primitive function used in `hist`, `barplot`, `legend`, etc.

## See Also

`box` for the “standard” box around the plot; `polygon` and `segments` for flexible line drawing.

`par` for how to specify colors.

## Examples

```
## set up the plot region:
op <- par(bg = "thistle")
plot(c(100, 250), c(300, 450), type = "n", xlab="", ylab="",
      main = "2x11 rectangles; 'rect(100+i,300+i,150+i,380+i)'" )
i <- 4*(0:10)
## draw rectangles with bottom left (100, 300)+i and top
## right (150, 380)+i
rect(100+i, 300+i, 150+i, 380+i,
     col=rainbow(11, start=.7,end=.1))
rect(240-i, 320+i, 250-i, 410+i,
     col=heat.colors(11), lwd=i/5)
## Background alternating (transparent / "bg"):
j <- 10*(0:5)
rect(125+j, 360+j, 141+j, 405+j/2, col = c(NA,0),
     border = "gold", lwd = 2)
rect(125+j, 296+j/2, 141+j, 331+j/5,
     col = c(NA,"midnightblue"))
mtext("+2x6 rect(*,col=c(NA,0)) and col=c(NA,\"m..blue\")")

## an example showing colouring and shading
plot(c(100, 200), c(300, 450), type= "n", xlab="", ylab="")
# transparent
rect(100, 300, 125, 350)
# coloured
rect(100, 400, 125, 450, col="green", border="blue")
rect(115, 375, 150, 425, col=par("bg"),
     border="transparent")
rect(150, 300, 175, 350, density=10, border="red")
rect(150, 400, 175, 450, density=30, col="blue",
```

```
angle=-30, border="transparent")

legend(180, 450, legend=1:4,
      fill=c(NA, "green", par("fg"), "blue"),
      density=c(NA, NA, 10, 30), angle=c(NA, NA, 30, -30))

par(op)
```

---

**rgb**     *RGB Color Specification*

---

**Description**

This function creates “colors” corresponding to the given intensities (between 0 and `max`) of the red, green and blue primaries. The `names` argument may be used to provide names for the colors.

The values returned by `rgb` can be used with a `col=` specification in graphics functions or in `par`.

**Usage**

```
rgb(red, green, blue, names=NULL, maxColorValue = 1)
```

**Arguments**

`red`, `blue`, `green`

vectors of same length with values in  $[0, M]$  where  $M$  is `maxColorValue`. When this is 255, the `red`, `blue` and `green` values are coerced to integers in 0:255 and the result is computed most efficiently.

`names`                      character. The names for the resulting vector.

`maxColorValue`            number giving the maximum of the color values range, see above.

**See Also**

`col2rgb` the “inverse” for translating R colors to RGB vectors; `rainbow`, `hsv`, `gray`.

**Examples**

```
rgb(0,1,0)
(u01 <- seq(0,1, length=11))
stopifnot(rgb(u01,u01,u01) == gray(u01))
reds <- rgb((0:15)/15, g=0,b=0,
            names=paste("red",0:15,sep="."))
reds

rgb(0, 0:12, 0, max = 255) # integer input
```

---

**rug**     *Add a Rug to a Plot*

---

## Description

Adds a *rug* representation (1-d plot) of the data to the plot.

## Usage

```
rug(x, ticksize=0.03, side=1, lwd=0.5, col,  
    quiet = getOption("warn") < 0, ...)
```

## Arguments

<b>x</b>	A numeric vector
<b>ticksize</b>	The length of the ticks making up the ‘rug’. Positive lengths give inwards ticks.
<b>side</b>	On which side of the plot box the rug will be plotted. Normally 1 (bottom) or 3 (top).
<b>lwd</b>	The line width of the ticks.
<b>col</b>	The colour the ticks are plotted in, default is black.
<b>quiet</b>	logical indicating if there should be a warning about clipped values.
<b>...</b>	further arguments, passed to <code>axis(...)</code> , such as <code>line</code> or <code>pos</code> for specifying the location of the rug.

## Details

Because of the way **rug** is implemented, only values of **x** that fall within the plot region are included. There will be a warning if any finite values are omitted, but non-finite values are omitted silently.

Because of the way colours are done the axis itself is coloured the same as the ticks. You can always replot the box in black if you don't like this feature.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

**See Also**

jitter which you may want for ties in x.

**Examples**

```
data(faithful)
with(faithful, {
  plot(density(eruptions, bw=0.15))
  rug(eruptions)
  rug(jitter(eruptions, amount = .01), side = 3,
      col = "light blue")
})
```

---

**screen**      *Creating and Controlling Multiple Screens on a Single Device*

---

## Description

`split.screen` defines a number of regions within the current device which can, to some extent, be treated as separate graphics devices. It is useful for generating multiple plots on a single device. Screens can themselves be split, allowing for quite complex arrangements of plots.

`screen` is used to select which screen to draw in.

`erase.screen` is used to clear a single screen, which it does by filling with the background colour.

`close.screen` removes the specified screen definition(s).

## Usage

```
split.screen(figs, screen = , erase = TRUE)
screen(n = , new = TRUE)
erase.screen(n = )
close.screen(n, all.screens = FALSE)
```

## Arguments

<b>figs</b>	A two-element vector describing the number of rows and the number of columns in a screen matrix <i>or</i> a matrix with 4 columns. If a matrix, then each row describes a screen with values for the left, right, bottom, and top of the screen (in that order) in NDC units.
<b>screen</b>	A number giving the screen to be split.
<b>erase</b>	logical: should selected screen be cleared?
<b>n</b>	A number indicating which screen to prepare for drawing ( <b>screen</b> ), erase ( <b>erase.screen</b> ), or close ( <b>close.screen</b> ).
<b>new</b>	A logical value indicating whether the screen should be erased as part of the preparation for drawing in the screen.
<b>all.screens</b>	A logical value indicating whether all of the screens should be closed.

## Details

The first call to `split.screen` places R into split-screen mode. The other split-screen functions only work within this mode. While in this mode, certain other commands should be avoided (see WARNINGS below). Split-screen mode is exited by the command `close.screen(all = TRUE)`

## Value

`split.screen` returns a vector of screen numbers for the newly-created screens. With no arguments, `split.screen` returns a vector of valid screen numbers.

`screen` invisibly returns the number of the selected screen. With no arguments, `screen` returns the number of the current screen.

`close.screen` returns a vector of valid screen numbers.

`screen`, `erase.screen`, and `close.screen` all return FALSE if R is not in split-screen mode.

## Warning

The recommended way to use these functions is to completely draw a plot and all additions (ie. points and lines) to the base plot, prior to selecting and plotting on another screen. The behavior associated with returning to a screen to add to an existing plot is unpredictable and may result in problems that are not readily visible.

These functions are totally incompatible with the other mechanisms for arranging plots on a device: `par(mfrow)`, `par(mfcol)`, and `layout()`.

The functions are also incompatible with some plotting functions, such as `coplot`, which make use of these other mechanisms.

The functions should not be used with multiple devices.

`erase.screen` will appear not to work if the background colour is transparent (as it is by default on most devices).

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

## See Also

`par`, `layout`, `Devices`, `dev.*`

## Examples

```
if (interactive()) {
  par(bg = "white") # default is likely to be transparent
  split.screen(c(2,1)) # split display into two screens
  # now split the bottom half into 3
  split.screen(c(1,3), screen = 2)
  screen(1) # prepare screen 1 for output
  plot(10:1)
  screen(4) # prepare screen 4 for output
  plot(10:1)
  close.screen(all = TRUE) # exit split-screen mode

  split.screen(c(2,1))      # split display into two screens
  split.screen(c(1,2),2)    # split bottom half in two
  plot(1:10)                # screen 3 is active, draw plot
  erase.screen()             # forgot label, erase and redraw
  plot(1:10, ylab= "ylab 3")
  screen(1)                 # prepare screen 1 for output
  plot(1:10)
  screen(4)                 # prepare screen 4 for output
  plot(1:10, ylab="ylab 4")
  screen(1, FALSE) # return to screen 1, but do not clear
  # overlay second plot
  plot(10:1, axes=FALSE, lty=2, ylab="")
  axis(4) # add tic marks to right-hand axis
  title("Plot 1")
  close.screen(all = TRUE) # exit split-screen mode
}
```



---

**segments**      *Add Line Segments to a Plot*

---

**Description**

Draw line segments between pairs of points.

**Usage**

```
segments(x0, y0, x1, y1,  
         col = par("fg"), lty = par("lty"), lwd = par("lwd"), ...)
```

**Arguments**

<code>x0,y0</code>	coordinates of points <b>from</b> which to draw.
<code>x1,y1</code>	coordinates of points <b>to</b> which to draw.
<code>col, lty, lwd</code>	usual graphical parameters as in <code>par</code> .
<code>...</code>	further graphical parameters (from <code>par</code> ).

**Details**

For each `i`, a line segment is drawn between the point `(x0[i], y0[i])` and the point `(x1[i], y1[i])`.

The graphical parameters `col` and `lty` can be used to specify a color and line texture for the line segments (`col` may be a vector).

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`arrows`, `polygon` for slightly easier and less flexible line drawing, and `lines` for the usual polygons.

**Examples**

```
x <- runif(12); y <- rnorm(12)  
i <- order(x,y); x <- x[i]; y <- y[i]  
plot(x,y, main="arrows(.) and segments(.)")  
## draw arrows from point to point :
```

```
s <- seq(length(x)-1) # one shorter than data
arrows(x[s], y[s], x[s+1], y[s+1], col= 1:3)
s <- s[-length(s)]
segments(x[s], y[s], x[s+2], y[s+2], col= 'pink')
```

---

**stars**      *Star (Spider/Radar) Plots and Segment Diagrams*


---

**Description**

Draw star plots or segment diagrams of a multivariate data set. With one single location, also draws “spider” (or “radar”) plots.

**Usage**

```
stars(x, full = TRUE, scale = TRUE, radius = TRUE,
      labels = dimnames(x)[[1]], locations = NULL,
      nrow = NULL, ncol = NULL, len = 1, key.loc = NULL,
      key.labels = dimnames(x)[[2]], key.xpd = TRUE,
      xlim = NULL, ylim = NULL, flip.labels = NULL,
      draw.segments = FALSE, col.segments = 1:n.seg,
      col.stars = NA, axes = FALSE, frame.plot = axes,
      main = NULL, sub = NULL, xlab = "", ylab = "",
      cex = 0.8, lwd = 0.25, lty = par("lty"), xpd = FALSE,
      mar = pmin(par("mar"), 1.1+ c(2*axes+ (xlab != ""),
                                     2*axes+ (ylab != ""), 1,0)),
      add=FALSE, plot=TRUE, ...)
```

**Arguments**

- |               |   |
|---------------|---|
| <b>x</b>      | matrix or data frame of data. One star or segment plot will be produced for each row of <b>x</b> . Missing values (NA) are allowed, but they are treated as if they were 0 (after scaling, if relevant).  |
| <b>full</b>   | logical flag: if <b>TRUE</b> , the segment plots will occupy a full circle. Otherwise, they occupy the (upper) semi-circle only.  |
| <b>scale</b>  | logical flag: if <b>TRUE</b> , the columns of the data matrix are scaled independently so that the maximum value in each column is 1 and the minimum is 0. If <b>FALSE</b> , the presumption is that the data have been scaled by some other algorithm to the range [0, 1]. |
| <b>radius</b> | logical flag: in <b>TRUE</b> , the radii corresponding to each variable in the data will be drawn.  |
| <b>labels</b> | vector of character strings for labeling the plots. Unlike the S function <b>stars</b> , no attempt is made to construct labels if <b>labels = NULL</b> .   |

<code>locations</code>	Either two column matrix with the x and y coordinates used to place each of the segment plots; or numeric of length 2 when all plots should be superimposed (for a “spider plot”). By default, <code>locations = NULL</code> , the segment plots will be placed in a rectangular grid.
<code>nrow, ncol</code>	integers giving the number of rows and columns to use when <code>locations</code> is <code>NULL</code> . By default, <code>nrow == ncol</code> , a square layout will be used.
<code>len</code>	scale factor for the length of radii or segments.
<code>key.loc</code>	vector with x and y coordinates of the unit key.
<code>key.labels</code>	vector of character strings for labeling the segments of the unit key. If omitted, the second component of <code>dimnames(x)</code> is used, if available.
<code>key.xpd</code>	clipping switch for the unit key (drawing and labeling), see <code>par("xpd")</code> .
<code>xlim</code>	vector with the range of x coordinates to plot.
<code>ylim</code>	vector with the range of y coordinates to plot.
<code>flip.labels</code>	logical indicating if the label locations should flip up and down from diagram to diagram. Defaults to a somewhat smart heuristic.
<code>draw.segments</code>	logical. If <code>TRUE</code> draw a segment diagram.
<code>col.segments</code>	color vector (integer or character, see <code>par</code> ), each specifying a color for one of the segments (variables). Ignored if <code>draw.segments = FALSE</code> .
<code>col.stars</code>	color vector (integer or character, see <code>par</code> ), each specifying a color for one of the stars (cases). Ignored if <code>draw.segments = TRUE</code> .
<code>axes</code>	logical flag: if <code>TRUE</code> axes are added to the plot.
<code>frame.plot</code>	logical flag: if <code>TRUE</code> , the plot region is framed.
<code>main</code>	a main title for the plot.
<code>sub</code>	a sub title for the plot.
<code>xlab</code>	a label for the x axis.
<code>ylab</code>	a label for the y axis.
<code>cex</code>	character expansion factor for the labels.
<code>lwd</code>	line width used for drawing.
<code>lty</code>	line type used for drawing.

<code>xpd</code>	logical or NA indicating if clipping should be done, see <code>par(xpd = .)</code> .
<code>mar</code>	argument to <code>par(mar = *)</code> , typically choosing smaller margins than by default.
<code>...</code>	further arguments, passed to the first call of <code>plot()</code> , see <code>plot.default</code> and to <code>box()</code> if <code>frame.plot</code> is true.
<code>add</code>	logical, if TRUE <i>add</i> stars to current plot.
<code>plot</code>	logical, if FALSE, nothing is plotted.

## Details

Missing values are treated as 0.

Each star plot or segment diagram represents one row of the input `x`. Variables (columns) start on the right and wind counterclockwise around the circle. The size of the (scaled) column is shown by the distance from the center to the point on the star or the radius of the segment representing the variable.

Only one page of output is produced.

## Note

This code started life as spatial star plots by David A. Andrews.

Prior to 1.4.1, scaling only shifted the maximum to 1, although documented as here.

## Author(s)

Thomas S. Dye

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## Examples

```
data(mtcars)
stars(mtcars[, 1:7], key.loc = c(14, 2),
      main = "Motor Trend Cars: stars(*, full=F)", full=FALSE)
stars(mtcars[, 1:7], key.loc = c(14, 1.5),
      main = "Motor Trend Cars: full stars()",
      flip.labels=FALSE)
```

```

## 'Spider' or 'Radar' plot:
stars(mtcars[, 1:7], locations = c(0,0), radius = FALSE,
      key.loc=c(0,0), main="Motor Trend Cars", lty = 2)

## Segment Diagrams:
palette(rainbow(12, s = 0.6, v = 0.75))
stars(mtcars[, 1:7], len = 0.8, key.loc = c(12, 1.5),
      main = "Motor Trend Cars", draw.segments = TRUE)
stars(mtcars[, 1:7], len = 0.6, key.loc = c(1.5, 0),
      main = "Motor Trend Cars", draw.segments = TRUE,
      frame.plot=TRUE, nrow = 4, cex = .7)

data(USJudgeRatings)
## scale linearly (not affinely) to [0, 1]
USJudge <- apply(USJudgeRatings, 2, function(x) x/max(x))
Jnam <- case.names(USJudgeRatings)
Snam <-
  abbreviate(substring(Jnam,1,regexpr("[,.] ",Jnam) - 1), 7)
stars(USJudge, labels = Jnam, scale = FALSE,
      key.loc = c(13, 1.5), main = "Judge not ...", len = 0.8)
stars(USJudge, labels = Snam, scale = FALSE,
      key.loc = c(13, 1.5), radius = FALSE)

loc <- stars(USJudge, labels = NULL, scale = FALSE,
            radius = FALSE, frame.plot = TRUE,
            key.loc = c(13, 1.5), main = "Judge not ...", len = 1.2)
text(loc, Snam, col = "blue", cex = 0.8, xpd = TRUE)

## 'Segments':
stars(USJudge, draw.segments = TRUE, scale = FALSE,
      key.loc = c(13,1.5))

## 'Spider':
stars(USJudgeRatings, locations=c(0,0), scale=FALSE,
      radius = FALSE, col.stars=1:10, key.loc = c(0,0),
      main="US Judges rated")
## 'Radar-Segments'
stars(USJudgeRatings[1:10,], locations = 0:1, scale=FALSE,
      draw.segments = TRUE, col.segments=0, col.stars=1:10,
      key.loc= 0:1, main="US Judges 1-10 ")
palette("default")
stars(cbind(1:16,10*(16:1)),draw.segments=TRUE,
      main = "A Joke -- do not use symbols on 2D data!")

```

---

**stripchart**     *1-D Scatter Plots*

---

**Description**

**stripchart** produces one dimensional scatter plots (or dot plots) of the given data. These plots are a good alternative to **boxplots** when sample sizes are small.

**Usage**

```
stripchart(x, method="overplot", jitter=0.1, offset=1/3,  
  vertical=FALSE, group.names, add = FALSE, at = NULL,  
  xlim=NULL, ylim=NULL, main="", ylab="", xlab="",  
  log="", pch=0, col=par("fg"), cex=par("cex"))
```

**Arguments**

<b>x</b>	the data from which the plots are to be produced. The data can be specified as a single vector, or as list of vectors, each corresponding to a component plot. Alternatively a symbolic specification of the form $\mathbf{x} \sim \mathbf{g}$ can be given, indicating the the observations in the vector <b>x</b> are to be grouped according to the levels of the factor <b>g</b> . NAs are allowed in the data.
<b>method</b>	the method to be used to separate coincident points. The default method <b>"overplot"</b> causes such points to be overplotted, but it is also possible to specify <b>"jitter"</b> to jitter the points, or <b>"stack"</b> have coincident points stacked. The last method only makes sense for very granular data.
<b>jitter</b>	when jittering is used, <b>jitter</b> gives the amount of jittering applied.
<b>offset</b>	when stacking is used, points are stacked this many line-heights (symbol widths) apart.
<b>vertical</b>	when vertical is <b>TRUE</b> the plots are drawn vertically rather than the default horizontal.
<b>group.names</b>	group labels which will be printed alongside (or underneath) each plot.
<b>add</b>	logical, if true <i>add</i> boxplot to current plot.

**at** numeric vector giving the locations where the boxplots should be drawn, particularly when `add = TRUE`; defaults to `1:n` where `n` is the number of boxes.

`xlim`, `ylim`, `main`, `ylab`, `xlab`, `log`, `pch`, `col`, `cex`  
Graphical parameters.

## Details

Extensive examples of the use of this kind of plot can be found in Box, Hunter and Hunter or Seber and Wild.

## Examples

```
x <- rnorm(50)
xr<- round(x, 1)
stripchart(x) ; m <- mean(par("usr")[1:2])
text(m, 1.04, "stripchart(x, \"overplot\")")
stripchart(xr, method = "stack", add = TRUE, at = 1.2)
text(m, 1.35, "stripchart(round(x,1), \"stack\")")
stripchart(xr, method = "jitter", add = TRUE, at = 0.7)
text(m, 0.85, "stripchart(round(x,1), \"jitter\")")

data(OrchardSprays)
with(OrchardSprays,
  stripchart(decrease ~ treatment,
    main = "stripchart(Orchardsprays)", ylab = "decrease",
    vertical = TRUE, log = "y"))

with(OrchardSprays,
  stripchart(decrease ~ treatment, at = c(1:8)^2,
    main = "stripchart(Orchardsprays)", ylab = "decrease",
    vertical = TRUE, log = "y"))
```



---

**strwidth**     *Plotting Dimensions of Character Strings and Math Expressions*

---

**Description**

These functions compute the width or height, respectively, of the given strings or mathematical expressions `s[i]` on the current plotting device in *user* coordinates, *inches* or as fraction of the figure width `par("fin")`.

**Usage**

```
strwidth(s, units = "user", cex = NULL)
strheight(s, units = "user", cex = NULL)
```

**Arguments**

- |              |  |
|--------------|--|
| <b>s</b>     | character vector or <b>expressions</b> whose string widths in plotting units are to be determined. An attempt is made to coerce other vectors to character, and other language objects to expressions. |
| <b>units</b> | character indicating in which units <b>s</b> is measured; should be one of "user", "inches", "figure"; partial matching is performed.  |
| <b>cex</b>   | character expansion to which it applies. By default, the current <code>par("cex")</code> is used.  |

**Value**

Numeric vector with the same length as **s**, giving the width or height for each `s[i]`. NA strings are given width and height 0 (as they are not plotted).

**See Also**

`text`, `nchar`

**Examples**

```

str.ex <- c("W","w","I",".","WwI.")
op <- par(pty='s'); plot(1:100,1:100, type="n")
sw <- strwidth(str.ex); sw
# since the last string contains the others
all.equal(sum(sw[1:4]), sw[5])

# width in [mm]
sw.i <- strwidth(str.ex, "inches"); 25.4 * sw.i
unique(sw / sw.i)
# constant factor: 1 value
mean(sw.i / strwidth(str.ex, "fig")) / par('fin')[1]

## See how letters fall in classes -- depending on graphics
## device and font!
all.lett <- c(letters, LETTERS)
# 'big points'
shL <- strheight(all.lett, units = "inches") * 72
table(shL) # all have same heights ...
mean(shL)/par("cin")[2] # around 0.6

# 'big points'
(swL <- strwidth(all.lett, units="inches") * 72)
split(all.lett, factor(round(swL, 2)))

sumex <- expression(sum(x[i], i=1,n), e^{i * pi} == -1)
strwidth(sumex)
strheight(sumex)

par(op) # reset to previous setting

```

---

**sunflowerplot**     *Produce a Sunflower Scatter Plot*


---

**Description**

Multiple points are plotted as “sunflowers” with multiple leaves (“petals”) such that overplotting is visualized instead of accidental and invisible.

**Usage**

```
sunflowerplot(x, y = NULL, number, log = "", digits = 6,
  xlab = NULL, ylab = NULL, xlim = NULL, ylim = NULL,
  add = FALSE, rotate = FALSE,
  pch = 16, cex = 0.8, cex.fact = 1.5,
  size = 1/8, seg.col = 2, seg.lwd = 1.5, ...)
```

**Arguments**

<b>x</b>	numeric vector of x-coordinates of length <b>n</b> , say, or another valid plotting structure, as for <code>plot.default</code> , see also <code>xy.coords</code> .
<b>y</b>	numeric vector of y-coordinates of length <b>n</b> .
<b>number</b>	integer vector of length <b>n</b> . <code>number[i]</code> = number of replicates for <code>(x[i],y[i])</code> , may be 0. Default: compute the exact multiplicity of the points <code>x[],y[]</code> .
<b>log</b>	character indicating log coordinate scale, see <code>plot.default</code> .
<b>digits</b>	when <code>number</code> is computed (i.e., not specified), <code>x</code> and <code>y</code> are rounded to <code>digits</code> significant digits before multiplicities are computed.
<b>xlab,ylab</b>	character label for x-, or y-axis, respectively.
<b>xlim,ylim</b>	<code>numeric(2)</code> limiting the extents of the x-, or y-axis.
<b>add</b>	logical; should the plot be added on a previous one ? Default is <b>FALSE</b> .
<b>rotate</b>	logical; if <b>TRUE</b> , randomly rotate the sunflowers (preventing artefacts).
<b>pch</b>	plotting character to be used for points ( <code>number[i]==1</code> ) and center of sunflowers.

<code>cex</code>	numeric; character size expansion of center points (s. <code>pch</code> ).
<code>cex.fact</code>	numeric <i>shrinking</i> factor to be used for the center points <i>when there are flower leaves</i> , i.e., <code>cex / cex.fact</code> is used for these.
<code>size</code>	of sunflower leaves in inches, <code>1[in] := 2.54[cm]</code> . Default: <code>1/8</code> ; approximately 3.2mm.
<code>seg.col</code>	color to be used for the <b>segments</b> which make the sunflowers leaves, see <code>par(col=)</code> ; <code>col = "gold"</code> reminds of real sunflowers.
<code>seg.lwd</code>	numeric; the line width for the leaves' segments.
<code>...</code>	further arguments to <code>plot</code> [if <code>add=FALSE</code> ].

## Details

For `number[i]==1`, a (slightly enlarged) usual plotting symbol (`pch`) is drawn. For `number[i] > 1`, a small plotting symbol is drawn and `number[i]` equi-angular “rays” emanate from it.

If `rotate=TRUE` and `number[i] >= 2`, a random direction is chosen (instead of the y-axis) for the first ray. The goal is to **jitter** the orientations of the sunflowers in order to prevent artefactual visual impressions.

## Value

A list with three components of same length,

<code>x</code>	x coordinates
<code>y</code>	y coordinates
<code>number</code>	number

## Side Effects

A scatter plot is drawn with “sunflowers” as symbols.

## Author(s)

Andreas Ruckstuhl, Werner Stahel, Martin Maechler, Tim Hesterberg, 1989–1993. Port to R by Martin Maechler.

## References

- Chambers, J. M., Cleveland, W. S., Kleiner, B. and Tukey, P. A. (1983) *Graphical Methods for Data Analysis*. Wadsworth.
- Schilling, M. F. and Watkins, A. E. (1994) A suggestion for sunflower plots. *The American Statistician*, **48**, 303–305.

## See Also

density

## Examples

```
data(iris)
## 'number' is computed automatically:
sunflowerplot(iris[, 3:4])
## Imitating Chambers et al., p.109, closely:
sunflowerplot(iris[, 3:4], cex=.2, cex.f=1, size=.035,
  seg.lwd=.8)

sunflowerplot(x=sort(2*round(rnorm(100))),
  y=round(rnorm(100),0),
  main = "Sunflower Plot of Rounded N(0,1)")

## A 'point process' {explicit 'number' argument}:
sunflowerplot(rnorm(100), rnorm(100),
  number=rpois(n=100, lambda=2),
  rotate=TRUE, main="Sunflower plot")
```

---

**symbols**     *Draw symbols on a plot*

---

## Description

This function draws symbols on a plot. One of six symbols; *circles*, *squares*, *rectangles*, *stars*, *thermometers*, and *boxplots*, can be plotted at a specified set of x and y coordinates. Specific aspects of the symbols, such as relative size, can be customized by additional parameters.

## Usage

```
symbols(x, y = NULL, circles, squares, rectangles, stars,
        thermometers, boxplots, inches = TRUE, add = FALSE,
        fg = 1, bg = NA, xlab = NULL, ylab = NULL, main = NULL,
        xlim = NULL, ylim = NULL, ...)
```

## Arguments

<b>x, y</b>	the x and y co-ordinates for the symbols. They can be specified in any way which is accepted by <code>xy.coords</code> .
<b>circles</b>	a vector giving the radii of the circles.
<b>squares</b>	a vector giving the length of the sides of the squares.
<b>rectangles</b>	a matrix with two columns. The first column gives widths and the second the heights of rectangle symbols.
<b>stars</b>	a matrix with three or more columns giving the lengths of the rays from the center of the stars. <code>NA</code> values are replaced by zeroes.
<b>thermometers</b>	a matrix with three or four columns. The first two columns give the width and height of the thermometer symbols. If there are three columns, the third is taken as a proportion. The thermometers are filled from their base to this proportion of their height. If there are four columns, the third and fourth columns are taken as proportions. The thermometers are filled between these two proportions of their heights.
<b>boxplots</b>	a matrix with five columns. The first two columns give the width and height of the boxes, the next two columns give the lengths of the lower and upper

	whiskers and the fifth the proportion (with a warning if not in $[0,1]$ ) of the way up the box that the median line is drawn.
<code>inches</code>	If <code>inches</code> is <code>FALSE</code> , the units are taken to be those of the x axis. If <code>inches</code> is <code>TRUE</code> , the symbols are scaled so that the largest symbol is one inch in height. If a number is given the symbols are scaled to make the largest symbol this height in inches.
<code>add</code>	if <code>add</code> is <code>TRUE</code> , the symbols are added to an existing plot, otherwise a new plot is created.
<code>fg</code>	colors the symbols are to be drawn in (the default is the value of the <code>col</code> graphics parameter).
<code>bg</code>	if specified, the symbols are filled with this color. The default is to leave the symbols unfilled.
<code>xlab</code>	the x label of the plot if <code>add</code> is not true; this applies to the following arguments as well. Defaults to the <code>deparsed</code> expression used for <code>x</code> .
<code>ylab</code>	the y label of the plot.
<code>main</code>	a main title for the plot.
<code>xlim</code>	numeric of length 2 giving the x limits for the plot.
<code>ylim</code>	numeric of length 2 giving the y limits for the plot.
<code>...</code>	graphics parameters can also be passed to this function.

## Details

Observations which have missing coordinates or missing size parameters are not plotted. The exception to this is *stars*. In that case, the length of any rays which are `NA` is reset to zero.

Circles of radius zero are plotted at radius one pixel (which is device-dependent).

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- W. S. Cleveland (1985) *The Elements of Graphing Data*. Monterey, California: Wadsworth.

## See Also

`stars` for drawing *stars* with a bit more flexibility; `sunflowerplot`.

## Examples

```
x <- 1:10
y <- sort(10*runif(10))
z <- runif(10)
z3 <- cbind(z, 2*runif(10), runif(10))
symbols(x, y, thermometers=cbind(.5, 1, z), inches=.5,
        fg = 1:10)
symbols(x, y, thermometers = z3, inches=FALSE)
text(x,y,
     apply(format(round(z3, dig=2)), 1, paste, collapse=","),
     adj = c(-.2,0), cex = .75, col = "purple", xpd=NA)

data(trees)
## Note that example(trees) shows more sensible plots!
N <- nrow(trees)
attach(trees)
## Girth is diameter in inches
symbols(Height, Volume, circles=Girth/24, inches=FALSE,
        main="Trees' Girth") # xlab and ylab automatically
## Colors too:
palette(rainbow(N, end = 0.9))
symbols(Height, Volume, circles=Girth/16, inches=FALSE,
        bg = 1:N, fg="gray30",
        main = "symbols(*, circles=Girth/16, bg = 1:N)")
palette("default"); detach()
```



---

**termplot**     *Plot regression terms*


---

**Description**

Plots regression terms against their predictors, optionally with standard errors and partial residuals added.

**Usage**

```
termplot(model, data=NULL,
  envir=environment(formula(model)),
  partial.resid=FALSE, rug=FALSE,
  terms=NULL, se=FALSE, xlabs=NULL, ylabs=NULL, main=NULL,
  col.term = 2, lwd.term = 1.5,
  col.se = "orange", lty.se = 2, lwd.se = 1,
  col.res = "gray", cex = 1, pch = par("pch"),
  ask = interactive() && nb.fig < n.tms
                                && .Device != "postscript",
  use.factor.levels=TRUE, ...)
```

**Arguments**

<b>model</b>	fitted model object
<b>data</b>	data frame in which variables in <b>model</b> can be found
<b>envir</b>	environment in which variables in <b>model</b> can be found
<b>partial.resid</b>	logical; should partial residuals be plotted?
<b>rug</b>	add rugplots (jittered 1-d histograms) to the axes?
<b>terms</b>	which terms to plot (default <b>NULL</b> means all terms)
<b>se</b>	plot pointwise standard errors?
<b>xlabs</b>	vector of labels for the x axes
<b>ylabs</b>	vector of labels for the y axes
<b>main</b>	logical, or vector of main titles; if <b>TRUE</b> , the model's call is taken as main title, <b>NULL</b> or <b>FALSE</b> mean no titles.
<b>col.term, lwd.term</b>	color and line width for the “term curve”, see <b>lines</b> .
<b>col.se, lty.se, lwd.se</b>	color, line type and line width for the “twice-standard-error curve” when <b>se = TRUE</b> .

<code>col.res</code> , <code>cex</code> , <code>pch</code>	color, plotting character expansion and type for partial residuals, when <code>partial.resid = TRUE</code> , see <code>points</code> .
<code>ask</code>	logical; if <code>TRUE</code> , the user is <i>asked</i> before each plot, see <code>par(ask=.)</code> .
<code>use.factor.levels</code>	Should x-axis ticks use factor levels or numbers for factor terms?
<code>...</code>	other graphical parameters

## Details

The model object must have a `predict` method that accepts `type=terms`, eg `glm` in the `base` package, `coxph` and `survreg` in the `survival` package.

For the `partial.resid=TRUE` option it must have a `residuals` method that accepts `type="partial"`, which `lm` and `glm` do.

The `data` argument should rarely be needed, but in some cases `termplot` may be unable to reconstruct the original data frame.

Nothing sensible happens for interaction terms.

## See Also

For (generalized) linear models, `plot.lm` and `predict.glm`.

## Examples

```
had.splines <- "package:splines" %in% search()
if(!had.splines) rs <- require(splines)
x <- 1:100
z <- factor(rep(LETTERS[1:4],25))
y <- rnorm(100,sin(x/10)+as.numeric(z))
model <- glm(y ~ ns(x,6) + z)

par(mfrow=c(2,2)) ## 2 x 2 plots for same model :
termplot(model,
  main = paste("termplot( ", deparse(model$call)," ...)")
termplot(model, rug=TRUE)
termplot(model, partial=TRUE, rug= TRUE,
  main="termplot(..., partial = TRUE, rug = TRUE)")
termplot(model, partial=TRUE, se = TRUE, main = TRUE)
if(!had.splines && rs) detach("package:splines")
```

---

<b>text</b>	<i>Add Text to a Plot</i>
-------------	---------------------------

---

## Description

**text** draws the strings given in the vector **labels** at the coordinates given by **x** and **y**. **y** may be missing since **xy.coords(x,y)** is used for construction of the coordinates.

## Usage

```
text(x, ...)
```

```
## Default S3 method:
```

```
text(x, y = NULL, labels = seq(along = x), adj = NULL,  
      pos = NULL, offset = 0.5, vfont = NULL,  
      cex = 1, col = NULL, font = NULL, xpd = NULL, ...)
```

## Arguments

<b>x, y</b>	numeric vectors of coordinates where the text <b>labels</b> should be written. If the length of <b>x</b> and <b>y</b> differs, the shorter one is recycled.
<b>labels</b>	one or more character strings or expressions specifying the <i>text</i> to be written. An attempt is made to coerce other vectors to character, and other language objects to expressions.
<b>adj</b>	one or two values in $[0, 1]$ which specify the x (and optionally y) adjustment of the labels. On most devices values outside that interval will also work.
<b>pos</b>	a position specifier for the text. If specified, this overrides any <b>adj</b> value given. Values of 1, 2, 3 and 4, respectively indicate positions below, to the left of, above and to the right of the specified coordinates.
<b>offset</b>	when <b>pos</b> is specified, this value gives the offset of the label from the specified coordinate in fractions of a character width.
<b>vfont</b>	if a character vector of length 2 is specified, then Hershey vector fonts are used. The first element of the vector selects a typeface and the second element selects a style.

<code>cex</code>	numeric character <b>exp</b> ansion factor; multiplied by <code>par("cex")</code> yields the final character size.
<code>col, font</code>	the color and font to be used; these default to the values of the global graphical parameters in <code>par()</code> .
<code>xpd</code>	(where) should clipping take place? Defaults to <code>par("xpd")</code> .
<code>...</code>	further graphical parameters (from <code>par</code> ).

## Details

`labels` must be of type **character** or **expression** (or be coercible to such a type). In the latter case, quite a bit of mathematical notation is available such as sub- and superscripts, greek letters, fractions, etc.

`adj` allows *adj*ustment of the text with respect to (`x,y`). Values of 0, 0.5, and 1 specify left/bottom, middle and right/top, respectively. The default is for centered text, i.e., `adj = c(0.5, 0.5)`. Accurate vertical centering needs character metric information on individual characters, which is only available on some devices.

The `pos` and `offset` arguments can be used in conjunction with values returned by `identify` to recreate an interactively labelled plot.

Text can be rotated by using graphical parameters `srt` (see `par`); this rotates about the centre set by `adj`.

Graphical parameters `col`, `cex` and `font` can be vectors and will then be applied cyclically to the `labels` (and extra values will be ignored).

Labels whose `x`, `y`, `labels`, `cex` or `col` value is `NA` are omitted from the plot.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`mtext`, `title`, `Hershey` for details on Hershey vector fonts, `plotmath` for details and more examples on mathematical annotation.

## Examples

```
plot(-1:1,-1:1, type = "n", xlab = "Re", ylab = "Im")
K <- 16; text(exp(1i * 2 * pi * (1:K) / K), col = 2)
```

```

plot(1:10, 1:10,
     main = "text(...) examples\n~~~~~")
points(c(6,2), c(2,1), pch = 3, cex = 4, col = "red")
text(6, 2,
     "the text is CENTERED around (x,y) = (6,2) by default",
     cex = .8)
text(2, 1,
     "or Left/Bottom - JUSTIFIED at (2,1) by 'adj = c(0,0)'",
     adj = c(0,0))
text(4, 9,
     expression(hat(beta) == (X^t * X)^{-1} * X^t * y))
text(4, 8.4,
     "expression(hat(beta) == (X^t * X)^{-1} * X^t * y)",
     cex = .75)
text(4, 7,
     expression(bar(x) == sum(frac(x[i], n), i==1, n))
)

```

---

<b>title</b>	<i>Plot Annotation</i>
--------------	------------------------

---

## Description

This function can be used to add labels to a plot. Its first four principal arguments can also be used as arguments in most high-level plotting functions. They must be of type **character** or **expression**. In the latter case, quite a bit of mathematical notation is available such as sub- and superscripts, greek letters, fractions, etc.

## Usage

```
title(main = NULL, sub = NULL, xlab = NULL, ylab = NULL,  
      line = NA, outer = FALSE, ...)
```

## Arguments

<b>main</b>	The main title (on top) using font and size (character expansion) <code>par("font.main")</code> and color <code>par("col.main")</code> .
<b>sub</b>	Sub-title (at bottom) using font and size <code>par("font.sub")</code> and color <code>par("col.sub")</code> .
<b>xlab</b>	X axis label using font and character expansion <code>par("font.axis")</code> and color <code>par("col.axis")</code> .
<b>ylab</b>	Y axis label, same font attributes as <b>xlab</b> .
<b>line</b>	specifying a value for <b>line</b> overrides the default placement of labels, and places them this many lines from the plot.
<b>outer</b>	a logical value. If <b>TRUE</b> , the titles are placed in the outer margins of the plot.
<b>...</b>	further graphical parameters from <b>par</b> . Use e.g., <code>col.main</code> or <code>cex.sub</code> instead of just <code>col</code> or <code>cex</code> .

## Details

The labels passed to **title** can be simple strings or expressions, or they can be a list containing the string to be plotted, and a selection of the optional modifying graphical parameters **cex=**, **col=**, **font=**.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`mtext`, `text`; `plotmath` for details on mathematical annotation.

## Examples

```
data(cars)
plot(cars, main = "") # here, could use main directly
title(main = "Stopping Distance versus Speed")

plot(cars, main = "")
title(main = list("Stopping Distance versus Speed",
                  cex=1.5, col="red", font=3))

## Specifying "...":
plot(1, col.axis = "sky blue", col.lab = "thistle")
title("Main Title", sub = "sub title",
      cex.main = 2, font.main = 4, col.main = "blue",
      cex.sub = 0.75, font.sub = 3, col.sub = "red")

x <- seq(-4, 4, len = 101)
y <- cbind(sin(x), cos(x))
matplot(x, y, type = "l", xaxt = "n",
        main = expression(paste(plain(sin) * phi, " and ",
                                plain(cos) * phi)),
        # only 1st is taken
        ylab = expression("sin" * phi, "cos" * phi),
        xlab = expression(paste("Phase Angle ", phi)),
        col.main = "blue")
axis(1, at = c(-pi, -pi/2, 0, pi/2, pi),
     lab = expression(-pi, -pi/2, 0, pi/2, pi))
abline(h = 0, v = pi/2 * c(-1,1), lty = 2, lwd = .1,
       col = "gray70")
```

---

**units**     *Graphical Units*

---

**Description**

`xinch` and `yinch` convert the specified number of inches given as their arguments into the correct units for plotting with graphics functions. Usually, this only makes sense when normal coordinates are used, i.e., *no log scale* (see the `log` argument to `par`).

`xyinch` does the same for a pair of numbers `xy`, simultaneously.

`cm` translates inches in to cm (centimeters).

**Usage**

```
xinch(x = 1, warn.log = TRUE)
yinch(y = 1, warn.log = TRUE)
xyinch(xy = 1, warn.log = TRUE)
cm(x)
```

**Arguments**

<code>x,y</code>	numeric vector
<code>xy</code>	numeric of length 1 or 2.
<code>warn.log</code>	logical; if TRUE, a warning is printed in case of active log scale.

**Examples**

```
all(c(xinch(),yinch()) == xyinch()) # TRUE
xyinch()
xyinch # to see that is really delta{"usr"} / "pin"

cm(1) # = 2.54

## plot labels offset 0.12 inches to the right of plotted
## symbols in a plot
data(mtcars)
with(mtcars, {
  plot(mpg, disp, pch=19, main= "Motor Trend Cars")
  text(mpg + xinch(0.12), disp, row.names(mtcars),
       adj = 0, cex = .7, col = 'blue')
})
```



---

**x11**     *X Window System Graphics*

---

**Description**

**X11** starts a graphics device driver for the X Window System (version 11). This can only be done on machines that run X. **x11** is recognized as a synonym for **X11**.

**Usage**

```
X11(display = "", width = 7, height = 7, pointsize = 12,  
     gamma = 1, colortype = getOption("X11colortype"),  
     maxcubysize = 256, canvas = "white")
```

**Arguments**

<b>display</b>	the display on which the graphics window will appear. The default is to use the value in the user's environment variable <b>DISPLAY</b> .
<b>width</b>	the width of the plotting window in inches.
<b>height</b>	the height of the plotting window in inches.
<b>pointsize</b>	the default pointsize to be used.
<b>gamma</b>	the gamma correction factor. This value is used to ensure that the colors displayed are linearly related to RGB values. A value of around 0.5 is appropriate for many PC displays. A value of 1.0 (no correction) is usually appropriate for high-end displays or Macintoshes.
<b>colortype</b>	the kind of color model to be used. The possibilities are <b>"mono"</b> , <b>"gray"</b> , <b>"pseudo"</b> , <b>"pseudo.cube"</b> and <b>"true"</b> . Ignored if an <b>X11</b> is already open.
<b>maxcubysize</b>	can be used to limit the size of color cube allocated for pseudocolor devices.
<b>canvas</b>	color. The color of the canvas, which is visible only when the background color is transparent.

## Details

By default, an X11 device will use the best color rendering strategy that it can. The choice can be overridden with the `colortype` parameter. A value of `"mono"` results in black and white graphics, `"gray"` in grayscale and `"true"` in truecolor graphics (if this is possible). The values `"pseudo"` and `"pseudo.cube"` provide color strategies for pseudocolor displays. The first strategy provides on-demand color allocation which produces exact colors until the color resources of the display are exhausted. The second causes a standard color cube to be set up, and requested colors are approximated by the closest value in the cube. The default strategy for pseudocolor displays is `"pseudo"`.

**Note:** All X11 devices share a `colortype` which is set by the first device to be opened. To change the `colortype` you need to close *all* open X11 devices then open one with the desired `colortype`.

With `colortype` equal to `"pseudo.cube"` or `"gray"` successively smaller palettes are tried until one is completely allocated. If allocation of the smallest attempt fails the device will revert to `"mono"`.

## See Also

Devices.

---

**xfig**     *XFig Graphics Device*


---

**Description**

**xfig** starts the graphics device driver for producing XFig (version 3.2) graphics.

The auxiliary function **ps.options** can be used to set and view (if called without arguments) default values for the arguments to **xfig** and **postscript**.

**Usage**

```
xfig(file = ifelse(onefile, "Rplots.fig", "Rplot%03d.fig"),
      onefile = FALSE, ...)
```

**Arguments**

- |                |  |
|----------------|--|
| <b>file</b>    | a character string giving the name of the file. If it is "", the output is piped to the command given by the argument <b>command</b> . For use with <b>onefile=FALSE</b> give a <b>printf</b> format such as "Rplot%d.fig" (the default in that case).   |
| <b>onefile</b> | logical: if true allow multiple figures in one file. If false, assume only one page per file and generate a file number containing the page number.  |
| <b>...</b>     | further arguments to <b>ps.options</b> accepted by <b>xfig()</b> :   |
|                | <b>paper</b> the size of paper in the printer. The choices are "A4", "Letter" and "Legal" (and these can be lowercase). A further choice is "default", which is the default. If this is selected, the papersize is taken from the option "papersize" if that is set and to "A4" if it is unset or empty. |
|                | <b>horizontal</b> the orientation of the printed image, a logical. Defaults to true, that is landscape orientation.  |
|                | <b>width, height</b> the width and height of the graphics region in inches. The default is to use the entire page less a 0.25 inch border.   |

`family` the font family to be used. This must be one of "AvantGarde", "Bookman", "Courier", "Helvetica", "Helvetica-Narrow", "NewCenturySchoolbook", "Palatino" or "Times".

`pointsize` the default point size to be used.

`bg` the default background color to be used.

`fg` the default foreground color to be used.

`pagecentre` logical: should the device region be centred on the page: defaults to TRUE.

## Details

Although xfig can produce multiple plots in one file, the XFig format does not say how to separate or view them. So `onefile=FALSE` is the default.

## Note

On some line textures ( $0 \leq \text{ity} < 4$ ) are used. Eventually this will be partially remedied, but the XFig file format does not allow as general line textures as the R model. Unimplemented line textures are displayed as *dash-double-dotted*.

There is a limit of 512 colours (plus white and black) per file.

## See Also

`Devices`, `postscript`, `ps.options`.

---

**xy.coords**     *Extracting Plotting Structures*

---

**Description**

`xy.coords` is used by many functions to obtain `x` and `y` coordinates for plotting. The use of this common mechanism across all R functions produces a measure of consistency.

**Usage**

```
xy.coords(x, y, xlab = NULL, ylab = NULL, log = NULL,  
          recycle = FALSE)
```

**Arguments**

<code>x, y</code>	the <code>x</code> and <code>y</code> coordinates of a set of points. Alternatively, a single argument <code>x</code> can be provided.
<code>xlab, ylab</code>	names for the <code>x</code> and <code>y</code> variables to be extracted.
<code>log</code>	character, " <code>x</code> ", " <code>y</code> " or both, as for <code>plot</code> . Sets negative values to <code>NA</code> and gives a warning.
<code>recycle</code>	logical; if <code>TRUE</code> , recycle ( <code>rep</code> ) the shorter of <code>x</code> or <code>y</code> if their lengths differ.

**Details**

An attempt is made to interpret the arguments `x` and `y` in a way suitable for plotting.

If `y` is missing and `x` is a

**formula:** of the form `yvar ~ xvar`. `xvar` and `yvar` are used as `x` and `y` variables.

**list:** containing components `x` and `y`, these are used to define plotting coordinates.

**time series:** the `x` values are taken to be `time(x)` and the `y` values to be the time series.

**matrix with two columns:** the first is assumed to contain the `x` values and the second the `y` values.

In any other case, the `x` argument is coerced to a vector and returned as `y` component where the resulting `x` is just the index vector `1:n`. In this case, the resulting `xlab` component is set to `"Index"`.

If `x` (after transformation as above) inherits from class `"POSIXt"` it is coerced to class `"POSIXct"`.

## Value

A list with the components

<code>x</code>	numeric (i.e., <code>"double"</code> ) vector of abscissa values.
<code>y</code>	numeric vector of the same length as <code>x</code> .
<code>xlab</code>	<code>character(1)</code> or <code>NULL</code> , the 'label' of <code>x</code> .
<code>ylab</code>	<code>character(1)</code> or <code>NULL</code> , the 'label' of <code>y</code> .

## See Also

`plot.default`, `lines`, `points` and `lowess` are examples of functions which use this mechanism.

## Examples

```
xy.coords(fft(c(1:10)), NULL)
data(cars) ; attach(cars)
xy.coords(dist ~ speed, NULL)$xlab # = "speed"

str(xy.coords(1:3, 1:2, recycle=TRUE))
str(xy.coords(-2:10, NULL, log="y"))
## warning: 3 y values <=0 omitted ..
detach()
```

---

**xyz.coords**      *Extracting Plotting Structures*


---

**Description**

Utility for obtaining consistent x, y and z coordinates and labels for three dimensional (3D) plots.

**Usage**

```
xyz.coords(x, y, z, xlab=NULL, ylab=NULL, zlab=NULL,
           log=NULL, recycle=FALSE)
```

**Arguments**

- |                         |   |
|-------------------------|---|
| <b>x, y, z</b>          | the x, y and z coordinates of a set of points. Alternatively, a single argument <b>x</b> can be provided. In this case, an attempt is made to interpret the argument in a way suitable for plotting.<br>If the argument is a formula <b>zvar ~ xvar + yvar</b> , <b>xvar</b> , <b>yvar</b> and <b>zvar</b> are used as x, y and z variables; if the argument is a list containing components <b>x</b> , <b>y</b> and <b>z</b> , these are assumed to define plotting coordinates; if the argument is a matrix with three columns, the first is assumed to contain the x values, etc.<br>Alternatively, two arguments <b>x</b> and <b>y</b> can be provided. One may be real, the other complex; in any other case, the arguments are coerced to vectors and the values plotted against their indices. |
| <b>xlab, ylab, zlab</b> | names for the x, y and z variables to be extracted.   |
| <b>log</b>              | character, "x", "y", "z" or combinations. Sets negative values to NA and gives a warning.   |
| <b>recycle</b>          | logical; if TRUE, recycle ( <b>rep</b> ) the shorter ones of <b>x</b> , <b>y</b> or <b>z</b> if their lengths differ.   |

**Value**

A list with the components

- |          |   |
|----------|---|
| <b>x</b> | numeric (i.e., <b>double</b> ) vector of abscissa values. |
| <b>y</b> | numeric vector of the same length as <b>x</b> .           |

<b>z</b>	numeric vector of the same length as <b>x</b> .
<b>xlab</b>	<code>character(1)</code> or <code>NULL</code> , the axis label of <b>x</b> .
<b>ylab</b>	<code>character(1)</code> or <code>NULL</code> , the axis label of <b>y</b> .
<b>zlab</b>	<code>character(1)</code> or <code>NULL</code> , the axis label of <b>z</b> .

### Author(s)

Uwe Ligges and Martin Maechler

### See Also

`xy.coords` for 2D.

### Examples

```
str(xyz.coords(data.frame(10*1:9, -4),y=NULL,z=NULL))

str(xyz.coords(1:6, fft(1:6),z=NULL,xlab="X", ylab="Y"))

y <- 2 * (x2 <- 10 + (x1 <- 1:10))
str(xyz.coords(y ~ x1 + x2,y=NULL,z=NULL))

str(xyz.coords(data.frame(x=-1:9,y=2:12,z=3:13),
  y=NULL, z=NULL, log="xy"))
## Warning message: 2 x values <= 0 omitted ...
```



## Chapter 2

# Base package — math

---

**abs**     *Miscellaneous Mathematical Functions*

---

**Description**

These functions compute miscellaneous mathematical functions. The naming follows the standard for computer languages such as C or Fortran.

**Usage**

```
abs(x)
sqrt(x)
```

**Arguments**

**x**                      a numeric vector

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

**Arithmetic** for simple, **log** for logarithmic, **sin** for trigonometric, and **Special** for special mathematical functions.

**Examples**

```
xx <- -9:9
plot(xx, sqrt(abs(xx)), col = "red")
lines(spline(xx, sqrt(abs(xx)), n=101), col = "pink")
```

---

**all.equal**      *Test if Two Objects are (Nearly) Equal*


---

## Description

`all.equal(x,y)` is a utility to compare R objects `x` and `y` testing “near equality”. If they are different, comparison is still made to some extent, and a report of the differences is returned. Don’t use `all.equal` directly in `if` expressions—either use `identical` or combine the two, as shown in the documentation for `identical`.

## Usage

```
all.equal(target, current, ...)

## S3 method for class 'numeric':
all.equal(target, current,
  tolerance= .Machine$double.eps ^ 0.5, scale=NULL, ...)
```

## Arguments

<b>target</b>	R object.
<b>current</b>	other R object, to be compared with <b>target</b> .
<b>...</b>	Further arguments for different methods, notably the following two, for numerical comparison:
<b>tolerance</b>	numeric $\geq 0$ . Differences smaller than <b>tolerance</b> are not considered.
<b>scale</b>	numeric scalar $> 0$ (or <code>NULL</code> ). See Details.

## Details

There are several methods available, most of which are dispatched by the default method, see `methods("all.equal")`. `all.equal.list` and `all.equal.language` provide comparison of recursive objects.

Numerical comparisons for `scale = NULL` (the default) are done by first computing the mean absolute difference of the two numerical vectors. If this is smaller than **tolerance** or not finite, absolute differences are used, otherwise relative differences scaled by the mean absolute difference.

If `scale` is positive, absolute comparisons are after scaling (dividing) by `scale`.

For complex arguments, the modulus `Mod` of the difference is used.

`attr.all.equal` is used for comparing `attributes`, returning `NULL` or `character`.

## Value

Either `TRUE` or a vector of mode `"character"` describing the differences between `target` and `current`.

Numerical differences are reported by relative error

## References

Chambers, J. M. (1998) *Programming with Data. A Guide to the S Language*. Springer (for `=`).

## See Also

`==`, and `all` for exact equality testing.

## Examples

```
# not precise enough (default tol) > relative error
all.equal(pi, 355/113)

d45 <- pi*(1/4 + 1:10)
stopifnot(
  all.equal(tan(d45), rep(1,10))) # TRUE, but
all      (tan(d45) == rep(1,10)) # FALSE, since not exactly
all.equal(tan(d45), rep(1,10), tol=0) # to see difference

all.equal(options(), .Options)
all.equal(options(), as.list(.Options)) # TRUE
.Options $ myopt <- TRUE
all.equal(options(), as.list(.Options))
rm(.Options)
```

---

**approxfun**     *Interpolation Functions*

---

**Description**

Return a list of points which linearly interpolate given data points, or a function performing the linear (or constant) interpolation.

**Usage**

```
approx    (x, y = NULL, xout, method="linear", n=50,  
           yleft, yright, rule = 1, f=0, ties = mean)  
  
approxfun(x, y = NULL,           method="linear",  
           yleft, yright, rule = 1, f=0, ties = mean)
```

**Arguments**

<b>x, y</b>	vectors giving the coordinates of the points to be interpolated. Alternatively a single plotting structure can be specified: see <code>xy.coords</code> .
<b>xout</b>	an optional set of values specifying where interpolation is to take place.
<b>method</b>	specifies the interpolation method to be used. Choices are "linear" or "constant".
<b>n</b>	If <b>xout</b> is not specified, interpolation takes place at <b>n</b> equally spaced points spanning the interval <code>[min(x), max(x)]</code> .
<b>yleft</b>	the value to be returned when input <b>x</b> values are less than <code>min(x)</code> . The default is defined by the value of <b>rule</b> given below.
<b>yright</b>	the value to be returned when input <b>x</b> values are greater than <code>max(x)</code> . The default is defined by the value of <b>rule</b> given below.
<b>rule</b>	an integer describing how interpolation is to take place outside the interval <code>[min(x), max(x)]</code> . If <b>rule</b> is 1 then NAs are returned for such points and if it is 2, the value at the closest data extreme is used.

<b>f</b>	For <b>method="constant"</b> a number between 0 and 1 inclusive, indicating a compromise between left- and right-continuous step functions. If <b>y0</b> and <b>y1</b> are the values to the left and right of the point then the value is $y0*(1-f)+y1*f$ so that <b>f=0</b> is right-continuous and <b>f=1</b> is left-continuous.
<b>ties</b>	Handling of tied <b>x</b> values. Either a function with a single vector argument returning a single number result or the string <b>"ordered"</b> .

## Details

The inputs can contain missing values which are deleted, so at least two complete (**x**, **y**) pairs are required. If there are duplicated (tied) **x** values and **ties** is a function it is applied to the **y** values for each distinct **x** value. Useful functions in this context include **mean**, **min**, and **max**. If **ties="ordered"** the **x** values are assumed to be already ordered. The first **y** value will be used for interpolation to the left and the last one for interpolation to the right.

## Value

**approx** returns a list with components **x** and **y**, containing **n** coordinates which interpolate the given data points according to the **method** (and **rule**) desired.

The function **approxfun** returns a function performing (linear or constant) interpolation of the given data points. For a given set of **x** values, this function will return the corresponding interpolated values. This is often more useful than **approx**.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**spline** and **splinefun** for spline interpolation.

## Examples

```
x <- 1:10
y <- rnorm(10)
par(mfrow = c(2,1))
```

```
plot(x, y, main = "approx(.) and approxfun(.)")
points(approx(x, y), col = 2, pch = "*")
points(approx(x, y, method = "constant"), col=4, pch="*")

f <- approxfun(x, y)
curve(f(x), 0, 10, col = "green")
points(x, y)
is.function(fc <- approxfun(x, y, method = "const")) # TRUE
curve(fc(x), 0, 10, col = "darkblue", add = TRUE)

## Show treatment of 'ties' :
x <- c(2,2:4,4,4,5,5,7,7,7)
y <- c(1:6, 5:4, 3:1)
approx(x,y, xout=x)$y # warning
(ay <- approx(x,y, xout=x, ties = "ordered")$y)
stopifnot(ay == c(2,2,3,6,6,6,4,4,1,1,1))
approx(x,y, xout=x, ties = min)$y
approx(x,y, xout=x, ties = max)$y
```

---

**Arithmetic**     *Arithmetic Operators*

---

**Description**

These binary operators perform arithmetic on vector objects.

**Usage**

```
x + y
x - y
x * y
x / y
x ^ y
x %% y
x %/% y
```

**Details**

$1 \wedge y$  and  $y \wedge 0$  are 1, *always*.  $x \wedge y$  should also give the proper “limit” result when either argument is infinite (i.e.,  $\pm\text{Inf}$ ).

Objects such as arrays or time-series can be operated on this way provided they are conformable.

**Value**

They return numeric vectors containing the result of the element by element operations. The elements of shorter vectors are recycled as necessary (with a **warning** when they are recycled only *fractionally*). The operators are + for addition, - for subtraction, \* for multiplication, / for division and ^ for exponentiation.

%% indicates  $x \bmod y$  and %/% indicates integer division. It is guaranteed that  $x == (x \% y) + y * (x \% y)$  unless  $y == 0$  where the result is NA or NaN (depending on the **typeof** of the arguments).

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.



**See Also**

`sqrt` for miscellaneous and `Special` for special mathematical functions.

`Syntax` for operator precedence.

**Examples**

```
x <- -1:12
x + 1
2 * x + 3
x %% 2 # is periodic
x %/% 5
```

---

**backsolve**     *Solve an Upper or Lower Triangular System*

---

## Description

Solves a system of linear equations where the coefficient matrix is upper or lower triangular.

## Usage

```
backsolve(r, x, k=ncol(r), upper.tri=TRUE, transpose=FALSE)
forwardsolve(l, x, k=ncol(l), upper.tri=FALSE,
             transpose=FALSE)
```

## Arguments

<code>r,l</code>	an upper (or lower) triangular matrix giving the coefficients for the system to be solved. Values below (above) the diagonal are ignored.
<code>x</code>	a matrix whose columns give “right-hand sides” for the equations.
<code>k</code>	The number of columns of <code>r</code> and rows of <code>x</code> to use.
<code>upper.tri</code>	logical; if TRUE (default), the <i>upper triangular</i> part of <code>r</code> is used. Otherwise, the lower one.
<code>transpose</code>	logical; if TRUE, solve $r' * y = x$ for $y$ , i.e., <code>t(r) %*% y == x</code> .

## Value

The solution of the triangular system. The result will be a vector if `x` is a vector and a matrix if `x` is a matrix.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Dongarra, J. J., Bunch, J. R., Moler, C. B. and Stewart, G. W. (1978) *LINPACK Users Guide*. Philadelphia: SIAM Publications.

## See Also

`chol`, `qr`, `solve`.

**Examples**

```
## upper triangular matrix 'r':  
r <- rbind(c(1,2,3),  
           c(0,1,1),  
           c(0,0,2))  
( y <- backsolve(r, x <- c(8,4,2)) ) # -1 3 1  
r %*% y # == x = (8,4,2)  
backsolve(r, x, transpose = TRUE) # 8 -12 -5
```

---

## Bessel *Bessel Functions*

---

### Description

Bessel Functions of integer and fractional order, of first and second kind,  $J_\nu$  and  $Y_\nu$ , and Modified Bessel functions (of first and third kind),  $I_\nu$  and  $K_\nu$ .

`gammaCody` is the ( $\Gamma$ ) function as from the `Specfun` package and originally used in the Bessel code.

### Usage

```
besselI(x, nu, expon.scaled = FALSE)
besselK(x, nu, expon.scaled = FALSE)
besselJ(x, nu)
besselY(x, nu)
gammaCody(x)
```

### Arguments

<code>x</code>	numeric, $\geq 0$ .
<code>nu</code>	numeric; The <i>order</i> (maybe fractional!) of the corresponding Bessel function.
<code>expon.scaled</code>	logical; if <code>TRUE</code> , the results are exponentially scaled in order to avoid overflow ( $I_\nu$ ) or underflow ( $K_\nu$ ), respectively.

### Details

The underlying C code stems from *Netlib* ([http://www.netlib.org/specfun/r\[ijky\]bes1](http://www.netlib.org/specfun/r[ijky]bes1)).

If `expon.scaled = TRUE`,  $e^{-x}I_\nu(x)$ , or  $e^xK_\nu(x)$  are returned.

`gammaCody` may be somewhat faster but less precise and/or robust than R's standard `gamma`. It is here for experimental purpose mainly, and *may be defunct very soon*.

For  $\nu < 0$ , formulae 9.1.2 and 9.6.2 from the reference below are applied (which is probably suboptimal), except for `besselK` which is symmetric in `nu`.

## Value

Numeric vector of the same length of `x` with the (scaled, if `expon.scale=TRUE`) values of the corresponding Bessel function.

## Author(s)

Original Fortran code: W. J. Cody, Argonne National Laboratory  
Translation to C and adaption to R: Martin Maechler.

## References

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. Dover, New York; Chapter 9: Bessel Functions of Integer Order.

## See Also

Other special mathematical functions, such as `gamma`,  $\Gamma(x)$ , and `beta`,  $B(x)$ .

## Examples

```
nus <- c(0:5,10,20)

x <- seq(0,4, len= 501)
plot(x,x, ylim = c(0,6), ylab="",type='n',
     main = "Bessel Functions I_nu(x)")
for(nu in nus) lines(x,besselI(x,nu=nu), col = nu+2)
legend(0,6, leg=paste("nu=",nus), col = nus+2, lwd=1)

x <- seq(0,40,len=801); y1 <- c(-.8,.8)
plot(x,x, ylim = y1, ylab="",type='n',
     main = "Bessel Functions J_nu(x)")
for(nu in nus) lines(x,besselJ(x,nu=nu), col = nu+2)
legend(32,-.18, leg=paste("nu=",nus), col = nus+2, lwd=1)

## Negative nu's :
xx <- 2:7
nu <- seq(-10,9, len = 2001)
op <- par(lab = c(16,5,7))
matplot(nu, t(outer(xx,nu, besselI)), type = 'l',
       ylim = c(-50,200),
       main = expression(paste("Bessel ",I[nu](x)," for fixed ",
                               x, ", as ",f(nu))),
```

```

      xlab = expression(nu))
abline(v=0, col = "light gray", lty = 3)
legend(5,200, leg = paste("x=",xx), col=seq(xx),
      lty=seq(xx))
par(op)

x0 <- 2^(-20:10)
plot(x0,x0^-8, log='xy', ylab="",type='n',
      main = "Bessel Functions J_nu(x) near 0\n log-log scale")
for(nu in sort(c(nus,nus+.5)))
  lines(x0,besselJ(x0,nu=nu), col = nu+2)
legend(3,1e50, leg=paste("nu=", paste(nus,nus+.5,sep=",")),
      col=nus+2, lwd=1)

plot(x0,x0^-8, log='xy', ylab="",type='n',
      main = "Bessel Functions K_nu(x) near 0\n log-log scale")
for(nu in sort(c(nus,nus+.5)))
  lines(x0,besselK(x0,nu=nu), col = nu+2)
legend(3,1e50,leg=paste("nu=", paste(nus,nus+.5, sep=",")),
      col=nus+2, lwd=1)

x <- x[x > 0]
plot(x,x, ylim=c(1e-18,1e11),log="y", ylab="",type='n',
      main = "Bessel Functions K_nu(x)")
for(nu in nus) lines(x,besselK(x,nu=nu), col = nu+2)
legend(0,1e-5, leg=paste("nu=",nus), col = nus+2, lwd=1)

yl <- c(-1.6, .6)
plot(x,x, ylim = yl, ylab="",type='n',
      main = "Bessel Functions Y_nu(x)")
for(nu in nus){
  xx <- x[x > .6*nus];
  lines(xx,besselY(xx,nu=nu), col = nu+2)
}
legend(25,-.5, leg=paste("nu=",nus), col = nus+2, lwd=1)

```

---

**chol**     *The Choleski Decomposition*


---

**Description**

Compute the Choleski factorization of a real symmetric positive-definite square matrix.

**Usage**

```
chol(x, pivot = FALSE, LINPACK = pivot)
La.chol(x)
```

**Arguments**

<b>x</b>	a real symmetric, positive-definite matrix
<b>pivot</b>	Should pivoting be used?
<b>LINPACK</b>	logical. Should LINPACK be used (for compatibility with R < 1.7.0)?

**Details**

`chol(pivot = TRUE)` provides an interface to the LINPACK routine DCHDC. `La.chol` provides an interface to the LAPACK routine DPOTRF.

Note that only the upper triangular part of **x** is used, so that  $R'R = x$  when **x** is symmetric.

If `pivot = FALSE` and **x** is not non-negative definite an error occurs. If **x** is positive semi-definite (i.e., some zero eigenvalues) an error will also occur, as a numerical tolerance is used.

If `pivot = TRUE`, then the Choleski decomposition of a positive semi-definite **x** can be computed. The rank of **x** is returned as `attr(Q, "rank")`, subject to numerical errors. The pivot is returned as `attr(Q, "pivot")`. It is no longer the case that `t(Q) %*% Q` equals **x**. However, setting `pivot <- attr(Q, "pivot")` and `oo <- order(pivot)`, it is true that `t(Q[, oo]) %*% Q[, oo]` equals **x**, or, alternatively, `t(Q) %*% Q` equals `x[pivot, pivot]`. See the examples.

## Value

The upper triangular factor of the Choleski decomposition, i.e., the matrix  $R$  such that  $R'R = x$  (see example).

If pivoting is used, then two additional attributes "pivot" and "rank" are also returned.

## Warning

The code does not check for symmetry.

If `pivot = TRUE` and `x` is not non-negative definite then there will be no error message but a meaningless result will occur. So only use `pivot = TRUE` when `x` is non-negative definite by construction.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Dongarra, J. J., Bunch, J. R., Moler, C. B. and Stewart, G. W. (1978) *LINPACK Users Guide*. Philadelphia: SIAM Publications.
- Anderson. E., et al. (1999) *LAPACK Users' Guide*. Third Edition. SIAM. ISBN 0-89871-447-8.

## See Also

`chol2inv` for its *inverse* (without pivoting), `backsolve` for solving linear systems with upper triangular left sides.

`qr`, `svd` for related matrix factorizations.

## Examples

```
( m <- matrix(c(5,1,1,3),2,2) )
( cm <- chol(m) )
t(cm) %*% cm  # = 'm'
crossprod(cm)  # = 'm'

# now for something positive semi-definite
x <- matrix(c(1:5, (1:5)^2), 5, 2)
x <- cbind(x, x[, 1] + 3*x[, 2])
m <- crossprod(x)
qr(m)$rank # is 2, as it should be

# chol() may fail, depending on numerical rounding:
```



```
# chol() unlike qr() does not use a tolerance.
try(chol(m))

# NB wrong rank here ... see Warning section.
(Q <- chol(m, pivot = TRUE))
## we can use this by
pivot <- attr(Q, "pivot")
oo <- order(pivot)
t(Q[, oo]) %*% Q[, oo] # recover m
```

---

**chol2inv**      *Inverse from Choleski Decomposition*


---

**Description**

Invert a symmetric, positive definite square matrix from its Choleski decomposition.

**Usage**

```
chol2inv(x, size = NCOL(x), LINPACK = FALSE)
La.chol2inv(x, size = ncol(x))
```

**Arguments**

<b>x</b>	a matrix. The first <code>nc</code> columns of the upper triangle contain the Choleski decomposition of the matrix to be inverted.
<b>size</b>	the number of columns of <code>x</code> containing the Choleski decomposition.
<b>LINPACK</b>	logical. Should LINPACK be used (for compatibility with R < 1.7.0)?

**Details**

`chol2inv(LINPACK=TRUE)` provides an interface to the LINPACK routine DPODI. `La.chol2inv` provides an interface to the LAPACK routine DPOTRI.

**Value**

The inverse of the decomposed matrix.

**References**

Dongarra, J. J., Bunch, J. R., Moler, C. B. and Stewart, G. W. (1978) *LINPACK Users Guide*. Philadelphia: SIAM Publications.

Anderson. E., et al. (1999) *LAPACK Users' Guide*. Third Edition. SIAM. ISBN 0-89871-447-8.

**See Also**

`chol`, `solve`.

**Examples**

```
cma <- chol(ma <- cbind(1, 1:3, c(1,3,7)))  
ma %*% chol2inv(cma)
```

---

`colSums`      *Form Row and Column Sums and Means*

---

## Description

Form row and column sums and means for numeric arrays.

## Usage

```
colSums(x, na.rm = FALSE, dims = 1)
rowSums(x, na.rm = FALSE, dims = 1)
colMeans(x, na.rm = FALSE, dims = 1)
rowMeans(x, na.rm = FALSE, dims = 1)
```

## Arguments

<code>x</code>	an array of two or more dimensions, containing numeric, complex, integer or logical values, or a numeric data frame.
<code>na.rm</code>	logical. Should missing values (including <code>NaN</code> ) be omitted from the calculations?
<code>dims</code>	Which dimensions are regarded as “rows” or “columns” to sum over. For <code>row*</code> , the sum or mean is over dimensions <code>dims+1, ...</code> ; for <code>col*</code> it is over dimensions <code>1:dims</code> .

## Details

These functions are equivalent to the use of `apply` with `FUN = mean` or `FUN = sum` with appropriate margins, but are a lot faster. As they are written for speed, they blur over some of the subtleties of `NaN` and `NA`. If `na.rm = FALSE` and either `NaN` or `NA` appears in a sum, the result will be one of `NaN` or `NA`, but which might be platform-dependent.

## Value

A numeric or complex array of suitable size, or a vector if the result is one-dimensional. The `dimnames` (or `names` for a vector result) are taken from the original array.

If there are no values in a range to be summed over (after removing missing values with `na.rm = TRUE`), that component of the output is set to 0 (`*Sums`) or `NA` (`*Means`), consistent with `sum` and `mean`.

## See Also

apply, rowsum

## Examples

```
## Compute row and column sums for a matrix:
x <- cbind(x1 = 3, x2 = c(4:1, 2:5))
rowSums(x); colSums(x)
dimnames(x)[[1]] <- letters[1:8]
rowSums(x); colSums(x); rowMeans(x); colMeans(x)
x[] <- as.integer(x)
rowSums(x); colSums(x)
x[] <- x < 3
rowSums(x); colSums(x)
x <- cbind(x1 = 3, x2 = c(4:1, 2:5))
x[3, ] <- NA; x[4, 2] <- NA
rowSums(x); colSums(x); rowMeans(x); colMeans(x)
rowSums(x, na.rm = TRUE); colSums(x, na.rm = TRUE)
rowMeans(x, na.rm = TRUE); colMeans(x, na.rm = TRUE)
```

```
## an array
data(UCBAdmissions)
dim(UCBAdmissions)
rowSums(UCBAdmissions); rowSums(UCBAdmissions, dims = 2)
colSums(UCBAdmissions); colSums(UCBAdmissions, dims = 2)
```

```
## complex case
x <- cbind(x1 = 3 + 2i, x2 = c(4:1, 2:5) - 5i)
x[3, ] <- NA; x[4, 2] <- NA
rowSums(x); colSums(x); rowMeans(x); colMeans(x)
rowSums(x, na.rm = TRUE); colSums(x, na.rm = TRUE)
rowMeans(x, na.rm = TRUE); colMeans(x, na.rm = TRUE)
```

---

**convolve**      *Fast Convolution*


---

**Description**

Use the Fast Fourier Transform to compute several kinds of convolutions of two sequences.

**Usage**

```
convolve(x, y, conj = TRUE, type = c("circular", "open",
                                     "filter"))
```

**Arguments**

<b>x,y</b>	numeric sequences <i>of the same length</i> to be convolved.
<b>conj</b>	logical; if TRUE, take the complex <i>conjugate</i> before back-transforming (default, and used for usual convolution).
<b>type</b>	character; one of "circular", "open", "filter" (beginning of word is ok). For <b>circular</b> , the two sequences are treated as <i>circular</i> , i.e., periodic. For <b>open</b> and <b>filter</b> , the sequences are padded with 0s (from left and right) first; "filter" returns the middle sub-vector of "open", namely, the result of running a weighted mean of <b>x</b> with weights <b>y</b> .

**Details**

The Fast Fourier Transform, **fft**, is used for efficiency.

The input sequences **x** and **y** must have the same length if **circular** is true.

Note that the usual definition of convolution of two sequences **x** and **y** is given by **convolve(x, rev(y), type = "o")**.

**Value**

If **r <- convolve(x,y, type = "open")** and **n <- length(x)**, **m <- length(y)**, then

$$r_k = \sum_i x_{k-m+i} y_i$$

where the sum is over all valid indices  $i$ , for  $k = 1, \dots, n + m - 1$

If `type == "circular"`,  $n = m$  is required, and the above is true for  $i, k = 1, \dots, n$  when  $x_j := x_{n+j}$  for  $j < 1$ .

## References

Brillinger, D. R. (1981) *Time Series: Data Analysis and Theory*, Second Edition. San Francisco: Holden-Day.

## See Also

`fft`, `nextn`, and particularly `filter` (from the `ts` package) which may be more appropriate.

## Examples

```
x <- c(0,0,0,100,0,0,0)
y <- c(0,0,1, 2 ,1,0,0)/4
zapsmall(convolve(x,y)) # NOT what you first thought.
zapsmall(convolve(x, y[3:5], type="f")) # rather
x <- rnorm(50)
y <- rnorm(50)
# Circular convolution has this symmetry:
all.equal(convolve(x,y, conj = FALSE),
          rev(convolve(rev(y),x)))

n <- length(x <- -20:24)
y <- (x-10)^2/1000 + rnorm(x)/8

Han <- function(y) # Hanning
  convolve(y, c(1,2,1)/4, type = "filter")

plot(x,y, main="Using convolve(.) for Hanning filters")
lines(x[-c(1,n)],Han(y),col="red")
lines(x[-c(1:2,(n-1):n)],Han(Han(y)),lwd=2,col="dark blue")
```

---

**crossprod**     *Matrix Crossproduct*

---

**Description**

Given matrices `x` and `y` as arguments, **crossprod** returns their matrix cross-product. This is formally equivalent to, but faster than, the call `t(x) %*% y`.

**Usage**

```
crossprod(x, y = NULL)
```

**Arguments**

`x`, `y`                    matrices: `y = NULL` is taken to be the same matrix as `x`.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`%*%` and outer product `%o%`.

**Examples**

```
(z <- crossprod(1:4))     # = sum(1 + 2^2 + 3^2 + 4^2)
drop(z)                   # scalar
```



---

**cumsum**      *Cumulative Sums, Products, and Extremes*

---

**Description**

Returns a vector whose elements are the cumulative sums, products, minima or maxima of the elements of the argument.

**Usage**

```
cumsum(x)
cumprod(x)
cummax(x)
cummin(x)
```

**Arguments**

**x**                      a numeric object.

**Details**

An NA value in **x** causes the corresponding and following elements of the return value to be NA.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (**cumsum** only.)

**Examples**

```
cumsum(1:10)
cumprod(1:10)
cummin(c(3:1, 2:0, 4:2))
cummax(c(3:1, 2:0, 4:2))
```

---

**deriv**     *Symbolic and Algorithmic Derivatives of Simple Expressions*

---

## Description

Compute derivatives of simple expressions, symbolically.

## Usage

```
D (expr, name)
deriv(expr, namevec, function.arg, tag = ".expr",
      hessian = FALSE)
deriv3(expr, namevec, function.arg, tag = ".expr",
      hessian = TRUE)
```

## Arguments

<b>expr</b>	expression or call to be differentiated.
<b>name, namevec</b>	character vector, giving the variable names (only one for <code>D()</code> ) with respect to which derivatives will be computed.
<b>function.arg</b>	If specified, a character vector of arguments for a function return, or a function (with empty body) or <code>TRUE</code> , the latter indicating that a function with argument names <code>namevec</code> should be used.
<b>tag</b>	character; the prefix to be used for the locally created variables in result.
<b>hessian</b>	a logical value indicating whether the second derivatives should be calculated and incorporated in the return value.

## Details

`D` is modelled after its `S` namesake for taking simple symbolic derivatives. `deriv` is a *generic* function with a default and a `formula` method. It returns a call for computing the `expr` and its (partial) derivatives, simultaneously. It uses so-called “*algorithmic derivatives*”. If `function.arg` is a function, its arguments can have default values, see the `fx` example below.

Currently, `deriv.formula` just calls `deriv.default` after extracting the expression to the right of `~`.

`deriv3` and its methods are equivalent to `deriv` and its methods except that `hessian` defaults to `TRUE` for `deriv3`.

## Value

`D` returns a call and therefore can easily be iterated for higher derivatives.

`deriv` and `deriv3` normally return an **expression** object whose evaluation returns the function values with a **"gradient"** attribute containing the gradient matrix. If `hessian` is `TRUE` the evaluation also returns a **"hessian"** attribute containing the Hessian array.

If `function.arg` is specified, `deriv` and `deriv3` return a function with those arguments rather than an expression.

## References

Griewank, A. and Corliss, G. F. (1991) *Automatic Differentiation of Algorithms: Theory, Implementation, and Application*. SIAM proceedings, Philadelphia.

Bates, D. M. and Chambers, J. M. (1992) *Nonlinear models*. Chapter 10 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`nlm` and `optim` for numeric minimization which could make use of derivatives, `nls` in package `nls`.

## Examples

```
## formula argument :
dx2x <- deriv(~ x^2, "x") ; dx2x
expression({
  .value <- x^2
  .grad <- array(0, c(length(.value),1), list(NULL,c("x")))
  .grad[, "x"] <- 2 * x
  attr(.value, "gradient") <- .grad
  .value
})
mode(dx2x)
x <- -1:2
eval(dx2x)
```

```

## Something 'tougher':
trig.exp <- expression(sin(cos(x + y^2)))
( D.sc <- D(trig.exp, "x") )
all.equal(D(trig.exp[[1]], "x"), D.sc)

( dxy <- deriv(trig.exp, c("x", "y")) )
y <- 1
eval(dxy)
eval(D.sc)

## function returned:
deriv((y ~ sin(cos(x) * y)), c("x","y"), func = TRUE)

## function with defaulted arguments:
(fx <- deriv(y ~ b0 + b1 * 2^(-x/th), c("b0", "b1", "th"),
            function(b0, b1, th, x = 1:7){} ) )
fx(2,3,4)

## Higher derivatives
deriv3(y ~ b0 + b1 * 2^(-x/th), c("b0", "b1", "th"),
      c("b0", "b1", "th", "x") )

## Higher derivatives:
DD <- function(expr,name, order = 1) {
  if(order < 1) stop("'order' must be >= 1")
  if(order == 1) D(expr,name)
  else DD(D(expr, name), name, order - 1)
}

DD(expression(sin(x^2)), "x", 3)
## showing the limits of the internal "simplify()" :
-sin(x^2) * (2 * x) * 2
+ ((cos(x^2) * (2 * x) * (2 * x) + sin(x^2) * 2) * (2 * x)
+ sin(x^2) * (2 * x) * 2)

```

---

**eigen**     *Spectral Decomposition of a Matrix*


---

**Description**

Computes eigenvalues and eigenvectors.

**Usage**

```
eigen(x, symmetric, only.values = FALSE, EISPACK = FALSE)
La.eigen(x, symmetric, only.values = FALSE,
         method = c("dsyevr", "dsyev"))
```

**Arguments**

<b>x</b>	a matrix whose spectral decomposition is to be computed.
<b>symmetric</b>	if <b>TRUE</b> , the matrix is assumed to be symmetric (or Hermitian if complex) and only its lower triangle is used. If <b>symmetric</b> is not specified, the matrix is inspected for symmetry.
<b>only.values</b>	if <b>TRUE</b> , only the eigenvalues are computed and returned, otherwise both eigenvalues and eigenvectors are returned.
<b>EISPACK</b>	logical. Should EISPACK be used (for compatibility with R < 1.7.0)?
<b>method</b>	The LAPACK routine to use in the real symmetric case.

**Details**

These functions use the LAPACK routines DSYEV/DSYEVR, DGEEV, ZHEEV and ZGEEV, and **eigen(EISPACK=TRUE)** provides an interface to the EISPACK routines RS, RG, CH and CG.

If **symmetric** is unspecified, the code attempts to determine if the matrix is symmetric up to plausible numerical inaccuracies. It is faster and surer to set the value yourself.

**eigen** is preferred to **eigen(EISPACK=TRUE)** for new projects, but its eigenvectors may differ in sign and (in the asymmetric case) in normalization. (They may also differ between methods and between platforms.)

The LAPACK routine DSYEVR is usually substantially faster than DSYEV. Most benefits are seen with an optimized BLAS system.

Using `method="dsyevr"` requires IEEE 754 arithmetic. Should this not be supported on your platform, `method="dsyev"` is used, with a warning.

Computing the eigenvectors is the slow part for large matrices.

## Value

The spectral decomposition of `x` is returned as components of a list.

- values**            a vector containing the  $p$  eigenvalues of `x`, sorted in *decreasing* order, according to `Mod(values)` if they are complex.
- vectors**        a  $p \times p$  matrix whose columns contain the eigenvectors of `x`, or NULL if `only.values` is TRUE.  
                     For `eigen(..., symmetric=FALSE, EISPACK=TRUE)` the choice of length of the eigenvectors is not defined by EISPACK. In all other cases the vectors are normalized to unit length.  
                     Recall that the eigenvectors are only defined up to a constant: even when the length is specified they are still only defined up to a scalar of modulus one (the sign for real matrices).

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Smith, B. T., Boyle, J. M., Dongarra, J. J., Garbow, B. S., Ikebe, Y., Klema, V., and Moler, C. B. (1976). *Matrix Eigensystems Routines – EISPACK Guide*. Springer-Verlag Lecture Notes in Computer Science.
- Anderson. E., et al. (1999) *LAPACK Users' Guide*. Third Edition. SIAM. ISBN 0-89871-447-8.

## See Also

- `svd`, a generalization of `eigen`; `qr`, and `chol` for related decompositions.
- To compute the determinant of a matrix, the `qr` decomposition is much more efficient: `det`.
- `capabilities` to test for IEEE 754 arithmetic.

**Examples**

```
eigen(cbind(c(1,-1),c(-1,1)))  
# same (different algorithm)  
eigen(cbind(c(1,-1),c(-1,1)), symmetric = FALSE)  
  
eigen(cbind(1,c(1,-1)), only.values = TRUE)  
eigen(cbind(-1,2:1)) # complex values  
# Hermitian, real eigenvalues  
eigen(print(cbind(c(0,1i), c(-1i,0))))  
## 3 x 3:  
eigen(cbind( 1,3:1,1:3))  
eigen(cbind(-1,c(1:2,0),0:2)) # complex values
```

---

**Extremes**      *Maxima and Minima*

---

**Description**

Returns the (parallel) maxima and minima of the input values.

**Usage**

```
max(..., na.rm=FALSE)
min(..., na.rm=FALSE)

pmax(..., na.rm=FALSE)
pmin(..., na.rm=FALSE)
```

**Arguments**

<code>...</code>	numeric arguments.
<code>na.rm</code>	a logical indicating whether missing values should be removed.

**Value**

`max` and `min` return the maximum or minimum of all the values present in their arguments, as **integer** if all are **integer**, or as **double** otherwise.

The minimum and maximum of an empty set are `+Inf` and `-Inf` (in this order!) which ensures *transitivity*, e.g., `min(x1, min(x2)) == min(x1,x2)`. In R versions before 1.5, `min(integer(0)) == .Machine$integer.max`, and analogously for `max`, preserving argument *type*, whereas from R version 1.5.0, `max(x) == -Inf` and `min(x) == +Inf` whenever `length(x) == 0` (after removing missing values if requested).

If `na.rm` is **FALSE** an NA value in any of the arguments will cause a value of NA to be returned, otherwise NA values are ignored.

`pmax` and `pmin` take several vectors (or matrices) as arguments and return a single vector giving the parallel maxima (or minima) of the vectors. The first element of the result is the maximum (minimum) of the first elements of all the arguments, the second element of the result is the maximum (minimum) of the second elements of all the arguments and so on. Shorter vectors are recycled if necessary. If `na.rm` is **FALSE**, NA values in the input vectors will produce NA values in the output. If



`na.rm` is TRUE, NA values are ignored. `attributes` (such as `names` or `dim`) are transferred from the first argument (if applicable).

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`range` (*both* min and max) and `which.min` (`which.max`) for the *arg min*, i.e., the location where an extreme value occurs.

## Examples

```
min(5:1,pi)
pmin(5:1, pi)
x <- sort(rnorm(100)); cH <- 1.35
pmin(cH, quantile(x)) # no names
pmin(quantile(x), cH) # has names
plot(x, pmin(cH, pmax(-cH, x)), type='b',
      main= "Huber's function")
```

---

**fft**     *Fast Discrete Fourier Transform*

---

**Description**

Performs the Fast Fourier Transform of an array.

**Usage**

```
fft(z, inverse = FALSE)
mvfft(z, inverse = FALSE)
```

**Arguments**

<b>z</b>	a real or complex array containing the values to be transformed.
<b>inverse</b>	if <b>TRUE</b> , the unnormalized inverse transform is computed (the inverse has a + in the exponent of $e$ , but here, we do <i>not</i> divide by $1/\text{length}(\mathbf{x})$ ).

**Value**

When **z** is a vector, the value computed and returned by **fft** is the unnormalized univariate Fourier transform of the sequence of values in **z**. When **z** contains an array, **fft** computes and returns the multivariate (spatial) transform. If **inverse** is **TRUE**, the (unnormalized) inverse Fourier transform is returned, i.e., if  $\mathbf{y} \leftarrow \text{fft}(\mathbf{z})$ , then  $\mathbf{z}$  is  $\text{fft}(\mathbf{y}, \text{inverse} = \text{TRUE}) / \text{length}(\mathbf{y})$ .

By contrast, **mvfft** takes a real or complex matrix as argument, and returns a similar shaped matrix, but with each column replaced by its discrete Fourier transform. This is useful for analyzing vector-valued series.

The FFT is fastest when the length of the series being transformed is highly composite (i.e., has many factors). If this is not the case, the transform may take a long time to compute and will use a large amount of memory.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Singleton, R. C. (1979) Mixed Radix Fast Fourier Transforms, in *Programs for Digital Signal Processing*, IEEE Digital Signal Processing Committee eds. IEEE Press.

### See Also

`convolve`, `nextn`.

### Examples

```
x <- 1:4
fft(x)
fft(fft(x), inverse = TRUE)/length(x)
```

---

**findInterval**      *Find Interval Numbers or Indices*


---

**Description**

Find the indices of **x** in **vec**, where **vec** must be sorted (non-decreasingly); i.e., if `i <- findInterval(x,v)`, we have  $v_{i_j} \leq x_j < v_{i_j+1}$  where  $v_0 := -\infty$ ,  $v_{N+1} := +\infty$ , and  $N <- \text{length}(\text{vec})$ . At the two boundaries, the returned index may differ by 1, depending on the optional arguments **rightmost.closed** and **all.inside**.

**Usage**

```
findInterval(x, vec, rightmost.closed = FALSE,
             all.inside = FALSE)
```

**Arguments**

<b>x</b>	numeric.
<b>vec</b>	numeric, sorted (weakly) increasingly, of length <b>N</b> , say.
<b>rightmost.closed</b>	logical; if true, the rightmost interval, <b>vec</b> [ <b>N</b> -1] .. <b>vec</b> [ <b>N</b> ] is treated as <i>closed</i> , see below.
<b>all.inside</b>	logical; if true, the returned indices are coerced into $\{1, \dots, N-1\}$ , i.e., 0 is mapped to 1 and <b>N</b> to <b>N</b> - 1.

**Details**

The function **findInterval** finds the index of one vector **x** in another, **vec**, where the latter must be non-decreasing. Where this is trivial, equivalent to `apply( outer(x, vec, ">="), 1, sum)`, as a matter of fact, the internal algorithm uses interval search ensuring  $O(n \log N)$  complexity where  $n <- \text{length}(\mathbf{x})$  (and  $N <- \text{length}(\text{vec})$ ). For (almost) sorted **x**, it will be even faster, basically  $O(n)$ .

This is the same computation as for the empirical distribution function, and indeed, `findInterval(t, sort(X))` is *identical* to  $nF_n(t; X_1, \dots, X_n)$  where  $F_n$  is the empirical distribution function of  $X_1, \dots, X_n$ .

When **rightmost.closed** = **TRUE**, the result for `x[j] = vec[N]` (= `max(vec)`), is **N** - 1 as for all other values in the last interval.

**Value**

vector of length `length(x)` with values in `0:N` where `N <- length(vec)`, or values coerced to `1:(N-1)` if `all.inside = TRUE` (equivalently coercing all `x` values *inside* the intervals).

**Author(s)**

Martin Maechler

**See Also**

`approx(*, method = "constant")` which is a generalization of `findInterval()`, `ecdf` for computing the empirical distribution function which is (up to a factor of  $n$ ) also basically the same as `findInterval()`.

**Examples**

```
N <- 100
X <- sort(round(rt(N, df=2), 2))
tt <- c(-100, seq(-2,2, len=201), +100)
it <- findInterval(tt, X)
# only first and last are outside range(X)
tt[it < 1 | it >= N]
```

---

**gl**     *Generate Factor Levels*

---

**Description**

Generate factors by specifying the pattern of their levels.

**Usage**

```
gl(n, k, length = n*k, labels = 1:n, ordered = FALSE)
```

**Arguments**

<b>n</b>	an integer giving the number of levels.
<b>k</b>	an integer giving the number of replications.
<b>length</b>	an integer giving the length of the result.
<b>labels</b>	an optional vector of labels for the resulting factor levels.
<b>ordered</b>	a logical indicating whether the result should be ordered or not.

**Value**

The result has levels from 1 to **n** with each value replicated in groups of length **k** out to a total length of **length**.

*gl* is modelled on the *GLIM* function of the same name.

**See Also**

The underlying `factor()`.

**Examples**

```
## First control, then treatment:
gl(2, 8, label = c("Control", "Treat"))
## 20 alternating 1s and 2s
gl(2, 1, 20)
## alternating pairs of 1s and 2s
gl(2, 2, 20)
```

---

**Hyperbolic**     *Hyperbolic Functions*

---

**Description**

These functions give the obvious hyperbolic functions. They respectively compute the hyperbolic cosine, sine, tangent, arc-cosine, arc-sine, arc-tangent.

**Usage**

```
cosh(x)
sinh(x)
tanh(x)
acosh(x)
asinh(x)
atanh(x)
```

**Arguments**

`x`                      a numeric vector

**See Also**

`cos`, `sin`, `tan`, `acos`, `asin`, `atan`.

---

**integrate**     *Integration of One-Dimensional Functions*


---

**Description**

Adaptive quadrature of functions of one variable over a finite or infinite interval.

**Usage**

```
integrate(f, lower, upper, subdivisions=100,
  rel.tol = .Machine$double.eps^0.25, abs.tol = rel.tol,
  stop.on.error = TRUE, keep.xy = FALSE, aux = NULL, ...)
```

**Arguments**

<b>f</b>	an R function taking a numeric first argument and returning a numeric vector of the same length. Returning a non-finite element will generate an error.
<b>lower, upper</b>	the limits of integration. Can be infinite.
<b>subdivisions</b>	the maximum number of subintervals.
<b>rel.tol</b>	relative accuracy requested.
<b>abs.tol</b>	absolute accuracy requested.
<b>stop.on.error</b>	logical. If true (the default) an error stops the function. If false some errors will give a result with a warning in the <b>message</b> component.
<b>keep.xy</b>	unused. For compatibility with S.
<b>aux</b>	unused. For compatibility with S.
<b>...</b>	additional arguments to be passed to <b>f</b> . Remember to use argument names <i>not</i> matching those of <b>integrate(.)</b> !

**Details**

If one or both limits are infinite, the infinite range is mapped onto a finite interval.

For a finite interval, globally adaptive interval subdivision is used in connection with extrapolation by the Epsilon algorithm.

**rel.tol** cannot be less than `max(50*.Machine$double.eps, 0.5e-28)` if **abs.tol** `<= 0`.



## Value

A list of class `"integrate"` with components

<code>value</code>	the final estimate of the integral.
<code>abs.error</code>	estimate of the modulus of the absolute error.
<code>subdivisions</code>	the number of subintervals produced in the subdivision process.
<code>message</code>	"OK" or a character string giving the error message.
<code>call</code>	the matched call.

## Note

Like all numerical integration routines, these evaluate the function on a finite set of points. If the function is approximately constant (in particular, zero) over nearly all its range it is possible that the result and error estimate may be seriously wrong.

When integrating over infinite intervals do so explicitly, rather than just using a large number as the endpoint. This increases the chance of a correct answer – any function whose integral over an infinite interval is finite must be near zero for most of that interval.

## References

Based on QUADPACK routines `dqags` and `dqagi` by R. Piessens and E. deDoncker-Kapenga, available from Netlib.

See

R. Piessens, E. deDoncker-Kapenga, C. Uberhuber, D. Kahaner (1983) *Quadpack: a Subroutine Package for Automatic Integration*; Springer Verlag.

## See Also

The function `adapt` in the `adapt` package on CRAN, for multivariate integration.

## Examples

```
integrate(dnorm, -1.96, 1.96)
integrate(dnorm, -Inf, Inf)

## a slowly-convergent integral
integrand <- function(x) {1/((x+1)*sqrt(x))}
integrate(integrand, lower = 0, upper = Inf)
```

```
## don't do this if you really want the integral from 0 to
## Inf
integrate(integrand, lower = 0, upper = 10)
integrate(integrand, lower = 0, upper = 100000)
integrate(integrand, lower = 0, upper = 1000000,
          stop.on.error = FALSE)

## no vectorizable function
try(integrate(function(x) 2, 0, 1))
integrate(function(x) rep(2, length(x)), 0, 1) ## correct

## integrate can fail if misused
integrate(dnorm,0,2)
integrate(dnorm,0,20)
integrate(dnorm,0,200)
integrate(dnorm,0,2000)
integrate(dnorm,0,20000) ## fails on many systems
integrate(dnorm,0,Inf)   ## works
```

---

**kappa**     *Estimate the Condition Number*

---

**Description**

An estimate of the condition number of a matrix or of the  $R$  matrix of a  $QR$  decomposition, perhaps of a linear fit. The condition number is defined as the ratio of the largest to the smallest *non-zero* singular value of the matrix.

**Usage**

```
kappa(z, ...)  
## S3 method for class 'lm':  
kappa(z, ...)  
## Default S3 method:  
kappa(z, exact = FALSE, ...)  
## S3 method for class 'qr':  
kappa(z, ...)  
  
kappa.tri(z, exact = FALSE, ...)
```

**Arguments**

<b>z</b>	A matrix or a the result of <b>qr</b> or a fit from a class inheriting from "lm".
<b>exact</b>	logical. Should the result be exact?
<b>...</b>	further arguments passed to or from other methods.

**Details**

If **exact** = **FALSE** (the default) the condition number is estimated by a cheap approximation. Following S, this uses the LINPACK routine 'dtrco.f'. However, in R (or S) the exact calculation is also likely to be quick enough.

**kappa.tri** is an internal function called by **kappa.qr**.

**Value**

The condition number, *kappa*, or an approximation if **exact** = **FALSE**.

## Author(s)

The design was inspired by (but differs considerably from) the `S` function of the same name described in Chambers (1992).

## References

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`svd` for the singular value decomposition and `qr` for the *QR* one.

## Examples

```
kappa(x1 <- cbind(1,1:10)) # 15.71
kappa(x1, exact = TRUE)    # 13.68
kappa(x2 <- cbind(x1,2:11)) # high! [x2 is singular!]

hilbert <- function(n) {
  i <- 1:n; 1 / outer(i - 1, i, "+")
}
sv9 <- svd(h9 <- hilbert(9))$ d
kappa(h9) # pretty high!
kappa(h9, exact = TRUE) == max(sv9) / min(sv9)
# .677 (i.e., rel.error = 32%)
kappa(h9, exact = TRUE) / kappa(h9)
```

---

**log**     *Logarithms and Exponentials*

---

**Description**

`log` computes natural logarithms, `log10` computes common (i.e., base 10) logarithms, and `log2` computes binary (i.e., base 2) logarithms. The general form `logb(x, base)` computes logarithms with base `base` (`log10` and `log2` are only special cases).

`log1p(x)` computes  $\log(1 + x)$  accurately also for  $|x| \ll 1$  (and less accurately when  $x \approx -1$ ).

`exp` computes the exponential function.

`expm1(x)` computes  $\exp(x) - 1$  accurately also for  $|x| \ll 1$ .

**Usage**

```
log(x, base = exp(1))
logb(x, base = exp(1))
log10(x)
log2(x)
exp(x)
expm1(x)
log1p(x)
```

**Arguments**

<code>x</code>	a numeric or complex vector.
<code>base</code>	positive number. The base with respect to which logarithms are computed. Defaults to <code>e=exp(1)</code> .

**Value**

A vector of the same length as `x` containing the transformed values. `log(0)` gives `-Inf` (when available).

**Note**

`log` and `logb` are the same thing in R, but `logb` is preferred if `base` is specified, for S-PLUS compatibility.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (for `log`, `log10` and `exp`.)

Chambers, J. M. (1998) *Programming with Data. A Guide to the S Language*. Springer. (for `logb`.)

## See Also

`Trig`, `sqrt`, `Arithmetic`.

## Examples

```
log(exp(3))
log10(1e7) # = 7
```

```
x <- 10^-(1+2*1:9)
cbind(x, log(1+x), log1p(x), exp(x)-1, expm1(x))
```

---

**matmult**     *Matrix Multiplication*

---

**Description**

Multiplies two matrices, if they are conformable. If one argument is a vector, it will be coerced to either a row or column matrix to make the two arguments conformable. If both are vectors it will return the inner product.

**Usage**

```
a %*% b
```

**Arguments**

a, b                    numeric or complex matrices or vectors.

**Value**

The matrix product. Use **drop** to get rid of dimensions which have only one level.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

matrix, Arithmetic, diag.

**Examples**

```
x <- 1:4
(z <- x %*% x)      # scalar ("inner") product (1 x 1 matrix)
drop(z)            # as scalar

y <- diag(x)
z <- matrix(1:12, ncol = 3, nrow = 4)
y %*% z
y %*% x
x %*% z
```

---

**matrix**     *Matrices*

---

**Description**

**matrix** creates a matrix from the given set of values.

**as.matrix** attempts to turn its argument into a matrix.

**is.matrix** tests if its argument is a (strict) matrix. It is generic: you can write methods to handle specific classes of objects, see Internal-Methods.

**Usage**

```
matrix(data = NA, nrow = 1, ncol = 1, byrow = FALSE,  
        dimnames = NULL)  
as.matrix(x)  
is.matrix(x)
```

**Arguments**

<b>data</b>	an optional data vector.
<b>nrow</b>	the desired number of rows
<b>ncol</b>	the desired number of columns
<b>byrow</b>	logical. If <b>FALSE</b> (the default) the matrix is filled by columns, otherwise the matrix is filled by rows.
<b>dimnames</b>	A <b>dimnames</b> attribute for the matrix: a list of length 2.
<b>x</b>	an R object.

**Details**

If either of **nrow** or **ncol** is not given, an attempt is made to infer it from the length of **data** and the other parameter.

**is.matrix** returns **TRUE** if **x** is a matrix (i.e., it is *not* a **data.frame** and has a **dim** attribute of length 2) and **FALSE** otherwise.

**as.matrix** is a generic function. The method for data frames will convert any non-numeric column into a character vector using **format** and so return a character matrix.



## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`data.matrix`, which attempts to convert to a numeric matrix.

## Examples

```
is.matrix(as.matrix(1:10))
data(warpbreaks)
!is.matrix(warpbreaks) # data.frame, NOT matrix!
str(warpbreaks)
# using as.matrix.data.frame(.) method
str(as.matrix(warpbreaks))
```

---

**nextn**     *Highly Composite Numbers*

---

**Description**

**nextn** returns the smallest integer, greater than or equal to **n**, which can be obtained as a product of powers of the values contained in **factors**. **nextn** is intended to be used to find a suitable length to zero-pad the argument of **fft** to so that the transform is computed quickly. The default value for **factors** ensures this.

**Usage**

```
nextn(n, factors=c(2,3,5))
```

**Arguments**

<b>n</b>	an integer.
<b>factors</b>	a vector of positive integer factors.

**See Also**

`convolve`, `fft`.

**Examples**

```
nextn(1001) # 1024
table(sapply(599:630, nextn))
```

---

**poly**     *Compute Orthogonal Polynomials*

---

**Description**

Returns or evaluates orthogonal polynomials of degree 1 to **degree** over the specified set of points **x**. These are all orthogonal to the constant polynomial of degree 0.

**Usage**

```
poly(x, ..., degree = 1, coefs = NULL)
polym(..., degree = 1)
```

```
## S3 method for class 'poly':
predict(object, newdata, ...)
```

**Arguments**

<b>x</b> , <b>newdata</b>	a numeric vector at which to evaluate the polynomial. <b>x</b> can also be a matrix.
<b>degree</b>	the degree of the polynomial
<b>coefs</b>	for prediction, coefficients from a previous fit.
<b>object</b>	an object inheriting from class " <b>poly</b> ", normally the result of a call to <b>poly</b> with a single vector argument.
<b>...</b>	<b>poly</b> , <b>polym</b> : further vectors. <b>predict.poly</b> : arguments to be passed to or from other methods.

**Details**

Although formally **degree** should be named (as it follows ...), an unnamed second argument of length 1 will be interpreted as the degree.

The orthogonal polynomial is summarized by the coefficients, which can be used to evaluate it via the three-term recursion given in Kennedy & Gentle (1980, pp. 343-4), and used in the “predict” part of the code.

## Value

For `poly` with a single vector argument:

A matrix with rows corresponding to points in `x` and columns corresponding to the degree, with attributes "`degree`" specifying the degrees of the columns and "`coefs`" which contains the centering and normalization constants used in constructing the orthogonal polynomials. The matrix is given class `c("poly", "matrix")` as from R 1.5.0.

Other cases of `poly` and `polym`, and `predict.poly`: a matrix.

## Note

This routine is intended for statistical purposes such as `contr.poly`: it does not attempt to orthogonalize to machine accuracy.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

Kennedy, W. J. Jr and Gentle, J. E. (1980) *Statistical Computing* Marcel Dekker.

## See Also

`contr.poly`

## Examples

```
(z <- poly(1:10, 3))
predict(z, seq(2, 4, 0.5))
poly(seq(4, 6, 0.5), 3, coefs = attr(z, "coefs"))

polym(1:4, c(1, 4:6), degree=3) # or just poly()
poly(cbind(1:4, c(1, 4:6)), degree=3)
```

---

**polyroot**     *Find Zeros of a Real or Complex Polynomial*

---

**Description**

Find zeros of a real or complex polynomial.

**Usage**

```
polyroot(z)
```

**Arguments**

**z**                      the vector of polynomial coefficients in increasing order.

**Details**

A polynomial of degree  $n - 1$ ,

$$p(x) = z_1 + z_2x + \cdots + z_nx^{n-1}$$

is given by its coefficient vector **z**[1:n]. **polyroot** returns the  $n - 1$  complex zeros of  $p(x)$  using the Jenkins-Traub algorithm.

If the coefficient vector **z** has zeroes for the highest powers, these are discarded.

**Value**

A complex vector of length  $n - 1$ , where  $n$  is the position of the largest non-zero element of **z**.

**References**

Jenkins and Traub (1972) TOMS Algorithm 419. *Comm. ACM*, **15**, 97–99.

**See Also**

**uniroot** for numerical root finding of arbitrary functions; **complex** and the **zero** example in the demos directory.

**Examples**

```
polyroot(c(1, 2, 1))  
round(polyroot(choose(8, 0:8)), 11) # guess what!  
for (n1 in 1:4) print(polyroot(1:n1), digits = 4)  
polyroot(c(1, 2, 1, 0, 0)) # same as the first
```

---

**prod**     *Product of Vector Elements*

---

**Description**

`prod` returns the product of all the values present in its arguments.

**Usage**

```
prod(..., na.rm = FALSE)
```

**Arguments**

<code>...</code>	numeric vectors.
<code>na.rm</code>	logical. Should missing values be removed?

**Details**

If `na.rm` is `FALSE` an `NA` value in any of the arguments will cause a value of `NA` to be returned, otherwise `NA` values are ignored.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`sum`, `cumprod`, `cumsum`.

**Examples**

```
print(prod(1:7)) == print(gamma(8))
```

---

**qr**     *The QR Decomposition of a Matrix*


---

**Description**

**qr** computes the QR decomposition of a matrix. It provides an interface to the techniques used in the LINPACK routine DQRDC or the LAPACK routines DGEQP3 and (for complex matrices) ZGEQP3.

**Usage**

```
qr(x, tol = 1e-07 , LAPACK = FALSE)
qr.coef(qr, y)
qr.qy(qr, y)
qr.qty(qr, y)
qr.resid(qr, y)
qr.fitted(qr, y, k = qr$rank)
qr.solve(a, b, tol = 1e-7)
## S3 method for class 'qr':
solve(a, b, ...)

is.qr(x)
as.qr(x)
```

**Arguments**

<b>x</b>	a matrix whose QR decomposition is to be computed.
<b>tol</b>	the tolerance for detecting linear dependencies in the columns of <b>x</b> . Only used by LINPACK.
<b>qr</b>	a QR decomposition of the type computed by <b>qr</b> .
<b>y, b</b>	a vector or matrix of right-hand sides of equations.
<b>a</b>	A QR decomposition or ( <b>qr.solve</b> only) a rectangular matrix.
<b>k</b>	effective rank.
<b>LAPACK</b>	logical. For real <b>x</b> , if true use LAPACK otherwise use LINPACK.
<b>...</b>	further arguments passed to or from other methods



## Details

The QR decomposition plays an important role in many statistical techniques. In particular it can be used to solve the equation  $\mathbf{Ax} = \mathbf{b}$  for given matrix  $\mathbf{A}$ , and vector  $\mathbf{b}$ . It is useful for computing regression coefficients and in applying the Newton-Raphson algorithm.

The functions `qr.coef`, `qr.resid`, and `qr.fitted` return the coefficients, residuals and fitted values obtained when fitting  $y$  to the matrix with QR decomposition `qr`. `qr.qy` and `qr.qty` return  $\mathbf{Q} \%*\% y$  and  $t(\mathbf{Q}) \%*\% y$ , where  $\mathbf{Q}$  is the  $\mathbf{Q}$  matrix.

All the above functions keep `dimnames` (and `names`) of  $x$  and  $y$  if there are.

`solve.qr` is the method for `solve` for `qr` objects. `qr.solve` solves systems of equations via the QR decomposition: if  $\mathbf{a}$  is a QR decomposition it is the same as `solve.qr`, but if  $\mathbf{a}$  is a rectangular matrix the QR decomposition is computed first. Either will handle over- and under-determined systems, providing a minimal-length solution or a least-squares fit if appropriate.

`is.qr` returns `TRUE` if  $x$  is a list with components named `qr`, `rank` and `graux` and `FALSE` otherwise.

It is not possible to coerce objects to mode `"qr"`. Objects either are QR decompositions or they are not.

## Value

The QR decomposition of the matrix as computed by LINPACK or LAPACK. The components in the returned value correspond directly to the values returned by DQRDC/DGEQP3/ZGEQP3.

<code>qr</code>	a matrix with the same dimensions as $x$ . The upper triangle contains the $\mathbf{R}$ of the decomposition and the lower triangle contains information on the $\mathbf{Q}$ of the decomposition (stored in compact form). Note that the storage used by DQRDC and DGEQP3 differs.
<code>graux</code>	a vector of length <code>ncol(x)</code> which contains additional information on $\mathbf{Q}$ .
<code>rank</code>	the rank of $x$ as computed by the decomposition: always full rank in the LAPACK case.
<code>pivot</code>	information on the pivoting strategy used during the decomposition.

Non-complex QR objects computed by LAPACK have the attribute `"useLAPACK"` with value `TRUE`.

## Note

To compute the determinant of a matrix (do you *really* need it?), the QR decomposition is much more efficient than using Eigen values (**eigen**). See **det**.

Using LAPACK (including in the complex case) uses column pivoting and does not attempt to detect rank-deficient matrices.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Dongarra, J. J., Bunch, J. R., Moler, C. B. and Stewart, G. W. (1978) *LINPACK Users Guide*. Philadelphia: SIAM Publications.
- Anderson. E., et al. (1999) *LAPACK Users' Guide*. Third Edition. SIAM. ISBN 0-89871-447-8.

## See Also

**qr.Q**, **qr.R**, **qr.X** for reconstruction of the matrices. **solve.qr**, **lsfit**, **eigen**, **svd**.

**det** (using **qr**) to compute the determinant of a matrix.

## Examples

```
hilbert <- function(n) {
  i <- 1:n; 1 / outer(i - 1, i, "+")
}
h9 <- hilbert(9); h9
qr(h9)$rank           # only 7
qrh9 <- qr(h9, tol = 1e-10)
qrh9$rank             # 9
## Solve linear equation system H %*% x = y :
y <- 1:9/10
x <- qr.solve(h9, y, tol = 1e-10) # or equivalently :
x <- qr.coef(qrh9, y) # is == but much better than
                        # solve(h9) %*% y
h9 %*% x              # = y
```

---

<b>QR.Auxiliaries</b>	<i>Reconstruct the <math>Q</math>, <math>R</math>, or <math>X</math> Matrices from a <math>QR</math> Object</i>
-----------------------	---

---

## Description

Returns the original matrix from which the object was constructed or the components of the decomposition.

## Usage

```
qr.X(qr, complete = FALSE, ncol =)
qr.Q(qr, complete = FALSE, Dvec =)
qr.R(qr, complete = FALSE)
```

## Arguments

<b>qr</b>	object representing a QR decomposition. This will typically have come from a previous call to <b>qr</b> or <b>lsfit</b> .
<b>complete</b>	logical expression of length 1. Indicates whether an arbitrary orthogonal completion of the $Q$ or $X$ matrices is to be made, or whether the $R$ matrix is to be completed by binding zero-value rows beneath the square upper triangle.
<b>ncol</b>	integer in the range $1:\text{ncol}(\text{qr}\$qr)$ . The number of columns to be in the reconstructed $X$ . The default when <b>complete</b> is <b>FALSE</b> is the first $\min(\text{ncol}(X), \text{nrow}(X))$ columns of the original $X$ from which the <b>qr</b> object was constructed. The default when <b>complete</b> is <b>TRUE</b> is a square matrix with the original $X$ in the first $\text{ncol}(X)$ columns and an arbitrary orthogonal completion (unitary completion in the complex case) in the remaining columns.
<b>Dvec</b>	vector (not matrix) of diagonal values. Each column of the returned $Q$ will be multiplied by the corresponding diagonal value. Defaults to all 1s.

## Value

**qr.X** returns  $X$ , the original matrix from which the **qr** object was constructed, provided  $\text{ncol}(X) \leq \text{nrow}(X)$ . If **complete** is **TRUE** or the

argument `ncol` is greater than `ncol(X)`, additional columns from an arbitrary orthogonal (unitary) completion of **X** are returned.

`qr.Q` returns **Q**, the order-`nrow(X)` orthogonal (unitary) transformation represented by `qr`. If `complete` is `TRUE`, **Q** has `nrow(X)` columns. If `complete` is `FALSE`, **Q** has `ncol(X)` columns. When `Dvec` is specified, each column of **Q** is multiplied by the corresponding value in `Dvec`.

`qr.R` returns **R**, the upper triangular matrix such that `X == Q %*% R`. The number of rows of **R** is `nrow(X)` or `ncol(X)`, depending on whether `complete` is `TRUE` or `FALSE`.

### See Also

`qr`, `qr.qy`.

### Examples

```
data(LifeCycleSavings)
p <- ncol(x <- LifeCycleSavings[, -1]) # not the 'sr'
qrstr <- qr(x) # dim(x) == c(n,p)
qrstr $ rank # = 4 = p
Q <- qr.Q(qrstr) # dim(Q) == dim(x)
R <- qr.R(qrstr) # dim(R) == ncol(x)
X <- qr.X(qrstr) # X == x
range(X - as.matrix(x)) # ~ < 6e-12
## X == Q %*% R :
Q %*% R
```

---

**range**     *Range of Values*

---

**Description**

**range** returns a vector containing the minimum and maximum of all the given arguments.

**Usage**

```
range(..., na.rm = FALSE)
```

```
## Default S3 method:
```

```
range(..., na.rm = FALSE, finite = FALSE)
```

**Arguments**

<code>...</code>	any <b>numeric</b> objects.
<code>na.rm</code>	logical, indicating if NA's should be omitted.
<code>finite</code>	logical, indicating if all non-finite elements should be omitted.

**Details**

This is a generic function; currently, it has only a default method (`range.default`).

It is also a member of the **Summary** group of functions, see **Methods**.

If `na.rm` is **FALSE**, NA and NaN values in any of the arguments will cause NA values to be returned, otherwise NA values are ignored.

If `finite` is **TRUE**, the minimum and maximum of all finite values is computed, i.e., `finite=TRUE` *includes* `na.rm=TRUE`.

A special situation occurs when there is no (after omission of NAs) nonempty argument left, see **min**.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`min`, `max`, **Methods**.

**Examples**

```
print(r.x <- range(rnorm(100)))  
diff(r.x) # the SAMPLE range
```

```
x <- c(NA, 1:3, -1:1/0); x  
range(x)  
range(x, na.rm = TRUE)  
range(x, finite = TRUE)
```

---

**Round**     *Rounding of Numbers*

---

**Description**

**ceiling** takes a single numeric argument **x** and returns a numeric vector containing the smallest integers not less than the corresponding elements of **x**.

**floor** takes a single numeric argument **x** and returns a numeric vector containing the largest integers not greater than the corresponding elements of **x**.

**round** rounds the values in its first argument to the specified number of decimal places (default 0). Note that for rounding off a 5, the IEEE standard is used, “*go to the even digit*”. Therefore **round**(0.5) is 0 and **round**(-1.5) is -2.

**signif** rounds the values in its first argument to the specified number of significant digits.

**trunc** takes a single numeric argument **x** and returns a numeric vector containing the integers by truncating the values in **x** toward 0.

**zapsmall** determines a **digits** argument **dr** for calling **round(x, digits = dr)** such that values “close to zero” (compared with the maximal absolute one) are “zapped”, i.e., treated as 0.

**Usage**

```
ceiling(x)
floor(x)
round(x, digits = 0)
signif(x, digits = 6)
trunc(x)
zapsmall(x, digits= getOption("digits"))
```

**Arguments**

<b>x</b>	a numeric vector.
<b>digits</b>	integer indicating the precision to be used.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (except `zapsmall`.)

Chambers, J. M. (1998) *Programming with Data. A Guide to the S Language*. Springer. (`zapsmall`.)

## See Also

`as.integer`.

## Examples

```
round(.5 + -2:4) # IEEE rounding: -2  0  0  2  2  4  4
( x1 <- seq(-2, 4, by = .5) )
round(x1) # IEEE rounding !
x1[trunc(x1) != floor(x1)]
x1[round(x1) != floor(x1 + .5)]
(non.int <- ceiling(x1) != floor(x1))

x2 <- pi * 100^(-1:3)
round(x2, 3)
signif(x2, 3)

print  (x2 / 1000, digits=4)
zapsmall(x2 / 1000, digits=4)
zapsmall(exp(1i*0:4*pi/2))
```



---

**sign**     *Sign Function*

---

### Description

**sign** returns a vector with the signs of the corresponding elements of **x** (the sign of a real number is 1, 0, or  $-1$  if the number is positive, zero, or negative, respectively).

Note that **sign** does not operate on complex vectors.

### Usage

```
sign(x)
```

### Arguments

**x**                      a numeric vector

### See Also

**abs**

### Examples

```
sign(pi) # == 1
sign(-2:3) # -1 -1 0 1 1 1
```

---

**solve**     *Solve a System of Equations*


---

**Description**

This generic function solves the equation `a %*% x = b` for `x`, where `b` can be either a vector or a matrix.

**Usage**

```
solve(a, b, ...)
```

```
## Default S3 method:
```

```
solve(a, b, tol, LINPACK = FALSE, ...)
```

**Arguments**

<code>a</code>	a square numeric or complex matrix containing the coefficients of the linear system.
<code>b</code>	a numeric or complex vector or matrix giving the right-hand side(s) of the linear system. If missing, <code>b</code> is taken to be an identity matrix and <code>solve</code> will return the inverse of <code>a</code> .
<code>tol</code>	the tolerance for detecting linear dependencies in the columns of <code>a</code> . If <code>LINPACK</code> is <code>TRUE</code> the default is <code>1e-7</code> , otherwise it is <code>.Machine\$double.eps</code> . Future versions of R may use a tighter tolerance. Not presently used with complex matrices <code>a</code> .
<code>LINPACK</code>	logical. Should LINPACK be used (for compatibility with R < 1.7.0)?
<code>...</code>	further arguments passed to or from other methods

**Details**

As from R 1.3.0, `a` or `b` can be complex, in which case LAPACK routine ZESV is used. This uses double complex arithmetic which might not be available on all platforms.

The row and column names of the result are taken from the column names of `a` and of `b` respectively. As from R 1.7.0 if `b` is missing the column names of the result are the row names of `a`. No check is made that the column names of `a` and the row names of `b` are equal.

For back-compatibility `a` can be a (real) QR decomposition, although `qr.solve` should be called in that case. `qr.solve` can handle non-square systems.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`solve.qr` for the `qr` method, `backsolve`, `qr.solve`.

## Examples

```
hilbert <- function(n) {
  i <- 1:n; 1 / outer(i - 1, i, "+")
}
h8 <- hilbert(8); h8
sh8 <- solve(h8)
round(sh8 %*% h8, 3)

A <- hilbert(4)
A[] <- as.complex(A)
## might not be supported on all platforms
try(solve(A))
```

---

**sort**     *Sorting or Ordering Vectors*

---

**Description**

Sort (or *order*) a numeric or complex vector (partially) into ascending (or descending) order.

**Usage**

```
sort(x, partial = NULL, na.last = NA, decreasing = FALSE,  
     method = c("shell", "quick"), index.return = FALSE)  
is.unsorted(x, na.rm = FALSE)
```

**Arguments**

<b>x</b>	a numeric or complex vector.
<b>partial</b>	a vector of indices for partial sorting.
<b>na.last</b>	for controlling the treatment of NAs. If <b>TRUE</b> , missing values in the data are put last; if <b>FALSE</b> , they are put first; if <b>NA</b> , they are removed.
<b>decreasing</b>	logical. Should the sort be increasing or decreasing?
<b>method</b>	character specifying the algorithm used.
<b>index.return</b>	logical indicating if the ordering index vector should be returned as well; this is only available for the default <b>na.last = NA</b> .
<b>na.rm</b>	logical. Should missing values be removed?

**Details**

If **partial** is not **NULL**, it is taken to contain indices of elements of **x** which are to be placed in their correct positions by partial sorting. After the sort, the values specified in **partial** are in their correct position in the sorted array. Any values smaller than these values are guaranteed to have a smaller index in the sorted array and any values which are greater are guaranteed to have a bigger index in the sorted array.

The sort order for character vectors will depend on the collating sequence of the locale in use: see [Comparison](#).

**is.unsorted** returns a logical indicating if **x** is sorted increasingly, i.e., **is.unsorted(x)** is true if **any(x != sort(x))** (and there are no NAs).

`method = "shell"` uses Shellsort (an  $O(n^{4/3})$  variant from Sedgewick (1996)). If `x` has names a stable sort is used, so ties are not reordered. (This only matters if names are present.)

Method `"quick"` uses Singleton's Quicksort implementation and is only available when `x` is numeric (double or integer) and `partial` is `NULL`. It is normally somewhat faster than Shellsort (perhaps twice as fast on vectors of length a million) but has poor performance in the rare worst case. (Peto's modification using a pseudo-random midpoint is used to make the worst case rarer.) This is not a stable sort, and ties may be reordered.

## Value

For `sort` the sorted vector unless `index.return` is true, when the result is a list with components named `x` and `ix` containing the sorted numbers and the ordering index vector. In the latter case, if `method == "quick"` ties may be reversed in the ordering, unlike `sort.list`, as quicksort is not stable.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Sedgewick, R. (1986) A new upper bound for Shell sort. *J. Algorithms* **7**, 159–173.
- Singleton, R. C. (1969) An efficient algorithm for sorting with minimal storage: Algorithm 347. *Communications of the ACM* **12**, 185–187.

## See Also

`order`, `rank`.

## Examples

```
data(swiss)
x <- swiss$Education[1:25]
x; sort(x); sort(x, partial = c(10, 15))
median # shows you another example for 'partial'

## illustrate 'stable' sorting (of ties):
sort(c(10:3,2:12), method = "sh", index=TRUE) # is stable
## $x : 2  3  3  4  4  5  5  6  6  7  7  8  8  9  9 10 10
## 11 12
## $ix: 9  8 10  7 11  6 12  5 13  4 14  3 15  2 16
```

```
## 1 17 18 19
sort(c(10:3,2:12), method = "qu", index=TRUE) # is not
## $x : 2 3 3 4 4 5 5 6 6 7 7 8 8 9 9 10 10
## 11 12
## $ix: 9 10 8 7 11 6 12 5 13 4 14 3 15 16 2
## 17 1 18 19

## Small speed comparison simulation:
N <- 2000
Sim <- 20
rep <- 50 # << adjust to your CPU
c1 <- c2 <- numeric(Sim)
for(is in 1:Sim) {
  x <- rnorm(N)
  gc() ## sort should not have to pay for gc
  c1[is] <- system.time(for(i in 1:rep)
                        sort(x, method = "shell"))[1]
  c2[is] <- system.time(for(i in 1:rep)
                        sort(x, method = "quick"))[1]
  stopifnot(sort(x, meth = "s") == sort(x, meth = "q"))
}
100 * rbind(ShellSort = c1, QuickSort = c2)
cat("Speedup factor of quick sort():\n")
summary({qq <- c1 / c2; qq[is.finite(qq)]})

## A larger test
x <- rnorm(1e6)
gc()
system.time(x1 <- sort(x, method = "shell"))
gc()
system.time(x2 <- sort(x, method = "quick"))
stopifnot(identical(x1, x2))
```

---

**Special**     *Special Functions of Mathematics*


---

**Description**

Special mathematical functions related to the beta and gamma functions.

**Usage**

```

beta(a, b)
lbeta(a, b)
gamma(x)
lgamma(x)
digamma(x)
trigamma(x)
tetragamma(x)
pentagamma(x)
choose(n, k)
lchoose(n, k)

```

**Arguments**

<code>a, b, x</code>	numeric vectors.
<code>n, k</code>	integer vectors.

**Details**

The functions `beta` and `lbeta` return the beta function and the natural logarithm of the beta function,

$$B(a, b) = \frac{\Gamma(a)\Gamma(b)}{\Gamma(a+b)}.$$

The functions `gamma` and `lgamma` return the gamma function  $\Gamma(x)$  and the natural logarithm of the absolute value of the gamma function.

The functions `digamma`, `trigamma`, `tetragamma` and `pentagamma` return the first, second, third and fourth derivatives of the logarithm of the gamma function.

$$\text{digamma}(x) = \psi(x) = \frac{d}{dx} \ln \Gamma(x) = \frac{\Gamma'(x)}{\Gamma(x)}$$

The functions `choose` and `lchoose` return binomial coefficients and their logarithms.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (for `gamma` and `lgamma`.)

Abramowitz, M. and Stegun, I. A. (1972) *Handbook of Mathematical Functions*. New York: Dover. Chapter 6: Gamma and Related Functions.

## See Also

Arithmetic for simple, `sqrt` for miscellaneous mathematical functions and `Bessel` for the real Bessel functions.

## Examples

```
choose(5, 2)
for (n in 0:10) print(choose(n, k = 0:n))

## gamma has discontinuities are 0, -1, -2, ... use plots
## of points to show this.
curve(gamma(x),-3,4, n=1001, ylim=c(-10,100),
      col="red", lwd=2, main="gamma(x)")
abline(h=0,v=0, lty=3, col="midnightblue")

x <- seq(.1, 4, length = 201); dx <- diff(x)[1]
par(mfrow = c(2, 3))
for (ch in c("", "l","di","tri","tetra","penta")) {
  is.deriv <- nchar(ch) >= 2
  if (is.deriv) dy <- diff(y) / dx
  nm <- paste(ch, "gamma", sep = "")
  y <- get(nm)(x)
  plot(x, y, type = "l", main = nm, col = "red")
  abline(h = 0, col = "lightgray")
  if (is.deriv) lines(x[-1], dy, col = "blue", lty = 2)
}
```



---

**splinefun**     *Interpolating Splines*

---

**Description**

Perform cubic spline interpolation of given data points, returning either a list of points obtained by the interpolation or a function performing the interpolation.

**Usage**

```
splinefun(x, y = NULL, method = "fmm")

spline(x, y = NULL, n = 3*length(x), method = "fmm",
       xmin = min(x), xmax = max(x))
```

**Arguments**

<b>x,y</b>	vectors giving the coordinates of the points to be interpolated. Alternatively a single plotting structure can be specified: see <code>xy.coords</code> .
<b>method</b>	specifies the type of spline to be used. Possible values are "fmm", "natural" and "periodic".
<b>n</b>	interpolation takes place at <b>n</b> equally spaced points spanning the interval <code>[xmin, xmax]</code> .
<b>xmin</b>	left-hand endpoint of the interpolation interval.
<b>xmax</b>	right-hand endpoint of the interpolation interval.

**Details**

If `method = "fmm"`, the spline used is that of Forsythe, Malcolm and Moler (an exact cubic is fitted through the four points at each end of the data, and this is used to determine the end conditions). Natural splines are used when `method = "natural"`, and periodic splines when `method = "periodic"`.

These interpolation splines can also be used for extrapolation, that is prediction at points outside the range of **x**. Extrapolation makes little sense for `method = "fmm"`; for natural splines it is linear using the slope of the interpolating curve at the nearest data point.

## Value

**spline** returns a list containing components **x** and **y** which give the ordinates where interpolation took place and the interpolated values.

**splinefun** returns a function which will perform cubic spline interpolation of the given data points. This is often more useful than **spline**.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Forsythe, G. E., Malcolm, M. A. and Moler, C. B. (1977) *Computer Methods for Mathematical Computations*.

## See Also

**approx** and **approxfun** for constant and linear interpolation.

Package **splines**, especially **interpSpline** and **periodicSpline** for interpolation splines. That package also generates spline bases that can be used for regression splines.

**smooth.spline** in package **modreg** for smoothing splines.

## Examples

```
op <- par(mfrow = c(2,1), mgp = c(2,.8,0),
          mar = .1+c(3,3,3,1))
n <- 9
x <- 1:n
y <- rnorm(n)
plot(x, y,
     main = paste("spline[fun](.) through", n, "points"))
lines(spline(x, y))
lines(spline(x, y, n = 201), col = 2)

y <- (x-6)^2
plot(x, y, main = "spline(.) -- 3 methods")
lines(spline(x, y, n = 201), col = 2)
lines(spline(x, y, n = 201, method = "natural"), col = 3)
lines(spline(x, y, n = 201, method = "periodic"), col = 4)
legend(6,25, c("fmm","natural","periodic"), col=2:4, lty=1)

f <- splinefun(x, y)
ls(envir = environment(f))
```

```
splinecoef <- eval(expression(z), envir = environment(f))  
curve(f(x), 1, 10, col = "green", lwd = 1.5)  
points(splinecoef, col = "purple", cex = 2)  
par(op)
```

---

**sum**     *Sum of Vector Elements*

---

**Description**

`sum` returns the sum of all the values present in its arguments. If `na.rm` is `FALSE` an NA value in any of the arguments will cause a value of NA to be returned, otherwise NA values are ignored.

**Usage**

```
sum(..., na.rm=FALSE)
```

**Arguments**

<code>...</code>	numeric or complex vectors.
<code>na.rm</code>	logical. Should missing values be removed?

**Value**

The sum. If all of `...` are of type integer, then so is the sum, and in that case the result will be NA (with a warning) if integer overflow occurs.

NB: the sum of an empty set is zero, by definition.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

---

**svd**     *Singular Value Decomposition of a Matrix*


---

**Description**

Compute the singular-value decomposition of a rectangular matrix.

**Usage**

```
svd(x, nu = min(n, p), nv = min(n, p), LINPACK = FALSE)
La.svd(x, nu = min(n, p), nv = min(n, p),
       method = c("dgesdd", "dgesvd"))
```

**Arguments**

<b>x</b>	a matrix whose SVD decomposition is to be computed.
<b>nu</b>	the number of left singular vectors to be computed. This must be one of 0, <code>nrow(x)</code> and <code>ncol(x)</code> , except for <code>method = "dgesdd"</code> .
<b>nv</b>	the number of right singular vectors to be computed. This must be one of 0 and <code>ncol(x)</code> .
<b>LINPACK</b>	logical. Should LINPACK be used (for compatibility with R < 1.7.0)?
<b>method</b>	The LAPACK routine to use in the real case.

**Details**

The singular value decomposition plays an important role in many statistical techniques. `svd` and `La.svd` provide two slightly different interfaces. The main functions used are the LAPACK routines DGESDD and ZGESVD; `svd(LINPACK=TRUE)` provides an interface to the LINPACK routine DSVDC, purely for backwards compatibility.

`La.svd` provides an interface to both the LAPACK routines DGESVD and DGESDD. The latter is usually substantially faster if singular vectors are required. Most benefit is seen with an optimized BLAS system. Using `method="dgesdd"` requires IEEE 754 arithmetic. Should this not be supported on your platform, `method="dgesvd"` is used, with a warning.

Computing the singular vectors is the slow part for large matrices.

## Value

The SVD decomposition of the matrix as computed by LINPACK,

$$\mathbf{X} = \mathbf{U} \mathbf{D} \mathbf{V}',$$

where  $\mathbf{U}$  and  $\mathbf{V}$  are orthogonal,  $\mathbf{V}'$  means  $\mathbf{V}$  *transposed*, and  $\mathbf{D}$  is a diagonal matrix with the singular values  $D_{ii}$ . Equivalently,  $\mathbf{D} = \mathbf{U}' \mathbf{X} \mathbf{V}$ , which is verified in the examples, below.

The returned value is a list with components

<code>d</code>	a vector containing the singular values of <code>x</code> .
<code>u</code>	a matrix whose columns contain the left singular vectors of <code>x</code> , present if <code>nu &gt; 0</code>
<code>v</code>	a matrix whose columns contain the right singular vectors of <code>x</code> , present if <code>nv &gt; 0</code> .

For `La.svd` the return value replaces `v` by `vt`, the (conjugated if complex) transpose of `v`.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Dongarra, J. J., Bunch, J. R., Moler, C. B. and Stewart, G. W. (1978) *LINPACK Users Guide*. Philadelphia: SIAM Publications.
- Anderson, E., et al. (1999) *LAPACK Users' Guide*. Third Edition. SIAM. ISBN 0-89871-447-8.

## See Also

`eigen`, `qr`.  
`capabilities` to test for IEEE 754 arithmetic.

## Examples

```
hilbert <- function(n) {
  i <- 1:n; 1 / outer(i - 1, i, "+")
}
str(X <- hilbert(9)[,1:6])
str(s <- svd(X))
D <- diag(s$d)
s$u %*% D %*% t(s$v) # X = U D V'
t(s$u) %*% X %*% s$v # D = U' X V
```

---

**tabulate**      *Tabulation for Vectors*

---

**Description**

**tabulate** takes the integer valued vector **bin** and counts the number of times each integer occurs in it. **tabulate** is used as the basis of the **table** function.

**Usage**

```
tabulate(bin, nbins = max(1, bin))
```

**Arguments**

<b>bin</b>	a vector of integers, or a factor.
<b>nbins</b>	the number of bins to be used.

**Details**

If **bin** is a factor, its internal integer representation is tabulated. If the elements of **bin** are not integers, they are rounded to the nearest integer. Elements outside the range  $1, \dots, \mathbf{nbins}$  are (silently) ignored in the tabulation.

**See Also**

**factor**, **table**.

**Examples**

```
tabulate(c(2,3,5))
tabulate(c(2,3,3,5), nb = 10)
tabulate(c(-2,0,2,3,3,5), nb = 3)
tabulate(factor(letters[1:10]))
```

---

**Trig**     *Trigonometric Functions*

---

**Description**

These functions give the obvious trigonometric functions. They respectively compute the cosine, sine, tangent, arc-cosine, arc-sine, arc-tangent, and the two-argument arc-tangent.

**Usage**

```
cos(x)
sin(x)
tan(x)
acos(x)
asin(x)
atan(x)
atan2(y, x)
```

**Arguments**

`x`, `y`                      numeric vector

**Details**

The arc-tangent of two arguments `atan2(y,x)` returns the angle between the x-axis and the vector from the origin to  $(x, y)$ , i.e., for positive arguments `atan2(y,x) == atan(y/x)`.

Angles are in radians, not degrees (i.e., a right angle is  $\pi/2$ ).

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.



## Chapter 3

# Base package — distributions and random numbers

---

**bandwidth**      *Bandwidth Selectors for Kernel Density Estimation*


---

**Description**

Bandwidth selectors for gaussian windows in **density**.

**Usage**

```
bw.nrd0(x)
bw.nrd(x)
bw.ucv(x, nb = 1000, lower, upper)
bw.bcv(x, nb = 1000, lower, upper)
bw.SJ(x, nb = 1000, lower, upper, method = c("ste", "dpi"))
```

**Arguments**

<b>x</b>	A data vector.
<b>nb</b>	number of bins to use.
<b>lower, upper</b>	Range over which to minimize. The default is almost always satisfactory.
<b>method</b>	Either <b>"ste"</b> ("solve-the-equation") or <b>"dpi"</b> ("direct plug-in").

**Details**

**bw.nrd0** implements a rule-of-thumb for choosing the bandwidth of a Gaussian kernel density estimator. It defaults to 0.9 times the minimum of the standard deviation and the interquartile range divided by 1.34 times the sample size to the negative one-fifth power (= Silverman's "rule of thumb", Silverman (1986, page 48, eqn (3.31)) *unless* the quartiles coincide when a positive result will be guaranteed.

**bw.nrd** is the more common variation given by Scott (1992), using factor 1.06.

**bw.ucv** and **bw.bcv** implement unbiased and biased cross-validation respectively.

**bw.SJ** implements the methods of Sheather & Jones (1991) to select the bandwidth using pilot estimation of derivatives.

**Value**

A bandwidth on a scale suitable for the **bw** argument of **density**.

## References

- Scott, D. W. (1992) *Multivariate Density Estimation: Theory, Practice, and Visualization*. Wiley.
- Sheather, S. J. and Jones, M. C. (1991) A reliable data-based bandwidth selection method for kernel density estimation. *Journal of the Royal Statistical Society series B*, **53**, 683–690.
- Silverman, B. W. (1986) *Density Estimation*. London: Chapman and Hall.
- Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. Springer.

## See Also

`density`.

`bandwidth.nrd`, `ucv`, `bcv` and `width.SJ` in package **MASS**, which are all scaled to the `width` argument of `density` and so give answers four times as large.

## Examples

```
data(precip)
plot(density(precip, n = 1000))
rug(precip)
lines(density(precip, bw="nrd"), col = 2)
lines(density(precip, bw="ucv"), col = 3)
lines(density(precip, bw="bcv"), col = 4)
lines(density(precip, bw="SJ-ste"), col = 5)
lines(density(precip, bw="SJ-dpi"), col = 6)
legend(55, 0.035,
      legend = c("nrd0", "nrd", "ucv", "bcv", "SJ-ste", "SJ-dpi"),
      col = 1:6, lty = 1)
```

---

## Beta      *The Beta Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the Beta distribution with parameters **shape1** and **shape2** (and optional non-centrality parameter **ncp**).

### Usage

```
dbeta(x, shape1, shape2, ncp=0, log=FALSE)
pbeta(q, shape1, shape2, ncp=0, lower.tail=TRUE, log.p=FALSE)
qbeta(p, shape1, shape2, lower.tail=TRUE, log.p=FALSE)
rbeta(n, shape1, shape2)
```

### Arguments

<b>x</b> , <b>q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>shape1</b> , <b>shape2</b>	positive parameters of the Beta distribution.
<b>ncp</b>	non-centrality parameter.
<b>log</b> , <b>log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

The Beta distribution with parameters **shape1** =  $a$  and **shape2** =  $b$  has density

$$f(x) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} x^a (1-x)^b$$

for  $a > 0$ ,  $b > 0$  and  $0 \leq x \leq 1$  where the boundary values at  $x = 0$  or  $x = 1$  are defined as by continuity (as limits).

### Value

**dbeta** gives the density, **pbeta** the distribution function, **qbeta** the quantile function, and **rbeta** generates random deviates.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`beta` for the Beta function, and `dgamma` for the Gamma distribution.

## Examples

```
x <- seq(0, 1, length=21)
dbeta(x, 1, 1)
pbeta(x, 1, 1)
```

---

## Binomial      *The Binomial Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the binomial distribution with parameters **size** and **prob**.

### Usage

```
dbinom(x, size, prob, log = FALSE)
pbinom(q, size, prob, lower.tail = TRUE, log.p = FALSE)
qbinom(p, size, prob, lower.tail = TRUE, log.p = FALSE)
rbinom(n, size, prob)
```

### Arguments

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>size</b>	number of trials.
<b>prob</b>	probability of success on each trial.
<b>log, log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

The binomial distribution with **size** =  $n$  and **prob** =  $p$  has density

$$p(x) = \binom{n}{x} p^x (1-p)^{n-x}$$

for  $x = 0, \dots, n$ .

If an element of **x** is not integer, the result of **dbinom** is zero, with a warning.  $p(x)$  is computed using Loader's algorithm, see the reference below.

The quantile is defined as the smallest value  $x$  such that  $F(x) \geq p$ , where  $F$  is the distribution function.

## Value

`dbinom` gives the density, `pbinom` gives the distribution function, `qbinom` gives the quantile function and `rbinom` generates random deviates.

If `size` is not an integer, `NaN` is returned.

## References

Catherine Loader (2000). *Fast and Accurate Computation of Binomial Probabilities*

## See Also

`dnbinom` for the negative binomial, and `dpois` for the Poisson distribution.

## Examples

```
# Compute P(45 < X < 55) for X Binomial(100,0.5)
sum(dbinom(46:54, 100, 0.5))

## Using "log = TRUE" for an extended range :
n <- 2000
k <- seq(0, n, by = 20)
plot(k, dbinom(k, n, pi/10, log=TRUE), type='l',
     ylab="log density",
     main="dbinom(*, log=TRUE) is better than log(dbinom(*))")
lines(k, log(dbinom(k, n, pi/10)), col='red', lwd=2)
## extreme points are omitted since dbinom gives 0.
mtext("dbinom(k, log=TRUE)", adj=0)
mtext("extended range", adj=0, line = -1, font=4)
mtext("log(dbinom(k))", col="red", adj=1)
```

---

<b>birthday</b>	<i>Probability of coincidences</i>
-----------------	------------------------------------

---

## Description

Computes approximate answers to a generalised “birthday paradox” problem. `pbirthday` computes the probability of a coincidence and `qbirthday` computes the number of observations needed to have a specified probability of coincidence.

## Usage

```
qbirthday(prob = 0.5, classes = 365, coincident = 2)
pbirthday(n, classes = 365, coincident = 2)
```

## Arguments

<code>classes</code>	How many distinct categories the people could fall into
<code>prob</code>	The desired probability of coincidence
<code>n</code>	The number of people
<code>coincident</code>	The number of people to fall in the same category

## Details

The birthday paradox is that a very small number of people, 23, suffices to have a 50-50 chance that two of them have the same birthday. This function generalises the calculation to probabilities other than 0.5, numbers of coincident events other than 2, and numbers of classes other than 365.

This formula is approximate, as the example below shows. For `coincident=2` the exact computation is straightforward and may be preferable.

## Value

<code>qbirthday</code>	Number of people needed for a probability <code>prob</code> that <code>k</code> of them have the same one out of <code>classes</code> equiprobable labels.
<code>pbirthday</code>	Probability of the specified coincidence



## References

Diaconis P, Mosteller F., “Methods for studying coincidences”. JASA 84:853-861

## Examples

```
## the standard version
qbirthday()
## same 4-digit PIN number
qbirthday(classes=10^4)
## 0.9 probability of three coincident birthdays
qbirthday(coincident=3,prob=0.9)
## Chance of 4 coincident birthdays in 150 people
pbirthday(150,coincident=4)
## Accuracy compared to exact calculation
x1<- sapply(10:100, pbirthday)
x2<-1-sapply(10:100, function(n)
               prod((365:(365-n+1))/rep(365,n)))
par(mfrow=c(2,2))
plot(x1,x2,xlab="approximate",ylab="exact")
abline(0,1)
plot(x1,x1-x2,xlab="approximate",ylab="error")
abline(h=0)
plot(x1,x2,log="xy",xlab="approximate",ylab="exact")
abline(0,1)
plot(1-x1,1-x2,log="xy",xlab="approximate",ylab="exact")
abline(0,1)
```

---

## Cauchy      *The Cauchy Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the Cauchy distribution with location parameter **location** and scale parameter **scale**.

### Usage

```
dcauchy(x, location=0, scale=1, log=FALSE)
pcauchy(q, location=0, scale=1, lower.tail=TRUE, log.p=FALSE)
qcauchy(p, location=0, scale=1, lower.tail=TRUE, log.p=FALSE)
rcauchy(n, location=0, scale=1)
```

### Arguments

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>location, scale</b>	location and scale parameters.
<b>log, log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

If **location** or **scale** are not specified, they assume the default values of 0 and 1 respectively.

The Cauchy distribution with location  $l$  and scale  $s$  has density

$$f(x) = \frac{1}{\pi s} \left( 1 + \left( \frac{x-l}{s} \right)^2 \right)^{-1}$$

for all  $x$ .

**Value**

`dcauchy`, `pcauchy`, and `qcauchy` are respectively the density, distribution function and quantile function of the Cauchy distribution. `rcauchy` generates random deviates from the Cauchy.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`dt` for the t distribution which generalizes `dcauchy(*, l = 0, s = 1)`.

**Examples**

```
dcauchy(-1:4)
```

---

**Chisquare**      *The (non-central) Chi-Squared Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the chi-squared ( $\chi^2$ ) distribution with **df** degrees of freedom and optional non-centrality parameter **ncp**.

**Usage**

```
dchisq(x, df, ncp=0, log = FALSE)
pchisq(q, df, ncp=0, lower.tail = TRUE, log.p = FALSE)
qchisq(p, df, ncp=0, lower.tail = TRUE, log.p = FALSE)
rchisq(n, df, ncp=0)
```

**Arguments**

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>df</b>	degrees of freedom (non-negative, but can be non-integer).
<b>ncp</b>	non-centrality parameter (non-negative). Note that <b>ncp</b> values larger than about 1417 are not allowed currently for <b>pchisq</b> and <b>qchisq</b> .
<b>log, log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The chi-squared distribution with **df**=  $n$  degrees of freedom has density

$$f_n(x) = \frac{1}{2^{n/2}\Gamma(n/2)} x^{n/2-1} e^{-x/2}$$

for  $x > 0$ . The mean and variance are  $n$  and  $2n$ .

The non-central chi-squared distribution with  $\mathbf{df} = n$  degrees of freedom and non-centrality parameter  $\mathbf{ncp} = \lambda$  has density

$$f(x) = e^{-\lambda/2} \sum_{r=0}^{\infty} \frac{(\lambda/2)^r}{r!} f_{n+2r}(x)$$

for  $x \geq 0$ . For integer  $n$ , this is the distribution of the sum of squares of  $n$  normals each with variance one,  $\lambda$  being the sum of squares of the normal means. Note that the degrees of freedom  $\mathbf{df} = n$ , can be non-integer, and for non-centrality  $\lambda > 0$ , even  $n = 0$ ; see the reference, chapter 29.

## Value

`dchisq` gives the density, `pchisq` gives the distribution function, `qchisq` gives the quantile function, and `rchisq` generates random deviates.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Johnson, Kotz and Balakrishnan (1995). *Continuous Univariate Distributions*, Vol 2; Wiley NY;

## See Also

`dgamma` for the Gamma distribution which generalizes the chi-squared one.

## Examples

```
dchisq(1, df=1:3)
pchisq(1, df= 3)
pchisq(1, df= 3, ncp = 0:4) # includes the above

x <- 1:10
## Chi-squared(df = 2) is a special exponential
## distribution
all.equal(dchisq(x, df=2), dexp(x, 1/2))
all.equal(pchisq(x, df=2), pexp(x, 1/2))

## non-central RNG -- df=0 is ok for ncp > 0: Z0 has point
## mass at 0!
Z0 <- rchisq(100, df = 0, ncp = 2.)
```

```

stem(Z0)

## visual testing do P-P plots for 1000 points at various
## degrees of freedom
L <- 1.2; n <- 1000; pp <- ppoints(n)
op <- par(mfrow = c(3,3), mar = c(3,3,1,1)+.1,
          mgp = c(1.5,.6,0), oma = c(0,0,3,0))
for(df in 2^(4*rnorm(9))) {
  plot(pp, sort(pchisq(rr <- rchisq(n,df=df, ncp=L),
    df=df, ncp=L)), ylab="pchisq(rchisq(.),.)", pch=".")
  mtext(paste("df = ",formatC(df, digits = 4)),
    line= -2, adj=0.05)
  abline(0,1,col=2)
}
mtext(expression("P-P plots : Noncentral " *
  chi^2 * "(n=1000, df=X, ncp= 1.2)"),
  cex = 1.5, font = 2, outer=TRUE)
par(op)

```

---

**density**     *Kernel Density Estimation*


---

**Description**

The function **density** computes kernel density estimates with the given kernel and bandwidth.

**Usage**

```
density(x, bw = "nrd0", adjust = 1,
        kernel = c("gaussian", "epanechnikov",
                    "rectangular", "triangular",
                    "biweight", "cosine", "optcosine"),
        window = kernel, width,
        give.Rkern = FALSE,
        n = 512, from, to, cut = 3, na.rm = FALSE)
```

**Arguments**

- |                       |  |
|-----------------------|--|
| <b>x</b>              | the data from which the estimate is to be computed.  |
| <b>bw</b>             | <p>the smoothing bandwidth to be used. The kernels are scaled such that this is the standard deviation of the smoothing kernel. (Note this differs from the reference books cited below, and from S-PLUS.)</p> <p><b>bw</b> can also be a character string giving a rule to choose the bandwidth. See <b>bw.nrd</b>.</p> <p>The specified (or computed) value of <b>bw</b> is multiplied by <b>adjust</b>.</p> |
| <b>adjust</b>         | the bandwidth used is actually <b>adjust*bw</b> . This makes it easy to specify values like “half the default” bandwidth.  |
| <b>kernel, window</b> | <p>a character string giving the smoothing kernel to be used. This must be one of "gaussian", "rectangular", "triangular", "epanechnikov", "biweight", "cosine" or "optcosine", with default "gaussian", and may be abbreviated to a unique prefix (single letter).</p> <p>"cosine" is smoother than "optcosine", which is the usual “cosine” kernel in the literature and almost</p>                          |

	MSE-efficient. However, "cosine" is the version used by S.
<code>width</code>	this exists for compatibility with S; if given, and <code>bw</code> is not, will set <code>bw</code> to <code>width</code> if this is a character string, or to a kernel-dependent multiple of <code>width</code> if this is numeric.
<code>give.Rkern</code>	logical; if true, <i>no</i> density is estimated, and the "canonical bandwidth" of the chosen <code>kernel</code> is returned instead.
<code>n</code>	the number of equally spaced points at which the density is to be estimated. When <code>n</code> > 512, it is rounded up to the next power of 2 for efficiency reasons ( <code>fft</code> ).
<code>from,to</code>	the left and right-most points of the grid at which the density is to be estimated.
<code>cut</code>	by default, the values of <code>left</code> and <code>right</code> are <code>cut</code> bandwidths beyond the extremes of the data. This allows the estimated density to drop to approximately zero at the extremes.
<code>na.rm</code>	logical; if TRUE, missing values are removed from <code>x</code> . If FALSE any missing values cause an error.

## Details

The algorithm used in `density` disperses the mass of the empirical distribution function over a regular grid of at least 512 points and then uses the fast Fourier transform to convolve this approximation with a discretized version of the kernel and then uses linear approximation to evaluate the density at the specified points.

The statistical properties of a kernel are determined by  $\sigma_K^2 = \int t^2 K(t) dt$  which is always = 1 for our kernels (and hence the bandwidth `bw` is the standard deviation of the kernel) and  $R(K) = \int K^2(t) dt$ .

MSE-equivalent bandwidths (for different kernels) are proportional to  $\sigma_K R(K)$  which is scale invariant and for our kernels equal to  $R(K)$ . This value is returned when `give.Rkern` = TRUE. See the examples for using exact equivalent bandwidths.

Infinite values in `x` are assumed to correspond to a point mass at `+/-Inf` and the density estimate is of the sub-density on `(-Inf, +Inf)`.

## Value

If `give.Rkern` is true, the number  $R(K)$ , otherwise an object with class "density" whose underlying structure is a list containing the following components.



<code>x</code>	the <code>n</code> coordinates of the points where the density is estimated.
<code>y</code>	the estimated density values.
<code>bw</code>	the bandwidth used.
<code>N</code>	the sample size after elimination of missing values.
<code>call</code>	the call which produced the result.
<code>data.name</code>	the deparsed name of the <code>x</code> argument.
<code>has.na</code>	logical, for compatibility (always <code>FALSE</code> ).

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole (for S version).
- Scott, D. W. (1992) *Multivariate Density Estimation. Theory, Practice and Visualization*. New York: Wiley.
- Sheather, S. J. and Jones M. C. (1991) A reliable data-based bandwidth selection method for kernel density estimation. *J. Roy. Statist. Soc. B*, 683–690.
- Silverman, B. W. (1986) *Density Estimation*. London: Chapman and Hall.
- Venables, W. N. and Ripley, B. D. (1999) *Modern Applied Statistics with S-PLUS*. New York: Springer.

## See Also

`bw.nrd`, `plot.density`, `hist`.

## Examples

```
plot(density(c(-20,rep(0,98),20)), xlim = c(-4,4)) # IQR = 0

# The Old Faithful geyser data
data(faithful)
d <- density(faithful$eruptions, bw = "sj")
d
plot(d)

plot(d, type = "n")
polygon(d, col = "wheat")

## Missing values:
```

```

x <- xx <- faithful$eruptions
x[i.out <- sample(length(x), 10)] <- NA
doR <- density(x, bw = 0.15, na.rm = TRUE)
lines(doR, col = "blue")
points(xx[i.out], rep(0.01, 10))

(kernels <- eval(formals(density)$kernel))

## show the kernels in the R parametrization
plot(density(0, bw = 1), xlab = "",
      main="R's density() kernels with bw = 1")
for(i in 2:length(kernels))
  lines(density(0, bw = 1, kern = kernels[i]), col = i)
legend(1.5,.4, legend = kernels, col = seq(kernels),
      lty = 1, cex = .8, y.int = 1)

## show the kernels in the S parametrization
plot(density(0, from=-1.2, to=1.2, width=2,
      kern="gaussian"), type="l", ylim = c(0, 1), xlab="",
      main="R's density() kernels with width = 1")
for(i in 2:length(kernels))
  lines(density(0, width=2, kern = kernels[i]), col = i)
legend(0.6, 1.0, legend = kernels, col = seq(kernels),
      lty = 1)

(RKs <- cbind(sapply(kernels, function(k)
      density(kern = k, give.Rkern = TRUE))))
100*round(RKs["epanechnikov",]/RKs, 4) ## Efficiencies

if(interactive()) {
data(precip)
bw <- bw.SJ(precip) ## sensible automatic choice
plot(density(precip, bw = bw, n = 2^13),
      main = "same sd bandwidths, 7 different kernels")
for(i in 2:length(kernels))
  lines(density(precip, bw = bw, kern = kernels[i],
      n = 2^13), col = i)

## Bandwidth Adjustment for "Exactly Equivalent Kernels"
h.f <- sapply(kernels, function(k)
      density(kern = k, give.Rkern = TRUE))
(h.f <- (h.f["gaussian"] / h.f)^.2)
## -> 1, 1.01, .995, 1.007,... close to 1 => adjustment

```

```
## barely visible..

plot(density(precip, bw = bw, n = 2^13),
     main = "equivalent bandwidths, 7 different kernels")
for(i in 2:length(kernels))
  lines(density(precip, bw = bw, adjust = h.f[i],
               kern = kernels[i], n = 2^13), col = i)
legend(55, 0.035, legend = kernels, col = seq(kernels),
      lty = 1)
}
```

---

## Exponential      *The Exponential Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the exponential distribution with rate **rate** (i.e., mean  $1/\text{rate}$ ).

### Usage

```
dexp(x, rate = 1, log = FALSE)
pexp(q, rate = 1, lower.tail = TRUE, log.p = FALSE)
qexp(p, rate = 1, lower.tail = TRUE, log.p = FALSE)
rexp(n, rate = 1)
```

### Arguments

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<b>rate</b>	vector of rates.
<b>log, log.p</b>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

If **rate** is not specified, it assumes the default value of 1.

The exponential distribution with rate  $\lambda$  has density

$$f(x) = \lambda e^{-\lambda x}$$

for  $x \geq 0$ .

### Value

**dexp** gives the density, **pexp** gives the distribution function, **qexp** gives the quantile function, and **rexp** generates random deviates.

## Note

The cumulative hazard  $H(t) = -\log(1 - F(t))$  is `-pexp(t, r, lower = FALSE, log = TRUE)`.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`exp` for the exponential function, `dgamma` for the gamma distribution and `dweibull` for the Weibull distribution, both of which generalize the exponential.

## Examples

```
dexp(1) - exp(-1) # 0
```

---

**F**Dist      *The F Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the F distribution with **df1** and **df2** degrees of freedom (and optional non-centrality parameter **ncp**).

**Usage**

```
df(x, df1, df2, log = FALSE)
pf(q, df1, df2, ncp=0, lower.tail = TRUE, log.p = FALSE)
qf(p, df1, df2,          lower.tail = TRUE, log.p = FALSE)
rf(n, df1, df2)
```

**Arguments**

<b>x</b> , <b>q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>df1</b> , <b>df2</b>	degrees of freedom.
<b>ncp</b>	non-centrality parameter.
<b>log</b> , <b>log.p</b>	logical; if TRUE, probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if TRUE (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The F distribution with **df1** =  $n_1$  and **df2** =  $n_2$  degrees of freedom has density

$$f(x) = \frac{\Gamma(n_1/2 + n_2/2)}{\Gamma(n_1/2)\Gamma(n_2/2)} \left(\frac{n_1}{n_2}\right)^{n_1/2} x^{n_1/2-1} \left(1 + \frac{n_1 x}{n_2}\right)^{-(n_1+n_2)/2}$$

for  $x > 0$ .

It is the distribution of the ratio of the mean squares of  $n_1$  and  $n_2$  independent standard normals, and hence of the ratio of two independent chi-squared variates each divided by its degrees of freedom. Since the ratio of a normal and the root mean-square of  $m$  independent normals

has a Student's  $t_m$  distribution, the square of a  $t_m$  variate has a F distribution on 1 and  $m$  degrees of freedom.

The non-central F distribution is again the ratio of mean squares of independent normals of unit variance, but those in the numerator are allowed to have non-zero means and `ncp` is the sum of squares of the means. See `Chisquare` for further details on non-central distributions.

## Value

`df` gives the density, `pf` gives the distribution function `qf` gives the quantile function, and `rf` generates random deviates.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`dchisq` for chi-squared and `dt` for Student's t distributions.

## Examples

```
## the density of the square of a t_m is 2*dt(x, m)/(2*x)
# check this is the same as the density of F_{1,m}
x <- seq(0.001, 5, len=100)
all.equal(df(x^2, 1, 5), dt(x, 5)/x)

## Identity: qf(2*p - 1, 1, df)) == qt(p, df)^2 for p
## >= 1/2
p <- seq(1/2, .99, length=50); df <- 10
rel.err <- function(x,y)
  ifelse(x==y, 0, abs(x-y)/mean(abs(c(x,y))))
quantile(rel.err(qf(2*p - 1, df1=1, df2=df), qt(p, df)^2),
  .90) # ~ = 7e-9
```

---

**GammaDist**     *The Gamma Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the Gamma distribution with parameters **shape** and **scale**.

**Usage**

```

dgamma(x,shape,rate=1,scale=1/rate,log=FALSE)
pgamma(q,shape,rate=1,scale=1/rate,lower.tail=TRUE,
       log.p=FALSE)
qgamma(p,shape,rate=1,scale=1/rate,lower.tail=TRUE,log.
       p=FALSE)
rgamma(n,shape,rate=1,scale=1/rate)

```

**Arguments**

<b>x</b> , <b>q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<b>rate</b>	an alternative way to specify the scale.
<b>shape</b> , <b>scale</b>	shape and scale parameters.
<b>log</b> , <b>log.p</b>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<b>lower.tail</b>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

If **scale** is omitted, it assumes the default value of 1.

The Gamma distribution with parameters **shape** =  $\alpha$  and **scale** =  $\sigma$  has density

$$f(x) = \frac{1}{\sigma^\alpha \Gamma(\alpha)} x^{\alpha-1} e^{-x/\sigma}$$

for  $x > 0$ ,  $\alpha > 0$  and  $\sigma > 0$ . The mean and variance are  $E(X) = \alpha\sigma$  and  $Var(X) = \alpha\sigma^2$ .



## Value

**dgamma** gives the density, **pgamma** gives the distribution function **qgamma** gives the quantile function, and **rgamma** generates random deviates.

## Note

The S parametrization is via **shape** and **rate**: S has no **scale** parameter. Prior to 1.4.0 R only had **scale**.

The cumulative hazard  $H(t) = -\log(1 - F(t))$  is **-pgamma(t, ..., lower = FALSE, log = TRUE)**.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**gamma** for the Gamma function, **dbeta** for the Beta distribution and **dchisq** for the chi-squared distribution which is a special case of the Gamma distribution.

## Examples

```
-log(dgamma(1:4, shape=1))
p <- (1:9)/10
pgamma(qgamma(p, shape=2), shape=2)
1 - 1/exp(qgamma(p, shape=1))
```

---

**Geometric**     *The Geometric Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the geometric distribution with parameter **prob**.

**Usage**

```
dgeom(x, prob, log = FALSE)
pgeom(q, prob, lower.tail = TRUE, log.p = FALSE)
qgeom(p, prob, lower.tail = TRUE, log.p = FALSE)
rgeom(n, prob)
```

**Arguments**

<b>x, q</b>	vector of quantiles representing the number of failures in a sequence of Bernoulli trials before success occurs.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>prob</b>	probability of success in each trial.
<b>log, log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The geometric distribution with **prob** =  $p$  has density

$$p(x) = p(1 - p)^x$$

for  $x = 0, 1, 2, \dots$

If an element of **x** is not integer, the result of **pgeom** is zero, with a warning.

The quantile is defined as the smallest value  $x$  such that  $F(x) \geq p$ , where  $F$  is the distribution function.

**Value**

`dgeom` gives the density, `pgeom` gives the distribution function, `qgeom` gives the quantile function, and `rgeom` generates random deviates.

**See Also**

`dnbinom` for the negative binomial which generalizes the geometric distribution.

**Examples**

```
qgeom((1:9)/10, prob = .2)
Ni <- rgeom(20, prob = 1/4); table(factor(Ni, 0:max(Ni)))
```

---

## Hypergeometric      *The Hypergeometric Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the hypergeometric distribution.

### Usage

```
dhyper(x, m, n, k, log = FALSE)
phyper(q, m, n, k, lower.tail = TRUE, log.p = FALSE)
qhyper(p, m, n, k, lower.tail = TRUE, log.p = FALSE)
rhyper(nn, m, n, k)
```

### Arguments

<code>x, q</code>	vector of quantiles representing the number of white balls drawn without replacement from an urn which contains both black and white balls.
<code>m</code>	the number of white balls in the urn.
<code>n</code>	the number of black balls in the urn.
<code>k</code>	the number of balls drawn from the urn.
<code>p</code>	probability, it must be between 0 and 1.
<code>nn</code>	number of observations. If <code>length(nn) &gt; 1</code> , the length is taken to be the number required.
<code>log, log.p</code>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<code>lower.tail</code>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

The hypergeometric distribution is used for sampling *without* replacement. The density of this distribution with parameters `m`, `n` and `k` (named  $Np$ ,  $N - Np$ , and  $n$ , respectively in the reference below) is given by

$$p(x) = \binom{m}{x} \binom{n}{k-x} / \binom{m+n}{k}$$

for  $x = 0, \dots, k$ .

## Value

**dhyper** gives the density, **phyper** gives the distribution function, **qhyper** gives the quantile function, and **rhyper** generates random deviates.

## References

Johnson, N. L., Kotz, S., and Kemp, A. W. (1992) *Univariate Discrete Distributions*, Second Edition. New York: Wiley.

## Examples

```
m <- 10; n <- 7; k <- 8
x <- 0:(k+1)
rbind(phyper(x, m, n, k), dhyper(x, m, n, k))
# FALSE
all(phyper(x, m, n, k) == cumsum(dhyper(x, m, n, k)))
## but error is very small:
signif(phyper(x, m, n, k) - cumsum(dhyper(x, m, n, k)),
      dig=3)
```

---

**Logistic**     *The Logistic Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the logistic distribution with parameters **location** and **scale**.

**Usage**

```
dlogis(x,location=0,scale=1,log=FALSE)
plogis(q,location=0,scale=1,lower.tail=TRUE,log.p=FALSE)
qlogis(p,location=0,scale=1,lower.tail=TRUE,log.p=FALSE)
rlogis(n,location=0,scale=1)
```

**Arguments**

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>location, scale</b>	location and scale parameters.
<b>log, log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

If **location** or **scale** are omitted, they assume the default values of 0 and 1 respectively.

The Logistic distribution with **location** =  $\mu$  and **scale** =  $\sigma$  has distribution function

$$F(x) = \frac{1}{1 + e^{-(x-\mu)/\sigma}}$$

and density

$$f(x) = \frac{1}{\sigma} \frac{e^{(x-\mu)/\sigma}}{(1 + e^{(x-\mu)/\sigma})^2}$$

It is a long-tailed distribution with mean  $\mu$  and variance  $\pi^2/3\sigma^2$ .

**Value**

**dlogis** gives the density, **plogis** gives the distribution function, **qlogis** gives the quantile function, and **rlogis** generates random deviates.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**Examples**

```
var(rlogis(4000, 0, s = 5)) # approximately (+/- 3)
pi^2/3 * 5^2
```

---

**Lognormal**      *The Log Normal Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the log normal distribution whose logarithm has mean equal to `meanlog` and standard deviation equal to `sdlog`.

**Usage**

```
dlnorm(x, meanlog=0, sdlog=1, log=false)
plnorm(q, meanlog=0, sdlog=1, lower.tail=true, log.p=false)
qlnorm(p, meanlog=0, sdlog=1, lower.tail=true, log.p=false)
rlnorm(n, meanlog=0, sdlog=1)
```

**Arguments**

<code>x, q</code>	vector of quantiles.
<code>p</code>	vector of probabilities.
<code>n</code>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<code>meanlog, sdlog</code>	mean and standard deviation of the distribution on the log scale with default values of 0 and 1 respectively.
<code>log, log.p</code>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<code>lower.tail</code>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The log normal distribution has density

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma x} e^{-(\log(x)-\mu)^2/2\sigma^2}$$

where  $\mu$  and  $\sigma$  are the mean and standard deviation of the logarithm. The mean is  $E(X) = \exp(\mu + 1/2\sigma^2)$ , and the variance  $Var(X) = \exp(2\mu + \sigma^2)(\exp(\sigma^2) - 1)$  and hence the coefficient of variation is  $\sqrt{\exp(\sigma^2) - 1}$  which is approximately  $\sigma$  when that is small (e.g.,  $\sigma < 1/2$ ).



## Value

**dlnorm** gives the density, **plnorm** gives the distribution function, **qlnorm** gives the quantile function, and **rlnorm** generates random deviates.

## Note

The cumulative hazard  $H(t) = -\log(1-F(t))$  is **-plnorm(t, r, lower = FALSE, log = TRUE)**.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**dnorm** for the normal distribution.

## Examples

```
dlnorm(1) == dnorm(0)
```

---

**Multinomial**      *The Multinomial Distribution*


---

**Description**

Generate multinomially distributed random number vectors and compute multinomial “density” probabilities.

**Usage**

```
rmultinom(n, size, prob)
dmultinom(x, size = NULL, prob, log = FALSE)
```

**Arguments**

<b>x</b>	vector of length $K$ of integers in $0:\text{size}$ .
<b>n</b>	number of random vectors to draw.
<b>size</b>	integer, say $N$ , specifying the total number of objects that are put into $K$ boxes in the typical multinomial experiment. For <code>dmultinom</code> , it defaults to <code>sum(x)</code> .
<b>prob</b>	numeric non-negative vector of length $K$ , specifying the probability for the $K$ classes; is internally normalized to sum 1.
<b>log</b>	logical; if <code>TRUE</code> , log probabilities are computed.

**Details**

If  $\mathbf{x}$  is a  $K$ -component vector, `dmultinom(x, prob)` is the probability

$$P(X_1 = x_1, \dots, X_K = x_k) = C \times \prod_{j=1}^K \pi_j^{x_j}$$

where  $C$  is the “multinomial coefficient”  $C = N!/(x_1! \cdots x_K!)$  and  $N = \sum_{j=1}^K x_j$ .

By definition, each component  $X_j$  is binomially distributed as `Bin(size, prob[j])` for  $j = 1, \dots, K$ .

The `rmultinom()` algorithm draws binomials from  $\text{Bin}(n_j, P_j)$  sequentially, where  $n_1 = N$  ( $N := \text{size}$ ),  $P_1 = \pi_1$  ( $\pi$  is `prob` scaled to sum 1), and for  $j \geq 2$ , recursively  $n_j = N - \sum_{k=1}^{j-1} n_k$  and  $P_j = \pi_j / (1 - \sum_{k=1}^{j-1} \pi_k)$ .

**Value**

For `rmultinom()`, an integer  $K \times n$  matrix where each column is a random vector generated according to the desired multinomial law, and hence summing to `size`. Whereas the *transposed* result would seem more natural at first, the returned matrix is more efficient because of columnwise storage.

**Note**

`dmultinom` is currently *not vectorized* at all and has no C interface (API); this may be amended in the future.

**See Also**

`rbinom` which is a special case conceptually.

**Examples**

```
rmultinom(10, size = 12, prob=c(0.1,0.2,0.8))

pr <-c(1,3,6,10) # normalization unnecessary for generation
rmultinom(10, 20, prob = pr)

## all possible outcomes of Multinom(N = 3, K = 3)
x <- t(as.matrix(expand.grid(0:3, 0:3)));
x <- x[, colsums(x) <= 3]
x <- rbind(x, 3:3 - colsums(x));
dimnames(x) <- list(letters[1:3], null)
X
round(apply(X, 2, function(x)
            dmultinom(x, prob = c(1,2,5))), 3)
```

---

## NegBinomial      *The Negative Binomial Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the negative binomial distribution with parameters **size** and **prob**.

### Usage

```
dnbinom(x, size, prob, mu, log=FALSE)
pnbinom(q, size, prob, mu, lower.tail=TRUE, log.p=FALSE)
qnbinom(p, size, prob, mu, lower.tail=TRUE, log.p=FALSE)
rnbinom(n, size, prob, mu)
```

### Arguments

<b>x</b>	vector of (non-negative integer) quantiles.
<b>q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<b>size</b>	target for number of successful trials, or dispersion parameter (the shape parameter of the gamma mixing distribution).
<b>prob</b>	probability of success in each trial.
<b>mu</b>	alternative parametrization via mean: see Details
<b>log, log.p</b>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<b>lower.tail</b>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

The negative binomial distribution with **size** =  $n$  and **prob** =  $p$  has density

$$p(x) = \frac{\Gamma(x+n)}{\Gamma(n)x!} p^n (1-p)^x$$

for  $x = 0, 1, 2, \dots$

This represents the number of failures which occur in a sequence of Bernoulli trials before a target number of successes is reached.

A negative binomial distribution can arise as a mixture of Poisson distributions with mean distributed as a gamma (`pgamma`) distribution with scale parameter  $(1 - \text{prob})/\text{prob}$  and shape parameter `size`. (This definition allows non-integer values of `size`.) In this model  $\text{prob} = \text{scale}/(1+\text{scale})$ , and the mean is  $\text{size} * (1 - \text{prob})/\text{prob}$

The alternative parametrization (often used in ecology) is by the *mean* `mu`, and *size*, the *dispersion parameter*, where  $\text{prob} = \text{size}/(\text{size}+\text{mu})$ . In this parametrization the variance is  $\text{mu} + \text{mu}^2/\text{size}$ .

If an element of `x` is not integer, the result of `dnbinom` is zero, with a warning.

The quantile is defined as the smallest value  $x$  such that  $F(x) \geq p$ , where  $F$  is the distribution function.

## Value

`dnbinom` gives the density, `pnbinom` gives the distribution function, `qpnbinom` gives the quantile function, and `rnbinom` generates random deviates.

## See Also

`dbinom` for the binomial, `dpois` for the Poisson and `dgeom` for the geometric distribution, which is a special case of the negative binomial.

## Examples

```
x <- 0:11
dnbinom(x, size = 1, prob = 1/2) * 2^(1 + x) # == 1
# theoretically integer
126 / dnbinom(0:8, size = 2, prob = 1/2)

## Cumulative ('p') = Sum of discrete prob.s ('d');
## Relative error :
summary(1 - cumsum(dnbinom(x, size = 2, prob = 1/2)) /
        pnbinom(x, size = 2, prob = 1/2))

x <- 0:15
size <- (1:20)/4
persp(x,size,
      dnb <- outer(x, size, function(x,s) dnbinom(x,s,pr=0.4)),
      xlab = "x", ylab = "s", zlab="density", theta = 150)
title(
  tit <- "negative binomial density(x,s,pr=0.4) vs. x,s"
```

)

```
image (x,size, log10(dnb), main= paste("log [",tit,""])\ncontour(x,size, log10(dnb),add=TRUE)
```

```
## Alternative parametrization\nx1 <- rnbinom(500, mu = 4, size = 1)\nx2 <- rnbinom(500, mu = 4, size = 10)\nx3 <- rnbinom(500, mu = 4, size = 100)\nh1 <- hist(x1, breaks = 20, plot = FALSE)\nh2 <- hist(x2, breaks = h1$breaks, plot = FALSE)\nh3 <- hist(x3, breaks = h1$breaks, plot = FALSE)\nbarplot(rbind(h1$counts, h2$counts, h3$counts),\n        beside = TRUE, col = c("red","blue","cyan"),\n        names.arg = round(h1$breaks[-length(h1$breaks)]))
```

---

## Normal      *The Normal Distribution*

---

### Description

Density, distribution function, quantile function and random generation for the normal distribution with mean equal to `mean` and standard deviation equal to `sd`.

### Usage

```
dnorm(x, mean=0, sd=1, log = FALSE)
pnorm(q, mean=0, sd=1, lower.tail = TRUE, log.p = FALSE)
qnorm(p, mean=0, sd=1, lower.tail = TRUE, log.p = FALSE)
rnorm(n, mean=0, sd=1)
```

### Arguments

<code>x, q</code>	vector of quantiles.
<code>p</code>	vector of probabilities.
<code>n</code>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<code>mean</code>	vector of means.
<code>sd</code>	vector of standard deviations.
<code>log, log.p</code>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<code>lower.tail</code>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

### Details

If `mean` or `sd` are not specified they assume the default values of 0 and 1, respectively.

The normal distribution has density

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/2\sigma^2}$$

where  $\mu$  is the mean of the distribution and  $\sigma$  the standard deviation.

`qnorm` is based on Wichura's algorithm AS 241 which provides precise results up to about 16 digits.

## Value

**dnorm** gives the density, **pnorm** gives the distribution function, **qnorm** gives the quantile function, and **rnorm** generates random deviates.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Wichura, M. J. (1988) Algorithm AS 241: The Percentage Points of the Normal Distribution. *Applied Statistics*, **37**, 477–484.

## See Also

**runif** and **.Random.seed** about random number generation, and **dlnorm** for the *Lognormal* distribution.

## Examples

```
dnorm(0) == 1/ sqrt(2*pi)
dnorm(1) == exp(-1/2)/ sqrt(2*pi)
dnorm(1) == 1/ sqrt(2*pi*exp(1))

## Using "log = TRUE" for an extended range :
par(mfrow=c(2,1))
plot(function(x)dnorm(x, log=TRUE), -60, 50,
     main = "log { Normal density }")
curve(log(dnorm(x)), add=TRUE, col="red",lwd=2)
mtext("dnorm(x, log=TRUE)", adj=0);
mtext("log(dnorm(x))", col="red", adj=1)

plot(function(x)pnorm(x, log=TRUE), -50, 10,
     main = "log { Normal Cumulative }")
curve(log(pnorm(x)), add=TRUE, col="red",lwd=2)
mtext("pnorm(x, log=TRUE)", adj=0);
mtext("log(pnorm(x))", col="red", adj=1)

## if you want the so-called 'error function'
erf <- function(x) 2 * pnorm(x * sqrt(2)) - 1
## and the so-called 'complementary error function'
erfc <- function(x) 2 * pnorm(x * sqrt(2), lower=FALSE)
```



---

**Poisson**     *The Poisson Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the Poisson distribution with parameter `lambda`.

**Usage**

```
dpois(x, lambda, log = FALSE)
ppois(q, lambda, lower.tail = TRUE, log.p = FALSE)
qpois(p, lambda, lower.tail = TRUE, log.p = FALSE)
rpois(n, lambda)
```

**Arguments**

<code>x</code>	vector of (non-negative integer) quantiles.
<code>q</code>	vector of quantiles.
<code>p</code>	vector of probabilities.
<code>n</code>	number of random values to return.
<code>lambda</code>	vector of positive means.
<code>log, log.p</code>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<code>lower.tail</code>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The Poisson distribution has density

$$p(x) = \frac{\lambda^x e^{-\lambda}}{x!}$$

for  $x = 0, 1, 2, \dots$ . The mean and variance are  $E(X) = Var(X) = \lambda$ .

If an element of `x` is not integer, the result of `dpois` is zero, with a warning.  $p(x)$  is computed using Loader's algorithm, see the reference in `dbinom`.

The quantile is left continuous: `qgeom(q, prob)` is the largest integer  $x$  such that  $P(X \leq x) < q$ .

Setting `lower.tail = FALSE` allows to get much more precise results when the default, `lower.tail = TRUE` would return 1, see the example below.

## Value

`dpois` gives the (log) density, `ppois` gives the (log) distribution function, `qpois` gives the quantile function, and `rpois` generates random deviates.

## See Also

`dbinom` for the binomial and `dnbinom` for the negative binomial distribution.

## Examples

```
-log(dpois(0:7, lambda=1) * gamma(1+ 0:7)) # == 1
Ni <- rpois(50, lam= 4); table(factor(Ni, 0:max(Ni)))

# becomes 0 (cancellation)
1 - ppois(10*(15:25), lambda=100)
# no cancellation
ppois(10*(15:25), lambda=100, lower=FALSE)

par(mfrow = c(2, 1))
x <- seq(-0.01, 5, 0.01)
plot(x, ppois(x, 1), type="s", ylab="F(x)",
      main="Poisson(1) CDF")
plot(x, pbinom(x, 100, 0.01), type="s", ylab="F(x)",
      main="Binomial(100, 0.01) CDF")
```

---

**r2dtable**      *Random 2-way Tables with Given Marginals*

---

**Description**

Generate random 2-way tables with given marginals using Patefield's algorithm.

**Usage**

```
r2dtable(n, r, c)
```

**Arguments**

- |          |  |
|----------|--|
| <b>n</b> | a non-negative numeric giving the number of tables to be drawn.  |
| <b>r</b> | a non-negative vector of length at least 2 giving the row totals, to be coerced to <b>integer</b> . Must sum to the same as <b>c</b> . |
| <b>c</b> | a non-negative vector of length at least 2 giving the column totals, to be coerced to <b>integer</b> .                                 |

**Value**

A list of length **n** containing the generated tables as its components.

**References**

Patefield, W. M. (1981) Algorithm AS159. An efficient method of generating  $r \times c$  tables with given row and column totals. *Applied Statistics* **30**, 91–97.

**Examples**

```
## Fisher's Tea Drinker data.
TeaTasting <-
matrix(c(3, 1, 1, 3),
      nr = 2,
      dimnames = list(Guess = c("Milk", "Tea"),
                       Truth = c("Milk", "Tea")))
## Simulate permutation test for independence based on the
## maximum Pearson residuals (rather than their sum).
rowTotals <- rowSums(TeaTasting)
```

```
colTotals <- colSums(TeaTasting)
nOfCases <- sum(rowTotals)
expected <- outer(rowTotals, colTotals, "*") / nOfCases
maxSqResid <-
  function(x) max((x - expected) ^ 2 / expected)
simMaxSqResid <-
  sapply(r2dtable(1000, rowTotals, colTotals), maxSqResid)
sum(simMaxSqResid >= maxSqResid(TeaTasting)) / 1000
## Fisher's exact test gives p = 0.4857 ...
```

---

**Random**      *Random Number Generation*

---

**Description**

`.Random.seed` is an integer vector, containing the random number generator (RNG) **state** for random number generation in R. It can be saved and restored, but should not be altered by the user.

`RNGkind` is a more friendly interface to query or set the kind of RNG in use.

`RNGversion` can be used to set the random generators as they were in an earlier R version (for reproducibility).

`set.seed` is the recommended way to specify seeds.

**Usage**

```
.Random.seed <- c(rng.kind, n1, n2, ...)  
save.seed <- .Random.seed
```

```
RNGkind(kind = NULL, normal.kind = NULL)  
RNGversion(vstr)  
set.seed(seed, kind = NULL)
```

**Arguments**

<code>kind</code>	character or NULL. If <code>kind</code> is a character string, set R's RNG to the kind desired. If it is NULL, return the currently used RNG. Use "default" to return to the R default.
<code>normal.kind</code>	character string or NULL. If it is a character string, set the method of Normal generation. Use "default" to return to the R default.
<code>seed</code>	a single value, interpreted as an integer.
<code>vstr</code>	a character string containing a version number, e.g., "1.6.2"
<code>rng.kind</code>	integer code in 0:k for the above <code>kind</code> .
<code>n1, n2, ...</code>	integers. See the details for how many are required (which depends on <code>rng.kind</code> ).

## Details

The currently available RNG kinds are given below. `kind` is partially matched to this list. The default is "Mersenne-Twister".

"Wichmann-Hill" The seed, `.Random.seed[-1] == r[1:3]` is an integer vector of length 3, where each `r[i]` is in  $1:(p[i] - 1)$ , where `p` is the length 3 vector of primes, `p = (30269, 30307, 30323)`. The Wichmann-Hill generator has a cycle length of  $6.9536 \times 10^{12}$  ( $= \text{prod}(p-1)/4$ , see *Applied Statistics* (1984) **33**, 123 which corrects the original article).

"Marsaglia-Multicarry": A *multiply-with-carry* RNG is used, as recommended by George Marsaglia in his post to the mailing list 'sci.stat.math'. It has a period of more than  $2^{60}$  and has passed all tests (according to Marsaglia). The seed is two integers (all values allowed).

"Super-Duper": Marsaglia's famous Super-Duper from the 70's. This is the original version which does *not* pass the MTUPLE test of the Diehard battery. It has a period of  $\approx 4.6 \times 10^{18}$  for most initial seeds. The seed is two integers (all values allowed for the first seed: the second must be odd).

We use the implementation by Reeds et al. (1982–84).

The two seeds are the Tausworthe and congruence long integers, respectively. A one-to-one mapping to S's `.Random.seed[1:12]` is possible but we will not publish one, not least as this generator is **not** exactly the same as that in recent versions of S-PLUS.

"Mersenne-Twister": From Matsumoto and Nishimura (1998). A twisted GFSR with period  $2^{19937} - 1$  and equidistribution in 623 consecutive dimensions (over the whole period). The "seed" is a 624-dimensional set of 32-bit integers plus a current position in that set.

"Knuth-TAOCP": From Knuth (1997). A GFSR using lagged Fibonacci sequences with subtraction. That is, the recurrence used is

$$X_j = (X_{j-100} - X_{j-37}) \bmod 2^{30}$$

and the "seed" is the set of the 100 last numbers (actually recorded as 101 numbers, the last being a cyclic shift of the buffer). The period is around  $2^{129}$ .

"Knuth-TAOCP-2002": The 2002 version which not backwards compatible with the earlier version: the initialization of the GFSR from the seed was altered. R did not allow you to choose consecutive seeds, the reported 'weakness', and already scrambled the seeds.

"user-supplied": Use a user-supplied generator. See `Random.user` for details.

`normal.kind` can be "Kinderman-Ramage", "Buggy Kinderman-Ramage", "Ahrens-Dieter", "Box-Muller", "Inversion" (the default), or "user-supplied". (For inversion, see the reference in `qnorm`.) The Kinderman-Ramage generator used in versions prior to 1.7.1 had several approximation errors and should only be used for reproduction of older results.

`set.seed` uses its single integer argument to set as many seeds as are required. It is intended as a simple way to get quite different seeds by specifying small integer arguments, and also as a way to get valid seed sets for the more complicated methods (especially "Mersenne-Twister" and "Knuth-TAOCP").

## Value

`.Random.seed` is an **integer** vector whose first element *codes* the kind of RNG and normal generator. The lowest two decimal digits are in  $0:(k-1)$  where  $k$  is the number of available RNGs. The hundreds represent the type of normal generator (starting at 0).

In the underlying C, `.Random.seed[-1]` is **unsigned**; therefore in R `.Random.seed[-1]` can be negative.

`RNGkind` returns a two-element character vector of the RNG and normal kinds in use *before* the call, invisibly if either argument is not `NULL`. `RNGversion` returns the same information.

`set.seed` returns `NULL`, invisibly.

## Note

Initially, there is no seed; a new one is created from the current time when one is required. Hence, different sessions will give different simulation results, by default.

`.Random.seed` saves the seed set for the uniform random-number generator, at least for the system generators. It does not necessarily save the state of other generators, and in particular does not save the state of the Box-Muller normal generator. If you want to reproduce work later, call `set.seed` rather than `set .Random.seed`.

As from R 1.8.0, `.Random.seed` is only looked for in the user's workspace.

## Author(s)

of RNGkind: Martin Maechler. Current implementation, B. D. Ripley

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (`set.seed`, storing in `.Random.seed`.)

Wichmann, B. A. and Hill, I. D. (1982) *Algorithm AS 183: An Efficient and Portable Pseudo-random Number Generator*, Applied Statistics, **31**, 188–190; Remarks: **34**, 198 and **35**, 89.

De Matteis, A. and Pagnutti, S. (1993) *Long-range Correlation Analysis of the Wichmann-Hill Random Number Generator*, Statist. Comput., **3**, 67–70.

Marsaglia, G. (1997) *A random number generator for C*. Discussion paper, posting on Usenet newsgroup `sci.stat.math` on September 29, 1997.

Reeds, J., Hubert, S. and Abrahams, M. (1982–4) C implementation of SuperDuper, University of California at Berkeley. (Personal communication from Jim Reeds to Ross Ihaka.)

Marsaglia, G. and Zaman, A. (1994) Some portable very-long-period random number generators. *Computers in Physics*, **8**, 117–121.

Matsumoto, M. and Nishimura, T. (1998) Mersenne Twister: A 623-dimensionally equidistributed uniform pseudo-random number generator, *ACM Transactions on Modeling and Computer Simulation*, **8**, 3–30. Source code at <http://www.math.keio.ac.jp/~matumoto/emt.html>.

Knuth, D. E. (1997) *The Art of Computer Programming*. Volume 2, third edition.

Source code at <http://www-cs-faculty.stanford.edu/~knuth/taocp.html>.

Knuth, D. E. (2002) *The Art of Computer Programming*. Volume 2, third edition, ninth printing.

Kinderman, A. J. and Ramage, J. G. (1976) Computer generation of normal random variables. *Journal of the American Statistical Association* **71**, 893–896.

Ahrens, J.H. and Dieter, U. (1973) Extensions of Forsythe’s method for random sampling from the normal distribution. *Mathematics of Computation* **27**, 927–937.



Box, G.E.P. and Muller, M.E. (1958) A note on the generation of normal random deviates. *Annals of Mathematical Statistics* **29**, 610–611.

## See Also

runif, rnorm, ....

## Examples

```
runif(1); .Random.seed; runif(1); .Random.seed
## If there is no seed, a "random" new one is created:
rm(.Random.seed); runif(1); .Random.seed

RNGkind("Wich") # (partial string matching on 'kind')

## This shows how 'runif(.)' works for Wichmann-Hill, using
## only R functions:
p.WH <- c(30269, 30307, 30323)
a.WH <- c( 171,  172,  170)
next.WHseed <- function(i.seed = .Random.seed[-1])
  { (a.WH * i.seed) %% p.WH }
my.runif1 <- function(i.seed = .Random.seed)
  { ns <- next.WHseed(i.seed[-1]); sum(ns / p.WH) %% 1 }
rs <- .Random.seed
(WHs <- next.WHseed(rs[-1]))
u <- runif(1)
stopifnot(
  next.WHseed(rs[-1]) == .Random.seed[-1],
  all.equal(u, my.runif1(rs))
)

## ----
.Random.seed
ok <- RNGkind()
RNGkind("Super") # matches "Super-Duper"
RNGkind()
.Random.seed # new, corresponding to Super-Duper

## Reset:
RNGkind(ok[1])
```

---

**Random.user**      *User-supplied Random Number Generation*

---

## Description

Function `RNGkind` allows user-coded uniform and normal random number generators to be supplied. The details are given here.

## Details

A user-specified uniform RNG is called from entry points in dynamically-loaded compiled code. The user must supply the entry point `user_unif_rand`, which takes no arguments and returns a *pointer to a double*. The example below will show the general pattern.

Optionally, the user can supply the entry point `user_unif_init`, which is called with an `unsigned int` argument when `RNGkind` (or `set.seed`) is called, and is intended to be used to initialize the user's RNG code. The argument is intended to be used to set the “seeds”; it is the `seed` argument to `set.seed` or an essentially random seed if `RNGkind` is called.

If only these functions are supplied, no information about the generator's state is recorded in `.Random.seed`. Optionally, functions `user_unif_nseed` and `user_unif_seedloc` can be supplied which are called with no arguments and should return pointers to the number of “seeds” and to an integer array of “seeds”. Calls to `GetRNGstate` and `PutRNGstate` will then copy this array to and from `.Random.seed`.

A user-specified normal RNG is specified by a single entry point `user_norm_rand`, which takes no arguments and returns a *pointer to a double*.

## Warning

As with all compiled code, mis-specifying these functions can crash R. Do include the ‘`R_ext/Random.h`’ header file for type checking.

## Examples

```
## Marsaglia's congruential PRNG
#include <R_ext/Random.h>

static Int32 seed;
static double res;
static int nseed = 1;
```

```

double * user_unif_rand()
{
    seed = 69069 * seed + 1;
    res = seed * 2.32830643653869e-10;
    return &res;
}

void user_unif_init(Int32 seed_in) { seed = seed_in; }
int * user_unif_nseed() { return &nseed; }
int * user_unif_seedloc() { return (int *) &seed; }

/* ratio-of-uniforms for normal */
#include <math.h>
static double x;

double * user_norm_rand()
{
    double u, v, z;
    do {
        u = unif_rand();
        v = 0.857764 * (2. * unif_rand() - 1);
        x = v/u; z = 0.25 * x * x;
        if (z < 1. - u) break;
        if (z > 0.259/u + 0.35) continue;
    } while (z > -log(u));
    return &x;
}

## Use under Unix:
R SHLIB urand.c
R
> dyn.load("urand.so")
> RNGkind("user")
> runif(10)
> .Random.seed
> RNGkind(, "user")
> rnorm(10)
> RNGkind()
[1] "user-supplied" "user-supplied"

```

---

**sample**     *Random Samples and Permutations*

---

**Description**

**sample** takes a sample of the specified size from the elements of **x** using either with or without replacement.

**Usage**

```
sample(x, size, replace = FALSE, prob = NULL)
```

**Arguments**

<b>x</b>	Either a (numeric, complex, character or logical) vector of more than one element from which to choose, or a positive integer.
<b>size</b>	non-negative integer giving the number of items to choose.
<b>replace</b>	Should sampling be with replacement?
<b>prob</b>	A vector of probability weights for obtaining the elements of the vector being sampled.

**Details**

If **x** has length 1, sampling takes place from **1:x**. *Note* that this convenience feature may lead to undesired behaviour when **x** is of varying length **sample(x)**. See the **resample()** example below.

By default **size** is equal to **length(x)** so that **sample(x)** generates a random permutation of the elements of **x** (or **1:x**).

The optional **prob** argument can be used to give a vector of weights for obtaining the elements of the vector being sampled. They need not sum to one, but they should be nonnegative and not all zero. If **replace** is false, these probabilities are applied sequentially, that is the probability of choosing the next item is proportional to the probabilities amongst the remaining items. The number of nonzero weights must be at least **size** in this case.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## Examples

```
x <- 1:12
# a random permutation
sample(x)
# bootstrap sampling -- only if length(x) > 1 !
sample(x,replace=TRUE)

# 100 Bernoulli trials
sample(c(0,1), 100, replace = TRUE)

## More careful bootstrapping -- Consider this when using
## sample() programmatically (i.e., in your function or
## simulation)!

# sample()'s surprise -- example
x <- 1:10
  sample(x[x > 8]) # length 2
  sample(x[x > 9]) # oops -- length 10!
try(sample(x[x > 10])) # error!

## This is safer:
resample <- function(x, size, ...)
  if(length(x) <= 1) {
    if(!missing(size) && size == 0) x[FALSE] else x
  } else sample(x, size, ...)

resample(x[x > 8]) # length 2
resample(x[x > 9]) # length 1
resample(x[x > 10]) # length 0
```

---

**SignRank**      *Distribution of the Wilcoxon Signed Rank Statistic*


---

**Description**

Density, distribution function, quantile function and random generation for the distribution of the Wilcoxon Signed Rank statistic obtained from a sample with size **n**.

**Usage**

```
dsignrank(x, n, log = FALSE)
psignrank(q, n, lower.tail = TRUE, log.p = FALSE)
qsignrank(p, n, lower.tail = TRUE, log.p = FALSE)
rsignrank(nn, n)
```

**Arguments**

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>nn</b>	number of observations. If <code>length(nn) &gt; 1</code> , the length is taken to be the number required.
<b>n</b>	numbers of observations in the sample. Must be positive integers less than 50.
<b>log, log.p</b>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<b>lower.tail</b>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

This distribution is obtained as follows. Let **x** be a sample of size **n** from a continuous distribution symmetric about the origin. Then the Wilcoxon signed rank statistic is the sum of the ranks of the absolute values **x[i]** for which **x[i]** is positive. This statistic takes values between 0 and  $n(n+1)/2$ , and its mean and variance are  $n(n+1)/4$  and  $n(n+1)(2n+1)/24$ , respectively.

**Value**

**dsignrank** gives the density, **psignrank** gives the distribution function, **qsignrank** gives the quantile function, and **rsignrank** generates random deviates.

**Author(s)**

Kurt Hornik

**See Also**

`dwilcox` etc, for the *two-sample* Wilcoxon rank sum statistic.

**Examples**

```
par(mfrow=c(2,2))
for(n in c(4:5,10,40)) {
  x <- seq(0, n*(n+1)/2, length=501)
  plot(x, dsignrank(x,n=n), type='l',
       main=paste("dsignrank(x,n=",n,""))
}
```

---

**TDist**     *The Student t Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the *t* distribution with **df** degrees of freedom (and optional noncentrality parameter **ncp**).

**Usage**

```
dt(x, df, ncp=0, log = FALSE)
pt(q, df, ncp=0, lower.tail = TRUE, log.p = FALSE)
qt(p, df,          lower.tail = TRUE, log.p = FALSE)
rt(n, df)
```

**Arguments**

<b>x</b> , <b>q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <b>length(n) &gt; 1</b> , the length is taken to be the number required.
<b>df</b>	degrees of freedom ( $> 0$ , maybe non-integer).
<b>ncp</b>	non-centrality parameter $\delta$ ; currently for <b>pt()</b> and <b>dt()</b> , only for <b>ncp</b> $\leq 37.62$ .
<b>log</b> , <b>log.p</b>	logical; if <b>TRUE</b> , probabilities <b>p</b> are given as $\log(p)$ .
<b>lower.tail</b>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The *t* distribution with **df** =  $\nu$  degrees of freedom has density

$$f(x) = \frac{\Gamma((\nu+1)/2)}{\sqrt{\pi\nu}\Gamma(\nu/2)} (1 + x^2/\nu)^{-(\nu+1)/2}$$

for all real  $x$ . It has mean 0 (for  $\nu > 1$ ) and variance  $\frac{\nu}{\nu-2}$  (for  $\nu > 2$ ).

The general *non-central t* with parameters  $(\nu, \delta) = (\mathbf{df}, \mathbf{ncp})$  is defined as the distribution of  $T_\nu(\delta) := \frac{U+\delta}{\chi_\nu/\sqrt{\nu}}$  where  $U$  and  $\chi_\nu$  are independent random variables,  $U \sim \mathcal{N}(0, 1)$ , and  $\chi_\nu^2$  is chi-squared, see **pchisq**.



The most used applications are power calculations for *t*-tests:

Let  $T = \frac{\bar{X} - \mu_0}{S/\sqrt{n}}$  where  $\bar{X}$  is the **mean** and  $S$  the sample standard deviation (**sd**) of  $X_1, X_2, \dots, X_n$  which are i.i.d.  $N(\mu, \sigma^2)$ . Then  $T$  is distributed as non-centrally *t* with **df** =  $n - 1$  degrees of freedom and non-centrality parameter **ncp** =  $(\mu - \mu_0)\sqrt{n}/\sigma$ .

## Value

**dt** gives the density, **pt** gives the distribution function, **qt** gives the quantile function, and **rt** generates random deviates.

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (except non-central versions.)
- Lenth, R. V. (1989). *Algorithm AS 243* — Cumulative distribution function of the non-central *t* distribution, *Appl. Statist.* **38**, 185–189.

## See Also

**df** for the F distribution.

## Examples

```
1 - pt(1:5, df = 1)
qt(.975, df = c(1:10,20,50,100,1000))

tt <- seq(0,10, len=21)
ncp <- seq(0,6, len=31)
ptn <- outer(tt,ncp, function(t,d) pt(t, df = 3, ncp=d))
image(tt,ncp,ptn, zlim=c(0,1),
      main=t.tit <- "Non-central t - Probabilities")
persp(tt,ncp,ptn, zlim=0:1, r=2, phi=20, theta=200,
      main=t.tit, xlab = "t",
      ylab = "noncentrality parameter",
      zlab = "Pr(T <= t)")

op <- par(yaxs="i")
plot(function(x) dt(x, df = 3, ncp = 2), -3, 11,
      ylim = c(0, 0.32), main="Non-central t - Density")
par(op)
```

---

**Tukey**     *The Studentized Range Distribution*


---

**Description**

Functions on the distribution of the studentized range,  $R/s$ , where  $R$  is the range of a standard normal sample of size  $n$  and  $s^2$  is independently distributed as chi-squared with  $df$  degrees of freedom, see `pchisq`.

**Usage**

```
ptukey(q,nmeans,df,nranges=1,lower.tail=TRUE,log.p=FALSE)
qtukey(p,nmeans,df,nranges=1,lower.tail=TRUE,log.p=FALSE)
```

**Arguments**

<code>q</code>	vector of quantiles.
<code>p</code>	vector of probabilities.
<code>nmeans</code>	sample size for range (same for each group).
<code>df</code>	degrees of freedom for $s$ (see below).
<code>nranges</code>	number of <i>groups</i> whose <b>maximum</b> range is considered.
<code>log.p</code>	logical; if <b>TRUE</b> , probabilities $p$ are given as $\log(p)$ .
<code>lower.tail</code>	logical; if <b>TRUE</b> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

If  $n_g = \text{nranges}$  is greater than one,  $R$  is the *maximum* of  $n_g$  groups of `nmeans` observations each.

**Value**

`ptukey` gives the distribution function and `qtukey` its inverse, the quantile function.

**Note**

A Legendre 16-point formula is used for the integral of `ptukey`. The computations are relatively expensive, especially for `qtukey` which uses a simple secant method for finding the inverse of `ptukey`. `qtukey` will be accurate to the 4th decimal place.

## References

Copenhaver, Margaret Diponzio and Holland, Burt S. (1988) Multiple comparisons of simple effects in the two-way analysis of variance with fixed effects. *Journal of Statistical Computation and Simulation*, **30**, 1–15.

## See Also

`pnorm` and `qnorm` for the corresponding functions for the normal distribution.

## Examples

```
if(interactive())
  curve(ptukey(x, nm=6, df=5), from=-1, to=8, n=101)
(ptt <- ptukey(0:10, 2, df= 5))
(qtt <- qtukey(.95, 2, df= 2:11))
## The precision may be not much more than about 8 digits:
summary(abs(.95 - ptukey(qtt,2, df = 2:11)))
```

---

**Uniform**     *The Uniform Distribution*


---

**Description**

These functions provide information about the uniform distribution on the interval from `min` to `max`. `dunif` gives the density, `punif` gives the distribution function `qunif` gives the quantile function and `runif` generates random deviates.

**Usage**

```
dunif(x, min=0, max=1, log = FALSE)
punif(q, min=0, max=1, lower.tail = TRUE, log.p = FALSE)
qunif(p, min=0, max=1, lower.tail = TRUE, log.p = FALSE)
runif(n, min=0, max=1)
```

**Arguments**

<code>x,q</code>	vector of quantiles.
<code>p</code>	vector of probabilities.
<code>n</code>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<code>min,max</code>	lower and upper limits of the distribution.
<code>log, log.p</code>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<code>lower.tail</code>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

If `min` or `max` are not specified they assume the default values of 0 and 1 respectively.

The uniform distribution has density

$$f(x) = \frac{1}{\max - \min}$$

for  $\min \leq x \leq \max$ .

For the case of  $u := \min == \max$ , the limit case of  $X \equiv u$  is assumed.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

.Random.seed about random number generation, `rnorm`, etc for other distributions.

## Examples

```
u <- runif(20)

## The following relations always hold :
punif(u) == u
dunif(u) == 1

var(runif(10000)) # ~ = 1/12 = .08333
```

---

**Weibull**     *The Weibull Distribution*


---

**Description**

Density, distribution function, quantile function and random generation for the Weibull distribution with parameters **shape** and **scale**.

**Usage**

```
dweibull(x, shape, scale=1, log=FALSE)
pweibull(q, shape, scale=1, lower.tail=TRUE, log.p=FALSE)
qweibull(p, shape, scale=1, lower.tail=TRUE, log.p=FALSE)
rweibull(n, shape, scale=1)
```

**Arguments**

<b>x, q</b>	vector of quantiles.
<b>p</b>	vector of probabilities.
<b>n</b>	number of observations. If <code>length(n) &gt; 1</code> , the length is taken to be the number required.
<b>shape, scale</b>	shape and scale parameters, the latter defaulting to 1.
<b>log, log.p</b>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as <code>log(p)</code> .
<b>lower.tail</b>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

The Weibull distribution with **shape** parameter  $a$  and **scale** parameter  $\sigma$  has density given by

$$f(x) = (a/\sigma)(x/\sigma)^{a-1} \exp(-(x/\sigma)^a)$$

for  $x > 0$ . The cumulative is  $F(x) = 1 - \exp(-(x/\sigma)^a)$ , the mean is  $E(X) = \sigma\Gamma(1+1/a)$ , and the  $Var(X) = \sigma^2(\Gamma(1+2/a) - (\Gamma(1+1/a))^2)$ .

**Value**

**dweibull** gives the density, **pweibull** gives the distribution function, **qweibull** gives the quantile function, and **rweibull** generates random deviates.

## Note

The cumulative hazard  $H(t) = -\log(1 - F(t))$  is `-pweibull(t, a, b, lower = FALSE, log = TRUE)` which is just  $H(t) = (t/b)^a$ .

## See Also

`dexp` for the Exponential which is a special case of a Weibull distribution.

## Examples

```
x <- c(0, rlnorm(50))
all.equal(dweibull(x, shape = 1), dexp(x))
all.equal(pweibull(x, shape = 1, scale = pi),
          pexp(x, rate = 1/pi))
## Cumulative hazard H():
all.equal(pweibull(x, 2.5, pi, lower=FALSE, log=TRUE),
          -(x/pi)^2.5, tol=1e-15)
all.equal(qweibull(x/11, shape = 1, scale = pi),
          qexp(x/11, rate = 1/pi))
```

---

**Wilcoxon**      *Distribution of the Wilcoxon Rank Sum Statistic*


---

**Description**

Density, distribution function, quantile function and random generation for the distribution of the Wilcoxon rank sum statistic obtained from samples with size `m` and `n`, respectively.

**Usage**

```
dwilcox(x, m, n, log = FALSE)
pwilcox(q, m, n, lower.tail = TRUE, log.p = FALSE)
qwilcox(p, m, n, lower.tail = TRUE, log.p = FALSE)
rwilcox(nn, m, n)
```

**Arguments**

<code>x, q</code>	vector of quantiles.
<code>p</code>	vector of probabilities.
<code>nn</code>	number of observations. If <code>length(nn) &gt; 1</code> , the length is taken to be the number required.
<code>m, n</code>	numbers of observations in the first and second sample, respectively.
<code>log, log.p</code>	logical; if <code>TRUE</code> , probabilities <code>p</code> are given as $\log(p)$ .
<code>lower.tail</code>	logical; if <code>TRUE</code> (default), probabilities are $P[X \leq x]$ , otherwise, $P[X > x]$ .

**Details**

This distribution is obtained as follows. Let `x` and `y` be two random, independent samples of size `m` and `n`. Then the Wilcoxon rank sum statistic is the number of all pairs (`x[i]`, `y[j]`) for which `y[j]` is not greater than `x[i]`. This statistic takes values between 0 and `m * n`, and its mean and variance are `m * n / 2` and `m * n * (m + n + 1) / 12`, respectively.

**Value**

`dwilcox` gives the density, `pwilcox` gives the distribution function, `qwilcox` gives the quantile function, and `rwilcox` generates random deviates.



**Note**

S-PLUS uses a different (but equivalent) definition of the Wilcoxon statistic.

**Author(s)**

Kurt Hornik

**See Also**

`dsignrank` etc, for the *one-sample* Wilcoxon rank statistic.

**Examples**

```
x <- -1:(4*6 + 1)
fx <- dwilcox(x, 4, 6)
Fx <- pwilcox(x, 4, 6)

layout(rbind(1,2),width=1,heights=c(3,2))
plot(x, fx,type='h', col="violet",
     main= "Prob (density) of Wilcoxon-Statist.(n=6,m=4)")
plot(x, Fx,type="s", col="blue",
     main= "Distribution of Wilcoxon-Statist.(n=6,m=4)")
abline(h=0:1, col="gray20",lty=2)
layout(1) # set back

N <- 200
hist(U <- rwilcox(N, m=4,n=6), breaks=0:25 - 1/2,
     border="red", col="pink", sub = paste("N =",N))
mtext("N * f(x), f() = true \"density\"", side=3,
     col="blue")
lines(x, N*fx, type='h', col='blue', lwd=2)
points(x, N*fx, cex=2)

## Better is a Quantile-Quantile Plot
qqplot(U, qw <- qwilcox((1:N - 1/2)/N, m=4,n=6),
     main = paste("QQ-Plot of empirical and theor. quantiles",
     "Wilcoxon Statistic, (m=4, n=6)",sep="\n"))
n <- as.numeric(names(print(tU <- table(U)))
text(n+.2, n+.5, labels=tU, col="red")
```



## Chapter 4

# Base package — models

---

**add1**     *Add or Drop All Possible Single Terms to a Model*

---

**Description**

Compute all the single terms in the `scope` argument that can be added to or dropped from the model, fit those models and compute a table of the changes in fit.

**Usage**

```
add1(object, scope, ...)

## Default S3 method:
add1(object, scope, scale = 0, test = c("none", "Chisq"),
      k = 2, trace = FALSE, ...)

## S3 method for class 'lm':
add1(object, scope, scale = 0,
      test = c("none", "Chisq", "F"),
      x = NULL, k = 2, ...)

## S3 method for class 'glm':
add1(object, scope, scale = 0,
      test = c("none", "Chisq", "F"),
      x = NULL, k = 2, ...)

drop1(object, scope, ...)

## Default S3 method:
drop1(object, scope, scale = 0, test = c("none", "Chisq"),
      k = 2, trace = FALSE, ...)

## S3 method for class 'lm':
drop1(object, scope, scale = 0, all.cols = TRUE,
      test=c("none", "Chisq", "F"),k = 2, ...)

## S3 method for class 'glm':
drop1(object, scope, scale = 0,
      test = c("none", "Chisq", "F"),
      k = 2, ...)
```

## Arguments

<b>object</b>	a fitted model object.
<b>scope</b>	a formula giving the terms to be considered for adding or dropping.
<b>scale</b>	an estimate of the residual mean square to be used in computing $C_p$ . Ignored if 0 or NULL.
<b>test</b>	should the results include a test statistic relative to the original model? The F test is only appropriate for <b>lm</b> and <b>aov</b> models or perhaps for <b>glm</b> fits with estimated dispersion. The $\chi^2$ test can be an exact test ( <b>lm</b> models with known scale) or a likelihood-ratio test or a test of the reduction in scaled deviance depending on the method.
<b>k</b>	the penalty constant in AIC / $C_p$ .
<b>trace</b>	if TRUE, print out progress reports.
<b>x</b>	a model matrix containing columns for the fitted model and all terms in the upper scope. Useful if <b>add1</b> is to be called repeatedly.
<b>all.cols</b>	(Provided for compatibility with S.) Logical to specify whether all columns of the design matrix should be used. If FALSE then non-estimable columns are dropped, but the result is not usually statistically meaningful.
<b>...</b>	further arguments passed to or from other methods.

## Details

For **drop1** methods, a missing **scope** is taken to be all terms in the model. The hierarchy is respected when considering terms to be added or dropped: all main effects contained in a second-order interaction must remain, and so on.

The methods for **lm** and **glm** are more efficient in that they do not recompute the model matrix and call the **fit** methods directly.

The default output table gives AIC, defined as minus twice log likelihood plus  $2p$  where  $p$  is the rank of the model (the number of effective parameters). This is only defined up to an additive constant (like log-likelihoods). For linear Gaussian models with fixed scale, the constant is chosen to give Mallows'  $C_p$ ,  $RSS/scale + 2p - n$ . Where  $C_p$  is used, the column is labelled as **Cp** rather than **AIC**.

## Value

An object of class "**anova**" summarizing the differences in fit between the models.

## Warning

The model fitting must apply the models to the same dataset. Most methods will attempt to use a subset of the data with no missing values for any of the variables if `na.action=na.omit`, but this may give biased results. Only use these functions with data containing missing values with great care.

## Note

These are not fully equivalent to the functions in S. There is no `keep` argument, and the methods used are not quite so computationally efficient.

Their authors' definitions of Mallows'  $C_p$  and Akaike's AIC are used, not those of the authors of the models chapter of S.

## Author(s)

The design was inspired by the S functions of the same names described in Chambers (1992).

## References

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`step`, `aov`, `lm`, `extractAIC`, `anova`

## Examples

```
example(step) # swiss
add1(lm1, ~ I(Education^2) + .^2)
drop1(lm1, test="F") # So called 'type II' anova
```

```
example(glm)
drop1(glm.D93, test="Chisq")
drop1(glm.D93, test="F")
```

---

**AIC**     *Akaike's An Information Criterion*

---

**Description**

Generic function calculating the Akaike information criterion for one or several fitted model objects for which a log-likelihood value can be obtained, according to the formula  $-2\log\text{-likelihood} + kn_{par}$ , where  $n_{par}$  represents the number of parameters in the fitted model, and  $k = 2$  for the usual AIC, or  $k = \log(n)$  ( $n$  the number of observations) for the so-called BIC or SBC (Schwarz's Bayesian criterion).

**Usage**

```
AIC(object, ..., k = 2)
```

**Arguments**

<b>object</b>	a fitted model object, for which there exists a <b>logLik</b> method to extract the corresponding log-likelihood, or an object inheriting from class <b>logLik</b> .
<b>...</b>	optionally more fitted model objects.
<b>k</b>	numeric, the “penalty” per parameter to be used; the default <b>k = 2</b> is the classical AIC.

**Details**

The default method for **AIC**, **AIC.default()** entirely relies on the existence of a **logLik** method computing the log-likelihood for the given class.

When comparing fitted objects, the smaller the AIC, the better the fit.

**Value**

If just one object is provided, returns a numeric value with the corresponding AIC (or BIC, or ..., depending on **k**); if more than one object is provided, returns a **data.frame** with rows corresponding to the objects and columns representing the number of parameters in the model (**df**) and the AIC.

**Author(s)**

Jose Pinheiro and Douglas Bates

## References

Sakamoto, Y., Ishiguro, M., and Kitagawa G. (1986). *Akaike Information Criterion Statistics*. D. Reidel Publishing Company.

## See Also

`extractAIC`, `logLik`.

## Examples

```
data(swiss)
lm1 <- lm(Fertility ~ . , data = swiss)
AIC(lm1)
stopifnot(all.equal(AIC(lm1),
                    AIC(logLik(lm1))))
## a version of BIC or Schwarz' BC :
AIC(lm1, k = log(nrow(swiss)))
```



---

**alias**      *Find Aliases (Dependencies) in a Model*


---

## Description

Find aliases (linearly dependent terms) in a linear model specified by a formula.

## Usage

```
alias(object, ...)

## S3 method for class 'formula':
alias(object, data, ...)

## S3 method for class 'lm':
alias(object, complete = TRUE, partial = FALSE,
      partial.pattern = FALSE, ...)
```

## Arguments

<b>object</b>	A fitted model object, for example from <code>lm</code> or <code>aov</code> , or a formula for <code>alias.formula</code> .
<b>data</b>	Optionally, a data frame to search for the objects in the formula.
<b>complete</b>	Should information on complete aliasing be included?
<b>partial</b>	Should information on partial aliasing be included?
<b>partial.pattern</b>	Should partial aliasing be presented in a schematic way? If this is done, the results are presented in a more compact way, usually giving the deciles of the coefficients.
<b>...</b>	further arguments passed to or from other methods.

## Details

Although the main method is for class "`lm`", `alias` is most useful for experimental designs and so is used with fits from `aov`. Complete aliasing refers to effects in linear models that cannot be estimated independently of the terms which occur earlier in the model and so have their coefficients omitted from the fit. Partial aliasing refers to effects that can be estimated less precisely because of correlations induced by the design.

**Value**

A list (of class "listof") containing components

<b>Model</b>	Description of the model; usually the formula.
<b>Complete</b>	A matrix with columns corresponding to effects that are linearly dependent on the rows; may be of class "mtable" which has its own <code>print</code> method.
<b>Partial</b>	The correlations of the estimable effects, with a zero diagonal.

**Note**

The aliasing pattern may depend on the contrasts in use: Helmert contrasts are probably most useful.

The defaults are different from those in S.

**Author(s)**

The design was inspired by the S function of the same name described in Chambers *et al.* (1992).

**References**

Chambers, J. M., Freeny, A and Heiberger, R. M. (1992) *Analysis of variance; designed experiments*. Chapter 5 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**Examples**

```
had.VR <- "package:MASS" %in% search()
## The next line is for fractions() which gives neater
## results
if(!had.VR) res <- require(MASS)
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)
npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)
```

```
op <- options(contrasts=c("contr.helmert", "contr.poly"))
npk.aov <- aov(yield ~ block + N*P*K, npk)
alias(npk.aov)
if(!had.VR && res) detach(package:MASS)
options(op) # reset
```

---

**anova**     *Anova Tables*

---

**Description**

Compute analysis of variance (or deviance) tables for one or more fitted model objects.

**Usage**

```
anova(object, ...)
```

**Arguments**

<code>object</code>	an object containing the results returned by a model fitting function (e.g., <code>lm</code> or <code>glm</code> ).
<code>...</code>	additional objects of the same type.

**Value**

This (generic) function returns an object of class **anova**. These objects represent analysis-of-variance and analysis-of-deviance tables. When given a single argument it produces a table which tests whether the model terms are significant.

When given a sequence of objects, **anova** tests the models against one another in the order specified.

The print method for **anova** objects prints tables in a “pretty” form.

**Warning**

The comparison between two or more models will only be valid if they are fitted to the same dataset. This may be a problem if there are missing values and R’s default of `na.action = na.omit` is used.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*, Wadsworth & Brooks/Cole.

**See Also**

`coefficients`, `effects`, `fitted.values`, `residuals`, `summary`, `drop1`, `add1`.

---

<code>anova.glm</code>	<i>Analysis of Deviance for Generalized Linear Model Fits</i>
------------------------	---

---

## Description

Compute an analysis of deviance table for one or more generalized linear model fits.

## Usage

```
## S3 method for class 'glm':  
anova(object, ..., dispersion = NULL, test = NULL)
```

## Arguments

<code>object, ...</code>	objects of class <code>glm</code> , typically the result of a call to <code>glm</code> , or a list of <code>objects</code> for the <code>"glmlist"</code> method.
<code>dispersion</code>	the dispersion parameter for the fitting family. By default it is obtained from <code>glm.obj</code> .
<code>test</code>	a character string, (partially) matching one of <code>"Chisq"</code> , <code>"F"</code> or <code>"Cp"</code> . See <code>stat.anova</code> .

## Details

Specifying a single object gives a sequential analysis of deviance table for that fit. That is, the reductions in the residual deviance as each term of the formula is added in turn are given in as the rows of a table, plus the residual deviances themselves.

If more than one object is specified, the table has a row for the residual degrees of freedom and deviance for each model. For all but the first model, the change in degrees of freedom and deviance is also given. (This only makes statistical sense if the models are nested.) It is conventional to list the models from smallest to largest, but this is up to the user.

The table will optionally contain test statistics (and P values) comparing the reduction in deviance for the row to the residuals. For models with known dispersion (e.g., binomial and Poisson fits) the chi-squared test is most appropriate, and for those with dispersion estimated by moments (e.g., `gaussian`, `quasibinomial` and `quasipoisson` fits) the F test is most appropriate. Mallows'  $C_p$  statistic is the residual deviance plus

twice the estimate of  $\sigma^2$  times the residual degrees of freedom, which is closely related to AIC (and a multiple of it if the dispersion is known).

## Value

An object of class `"anova"` inheriting from class `"data.frame"`.

## Warning

The comparison between two or more models by `anova` or `anova.glm` will only be valid if they are fitted to the same dataset. This may be a problem if there are missing values and R's default of `na.action = na.omit` is used, and `anova.glm` will detect this with an error.

## References

Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`glm`, `anova`.

`drop1` for so-called 'type II' anova where each term is dropped one at a time respecting their hierarchy.

## Examples

```
# Continuing the Example from '?glm':
anova(glm.D93)
anova(glm.D93, test = "Cp")
anova(glm.D93, test = "Chisq")
```

---

**anova.lm**      *ANOVA for Linear Model Fits*

---

**Description**

Compute an analysis of variance table for one or more linear model fits.

**Usage**

```
## S3 method for class 'lm':  
anova(object, ...)  
  
anova.lmlist(object, ..., scale = 0, test = "F")
```

**Arguments**

<b>object, ...</b>	objects of class <code>lm</code> , usually, a result of a call to <code>lm</code> .
<b>test</b>	a character string specifying the test statistic to be used. Can be one of "F", "Chisq" or "Cp", with partial matching allowed, or NULL for no test.
<b>scale</b>	numeric. An estimate of the noise variance $\sigma^2$ . If zero this will be estimated from the largest model considered.

**Details**

Specifying a single object gives a sequential analysis of variance table for that fit. That is, the reductions in the residual sum of squares as each term of the formula is added in turn are given in as the rows of a table, plus the residual sum of squares.

The table will contain F statistics (and P values) comparing the mean square for the row to the residual mean square.

If more than one object is specified, the table has a row for the residual degrees of freedom and sum of squares for each model. For all but the first model, the change in degrees of freedom and sum of squares is also given. (This only makes statistical sense if the models are nested.) It is conventional to list the models from smallest to largest, but this is up to the user.

Optionally the table can include test statistics. Normally the F statistic is most appropriate, which compares the mean square for a row to the residual sum of squares for the largest model considered. If `scale` is

specified chi-squared tests can be used. Mallows'  $C_p$  statistic is the residual sum of squares plus twice the estimate of  $\sigma^2$  times the residual degrees of freedom.

## Value

An object of class `"anova"` inheriting from class `"data.frame"`.

## Warning

The comparison between two or more models will only be valid if they are fitted to the same dataset. This may be a problem if there are missing values and R's default of `na.action = na.omit` is used, and `anova.lm` will detect this with an error.

## Note

Versions of R prior to 1.2.0 based F tests on pairwise comparisons, and this behaviour can still be obtained by a direct call to `anova.lm`.

## References

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

The model fitting function `lm`, `anova`.

`drop1` for so-called 'type II' anova where each term is dropped one at a time respecting their hierarchy.

## Examples

```
## sequential table
data(LifeCycleSavings)
fit <- lm(sr ~ ., data = LifeCycleSavings)
anova(fit)
## same effect via separate models
fit0 <- lm(sr ~ 1, data = LifeCycleSavings)
fit1 <- update(fit0, . ~ . + pop15)
fit2 <- update(fit1, . ~ . + pop75)
fit3 <- update(fit2, . ~ . + dpi)
fit4 <- update(fit3, . ~ . + ddpi)
anova(fit0, fit1, fit2, fit3, fit4, test="F")
anova(fit4, fit2, fit0, test="F") # unconventional order
```



---

**aov**     *Fit an Analysis of Variance Model*

---

**Description**

Fit an analysis of variance model by a call to `lm` for each stratum.

**Usage**

```
aov(formula, data = NULL, projections = FALSE, qr = TRUE,  
     contrasts = NULL, ...)
```

**Arguments**

<b>formula</b>	A formula specifying the model.
<b>data</b>	A data frame in which the variables specified in the formula will be found. If missing, the variables are searched for in the standard way.
<b>projections</b>	Logical flag: should the projections be returned?
<b>qr</b>	Logical flag: should the QR decomposition be returned?
<b>contrasts</b>	A list of contrasts to be used for some of the factors in the formula. These are not used for any <b>Error</b> term, and supplying contrasts for factors only in the <b>Error</b> term will give a warning.
<b>...</b>	Arguments to be passed to <code>lm</code> , such as <code>subset</code> or <code>na.action</code> .

**Details**

This provides a wrapper to `lm` for fitting linear models to balanced or unbalanced experimental designs.

The main difference from `lm` is in the way `print`, `summary` and so on handle the fit: this is expressed in the traditional language of the analysis of variance rather than of linear models.

If the formula contains a single **Error** term, this is used to specify error strata, and appropriate models are fitted within each error stratum.

The formula can specify multiple responses.

Weights can be specified by a `weights` argument, but should not be used with an **Error** term, and are incompletely supported (e.g., not by `model.tables`).

## Value

An object of class `c("aov", "lm")` or for multiple responses of class `c("maov", "aov", "mlm", "lm")` or for multiple error strata of class `"aovlist"`. There are `print` and `summary` methods available for these.

## Author(s)

The design was inspired by the S function of the same name described in Chambers *et al.* (1992).

## References

Chambers, J. M., Freeny, A and Heiberger, R. M. (1992) *Analysis of variance; designed experiments*. Chapter 5 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`lm`, `summary.aov`, `alias`, `proj`, `model.tables`, `TukeyHSD`

## Examples

```
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)
npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)

( npk.aov <- aov(yield ~ block + N*P*K, npk) )
summary(npk.aov)
coefficients(npk.aov)

## as a test, not particularly sensible statistically
op <- options(contrasts = c("contr.helmert",
                           "contr.treatment"))
npk.aovE <- aov(yield ~ N*P*K + Error(block), npk)
npk.aovE
summary(npk.aovE)
options(op) # reset to previous
```

---

**AsIs**     *Inhibit Interpretation/Conversion of Objects*

---

**Description**

Change the class of an object to indicate that it should be treated “as is”.

**Usage**

`I(x)`

**Arguments**

`x`                      an object

**Details**

Function `I` has two main uses.

- In function `data.frame`. Protecting an object by enclosing it in `I()` in a call to `data.frame` inhibits the conversion of character vectors to factors. `I` can also be used to protect objects which are to be added to a data frame, or converted to a data frame *via* `as.data.frame`.

It achieves this by prepending the class `"AsIs"` to the object's classes. Class `"AsIs"` has a few of its own methods, including for `[, as.data.frame, print` and `format`.

- In function `formula`. There it is used to inhibit the interpretation of operators such as `"+"`, `"-"`, `"*"` and `"^"` as formula operators, so they are used as arithmetical operators. This is interpreted as a symbol by `terms.formula`.

**Value**

A copy of the object with class `"AsIs"` prepended to the class(es).

**References**

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

`data.frame`, `formula`

---

**C**     *Sets Contrasts for a Factor*

---

**Description**

Sets the "contrasts" attribute for the factor.

**Usage**

```
C(object, contr, how.many, ...)
```

**Arguments**

<code>object</code>	a factor or ordered factor
<code>contr</code>	which contrasts to use. Can be a matrix with one row for each level of the factor or a suitable function like <code>contr.poly</code> or a character string giving the name of the function
<code>how.many</code>	the number of contrasts to set, by default one less than <code>nlevels(object)</code> .
<code>...</code>	additional arguments for the function <code>contr</code> .

**Details**

For compatibility with S, `contr` can be `treatment`, `helmert`, `sum` or `poly` (without quotes) as shorthand for `contr.treatment` and so on.

**Value**

The factor `object` with the "contrasts" attribute set.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical models*. Chapter 2 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

`contrasts`, `contr.sum`, etc.

## Examples

```
## reset contrasts to defaults
options(contrasts=c("contr.treatment", "contr.poly"))
data(warpbreaks)
attach(warpbreaks)
tens <- C(tension, poly, 1)
attributes(tens)
detach()
## tension SHOULD be an ordered factor, but as it is not we
## can use
aov(breaks ~ wool + tens + tension, data=warpbreaks)

## show the use of ... The default contrast is
## contr.treatment here
summary(lm(breaks ~ wool + C(tension, base=2),
           data=warpbreaks))

data(esoph) # following on from help(esoph)
model3 <-
  glm(cbind(ncases, ncontrols) ~ agegp + C(tobgp, , 1) +
       C(alcgp, , 1), data = esoph, family = binomial())
summary(model3)
```

---

<code>case/variable.names</code>	<i>Case and Variable Names of Fitted Models</i>
----------------------------------	---

---

## Description

Simple utilities returning (non-missing) case names, and (non-eliminated) variable names.

## Usage

```
case.names(object, ...)
## S3 method for class 'lm':
case.names(object, full = FALSE, ...)

variable.names(object, ...)
## S3 method for class 'lm':
variable.names(object, full = FALSE, ...)
```

## Arguments

<code>object</code>	an R object, typically a fitted model.
<code>full</code>	logical; if TRUE, all names (including zero weights, ...) are returned.
<code>...</code>	further arguments passed to or from other methods.

## Value

A character vector.

## See Also

`lm`

## Examples

```
x <- 1:20
y <- x + (x/4 - 2)^3 + rnorm(20, s=3)
names(y) <- paste("0",x,sep=".")
ww <- rep(1,20); ww[13] <- 0
summary(lmxy <- lm(y ~ x + I(x^2)+I(x^3) + I((x-10)^2),
                  weights = ww), cor = TRUE)
variable.names(lmxy)
```

```
variable.names(lmxy, full= TRUE) # includes the last  
case.names(lmxy)  
case.names(lmxy, full = TRUE) # includes the 0-weight case
```



---

**coef**     *Extract Model Coefficients*

---

**Description**

`coef` is a generic function which extracts model coefficients from objects returned by modeling functions. `coefficients` is an *alias* for it.

**Usage**

```
coef(object, ...)  
coefficients(object, ...)
```

**Arguments**

<code>object</code>	an object for which the extraction of model coefficients is meaningful.
<code>...</code>	other arguments.

**Details**

All object classes which are returned by model fitting functions should provide a `coef` method. (Note that the method is `coef` and not `coefficients`.)

**Value**

Coefficients extracted from the model object `object`.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

**See Also**

`fitted.values` and `residuals` for related methods; `glm`, `lm` for model fitting.

**Examples**

```
x <- 1:5; coef(lm(c(1:3,7,6) ~ x))
```

---

**confint**      *Confidence Intervals for Model Parameters*

---

**Description**

Computes confidence intervals for one or more parameters in a fitted model. Base has a method for objects inheriting from class "lm".

**Usage**

```
confint(object, parm, level = 0.95, ...)
```

**Arguments**

<b>object</b>	a fitted model object.
<b>parm</b>	a specification of which parameters are to be given confidence intervals, either a vector of numbers or a vector of names. If missing, all parameters are considered.
<b>level</b>	the confidence level required.
<b>...</b>	additional argument(s) for methods

**Details**

**confint** is a generic function with no default method. For objects of class "lm" the direct formulae based on t values are used.

Package **MASS** contains methods for "glm" and "nls" objects.

**Value**

A matrix (or vector) with columns giving lower and upper confidence limits for each parameter. These will be labelled as (1-level)/2 and 1 - (1-level)/2 in % (by default 2.5% and 97.5%).

**See Also**

`confint.nls`

**Examples**

```
data(mtcars)
fit <- lm(100/mpg ~ disp + hp + wt + am, data=mtcars)
confint(fit)
confint(fit, "wt")
```

---

**constrOptim**     *Linearly constrained optimisation*


---

**Description**

Minimise a function subject to linear inequality constraints using an adaptive barrier algorithm.

**Usage**

```
constrOptim(theta, f, grad, ui, ci, mu = 1e-04,
  control = list(),
  method = if(is.null(grad)) "Nelder-Mead" else "BFGS",
  outer.iterations = 100, outer.eps = 1e-05, ...)
```

**Arguments**

<b>theta</b>	Starting value: must be in the feasible region.
<b>f</b>	Function to minimise.
<b>grad</b>	Gradient of <b>f</b> .
<b>ui</b>	Constraints (see below).
<b>ci</b>	Constraints (see below).
<b>mu</b>	(Small) tuning parameter.
<b>control</b>	Passed to <b>optim</b> .
<b>method</b>	Passed to <b>optim</b> .
<b>outer.iterations</b>	Iterations of the barrier algorithm.
<b>outer.eps</b>	Criterion for relative convergence of the barrier algorithm.
<b>...</b>	Other arguments passed to <b>optim</b>

**Details**

The feasible region is defined by `ui %% theta - ci >= 0`. The starting value must be in the interior of the feasible region, but the minimum may be on the boundary.

A logarithmic barrier is added to enforce the constraints and then **optim** is called. The barrier function is chosen so that the objective function should decrease at each outer iteration. Minima in the interior of the

feasible region are typically found quite quickly, but a substantial number of outer iterations may be needed for a minimum on the boundary.

The tuning parameter `mu` multiplies the barrier term. Its precise value is often relatively unimportant. As `mu` increases the augmented objective function becomes closer to the original objective function but also less smooth near the boundary of the feasible region.

Any `optim` method that permits infinite values for the objective function may be used (currently all but "L-BFGS-B"). The gradient function must be supplied except with `method="Nelder-Mead"`.

As with `optim`, the default is to minimise and maximisation can be performed by setting `control$fnscale` to a negative value.

## Value

As for `optim`, but with two extra components: `barrier.value` giving the value of the barrier function at the optimum and `outer.iterations` gives the number of outer iterations (calls to `optim`)

## References

K. Lange *Numerical Analysis for Statisticians*. Springer 2001, p185ff

## See Also

`optim`, especially `method="L-BGFS-B"` which does box-constrained optimisation.

## Examples

```
## from optim
fr <- function(x) { ## Rosenbrock Banana function
  x1 <- x[1]
  x2 <- x[2]
  100 * (x2 - x1 * x1)^2 + (1 - x1)^2
}
grr <- function(x) { ## Gradient of 'fr'
  x1 <- x[1]
  x2 <- x[2]
  c(-400 * x1 * (x2 - x1 * x1) - 2 * (1 - x1),
    200 * (x2 - x1 * x1))
}

optim(c(-1.2,1), fr, grr)
# Box-constraint, optimum on the boundary
```

```

constrOptim(c(-1.2,0.9), fr, grr,
            ui=rbind(c(-1,0),c(0,-1)), ci=c(-1,-1))
# x<=0.9, y-x>0.1
constrOptim(c(.5,0), fr, grr,
            ui=rbind(c(-1,0),c(1,-1)), ci=c(-0.9,0.1))

## Solves linear and quadratic programming problems but
## needs a feasible starting value
#
# from example(solve.QP) in 'quadprog'
# no derivative
fQP <- function(b) {-sum(c(0,5,0)*b)+0.5*sum(b*b)}
Amat <- matrix(c(-4,-3,0,2,1,0,0,-2,1),3,3)
bvec <- c(-8,2,0)
constrOptim(c(2,-1,-1), fQP, NULL, ui=t(Amat),ci=bvec)
# derivative
gQP <- function(b) {-c(0,5,0)+b}
constrOptim(c(2,-1,-1), fQP, gQP, ui=t(Amat), ci=bvec)

## Now with maximisation instead of minimisation
hQP <- function(b) {sum(c(0,5,0)*b)-0.5*sum(b*b)}
constrOptim(c(2,-1,-1), hQP, NULL, ui=t(Amat), ci=bvec,
            control=list(fnscale=-1))

```

---

**contrast**      *Contrast Matrices*

---

**Description**

Return a matrix of contrasts.

**Usage**

```
contr.helmert(n, contrasts = TRUE)
contr.poly(n, scores = 1:n, contrasts = TRUE)
contr.sum(n, contrasts = TRUE)
contr.treatment(n, base = 1, contrasts = TRUE)
```

**Arguments**

<b>n</b>	a vector of levels for a factor, or the number of levels.
<b>contrasts</b>	a logical indicating whether contrasts should be computed.
<b>scores</b>	the set of values over which orthogonal polynomials are to be computed.
<b>base</b>	an integer specifying which group is considered the baseline group. Ignored if <b>contrasts</b> is <b>FALSE</b> .

**Details**

These functions are used for creating contrast matrices for use in fitting analysis of variance and regression models. The columns of the resulting matrices contain contrasts which can be used for coding a factor with **n** levels. The returned value contains the computed contrasts. If the argument **contrasts** is **FALSE** then a square indicator matrix is returned.

**contr.helmert** returns Helmert contrasts, which contrast the second level with the first, the third with the average of the first two, and so on. **contr.poly** returns contrasts based on orthogonal polynomials. **contr.sum** uses “sum to zero contrasts”.

**contr.treatment** contrasts each level with the baseline level (specified by **base**): the baseline level is omitted. Note that this does not produce “contrasts” as defined in the standard theory for linear models as they are not orthogonal to the constant.

## Value

A matrix with `n` rows and `k` columns, with `k=n-1` if `contrasts` is `TRUE` and `k=n` if `contrasts` is `FALSE`.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical models*. Chapter 2 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`contrasts`, `C`, and `aov`, `glm`, `lm`.

## Examples

```
(cH <- contr.helmert(4))
apply(cH, 2,sum) # column sums are 0!
crossprod(cH) # diagonal -- columns are orthogonal
# just the 4 x 4 identity matrix
contr.helmert(4, contrasts = FALSE)

(cT <- contr.treatment(5))
all(crossprod(cT) == diag(4)) # TRUE: even orthonormal

(cP <- contr.poly(3)) # Linear and Quadratic
zapsmall(crossprod(cP), dig=15) # orthonormal up to fuzz
```



---

**contrasts**      *Get and Set Contrast Matrices*

---

**Description**

Set and view the contrasts associated with a factor.

**Usage**

```
contrasts(x, contrasts = TRUE)
contrasts(x, how.many) <- value
```

**Arguments**

<b>x</b>	a factor.
<b>contrasts</b>	logical. See Details.
<b>how.many</b>	How many contrasts should be made. Defaults to one less than the number of levels of <b>x</b> . This need not be the same as the number of columns of <b>ctr</b> .
<b>value</b>	either a matrix whose columns give coefficients for contrasts in the levels of <b>x</b> , or the (quoted) name of a function which computes such matrices.

**Details**

If contrasts are not set for a factor the default functions from `options("contrasts")` are used.

The argument **contrasts** is ignored if **x** has a matrix **contrasts** attribute set. Otherwise if **contrasts** = **TRUE** it is passed to a contrasts function such as **contr.treatment** and if **contrasts** = **FALSE** an identity matrix is returned.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical models*. Chapter 2 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

**C**, **contr.helmert**, **contr.poly**, **contr.sum**, **contr.treatment**; **glm**, **aov**, **lm**.

## Examples

```
example(factor)
fff <- ff[, drop=TRUE] # reduce to 5 levels.
contrasts(fff) # treatment contrasts by default
contrasts(C(fff, sum))
contrasts(fff, contrasts = FALSE) # the 5x5 identity matrix

# set sum contrasts
contrasts(fff) <- contr.sum(5); contrasts(fff)
# set 2 contrasts
contrasts(fff, 2) <- contr.sum(5); contrasts(fff)
# supply 2 contrasts, compute 2 more to make full set of 4.
contrasts(fff) <- contr.sum(5)[,1:2]; contrasts(fff)
```

---

deviance	<i>Model Deviance</i>
----------	-----------------------

---

## Description

Returns the deviance of a fitted model object.

## Usage

```
deviance(object, ...)
```

## Arguments

<code>object</code>	an object for which the deviance is desired.
<code>...</code>	additional optional argument.

## Details

This is a generic function which can be used to extract deviances for fitted models. Consult the individual modeling functions for details on how to use this function.

## Value

The value of the deviance extracted from the object `object`.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

## See Also

`df.residual`, `extractAIC`, `glm`, `lm`.

---

**df.residual**      *Residual Degrees-of-Freedom*

---

**Description**

Returns the residual degrees-of-freedom extracted from a fitted model object.

**Usage**

```
df.residual(object, ...)
```

**Arguments**

<b>object</b>	an object for which the degrees-of-freedom are desired.
<b>...</b>	additional optional arguments.

**Details**

This is a generic function which can be used to extract residual degrees-of-freedom for fitted models. Consult the individual modeling functions for details on how to use this function.

The default method just extracts the **df.residual** component.

**Value**

The value of the residual degrees-of-freedom extracted from the object **x**.

**See Also**

**deviance**, **glm**, **lm**.

---

**dummy.coef**      *Extract Coefficients in Original Coding*

---

**Description**

This extracts coefficients in terms of the original levels of the coefficients rather than the coded variables.

**Usage**

```
dummy.coef(object, ...)  
  
## S3 method for class 'lm':  
dummy.coef(object, use.na = FALSE, ...)  
  
## S3 method for class 'aovlist':  
dummy.coef(object, use.na = FALSE, ...)
```

**Arguments**

<b>object</b>	a linear model fit.
<b>use.na</b>	logical flag for coefficients in a singular model. If <b>use.na</b> is true, undetermined coefficients will be missing; if false they will get one possible value.
<b>...</b>	arguments passed to or from other methods.

**Details**

A fitted linear model has coefficients for the contrasts of the factor terms, usually one less in number than the number of levels. This function re-expresses the coefficients in the original coding; as the coefficients will have been fitted in the reduced basis, any implied constraints (e.g., zero sum for **contr.helmert** or **contr.sum** will be respected. There will be little point in using **dummy.coef** for **contr.treatment** contrasts, as the missing coefficients are by definition zero.

The method used has some limitations, and will give incomplete results for terms such as **poly(x, 2)**). However, it is adequate for its main purpose, **aov** models.

**Value**

A list giving for each term the values of the coefficients. For a multi-stratum **aov** model, such a list for each stratum.

**Warning**

This function is intended for human inspection of the output: it should not be used for calculations. Use coded variables for all calculations.

The results differ from *S* for singular values, where *S* can be incorrect.

**See Also**

`aov`, `model.tables`

**Examples**

```
options(contrasts=c("contr.helmert", "contr.poly"))
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)

npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)
npk.aov <- aov(yield ~ block + N*P*K, npk)
dummy.coef(npk.aov)

npk.aovE <- aov(yield ~ N*P*K + Error(block), npk)
dummy.coef(npk.aovE)
```

---

<b>eff.aovlist</b>	<i>Compute Efficiencies of Multistratum Analysis of Variance</i>
--------------------	--

---

## Description

Computes the efficiencies of fixed-effect terms in an analysis of variance model with multiple strata.

## Usage

```
eff.aovlist(aovlist)
```

## Arguments

**aovlist**                The result of a call to **aov** with an **Error** term.

## Details

Fixed-effect terms in an analysis of variance model with multiple strata may be estimable in more than one stratum, in which case there is less than complete information in each. The efficiency is the fraction of the maximum possible precision (inverse variance) obtainable by estimating in just that stratum.

This is used to pick strata in which to estimate terms in **model.tables.aovlist** and elsewhere.

## Value

A matrix giving for each non-pure-error stratum (row) the efficiencies for each fixed-effect term in the model.

## See Also

**aov**, **model.tables.aovlist**, **se.contrast.aovlist**

## Examples

```
## for balanced designs all efficiencies are zero or one.
## so as a statistically meaningless test:
options(contrasts=c("contr.helmert", "contr.poly"))
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
```

```
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)

npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)
npk.aovE <- aov(yield ~ N*P*K + Error(block), npk)
eff.aovlist(npk.aovE)
```



---

**effects**     *Effects from Fitted Model*

---

**Description**

Returns (orthogonal) effects from a fitted model, usually a linear model. This is a generic function, but currently only has a methods for objects inheriting from classes "lm" and "glm".

**Usage**

```
effects(object, ...)  
  
## S3 method for class 'lm':  
effects(object, set.sign=FALSE, ...)
```

**Arguments**

<b>object</b>	an R object; typically, the result of a model fitting function such as <code>lm</code> .
<b>set.sign</b>	logical. If <code>TRUE</code> , the sign of the effects corresponding to coefficients in the model will be set to agree with the signs of the corresponding coefficients, otherwise the sign is arbitrary.
<b>...</b>	arguments passed to or from other methods.

**Details**

For a linear model fitted by `lm` or `aov`, the effects are the uncorrelated single-degree-of-freedom values obtained by projecting the data onto the successive orthogonal subspaces generated by the QR decomposition during the fitting process. The first  $r$  (the rank of the model) are associated with coefficients and the remainder span the space of residuals (but are not associated with particular residuals).

Empty models do not have effects.

**Value**

A (named) numeric vector of the same length as `residuals`, or a matrix if there were multiple responses in the fitted model, in either case of class `"coef"`.

The first  $r$  rows are labelled by the corresponding coefficients, and the remaining rows are unlabelled. Note that in rank-deficient models the “corresponding” coefficients will be in a different order if pivoting occurred.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

## See Also

`coef`

## Examples

```
y <- c(1:3,7,5)
x <- c(1:3,6:7)
( ee <- effects(lm(y ~ x)) )
# just the first is different
c(round(ee - effects(lm(y+10 ~ I(x-3.8))),3))
```

---

**expand.grid**     *Create a Data Frame from All Combinations of Factors*

---

## Description

Create a data frame from all combinations of the supplied vectors or factors. See the description of the return value for precise details of the way this is done.

## Usage

```
expand.grid(...)
```

## Arguments

...                      Vectors, factors or a list containing these.

## Value

A data frame containing one row for each combination of the supplied factors. The first factors vary fastest. The columns are labelled by the factors if these are supplied as named arguments or named components of a list.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

## Examples

```
expand.grid(height = seq(60, 80, 5),  
             weight = seq(100, 300, 50),  
             sex = c("Male", "Female"))
```

---

<code>expand.model.frame</code>	<i>Add new variables to a model frame</i>
---------------------------------	---

---

## Description

Evaluates new variables as if they had been part of the formula of the specified model. This ensures that the same **na.action** and **subset** arguments are applied and allows, for example, **x** to be recovered for a model using **sin(x)** as a predictor.

## Usage

```
expand.model.frame(model, extras,
                   envir=environment(formula(model)),
                   na.expand = FALSE)
```

## Arguments

<b>model</b>	a fitted model
<b>extras</b>	one-sided formula or vector of character strings describing new variables to be added
<b>envir</b>	an environment to evaluate things in
<b>na.expand</b>	logical; see below

## Details

If **na.expand=FALSE** then NA values in the extra variables will be passed to the **na.action** function used in **model**. This may result in a shorter data frame (with **na.omit**) or an error (with **na.fail**). If **na.expand=TRUE** the returned data frame will have precisely the same rows as **model.frame(model)**, but the columns corresponding to the extra variables may contain NA.

## Value

A data frame.

## See Also

`model.frame`, `predict`

**Examples**

```
data(trees)
model <- lm(log(Volume) ~ log(Girth) + log(Height),
            data=trees)
expand.model.frame(model, ~ Girth) # prints data.frame like

dd <- data.frame(x=1:5, y=rnorm(5), z=c(1,2,NA,4,5))
model <- glm(y ~ x, data=dd, subset=1:4, na.action=na.omit)
expand.model.frame(model, "z", na.expand=FALSE) # = default
expand.model.frame(model, "z", na.expand=TRUE)
```

---

<b>extractAIC</b>	<i>Extract AIC from a Fitted Model</i>
-------------------	--

---

## Description

Computes the (generalized) Akaike **A**n **I**nformation **C**riterion for a fitted parametric model.

## Usage

```
extractAIC(fit, scale,      k = 2, ...)
```

## Arguments

<b>fit</b>	fitted model, usually the result of a fitter like <b>lm</b> .
<b>scale</b>	optional numeric specifying the scale parameter of the model, see <b>scale</b> in <b>step</b> .
<b>k</b>	numeric specifying the “weight” of the <i>equivalent degrees of freedom</i> ( $\equiv \mathbf{edf}$ ) part in the AIC formula.
<b>...</b>	further arguments (currently unused in base R).

## Details

This is a generic function, with methods in base R for "**aov**", "**coxph**", "**glm**", "**lm**", "**negbin**" and "**survreg**" classes.

The criterion used is

$$AIC = -2 \log L + k \times \mathbf{edf},$$

where  $L$  is the likelihood and **edf** the equivalent degrees of freedom (i.e., the number of parameters for usual parametric models) of **fit**.

For linear models with unknown scale (i.e., for **lm** and **aov**),  $-2 \log L$  is computed from the *deviance* and uses a different additive constant to **AIC**.

**k** = 2 corresponds to the traditional AIC, using **k** = **log(n)** provides the BIC (Bayes IC) instead.

For further information, particularly about **scale**, see **step**.

**Value**

A numeric vector of length 2, giving

<b>edf</b>	the “equivalent <b>d</b> egrees of <b>f</b> reedom” of the fitted model <b>fit</b> .
<b>AIC</b>	the (generalized) Akaike Information Criterion for <b>fit</b> .

**Note**

These functions are used in **add1**, **drop1** and **step** and that may be their main use.

**Author(s)**

B. D. Ripley

**References**

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. New York: Springer (4th ed).

**See Also**

AIC, deviance, **add1**, **step**

**Examples**

```
example(glm)
extractAIC(glm.D93) # 5 15.129
```

---

**factor.scope**     *Compute Allowed Changes in Adding to or Dropping from a Formula*

---

## Description

`add.scope` and `drop.scope` compute those terms that can be individually added to or dropped from a model while respecting the hierarchy of terms.

## Usage

```
add.scope(terms1, terms2)
drop.scope(terms1, terms2)
factor.scope(factor, scope)
```

## Arguments

<code>terms1</code>	the terms or formula for the base model.
<code>terms2</code>	the terms or formula for the upper ( <code>add.scope</code> ) or lower ( <code>drop.scope</code> ) scope. If missing for <code>drop.scope</code> it is taken to be the null formula, so all terms (except any intercept) are candidates to be dropped.
<code>factor</code>	the " <b>factor</b> " attribute of the terms of the base object.
<code>scope</code>	a list with one or both components <code>drop</code> and <code>add</code> giving the " <b>factor</b> " attribute of the lower and upper scopes respectively.

## Details

`factor.scope` is not intended to be called directly by users.

## Value

For `add.scope` and `drop.scope` a character vector of terms labels. For `factor.scope`, a list with components `drop` and `add`, character vectors of terms labels.

## See Also

`add1`, `drop1`, `aov`, `lm`



**Examples**

```
add.scope( ~ a + b + c + a:b, ~ (a + b + c)^3)
# [1] "a:c" "b:c"
drop.scope( ~ a + b + c + a:b)
# [1] "c"   "a:b"
```

---

**family**     *Family Objects for Models*


---

**Description**

Family objects provide a convenient way to specify the details of the models used by functions such as `glm`. See the documentation for `glm` for the details on how such model fitting takes place.

**Usage**

```
family(object, ...)

binomial(link = "logit")
gaussian(link = "identity")
Gamma(link = "inverse")
inverse.gaussian(link = "1/mu^2")
poisson(link = "log")
quasi(link = "identity", variance = "constant")
quasibinomial(link = "logit")
quasipoisson(link = "log")
```

**Arguments**

<b>link</b>	<p>a specification for the model link function. The <code>gaussian</code> family accepts the links "identity", "log" and "inverse"; the <code>binomial</code> family the links "logit", "probit", "log" and "cloglog" (complementary log-log); the <code>Gamma</code> family the links "inverse", "identity" and "log"; the <code>poisson</code> family the links "log", "identity", and "sqrt" and the <code>inverse.gaussian</code> family the links "1/mu^2", "inverse", "inverse" and "log".</p> <p>The <code>quasi</code> family allows the links "logit", "probit", "cloglog", "identity", "inverse", "log", "1/mu^2" and "sqrt". The function <code>power</code> can also be used to create a power link function for the <code>quasi</code> family.</p>
<b>variance</b>	<p>for all families, other than <code>quasi</code>, the variance function is determined by the family. The <code>quasi</code> family will accept the specifications "constant", "mu(1-</p>

	<code>mu)</code> ", <code>"mu"</code> , <code>"mu^2"</code> and <code>"mu^3"</code> for the variance function.
<code>object</code>	the function <code>family</code> accesses the <code>family</code> objects which are stored within objects created by modelling functions (e.g., <code>glm</code> ).
<code>...</code>	further arguments passed to methods.

## Details

The `quasibinomial` and `quasipoisson` families differ from the `binomial` and `poisson` families only in that the dispersion parameter is not fixed at one, so they can “model” over-dispersion. For the binomial case see McCullagh and Nelder (1989, pp. 124–8). Although they show that there is (under some restrictions) a model with variance proportional to mean as in the quasi-binomial model, note that `glm` does not compute maximum-likelihood estimates in that model. The behaviour of `S` is closer to the quasi- variants.

## Author(s)

The design was inspired by `S` functions of the same names described in Hastie & Pregibon (1992).

## References

- McCullagh P. and Nelder, J. A. (1989) *Generalized Linear Models*. London: Chapman and Hall.
- Dobson, A. J. (1983) *An Introduction to Statistical Modelling*. London: Chapman and Hall.
- Cox, D. R. and Snell, E. J. (1981). *Applied Statistics; Principles and Examples*. London: Chapman and Hall.
- Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`glm`, `power`.

## Examples

```
nf <- gaussian() # Normal family
nf
```

```

str(nf) # internal STRucture

gf <- Gamma()
gf
str(gf)
gf$linkinv
gf$variance(-3:4) # == (.)^2

## quasipoisson. compare with example(glm)
counts <- c(18,17,15,20,10,20,25,13,12)
outcome <- gl(3,1,9)
treatment <- gl(3,3)
d.AD <- data.frame(treatment, outcome, counts)
glm.qD93 <- glm(counts ~ outcome + treatment,
                family=quasipoisson())

glm.qD93
anova(glm.qD93, test="F")
summary(glm.qD93)
## for Poisson results use
anova(glm.qD93, dispersion = 1, test="Chisq")
summary(glm.qD93, dispersion = 1)

## tests of quasi
x <- rnorm(100)
y <- rpois(100, exp(1+x))
glm(y ~x, family=quasi(var="mu", link="log"))
# which is the same as
glm(y ~x, family=poisson)
glm(y ~x, family=quasi(var="mu^2", link="log"))
# should fail
glm(y ~x, family=quasi(var="mu^3", link="log"))
y <- rbinom(100, 1, plogis(x))
# needs to set a starting value for the next fit
glm(y ~x, family=quasi(var="mu(1-mu)", link="logit"),
    start=c(0,1))

```

---

**fitted**     *Extract Model Fitted Values*

---

**Description**

**fitted** is a generic function which extracts fitted values from objects returned by modeling functions. **fitted.values** is an alias for it.

All object classes which are returned by model fitting functions should provide a **fitted** method. (Note that the generic is **fitted** and not **fitted.values**.)

Methods can make use of **napredict** methods to compensate for the omission of missing values. The default, **lm** and **glm** methods do.

**Usage**

```
fitted(object, ...)  
fitted.values(object, ...)
```

**Arguments**

<b>object</b>	an object for which the extraction of model fitted values is meaningful.
<b>...</b>	other arguments.

**Value**

Fitted values extracted from the object **x**.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

**See Also**

**coefficients**, **glm**, **lm**, **residuals**.

---

formula	<i>Model Formulae</i>
---------	-----------------------

---

## Description

The generic function `formula` and its specific methods provide a way of extracting formulae which have been included in other objects.

`as.formula` is almost identical, additionally preserving attributes when `object` already inherits from `"formula"`. The default value of the `env` argument is used only when the formula would otherwise lack an environment.

## Usage

```
y ~ model
formula(x, ...)
as.formula(object, env = parent.frame())
```

## Arguments

<code>x</code> , <code>object</code>	an object
<code>...</code>	further arguments passed to or from other methods.
<code>env</code>	the environment to associate with the result.

## Details

The models fit by, e.g., the `lm` and `glm` functions are specified in a compact symbolic form. The `~` operator is basic in the formation of such models. An expression of the form `y ~ model` is interpreted as a specification that the response `y` is modelled by a linear predictor specified symbolically by `model`. Such a model consists of a series of terms separated by `+` operators. The terms themselves consist of variable and factor names separated by `:` operators. Such a term is interpreted as the interaction of all the variables and factors appearing in the term.

In addition to `+` and `:`, a number of other operators are useful in model formulae. The `*` operator denotes factor crossing: `a*b` interpreted as `a+b+a:b`. The `^` operator indicates crossing to the specified degree. For example `(a+b+c)^2` is identical to `(a+b+c)*(a+b+c)` which in turn expands to a formula containing the main effects for `a`, `b` and `c` together with their second-order interactions. The `%in%` operator indicates that the terms on its left are nested within those on the right. For example

`a+b%in%a` expands to the formula `a+a:b`. The `-` operator removes the specified terms, so that `(a+b+c)^2 - a:b` is identical to `a + b + c + b:c + a:c`. It can also be used to remove the intercept term: `y~x - 1` is a line through the origin. A model with no intercept can be also specified as `y~x + 0` or `0 + y~x`.

While formulae usually involve just variable and factor names, they can also involve arithmetic expressions. The formula `log(y) ~ a + log(x)` is quite legal. When such arithmetic expressions involve operators which are also used symbolically in model formulae, there can be confusion between arithmetic and symbolic operator use.

To avoid this confusion, the function `I()` can be used to bracket those portions of a model formula where the operators are used in their arithmetic sense. For example, in the formula `y ~ a + I(b+c)`, the term `b+c` is to be interpreted as the sum of `b` and `c`.

As from R 1.8.0 variable names can be quoted by backticks ‘`like this`’ in formulae, although there is no guarantee that all code using formulae will accept such non-syntactic names.

## Value

All the functions above produce an object of class “`formula`” which contains a symbolic model formula.

## Environments

A formula object has an associated environment, and this environment (rather than the parent environment) is used by `model.frame` to evaluate variables that are not found in the supplied `data` argument.

Formulas created with the `~` operator use the environment in which they were created. Formulas created with `as.formula` will use the `env` argument for their environment. Pre-existing formulas extracted with `as.formula` will only have their environment changed if `env` is explicitly given.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical models*. Chapter 2 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

I.

For formula manipulation: `terms`, and `all.vars`; for typical use: `lm`, `glm`, and `coplot`.

## Examples

```
class(fo <- y ~ x1*x2) # "formula"
fo
typeof(fo) # R internal : "language"
terms(fo)

environment(fo)
environment(as.formula("y ~ x"))
environment(as.formula("y ~ x",env=new.env()))

## Create a formula for a model with a large number of
## variables:
xnam <- paste("x", 1:25, sep="")
(fmla <-
  as.formula(paste("y ~ ", paste(xnam, collapse= "+"))))
```



---

**glm**     *Fitting Generalized Linear Models*

---

**Description**

`glm` is used to fit generalized linear models, specified by giving a symbolic description of the linear predictor and a description of the error distribution.

**Usage**

```
glm(formula, family = gaussian, data, weights = NULL,
     subset = NULL, na.action, start = NULL, etastart = NULL,
     mustart = NULL, offset = NULL,
     control = glm.control(...), model = TRUE,
     method = "glm.fit", x = FALSE, y = TRUE,
     contrasts = NULL, ...)
```

```
glm.fit(x, y, weights = rep(1, nobs),
        start = NULL, etastart = NULL, mustart = NULL,
        offset = rep(0, nobs), family = gaussian(),
        control = glm.control(), intercept = TRUE)
```

```
## S3 method for class 'glm':
weights(object, type = c("prior", "working"), ...)
```

**Arguments**

<b>formula</b>	a symbolic description of the model to be fit. The details of model specification are given below.
<b>family</b>	a description of the error distribution and link function to be used in the model. This can be a character string naming a family function, a family function or the result of a call to a family function. (See <b>family</b> for details of family functions.)
<b>data</b>	an optional data frame containing the variables in the model. By default the variables are taken from <code>environment(formula)</code> , typically the environment from which <code>glm</code> is called.
<b>weights</b>	an optional vector of weights to be used in the fitting process.

<code>subset</code>	an optional vector specifying a subset of observations to be used in the fitting process.
<code>na.action</code>	a function which indicates what should happen when the data contain NAs. The default is set by the <code>na.action</code> setting of <code>options</code> , and is <code>na.fail</code> if that is unset. The “factory-fresh” default is <code>na.omit</code> .
<code>start</code>	starting values for the parameters in the linear predictor.
<code>etastart</code>	starting values for the linear predictor.
<code>mustart</code>	starting values for the vector of means.
<code>offset</code>	this can be used to specify an <i>a priori</i> known component to be included in the linear predictor during fitting.
<code>control</code>	a list of parameters for controlling the fitting process. See the documentation for <code>glm.control</code> for details.
<code>model</code>	a logical value indicating whether <i>model frame</i> should be included as a component of the returned value.
<code>method</code>	the method to be used in fitting the model. The default method <code>"glm.fit"</code> uses iteratively reweighted least squares (IWLS). The only current alternative is <code>"model.frame"</code> which returns the model frame and does no fitting.
<code>x, y</code>	For <code>glm</code> : logical values indicating whether the response vector and model matrix used in the fitting process should be returned as components of the returned value.  For <code>glm.fit</code> : <code>x</code> is a design matrix of dimension <code>n * p</code> , and <code>y</code> is a vector of observations of length <code>n</code> .
<code>contrasts</code>	an optional list. See the <code>contrasts.arg</code> of <code>model.matrix.default</code> .
<code>object</code>	an object inheriting from class <code>"glm"</code> .
<code>type</code>	character, partial matching allowed. Type of weights to extract from the fitted model object.
<code>intercept</code>	logical. Should an intercept be included?
<code>...</code>	further arguments passed to or from other methods.

## Details

A typical predictor has the form `response ~ terms` where `response` is the (numeric) response vector and `terms` is a series of terms which specifies a linear predictor for `response`. For `binomial` models the response

can also be specified as a **factor** (when the first level denotes failure and all others success) or as a two-column matrix with the columns giving the numbers of successes and failures. A terms specification of the form **first + second** indicates all the terms in **first** together with all the terms in **second** with duplicates removed.

A specification of the form **first:second** indicates the the set of terms obtained by taking the interactions of all terms in **first** with all terms in **second**. The specification **first\*second** indicates the *cross* of **first** and **second**. This is the same as **first + second + first:second**.

**glm.fit** and **glm.fit.null** are the workhorse functions: the former calls the latter for a null model (with no intercept).

If more than one of **etastart**, **start** and **mustart** is specified, the first in the list will be used.

## Value

**glm** returns an object of class inheriting from "**glm**" which inherits from the class "**lm**". See later in this section.

The function **summary** (i.e., **summary.glm**) can be used to obtain or print a summary of the results and the function **anova** (i.e., **anova.glm**) to produce an analysis of variance table.

The generic accessor functions **coefficients**, **effects**, **fitted.values** and **residuals** can be used to extract various useful features of the value returned by **glm**.

**weights** extracts a vector of weights, one for each case in the fit (after subsetting and **na.action**).

An object of class "**glm**" is a list containing at least the following components:

<b>coefficients</b>	a named vector of coefficients
<b>residuals</b>	the <i>working</i> residuals, that is the residuals in the final iteration of the IWLS fit.
<b>fitted.values</b>	the fitted mean values, obtained by transforming the linear predictors by the inverse of the link function.
<b>rank</b>	the numeric rank of the fitted linear model.
<b>family</b>	the <b>family</b> object used.
<b>linear.predictors</b>	the linear fit on link scale.
<b>deviance</b>	up to a constant, minus twice the maximized log-likelihood. Where sensible, the constant is chosen so that a saturated model has deviance zero.

<code>aic</code>	Akaike's <i>An Information Criterion</i> , minus twice the maximized log-likelihood plus twice the number of coefficients (so assuming that the dispersion is known).
<code>null.deviance</code>	The deviance for the null model, comparable with <b>deviance</b> . The null model will include the offset, and an intercept if there is one in the model
<code>iter</code>	the number of iterations of IWLS used.
<code>weights</code>	the <i>working</i> weights, that is the weights in the final iteration of the IWLS fit.
<code>prior.weights</code>	the case weights initially supplied.
<code>df.residual</code>	the residual degrees of freedom.
<code>df.null</code>	the residual degrees of freedom for the null model.
<code>y</code>	the <code>y</code> vector used. (It is a vector even for a binomial model.)
<code>converged</code>	logical. Was the IWLS algorithm judged to have converged?
<code>boundary</code>	logical. Is the fitted value on the boundary of the attainable values?
<code>call</code>	the matched call.
<code>formula</code>	the formula supplied.
<code>terms</code>	the <b>terms</b> object used.
<code>data</code>	the <b>data argument</b> .
<code>offset</code>	the offset vector used.
<code>control</code>	the value of the <b>control</b> argument used.
<code>method</code>	the name of the fitter function used, in R always <b>"glm.fit"</b> .
<code>contrasts</code>	(where relevant) the contrasts used.
<code>xlevels</code>	(where relevant) a record of the levels of the factors used in fitting.

In addition, non-empty fits will have components **qr**, **R** and **effects** relating to the final weighted linear fit.

Objects of class **"glm"** are normally of class **c("glm", "lm")**, that is inherit from class **"lm"**, and well-designed methods for class **"lm"** will be applied to the weighted linear model at the final iteration of IWLS. However, care is needed, as extractor functions for class **"glm"** such as **residuals** and **weights** do **not** just pick out the component of the fit with the same name.

If a **binomial glm** model is specified by giving a two-column response, the weights returned by `prior.weights` are the total numbers of cases (factored by the supplied case weights) and the component `y` of the result is the proportion of successes.

## Author(s)

The original R implementation of `glm` was written by Simon Davies working for Ross Ihaka at the University of Auckland, but has since been extensively re-written by members of the R Core team.

The design was inspired by the S function of the same name described in Hastie & Pregibon (1992).

## References

- Dobson, A. J. (1990) *An Introduction to Generalized Linear Models*. London: Chapman and Hall.
- Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.
- McCullagh P. and Nelder, J. A. (1989) *Generalized Linear Models*. London: Chapman and Hall.
- Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. New York: Springer.

## See Also

`anova.glm`, `summary.glm`, etc. for `glm` methods, and the generic functions `anova`, `summary`, `effects`, `fitted.values`, and `residuals`. Further, `lm` for non-generalized *linear* models.

`esoph`, `infert` and `predict.glm` have examples of fitting binomial glms.

## Examples

```
## Dobson (1990) Page 93: Randomized Controlled Trial :
counts <- c(18,17,15,20,10,20,25,13,12)
outcome <- gl(3,1,9)
treatment <- gl(3,3)
print(d.AD <- data.frame(treatment, outcome, counts))
glm.D93 <- glm(counts ~ outcome + treatment,
               family=poisson())
anova(glm.D93)
summary(glm.D93)
```

```
## an example with offsets from Venables & Ripley (2002,
## p.189)

## Need the anorexia data from a recent version of the
## package 'MASS':
library(MASS)
data(anorexia)

anorex.1 <- glm(Postwt ~ Prewt + Treat + offset(Prewt),
               family = gaussian, data = anorexia)
summary(anorex.1)

# A Gamma example, from McCullagh & Nelder (1989, pp.300-2)
clotting <- data.frame(
  u = c(5,10,15,20,30,40,60,80,100),
  lot1 = c(118,58,42,35,27,25,21,19,18),
  lot2 = c(69,35,26,21,18,16,13,12,12))
summary(glm(lot1 ~ log(u), data=clotting, family=Gamma))
summary(glm(lot2 ~ log(u), data=clotting, family=Gamma))
```

---

**glm.control**      *Auxiliary for Controlling GLM Fitting*

---

**Description**

Auxiliary function as user interface for `glm` fitting. Typically only used when calling `glm` or `glm.fit`.

**Usage**

```
glm.control(epsilon=1e-8, maxit=25, trace=FALSE)
```

**Arguments**

<code>epsilon</code>	positive convergence tolerance <i>epsilon</i> ; the iterations converge when $ dev - devold /( dev  + 0.1) < epsilon$ .
<code>maxit</code>	integer giving the maximal number of IWLS iterations.
<code>trace</code>	logical indicating if output should be produced for each iteration.

**Details**

If `epsilon` is small, it is also used as the tolerance for the least squares solution.

When `trace` is true, calls to `cat` produce the output for each IWLS iteration. Hence, `options(digits = *)` can be used to increase the precision, see the example.

**Value**

A list with the arguments as components.

**References**

Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

`glm.fit`, the fitting procedure used by `glm`.

## Examples

```
### A variation on example(glm) :
## Annette Dobson's example ...
counts <- c(18,17,15,20,10,20,25,13,12)
outcome <- gl(3,1,9)
treatment <- gl(3,3)
oo <- options(digits = 12) # to see more when tracing :
glm.D93X <-
  glm(counts ~ outcome + treatment, family=poisson(),
      trace = TRUE, epsilon = 1e-14)
options(oo)
# the last two are closer to 0 than in ?glm's glm.D93
coef(glm.D93X)
# put less so than in R < 1.8.0 when the default was 1e-4
```



---

**glm.summaries**      *Accessing Generalized Linear Model Fits*

---

## Description

These functions are all **methods** for class **glm** or **summary.glm** objects.

## Usage

```
## S3 method for class 'glm':  
family(object, ...)  
  
## S3 method for class 'glm':  
residuals(object,  
           type = c("deviance", "pearson", "working",  
                   "response", "partial"), ...)
```

## Arguments

<b>object</b>	an object of class <b>glm</b> , typically the result of a call to <b>glm</b> .
<b>type</b>	the type of residuals which should be returned. The alternatives are: " <b>deviance</b> " (default), " <b>pearson</b> ", " <b>working</b> ", " <b>response</b> ", and " <b>partial</b> ".
<b>...</b>	further arguments passed to or from other methods.

## Details

The references define the types of residuals: Davison & Snell is a good reference for the usages of each.

The partial residuals are a matrix of working residuals, with each column formed by omitting a term from the model.

## References

Davison, A. C. and Snell, E. J. (1991) *Residuals and diagnostics*. In: Statistical Theory and Modelling. In Honour of Sir David Cox, FRS, eds. Hinkley, D. V., Reid, N. and Snell, E. J., Chapman & Hall.

Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

McCullagh P. and Nelder, J. A. (1989) *Generalized Linear Models*. London: Chapman and Hall.

### See Also

`glm` for computing `glm.obj`, `anova.glm`; the corresponding *generic* functions, `summary.glm`, `coef`, `deviance`, `df.residual`, `effects`, `fitted`, `residuals`.

---

**influence.measures**      *Regression Deletion Diagnostics*

---

**Description**

This suite of functions can be used to compute some of the regression (leave-one-out deletion) diagnostics for linear and generalized linear models discussed in Belsley, Kuh and Welsch (1980), Cook and Weisberg (1982), etc.

**Usage**

```
influence.measures(model)

rstandard(model, ...)
## S3 method for class 'lm':
rstandard(model, infl = lm.influence(model, do.coef=FALSE),
  sd = sqrt(deviance(model)/df.residual(model)), ...)
## S3 method for class 'glm':
rstandard(model, infl = lm.influence(model, do.coef=FALSE),
  ...)

rstudent(model, ...)
## S3 method for class 'lm':
rstudent(model, infl = lm.influence(model, do.coef=FALSE),
  res = infl$wt.res, ...)
## S3 method for class 'glm':
rstudent(model, infl = influence(model, do.coef=FALSE),
  ...)

dffits(model, infl = , res = )

dfbeta(model, ...)
## S3 method for class 'lm':
dfbeta(model, infl = lm.influence(model, do.coef=TRUE),
  ...)

dfbetas(model, ...)
## S3 method for class 'lm':
dfbetas(model, infl = lm.influence(model, do.coef=TRUE),
  ...)
```

```

covratio(model, infl = lm.influence(model, do.coef=FALSE),
          res = weighted.residuals(model))

cooks.distance(model, ...)
## S3 method for class 'lm':
cooks.distance(model,
  infl = lm.influence(model, do.coef=FALSE),
  res = weighted.residuals(model),
  sd = sqrt(deviance(model)/df.residual(model)),
  hat = infl$hat, ...)
## S3 method for class 'glm':
cooks.distance(model,
  infl = influence(model, do.coef=FALSE),
  res = infl$pear.res,
  dispersion = summary(model)$dispersion,
  hat = infl$hat, ...)

hatvalues(model, ...)
## S3 method for class 'lm':
hatvalues(model,
  infl = lm.influence(model, do.coef=FALSE), ...)

hat(x, intercept = TRUE)

```

## Arguments

<code>model</code>	an R object, typically returned by <code>lm</code> or <code>glm</code> .
<code>infl</code>	influence structure as returned by <code>lm.influence</code> or <code>influence</code> (the latter only for the <code>glm</code> method of <code>rstudent</code> and <code>cooks.distance</code> ).
<code>res</code>	(possibly weighted) residuals, with proper default.
<code>sd</code>	standard deviation to use, see default.
<code>dispersion</code>	dispersion (for <code>glm</code> objects) to use, see default.
<code>hat</code>	hat values $H_{ii}$ , see default.
<code>x</code>	the $X$ or design matrix.
<code>intercept</code>	should an intercept column be pre-pended to <code>x</code> ?
<code>...</code>	further arguments passed to or from other methods.

## Details

The primary high-level function is `influence.measures` which produces a class "`infl`" object tabular display showing the DFBETAS for

each model variable, DFFITS, covariance ratios, Cook's distances and the diagonal elements of the hat matrix. Cases which are influential with respect to any of these measures are marked with an asterisk.

The functions `dfbetas`, `dffits`, `covratio` and `cooks.distance` provide direct access to the corresponding diagnostic quantities. Functions `rstandard` and `rstudent` give the standardized and Studentized residuals respectively. (These re-normalize the residuals to have unit variance, using an overall and leave-one-out measure of the error variance respectively.)

Values for generalized linear models are approximations, as described in Williams (1987) (except that Cook's distances are scaled as  $F$  rather than as chi-square values).

The optional `infl`, `res` and `sd` arguments are there to encourage the use of these direct access functions, in situations where, e.g., the underlying basic influence measures (from `lm.influence` or the generic `influence`) are already available.

Note that cases with `weights == 0` are *dropped* from all these functions, but that if a linear model has been fitted with `na.action = na.exclude`, suitable values are filled in for the cases excluded during fitting.

The function `hat()` exists mainly for S (version 2) compatibility; we recommend using `hatvalues()` instead.

## Note

For `hatvalues`, `dfbeta`, and `dfbetas`, the method for linear models also works for generalized linear models.

## Author(s)

Several R core team members and John Fox, originally in his 'car' package.

## References

- Belsley, D. A., Kuh, E. and Welsch, R. E. (1980) *Regression Diagnostics*. New York: Wiley.
- Cook, R. D. and Weisberg, S. (1982) *Residuals and Influence in Regression*. London: Chapman and Hall.
- Williams, D. A. (1987) Generalized linear model diagnostics using the deviance and single case deletions. *Applied Statistics* **36**, 181–191.

Fox, J. (1997) *Applied Regression, Linear Models, and Related Methods*. Sage.

Fox, J. (2002) *An R and S-Plus Companion to Applied Regression*. Sage Publ.

## See Also

influence (containing lm.influence).

## Examples

```
## Analysis of the life-cycle savings data given in
## Belsley, Kuh and Welsch.
data(LifeCycleSavings)
lm.SR <- lm(sr ~ pop15 + pop75 + dpi + ddpi,
            data = LifeCycleSavings)

inflm.SR <- influence.measures(lm.SR)
# which observations 'are' influential
which(apply(inflm.SR$is.inf, 1, any))
summary(inflm.SR) # only these
inflm.SR          # all
# recommended by some
plot(rstudent(lm.SR) ~ hatvalues(lm.SR))

## The 'infl' argument is not needed, but avoids
## recomputation:
rs <- rstandard(lm.SR)
iflSR <- influence(lm.SR)
identical(rs, rstandard(lm.SR, infl = iflSR))
## to "see" the larger values:
1000 * round(dfbetas(lm.SR, infl = iflSR), 3)

## Huber's data [Atkinson 1985]
xh <- c(-4:0, 10)
yh <- c(2.48, .73, -.04, -1.44, -1.32, 0)
summary(lmH <- lm(yh ~ xh))
(im <- influence.measures(lmH))
plot(xh, yh,
     main = "Huber's data: L.S. line and influential obs.")
abline(lmH);
points(xh[im$is.inf], yh[im$is.inf], pch=20, col=2)
```

---

<code>is.empty.model</code>	<i>Check if a Model is Empty</i>
-----------------------------	----------------------------------

---

## Description

R model notation allows models with no intercept and no predictors. These require special handling internally. `is.empty.model()` checks whether an object describes an empty model.

## Usage

```
is.empty.model(x)
```

## Arguments

`x`                    A **terms** object or an object with a **terms** method.

## Value

TRUE if the model is empty

## See Also

`lm`, `glm`

## Examples

```
y <- rnorm(20)
is.empty.model(y ~ 0)
is.empty.model(y ~ -1)
is.empty.model(lm(y ~ 0))
```

---

**labels**     *Find Labels from Object*

---

**Description**

Find a suitable set of labels from an object for use in printing or plotting, for example. A generic function.

**Usage**

```
labels(object, ...)
```

**Arguments**

<b>object</b>	Any R object: the function is generic.
<b>...</b>	further arguments passed to or from other methods.

**Value**

A character vector or list of such vectors. For a vector the results is the names or `seq(along=x)`, for a data frame or array it is the `dimnames` (with `NULL` expanded to `seq(len=d[i])`), for a `terms` object it is the term labels and for an `lm` object it is the term labels for estimable terms.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.



---

**lm**     *Fitting Linear Models*

---

**Description**

`lm` is used to fit linear models. It can be used to carry out regression, single stratum analysis of variance and analysis of covariance (although `aov` may provide a more convenient interface for these).

**Usage**

```
lm(formula, data, subset, weights, na.action,  
   method = "qr", model = TRUE, x = FALSE, y = FALSE,  
   qr = TRUE, singular.ok = TRUE, contrasts = NULL,  
   offset = NULL, ...)
```

**Arguments**

<code>formula</code>	a symbolic description of the model to be fit. The details of model specification are given below.
<code>data</code>	an optional data frame containing the variables in the model. By default the variables are taken from <code>environment(formula)</code> , typically the environment from which <code>lm</code> is called.
<code>subset</code>	an optional vector specifying a subset of observations to be used in the fitting process.
<code>weights</code>	an optional vector of weights to be used in the fitting process. If specified, weighted least squares is used with weights <code>weights</code> (that is, minimizing $\sum(w \cdot e^2)$ ); otherwise ordinary least squares is used.
<code>na.action</code>	a function which indicates what should happen when the data contain NAs. The default is set by the <code>na.action</code> setting of <code>options</code> , and is <code>na.fail</code> if that is unset. The “factory-fresh” default is <code>na.omit</code> .
<code>method</code>	the method to be used; for fitting, currently only <code>method="qr"</code> is supported; <code>method="model.frame"</code> returns the model frame (the same as with <code>model = TRUE</code> , see below).
<code>model, x, y, qr</code>	logicals. If <code>TRUE</code> the corresponding components of the fit (the model frame, the model matrix, the response, the QR decomposition) are returned.

<b>singular.ok</b>	logical. If <b>FALSE</b> (the default in S but not in R) a singular fit is an error.
<b>contrasts</b>	an optional list. See the <b>contrasts.arg</b> of <b>model.matrix.default</b> .
<b>offset</b>	this can be used to specify an <i>a priori</i> known component to be included in the linear predictor during fitting. An <b>offset</b> term can be included in the formula instead or as well, and if both are specified their sum is used.
<b>...</b>	additional arguments to be passed to the low level regression fitting functions (see below).

## Details

Models for **lm** are specified symbolically. A typical model has the form **response ~ terms** where **response** is the (numeric) response vector and **terms** is a series of terms which specifies a linear predictor for **response**. A terms specification of the form **first + second** indicates all the terms in **first** together with all the terms in **second** with duplicates removed. A specification of the form **first:second** indicates the set of terms obtained by taking the interactions of all terms in **first** with all terms in **second**. The specification **first\*second** indicates the *cross* of **first** and **second**. This is the same as **first + second + first:second**. If **response** is a matrix, a linear model is fitted to each column of the matrix. See **model.matrix** for some further details.

**lm** calls the lower level functions **lm.fit**, etc, see below, for the actual numerical computations. For programming only, you may consider doing likewise.

## Value

**lm** returns an object of class **"lm"** or for multiple responses of class **c("mlm", "lm")**.

The functions **summary** and **anova** are used to obtain and print a summary and analysis of variance table of the results. The generic accessor functions **coefficients**, **effects**, **fitted.values** and **residuals** extract various useful features of the value returned by **lm**.

An object of class **"lm"** is a list containing at least the following components:

<b>coefficients</b>	a named vector of coefficients
<b>residuals</b>	the residuals, that is response minus fitted values.

<code>fitted.values</code>	the fitted mean values.
<code>rank</code>	the numeric rank of the fitted linear model.
<code>weights</code>	(only for weighted fits) the specified weights.
<code>df.residual</code>	the residual degrees of freedom.
<code>call</code>	the matched call.
<code>terms</code>	the <code>terms</code> object used.
<code>contrasts</code>	(only where relevant) the contrasts used.
<code>xlevels</code>	(only where relevant) a record of the levels of the factors used in fitting.
<code>y</code>	if requested, the response used.
<code>x</code>	if requested, the model matrix used.
<code>model</code>	if requested (the default), the model frame used.

In addition, non-null fits will have components `assign`, `effects` and (unless not requested) `qr` relating to the linear fit, for use by extractor functions such as `summary` and `effects`.

## Note

Offsets specified by `offset` will not be included in predictions by `predict.lm`, whereas those specified by an offset term in the formula will be.

## Author(s)

The design was inspired by the S function of the same name described in Chambers (1992). The implementation of model formula by Ross Ihaka was based on Wilkinson & Rogers (1973).

## References

- Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.
- Wilkinson, G. N. and Rogers, C. E. (1973) Symbolic descriptions of factorial models for analysis of variance. *Applied Statistics*, **22**, 392–9.

## See Also

`summary.lm` for summaries and `anova.lm` for the ANOVA table; `aov` for a different interface.

The generic functions `coefficients`, `effects`, `residuals`, `fitted.values`.

`predict.lm` (via `predict`) for prediction, including confidence and prediction intervals.

`lm.influence` for regression diagnostics, and `glm` for **generalized** linear models.

The underlying low level functions, `lm.fit` for plain, and `lm.wfit` for weighted regression fitting.

## Examples

```
## Annette Dobson (1990) "An Introduction to Generalized
## Linear Models". Page 9: Plant Weight Data.
ctl <- c(4.17,5.58,5.18,6.11,4.50,4.61,5.17,4.53,5.33,5.14)
trt <- c(4.81,4.17,4.41,3.59,5.87,3.83,6.03,4.89,4.32,4.69)
group <- gl(2,10,20, labels=c("Ctl","Trt"))
weight <- c(ctl, trt)
anova(lm.D9 <- lm(weight ~ group))
# omitting intercept
summary(lm.D90 <- lm(weight ~ group - 1))
# residuals almost identical
summary(resid(lm.D9) - resid(lm.D90))

opar <- par(mfrow = c(2,2), oma = c(0, 0, 1.1, 0))
plot(lm.D9, las = 1)      # Residuals, Fitted, ...
par(opar)

## model frame :
stopifnot(
  identical(lm(weight ~ group, method = "model.frame"),
    model.frame(lm.D9))
)
```

---

**lm.fit**      *Fitter Functions for Linear Models*

---

**Description**

These are the basic computing engines called by **lm** used to fit linear models. These should usually *not* be used directly unless by experienced users.

**Usage**

```
lm.fit (x, y,      offset = NULL, method = "qr", tol = 1e-7,  
        singular.ok = TRUE, ...)
```

```
lm.wfit(x, y, w, offset = NULL, method = "qr", tol = 1e-7,  
        singular.ok = TRUE, ...)
```

**Arguments**

<b>x</b>	design matrix of dimension <b>n * p</b> .
<b>y</b>	vector of observations of length <b>n</b> .
<b>w</b>	vector of weights (length <b>n</b> ) to be used in the fitting process for the <b>wfit</b> functions. Weighted least squares is used with weights <b>w</b> , i.e., $\text{sum}(w * e^2)$ is minimized.
<b>offset</b>	numeric of length <b>n</b> ). This can be used to specify an <i>a priori</i> known component to be included in the linear predictor during fitting.
<b>method</b>	currently, only <b>method="qr"</b> is supported.
<b>tol</b>	tolerance for the <b>qr</b> decomposition. Default is 1e-7.
<b>singular.ok</b>	logical. If <b>FALSE</b> , a singular model is an error.
<b>...</b>	currently disregarded.

**Details**

The functions **lm.{w}fit.null** are called by **lm.fit** or **lm.wfit** respectively, when **x** has zero columns.

**Value**

a list with components

<code>coefficients</code>	<code>p</code> vector
<code>residuals</code>	<code>n</code> vector
<code>fitted.values</code>	<code>n</code> vector
<code>effects</code>	(not null fits) <code>n</code> vector of orthogonal single-df effects. The first <b>rank</b> of them correspond to non-aliased coefficients, and are named accordingly.
<code>weights</code>	<code>n</code> vector — <i>only</i> for the <i>*wfit*</i> functions.
<code>rank</code>	integer, giving the rank
<code>df.residual</code>	degrees of freedom of residuals
<code>qr</code>	(not null fits) the QR decomposition, see <code>qr</code> .

**See Also**

`lm` which you should use for linear least squares regression, unless you know better.

**Examples**

```
set.seed(129)
n <- 7 ; p <- 2
X <- matrix(rnorm(n * p), n,p) # no intercept!
y <- rnorm(n)
w <- rnorm(n)^2

str(lmw <- lm.wfit(x=X, y=y, w=w))

str(lm. <- lm.fit (x=X, y=y))
```

---

**lm.influence**      *Regression Diagnostics*

---

**Description**

This function provides the basic quantities which are used in forming a wide variety of diagnostics for checking the quality of regression fits.

**Usage**

```
influence(model, ...)
## S3 method for class 'lm':
influence(model, do.coef = TRUE, ...)
## S3 method for class 'glm':
influence(model, do.coef = TRUE, ...)

lm.influence(model, do.coef = TRUE)
```

**Arguments**

<code>model</code>	an object as returned by <code>lm</code> .
<code>do.coef</code>	logical indicating if the changed <b>coefficients</b> (see below) are desired. These need $O(n^2p)$ computing time.
<code>...</code>	further arguments passed to or from other methods.

**Details**

The `influence.measures()` and other functions listed in **See Also** provide a more user oriented way of computing a variety of regression diagnostics. These all build on `lm.influence`.

An attempt is made to ensure that computed hat values that are probably one are treated as one, and the corresponding rows in **sigma** and **coefficients** are NaN. (Dropping such a case would normally result in a variable being dropped, so it is not possible to give simple drop-one diagnostics.)

**Value**

A list containing the following components of the same length or number of rows  $n$ , which is the number of non-zero weights. Cases omitted in the fit are omitted unless a **na.action** method was used (such as **na.exclude**) which restores them.

<b>hat</b>	a vector containing the diagonal of the “hat” matrix.
<b>coefficients</b>	(unless <code>do.coef</code> is false) a matrix whose i-th row contains the change in the estimated coefficients which results when the i-th case is dropped from the regression. Note that aliased coefficients are not included in the matrix.
<b>sigma</b>	a vector whose i-th element contains the estimate of the residual standard deviation obtained when the i-th case is dropped from the regression.
<b>wt.res</b>	a vector of <i>weighted</i> (or for class <code>glm</code> rather <i>deviance</i> ) residuals.

## Note

The `coefficients` returned by the R version of `lm.influence` differ from those computed by S. Rather than returning the coefficients which result from dropping each case, we return the changes in the coefficients. This is more directly useful in many diagnostic measures.

Since these need  $O(n^2p)$  computing time, they can be omitted by `do.coef = FALSE`.

Note that cases with `weights == 0` are *dropped* (contrary to the situation in S).

If a model has been fitted with `na.action=na.exclude` (see `na.exclude`), cases excluded in the fit *are* considered here.

## References

See the list in the documentation for `influence.measures`.

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`summary.lm` for `summary` and related methods;  
`influence.measures`,  
`hat` for the hat matrix diagonals,  
`dfbetas`, `dfhits`, `covratio`, `cooks.distance`, `lm`.

## Examples

```
## Analysis of the life-cycle savings data given in
## Belsley, Kuh and Welsch.
data(LifeCycleSavings)
```



```
summary(lm.SR <- lm(sr ~ pop15 + pop75 + dpi + ddpi,
                    data = LifeCycleSavings),
        corr = TRUE)
str(lmI <- lm.influence(lm.SR))

## For more "user level" examples, use
## example(influence.measures)
```

---

**lm.summaries**      *Accessing Linear Model Fits*

---

## Description

All these functions are **methods** for class "lm" objects.

## Usage

```
## S3 method for class 'lm':  
family(object, ...)  
  
## S3 method for class 'lm':  
formula(x, ...)  
  
## S3 method for class 'lm':  
residuals(object, type = c("working", "response",  
    "deviance", "pearson", "partial"), ...)  
  
weights(object, ...)
```

## Arguments

<b>object, x</b>	an object inheriting from class <b>lm</b> , usually the result of a call to <b>lm</b> or <b>aov</b> .
<b>...</b>	further arguments passed to or from other methods.
<b>type</b>	the type of residuals which should be returned.

## Details

The generic accessor functions **coef**, **effects**, **fitted** and **residuals** can be used to extract various useful features of the value returned by **lm**.

The working and response residuals are “observed - fitted”. The deviance and pearson residuals are weighted residuals, scaled by the square root of the weights used in fitting. The partial residuals are a matrix with each column formed by omitting a term from the model. In all these, zero weight cases are never omitted (as opposed to the standardized **rstudent** residuals).

## References

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

The model fitting function `lm`, `anova.lm`.

`coef`, `deviance`, `df.residual`, `effects`, `fitted`, `glm` for **generalized** linear models, `influence` (etc on that page) for regression diagnostics, `weighted.residuals`, `residuals`, `residuals.glm`, `summary.lm`.

## Examples

```
## Continuing the lm(.) example:
coef(lm.D90) # the bare coefficients

## The 2 basic regression diagnostic plots [plot.lm(.) is
## preferred]
plot(resid(lm.D90), fitted(lm.D90)) # Tukey-Anscombe's
abline(h=0, lty=2, col = 'gray')

qqnorm(residuals(lm.D90))
```

---

**logLik**     *Extract Log-Likelihood*

---

**Description**

This function is generic; method functions can be written to handle specific classes of objects. Classes which already have methods for this function include: **glm**, **lm**, **nls** in package **nls** and **gls**, **lme** and others in package **nlme**.

**Usage**

```
logLik(object, ...)
```

```
## S3 method for class 'logLik':  
as.data.frame(x, row.names = NULL, optional = FALSE)
```

**Arguments**

<b>object</b>	any object from which a log-likelihood value, or a contribution to a log-likelihood value, can be extracted.
<b>...</b>	some methods for this generic function require additional arguments.
<b>x</b>	an object of class <b>logLik</b> .
<b>row.names</b> , <b>optional</b>	arguments to the <b>as.data.frame</b> method; see its documentation.

**Value**

Returns an object, say **r**, of class **logLik** which is a number with attributes, **attr(r, "df")** (**d**egrees of **f**reedom) giving the number of parameters in the model. There's a simple **print** method for **logLik** objects.

The details depend on the method function used; see the appropriate documentation.

**Author(s)**

Jose Pinheiro and Douglas Bates

**See Also**

`logLik.lm`, `logLik.glm`, `logLik.gls`, `logLik.lme`, etc.

**Examples**

```
## see the method function documentation
x <- 1:5
lmx <- lm(x ~ 1)
logLik(lmx) # using print.logLik() method
str(logLik(lmx))
```

---

<code>logLik.glm</code>	<i>Extract Log-Likelihood from an glm Object</i>
-------------------------	--

---

## Description

Returns the log-likelihood value of the generalized linear model represented by `object` evaluated at the estimated coefficients.

## Usage

```
## S3 method for class 'glm':  
logLik(object, ...)
```

## Arguments

<code>object</code>	an object inheriting from class "glm".
<code>...</code>	further arguments to be passed to or from methods.

## Details

As a `family` does not have to specify how to calculate the log-likelihood, this is based on the family's function to compute the AIC. For `gaussian`, `Gamma` and `inverse.gaussian` families it assumed that the dispersion of the GLM is estimated and has been included in the AIC, and for all other families it is assumed that the dispersion is known.

Not that this procedure is not completely accurate for the gamma and inverse gaussian families, as the estimate of dispersion used is not the MLE.

## Value

the log-likelihood of the linear model represented by `object` evaluated at the estimated coefficients.

## See Also

`glm`, `logLik.lm`

---

**logLik.lm**     *Extract Log-Likelihood from an lm Object*

---

**Description**

If `REML = FALSE`, returns the log-likelihood value of the linear model represented by `object` evaluated at the estimated coefficients; else, the restricted log-likelihood evaluated at the estimated coefficients is returned.

**Usage**

```
## S3 method for class 'lm':  
logLik(object, REML = FALSE, ...)
```

**Arguments**

<code>object</code>	an object inheriting from class "lm".
<code>REML</code>	an optional logical value. If <code>TRUE</code> the restricted log-likelihood is returned, else, if <code>FALSE</code> , the log-likelihood is returned. Defaults to <code>FALSE</code> .
<code>...</code>	further arguments to be passed to or from methods.

**Value**

an object of class `logLik`, the (restricted) log-likelihood of the linear model represented by `object` evaluated at the estimated coefficients. Note that error variance  $\sigma^2$  is estimated in `lm()` and hence counted as well.

**Author(s)**

Jose Pinheiro and Douglas Bates

**References**

Harville, D.A. (1974). Bayesian inference for variance components using only error contrasts. *Biometrika*, **61**, 383–385.

**See Also**

`lm`

**Examples**

```
data(attitude)
(fm1 <- lm(rating ~ ., data = attitude))
logLik(fm1)
logLik(fm1, REML = TRUE)

res <- try(data(Orthodont, package="nlme"))
if(!inherits(res, "try-error")) {
  fm1 <- lm(distance ~ Sex * age, Orthodont)
  print(logLik(fm1))
  print(logLik(fm1, REML = TRUE))
}
```



---

**loglin**     *Fitting Log-Linear Models*

---

**Description**

**loglin** is used to fit log-linear models to multidimensional contingency tables by Iterative Proportional Fitting.

**Usage**

```
loglin(table, margin, start = rep(1, length(table)),  
       fit = FALSE, eps = 0.1, iter = 20, param = FALSE,  
       print = TRUE)
```

**Arguments**

- |               |  |
|---------------|--|
| <b>table</b>  | a contingency table to be fit, typically the output from <b>table</b> .  |
| <b>margin</b> | <p>a list of vectors with the marginal totals to be fit.</p> <p>(Hierarchical) log-linear models can be specified in terms of these marginal totals which give the “maximal” factor subsets contained in the model. For example, in a three-factor model, <code>list(c(1, 2), c(1, 3))</code> specifies a model which contains parameters for the grand mean, each factor, and the 1-2 and 1-3 interactions, respectively (but no 2-3 or 1-2-3 interaction), i.e., a model where factors 2 and 3 are independent conditional on factor 1 (sometimes represented as ‘[12][13]’).</p> <p>The names of factors (i.e., <code>names(dimnames(table))</code>) may be used rather than numeric indices.</p> |
| <b>start</b>  | a starting estimate for the fitted table. This optional argument is important for incomplete tables with structural zeros in <b>table</b> which should be preserved in the fit. In this case, the corresponding entries in <b>start</b> should be zero and the others can be taken as one.   |
| <b>fit</b>    | a logical indicating whether the fitted values should be returned.   |
| <b>eps</b>    | maximum deviation allowed between observed and fitted margins.   |

<b>iter</b>	maximum number of iterations.
<b>param</b>	a logical indicating whether the parameter values should be returned.
<b>print</b>	a logical. If <b>TRUE</b> , the number of iterations and the final deviation are printed.

## Details

The Iterative Proportional Fitting algorithm as presented in Haberman (1972) is used for fitting the model. At most **iter** iterations are performed, convergence is taken to occur when the maximum deviation between observed and fitted margins is less than **eps**. All internal computations are done in double precision; there is no limit on the number of factors (the dimension of the table) in the model.

Assuming that there are no structural zeros, both the Likelihood Ratio Test and Pearson test statistics have an asymptotic chi-squared distribution with **df** degrees of freedom.

Package **MASS** contains **loglm**, a front-end to **loglin** which allows the log-linear model to be specified and fitted in a formula-based manner similar to that of other fitting functions such as **lm** or **glm**.

## Value

A list with the following components.

<b>lrt</b>	the Likelihood Ratio Test statistic.
<b>pearson</b>	the Pearson test statistic (X-squared).
<b>df</b>	the degrees of freedom for the fitted model. There is no adjustment for structural zeros.
<b>margin</b>	list of the margins that were fit. Basically the same as the input <b>margin</b> , but with numbers replaced by names where possible.
<b>fit</b>	An array like <b>table</b> containing the fitted values. Only returned if <b>fit</b> is <b>TRUE</b> .
<b>param</b>	A list containing the estimated parameters of the model. The “standard” constraints of zero marginal sums (e.g., zero row and column sums for a two factor parameter) are employed. Only returned if <b>param</b> is <b>TRUE</b> .

## Author(s)

Kurt Hornik

## References

- Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.
- Haberman, S. J. (1972) Log-linear fit for contingency tables—Algorithm AS51. *Applied Statistics*, **21**, 218–225.
- Agresti, A. (1990) *Categorical data analysis*. New York: Wiley.

## See Also

table

## Examples

```
data(HairEyeColor)
## Model of joint independence of sex from hair and eye
## color.
fm <- loglin(HairEyeColor, list(c(1, 2), c(1, 3), c(2, 3)))
fm
1 - pchisq(fm$lrt, fm$df)
## Model with no three-factor interactions fits well.
```

---

`ls.diag`*Compute Diagnostics for ‘lsfit’ Regression Results*

---

**Description**

Computes basic statistics, including standard errors, t- and p-values for the regression coefficients.

**Usage**

`ls.diag(ls.out)`

**Arguments**

`ls.out`                Typically the result of `lsfit()`

**Value**

A list with the following numeric components.

<code>std.dev</code>	The standard deviation of the errors, an estimate of $\sigma$ .
<code>hat</code>	diagonal entries $h_{ii}$ of the hat matrix $H$
<code>std.res</code>	standardized residuals
<code>stud.res</code>	studentized residuals
<code>cooks</code>	Cook’s distances
<code>dfits</code>	DFITS statistics
<code>correlation</code>	correlation matrix
<code>std.err</code>	standard errors of the regression coefficients
<code>cov.scaled</code>	Scaled covariance matrix of the coefficients
<code>cov.unscaled</code>	Unscaled covariance matrix of the coefficients

**References**

Belsley, D. A., Kuh, E. and Welsch, R. E. (1980) *Regression Diagnostics*. New York: Wiley.

**See Also**

`hat` for the hat matrix diagonals, `ls.print`, `lm.influence`, `summary.lm`, `anova`.

**Examples**

```
## Using the same data as the lm(.) example:
lsD9 <- lsfit(x = as.numeric(gl(2, 10, 20)), y = weight)
dlsD9 <- ls.diag(lsD9)
str(dlsD9, give.attr=FALSE)
# sum(h.ii) = p
abs(1 - sum(dlsD9$hat) / 2) < 10*.Machine$double.eps
plot(dlsD9$hat, dlsD9$stud.res, xlim=c(0,0.11))
abline(h = 0, lty = 2, col = "lightgray")
```

---

<b>ls.print</b>	<i>Print 'lsfit' Regression Results</i>
-----------------	---

---

### Description

Computes basic statistics, including standard errors, t- and p-values for the regression coefficients and prints them if `print.it` is TRUE.

### Usage

```
ls.print(ls.out, digits = 4, print.it = TRUE)
```

### Arguments

<code>ls.out</code>	Typically the result of <code>lsfit()</code>
<code>digits</code>	The number of significant digits used for printing
<code>print.it</code>	a logical indicating whether the result should also be printed

### Value

A list with the components

<code>summary</code>	The ANOVA table of the regression
<code>coef.table</code>	matrix with regression coefficients, standard errors, t- and p-values

### Note

Usually, you'd rather use `summary(lm(...))` and `anova(lm(...))` for obtaining similar output.

### See Also

`ls.diag`, `lsfit`, also for examples; `lm`, `lm.influence` which usually are preferable.

---

**lsfit**     *Find the Least Squares Fit*

---

**Description**

The least squares estimate of  $\beta$  in the model

$$\mathbf{Y} = \mathbf{X}\beta + \epsilon$$

is found.

**Usage**

```
lsfit(x, y, wt=NULL, intercept=TRUE, tolerance=1e-07,  
      yname=NULL)
```

**Arguments**

<b>x</b>	a matrix whose rows correspond to cases and whose columns correspond to variables.
<b>y</b>	the responses, possibly a matrix if you want to fit multiple left hand sides.
<b>wt</b>	an optional vector of weights for performing weighted least squares.
<b>intercept</b>	whether or not an intercept term should be used.
<b>tolerance</b>	the tolerance to be used in the matrix decomposition.
<b>yname</b>	names to be used for the response variables.

**Details**

If weights are specified then a weighted least squares is performed with the weight given to the  $j$ th case specified by the  $j$ th entry in **wt**.

If any observation has a missing value in any field, that observation is removed before the analysis is carried out. This can be quite inefficient if there is a lot of missing data.

The implementation is via a modification of the LINPACK subroutines which allow for multiple left-hand sides.

## Value

A list with the following named components:

<code>coef</code>	the least squares estimates of the coefficients in the model ( $\beta$ as stated above).
<code>residuals</code>	residuals from the fit.
<code>intercept</code>	indicates whether an intercept was fitted.
<code>qr</code>	the QR decomposition of the design matrix.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`lm` which usually is preferable; `ls.print`, `ls.diag`.

## Examples

```
## Using the same data as the lm(.) example:
lsD9 <- lsfit(x = unclass(gl(2,10)), y = weight)
ls.print(lsD9)
```



---

**make.link**     *Create a Link for GLM families*

---

## Description

This function is used with the **family** functions in **glm()**. Given a link, it returns a link function, an inverse link function, the derivative  $d\mu/d\eta$  and a function for domain checking.

## Usage

```
make.link(link)
```

## Arguments

<b>link</b>	character or numeric; one of "logit", "probit", "cloglog", "identity", "log", "sqrt", "1/mu^2", "inverse", or number, say $\lambda$ resulting in power link $= \mu^\lambda$ .
-------------	---

## Value

A list with components

<b>linkfun</b>	Link function <b>function(mu)</b>
<b>linkinv</b>	Inverse link function <b>function(eta)</b>
<b>mu.eta</b>	Derivative <b>function(eta)</b> $d\mu/d\eta$
<b>valideta</b>	<b>function(eta)</b> { TRUE if all of <b>eta</b> is in the domain of <b>linkinv</b> }.

## See Also

**glm**, **family**.

## Examples

```
str(make.link("logit"))

l2 <- make.link(2)
l2$linkfun(0:3) # 0 1 4 9
l2$mu.eta(eta= 1:2) # = 1/(2*sqrt(eta))
```

---

**makepredictcall**      *Utility Function for Safe Prediction*

---

## Description

A utility to help `model.frame.default` create the right matrices when predicting from models with terms like `poly` or `ns`.

## Usage

```
makepredictcall(var, call)
```

## Arguments

<code>var</code>	A variable.
<code>call</code>	The term in the formula, as a call.

## Details

This is a generic function with methods for `poly`, `bs` and `ns`: the default method handles `scale`. If `model.frame.default` encounters such a term when creating a model frame, it modifies the `predvars` attribute of the terms supplied to replace the term with one that will work for predicting new data. For example `makepredictcall.ns` adds arguments for the knots and intercept.

To make use of this, have your model-fitting function return the `terms` attribute of the model frame, or copy the `predvars` attribute of the `terms` attribute of the model frame to your `terms` object.

To extend this, make sure the term creates variables with a class, and write a suitable method for that class.

## Value

A replacement for `call` for the `predvars` attribute of the terms.

## See Also

`model.frame`, `poly`, `scale`, `bs`, `ns`, `cars`

**Examples**

```
## using poly: this did not work in R < 1.5.0
data(women)
fm <- lm(weight ~ poly(height, 2), data = women)
plot(women, xlab = "Height (in)", ylab = "Weight (lb)")
ht <- seq(57, 73, len = 200)
lines(ht, predict(fm, data.frame(height=ht)))

## see also example(cars)
## see bs and ns for spline examples.
```

---

**manova**     *Multivariate Analysis of Variance*

---

**Description**

A class for the multivariate analysis of variance.

**Usage**

```
manova(...)
```

**Arguments**

...                      Arguments to be passed to `aov`.

**Details**

Class "manova" differs from class "aov" in selecting a different `summary` method. Function `manova` calls `aov` and then add class "manova" to the result object for each stratum.

**Value**

See `aov` and the comments in Details here.

**Note**

`manova` does not support multistratum analysis of variance, so the formula should not include an **Error** term.

**References**

- Krzanowski, W. J. (1988) *Principles of Multivariate Analysis. A User's Perspective*. Oxford.
- Hand, D. J. and Taylor, C. C. (1987) *Multivariate Analysis of Variance and Repeated Measures*. Chapman and Hall.

**See Also**

`aov`, `summary.manova`, the latter containing examples.

---

<code>model.extract</code>	<i>Extract Components from a Model Frame</i>
----------------------------	--

---

## Description

Returns the response, offset, subset, weights or other special components of a model frame passed as optional arguments to `model.frame`.

## Usage

```
model.extract(frame, component)
model.offset(x)
model.response(data, type = "any")
model.weights(x)
```

## Arguments

<code>frame</code> , <code>x</code> , <code>data</code>	A model frame.
<code>component</code>	literal character string or name. The name of a component to extract, such as <code>"weights"</code> , <code>"subset"</code> .
<code>type</code>	One of <code>"any"</code> , <code>"numeric"</code> , <code>"double"</code> . Using the either of latter two coerces the result to have storage mode <code>"double"</code> .

## Details

`model.extract` is provided for compatibility with S, which does not have the more specific functions.

`model.offset` and `model.response` are equivalent to `model.frame(, "offset")` and `model.frame(, "response")` respectively.

`model.weights` is slightly different from `model.frame(, "weights")` in not naming the vector it returns.

## Value

The specified component of the model frame, usually a vector.

## See Also

`model.frame`, `offset`

**Examples**

```
data(esoph)
a <-
  model.frame(cbind(ncases,ncontrols) ~ agegp+tobgp+alcgp,
              data=esoph)
model.extract(a, "response")
stopifnot(
  model.extract(a, "response") == model.response(a)
)

a <-
  model.frame(ncases/(ncases+ncontrols) ~ agegp+tobgp+alcgp,
              data = esoph, weights = ncases+ncontrols)
model.response(a)
model.extract(a, "weights")

a <- model.frame(cbind(ncases,ncontrols) ~ agegp,
                 something = tobgp, data = esoph)
names(a)
stopifnot(model.extract(a, "something") == esoph$tobgp)
```

---

<code>model.frame</code>	<i>Extracting the “Environment” of a Model Formula</i>
--------------------------	--

---

## Description

`model.frame` (a generic function) and its methods return a `data.frame` with the variables needed to use `formula` and any ... arguments.

## Usage

```
model.frame(formula, ...)  
  
## Default S3 method:  
model.frame(formula, data = NULL,  
             subset = NULL, na.action = na.fail,  
             drop.unused.levels = FALSE, xlev = NULL, ...)  
  
## S3 method for class 'aovlist':  
model.frame(formula, data = NULL, ...)  
  
## S3 method for class 'glm':  
model.frame(formula, data, na.action, ...)  
  
## S3 method for class 'lm':  
model.frame(formula, data, na.action, ...)
```

## Arguments

<code>formula</code>	a model formula
<code>data</code>	<code>data.frame</code> , list, <code>environment</code> or object coercible to <code>data.frame</code> containing the variables in <code>formula</code> .
<code>subset</code>	a specification of the rows to be used: defaults to all rows. This can be any valid indexing vector (see <code>\$.data.frame</code> for the rows of <code>data</code> or if that is not supplied, a data frame made up of the variables used in <code>formula</code> ).
<code>na.action</code>	how NAs are treated. The default is first, any <code>na.action</code> attribute of <code>data</code> , second a <code>na.action</code> setting of <code>options</code> , and third <code>na.fail</code> if that is unset. The “factory-fresh” default is <code>na.omit</code> .

<code>drop.unused.levels</code>	should factors have unused levels dropped? Defaults to <b>FALSE</b> .
<code>xlev</code>	a named list of character vectors giving the full set of levels to be assumed for each factor.
<code>...</code>	further arguments such as <b>subset</b> , <b>offset</b> and <b>weights</b> . <b>NULL</b> arguments are treated as missing.

## Details

Variables in the formula, **subset** and in `...` are looked for first in **data** and then in the environment of **formula**: see the help for **formula()** for further details.

First all the variables needed are collected into a data frame. Then **subset** expression is evaluated, and it is used as a row index to the data frame. Then the **na.action** function is applied to the data frame (and may well add attributes). The levels of any factors in the data frame are adjusted according to the **drop.unused.levels** and **xlev** arguments.

## Value

A **data.frame** containing the variables used in **formula** plus those specified `....`

## References

Chambers, J. M. (1992) *Data for models*. Chapter 3 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

**model.matrix** for the “design matrix”, **formula** for formulas and **expand.model.frame** for **model.frame** manipulation.

## Examples

```
data(cars)
data.class(model.frame(dist ~ speed, data = cars))
```



---

**model.matrix**     *Construct Design Matrices*

---

**Description**

`model.matrix` creates a design matrix.

**Usage**

```
model.matrix(object, ...)
```

```
## Default S3 method:
```

```
model.matrix(object, data = environment(object),  
             contrasts.arg = NULL, xlev = NULL, ...)
```

**Arguments**

<code>object</code>	an object of an appropriate class. For the default method, a model formula or terms object.
<code>data</code>	a data frame created with <code>model.frame</code> .
<code>contrasts.arg</code>	A list, whose entries are contrasts suitable for input to the <code>contrasts</code> replacement function and whose names are the names of columns of <code>data</code> containing <b>factors</b> .
<code>xlev</code>	to be used as argument of <code>model.frame</code> if <code>data</code> has no <b>"terms"</b> attribute.
<code>...</code>	further arguments passed to or from other methods.

**Details**

`model.matrix` creates a design matrix from the description given in `terms(formula)`, using the data in `data` which must contain columns with the same names as would be created by a call to `model.frame(formula)` or, more precisely, by evaluating `attr(terms(formula), "variables")`. There may be other columns and the order is not important. If `contrasts` is specified it overrides the default factor coding for that variable.

In interactions, the variable whose levels vary fastest is the first one to appear in the formula (and not in the term), so in `~ a + b + b:a` the interaction will have **a** varying fastest.

By convention, if the response variable also appears on the right-hand side of the formula it is dropped (with a warning), although interactions involving the term are retained.

**Value**

The design matrix for a regression model with the specified formula and data.

**References**

Chambers, J. M. (1992) *Data for models*. Chapter 3 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

`model.frame`, `model.extract`, `terms`

**Examples**

```
data(trees)
ff <- log(Volume) ~ log(Height) + log(Girth)
str(m <- model.frame(ff, trees))
mat <- model.matrix(ff, m)

# balanced 2-way
dd <- data.frame(a = gl(3,4), b = gl(4,1,12))
options("contrasts")
model.matrix(~ a + b, dd)
model.matrix(~ a + b, dd, contrasts = list(a="contr.sum"))
model.matrix(~ a + b, dd, contrasts = list(a="contr.sum",
                                          b="contr.poly"))
m.orth <- model.matrix(~a+b, dd,
  contrasts = list(a="contr.helmert"))
crossprod(m.orth) # m.orth is ALMOST orthogonal
```

---

<b>model.tables</b>	<i>Compute Tables of Results from an Aov Model Fit</i>
---------------------	--

---

## Description

Computes summary tables for model fits, especially complex aov fits.

## Usage

```
model.tables(x, ...)  
  
## S3 method for class 'aov':  
model.tables(x, type = "effects", se = FALSE, cterms, ...)  
  
## S3 method for class 'aovlist':  
model.tables(x, type = "effects", se = FALSE, ...)
```

## Arguments

<b>x</b>	a model object, usually produced by <code>aov</code>
<b>type</b>	type of table: currently only "effects" and "means" are implemented.
<b>se</b>	should standard errors be computed?
<b>cterm</b>	A character vector giving the names of the terms for which tables should be computed. The default is all tables.
<b>...</b>	further arguments passed to or from other methods.

## Details

For `type = "effects"` give tables of the coefficients for each term, optionally with standard errors.

For `type = "means"` give tables of the mean response for each combinations of levels of the factors in a term.

## Value

An object of class `"tables.aov"`, as list which may contain components

<b>tables</b>	A list of tables for each requested term.
<b>n</b>	The replication information for each term.
<b>se</b>	Standard error information.

**Warning**

The implementation is incomplete, and only the simpler cases have been tested thoroughly.

Weighted aov fits are not supported.

**See Also**

`aov`, `proj`, `replications`, `TukeyHSD`, `se.contrast`

**Examples**

```
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)

npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)
npk.aov <- aov(yield ~ block + N*P*K, npk)
model.tables(npk.aov, "means", se=TRUE)

## as a test, not particularly sensible statistically
options(contrasts=c("contr.helmert", "contr.treatment"))
npk.aovE <- aov(yield ~ N*P*K + Error(block), npk)
model.tables(npk.aovE, se=TRUE)
model.tables(npk.aovE, "means")
```

---

**naprint**     *Adjust for Missing Values*

---

**Description**

Use missing value information to report the effects of an **na.action**.

**Usage**

```
naprint(x, ...)
```

**Arguments**

<code>x</code>	An object produced by an <b>na.action</b> function.
<code>...</code>	further arguments passed to or from other methods.

**Details**

This is a generic function, and the exact information differs by method. **naprint.omit** reports the number of rows omitted: **naprint.default** reports an empty string.

**Value**

A character string providing information on missing values, for example the number.

---

**naresid**     *Adjust for Missing Values*

---

**Description**

Use missing value information to adjust residuals and predictions.

**Usage**

```
naresid(omit, x, ...)  
napredict(omit, x, ...)
```

**Arguments**

<code>omit</code>	An object produced by an <code>na.action</code> function.
<code>x</code>	A vector, data frame, or matrix to be adjusted based upon the missing value information.
<code>...</code>	further arguments passed to or from other methods.

**Details**

These are utility functions used to allow `predict` and `resid` methods for modelling functions to compensate for the removal of NAs in the fitting process. These are used by the default `"lm"` and `"glm"` methods, and by further methods in packages **MASS**, **rpart** and **survival**.

The default methods do nothing. The default method for the `na.exclude` action is to pad the object with NAs in the correct positions to have the same number of rows as the original data frame.

Currently `naresid` and `napredict` are identical, but future methods need not be. `naresid` is used for residuals, and `napredict` for fitted values and predictions.

**Value**

These return a similar object to `x`.

**Note**

Packages **rpart** and **survival5** used to contain versions of these functions that had an `na.omit` action equivalent to that now used for `na.exclude`.

---

**nlm**     *Non-Linear Minimization*

---

**Description**

This function carries out a minimization of the function **f** using a Newton-type algorithm. See the references for details.

**Usage**

```
nlm(f, p, hessian = FALSE, typsize = rep(1, length(p)),  
    fscale = 1, print.level = 0, ndigit = 12, gradtol = 1e-6,  
    stepmax = max(1000 * sqrt(sum((p/typsize)^2)), 1000),  
    steptol = 1e-6, iterlim = 100, check.analyticals = TRUE,  
    ...)
```

**Arguments**

<b>f</b>	the function to be minimized. If the function value has an attribute called <b>gradient</b> or both <b>gradient</b> and <b>hessian</b> attributes, these will be used in the calculation of updated parameter values. Otherwise, numerical derivatives are used. <b>deriv</b> returns a function with suitable <b>gradient</b> attribute. This should be a function of a vector of the length of <b>p</b> followed by any other arguments specified in <b>dots</b> .
<b>p</b>	starting parameter values for the minimization.
<b>hessian</b>	if <b>TRUE</b> , the hessian of <b>f</b> at the minimum is returned.
<b>typsize</b>	an estimate of the size of each parameter at the minimum.
<b>fscale</b>	an estimate of the size of <b>f</b> at the minimum.
<b>print.level</b>	this argument determines the level of printing which is done during the minimization process. The default value of 0 means that no printing occurs, a value of 1 means that initial and final details are printed and a value of 2 means that full tracing information is printed.
<b>ndigit</b>	the number of significant digits in the function <b>f</b> .
<b>gradtol</b>	a positive scalar giving the tolerance at which the scaled gradient is considered close enough to zero to

terminate the algorithm. The scaled gradient is a measure of the relative change in **f** in each direction **p[i]** divided by the relative change in **p[i]**.

<b>stepmax</b>	a positive scalar which gives the maximum allowable scaled step length. <b>stepmax</b> is used to prevent steps which would cause the optimization function to overflow, to prevent the algorithm from leaving the area of interest in parameter space, or to detect divergence in the algorithm. <b>stepmax</b> would be chosen small enough to prevent the first two of these occurrences, but should be larger than any anticipated reasonable step.
<b>steptol</b>	A positive scalar providing the minimum allowable relative step length.
<b>iterlim</b>	a positive integer specifying the maximum number of iterations to be performed before the program is terminated.
<b>check.analyticals</b>	a logical scalar specifying whether the analytic gradients and Hessians, if they are supplied, should be checked against numerical derivatives at the initial parameter values. This can help detect incorrectly formulated gradients or Hessians.
<b>...</b>	additional arguments to <b>f</b> .

## Details

If a gradient or hessian is supplied but evaluates to the wrong mode or length, it will be ignored if **check.analyticals = TRUE** (the default) with a warning. The hessian is not even checked unless the gradient is present and passes the sanity checks.

From the three methods available in the original source, we always use method “1” which is line search.

## Value

A list containing the following components:

<b>minimum</b>	the value of the estimated minimum of <b>f</b> .
<b>estimate</b>	the point at which the minimum value of <b>f</b> is obtained.
<b>gradient</b>	the gradient at the estimated minimum of <b>f</b> .



<b>hessian</b>	the hessian at the estimated minimum of <b>f</b> (if requested).
<b>code</b>	an integer indicating why the optimization process terminated. <ol style="list-style-type: none"> <li><b>1:</b> relative gradient is close to zero, current iterate is probably solution.</li> <li><b>2:</b> successive iterates within tolerance, current iterate is probably solution.</li> <li><b>3:</b> last global step failed to locate a point lower than <b>estimate</b>. Either <b>estimate</b> is an approximate local minimum of the function or <b>steptol</b> is too small.</li> <li><b>4:</b> iteration limit exceeded.</li> <li><b>5:</b> maximum step size <b>stepmax</b> exceeded five consecutive times. Either the function is unbounded below, becomes asymptotic to a finite value from above in some direction or <b>stepmax</b> is too small.</li> </ol>
<b>iterations</b>	the number of iterations performed.

## References

Dennis, J. E. and Schnabel, R. B. (1983) *Numerical Methods for Unconstrained Optimization and Nonlinear Equations*. Prentice-Hall, Englewood Cliffs, NJ.

Schnabel, R. B., Koontz, J. E. and Weiss, B. E. (1985) A modular system of algorithms for unconstrained minimization. *ACM Trans. Math. Software*, **11**, 419–440.

## See Also

**optim**, **optimize** for one-dimensional minimization and **uniroot** for root finding. **deriv** to calculate analytical derivatives.

For nonlinear regression, **nls** (in package **nls**), may be of better use.

## Examples

```
f <- function(x) sum((x-1:length(x))^2)
nlm(f, c(10,10))
nlm(f, c(10,10), print.level = 2)
str(nlm(f, c(5), hessian = TRUE))

f <- function(x, a) sum((x-a)^2)
```

```
nlm(f, c(10,10), a=c(3,5))
f <- function(x, a)
{
  res <- sum((x-a)^2)
  attr(res, "gradient") <- 2*(x-a)
  res
}
nlm(f, c(10,10), a=c(3,5))

## more examples, including the use of derivatives.
demo(nlm)
```

---

<b>offset</b>	<i>Include an Offset in a Model Formula</i>
---------------	---

---

## Description

An offset is a term to be added to a linear predictor, such as in a generalised linear model, with known coefficient 1 rather than an estimated coefficient.

## Usage

```
offset(object)
```

## Arguments

<b>object</b>	An offset to be included in a model frame
---------------	---

## Value

The input value.

## See Also

`model.offset`, `model.frame`.

For examples see `glm`, `Insurance`.

---

**optim**     *General-purpose Optimization*


---

**Description**

General-purpose optimization based on Nelder–Mead, quasi-Newton and conjugate-gradient algorithms. It includes an option for box-constrained optimization and simulated annealing.

**Usage**

```
optim(par, fn, gr = NULL,
      method = c("Nelder-Mead", "BFGS", "CG", "L-BFGS-B", "SANN"),
      lower = -Inf, upper = Inf,
      control = list(), hessian = FALSE, ...)
```

**Arguments**

<b>par</b>	Initial values for the parameters to be optimized over.
<b>fn</b>	A function to be minimized (or maximized), with first argument the vector of parameters over which minimization is to take place. It should return a scalar result.
<b>gr</b>	A function to return the gradient for the "BFGS", "CG" and "L-BFGS-B" methods. If it is NULL, a finite-difference approximation will be used.  For the "SANN" method it specifies a function to generate a new candidate point. If it is NULL a default Gaussian Markov kernel is used.
<b>method</b>	The method to be used. See <b>Details</b> .
<b>lower, upper</b>	Bounds on the variables for the "L-BFGS-B" method.
<b>control</b>	A list of control parameters. See <b>Details</b> .
<b>hessian</b>	Logical. Should a numerically differentiated Hessian matrix be returned?
<b>...</b>	Further arguments to be passed to <b>fn</b> and <b>gr</b> .

## Details

By default this function performs minimization, but it will maximize if `control$fnscale` is negative.

The default method is an implementation of that of Nelder and Mead (1965), that uses only function values and is robust but relatively slow. It will work reasonably well for non-differentiable functions.

Method "BFGS" is a quasi-Newton method (also known as a variable metric algorithm), specifically that published simultaneously in 1970 by Broyden, Fletcher, Goldfarb and Shanno. This uses function values and gradients to build up a picture of the surface to be optimized.

Method "CG" is a conjugate gradients method based on that by Fletcher and Reeves (1964) (but with the option of Polak–Ribiere or Beale–Sorenson updates). Conjugate gradient methods will generally be more fragile than the BFGS method, but as they do not store a matrix they may be successful in much larger optimization problems.

Method "L-BFGS-B" is that of Byrd *et al.* (1994) which allows *box constraints*, that is each variable can be given a lower and/or upper bound. The initial value must satisfy the constraints. This uses a limited-memory modification of the BFGS quasi-Newton method. If non-trivial bounds are supplied, this method will be selected, with a warning.

Nocedal and Wright (1999) is a comprehensive reference for the previous three methods.

Method "SANN" is by default a variant of simulated annealing given in Belisle (1992). Simulated-annealing belongs to the class of stochastic global optimization methods. It uses only function values but is relatively slow. It will also work for non-differentiable functions. This implementation uses the Metropolis function for the acceptance probability. By default the next candidate point is generated from a Gaussian Markov kernel with scale proportional to the actual temperature. If a function to generate a new candidate point is given, method "SANN" can also be used to solve combinatorial optimization problems. Temperatures are decreased according to the logarithmic cooling schedule as given in Belisle (1992, p. 890). Note that the "SANN" method depends critically on the settings of the control parameters. It is not a general-purpose method but can be very useful in getting to a good value on a very rough surface.

Function `fn` can return `NA` or `Inf` if the function cannot be evaluated at the supplied value, but the initial value must have a computable finite value of `fn`. (Except for method "L-BFGS-B" where the values should always be finite.)

**optim** can be used recursively, and for a single parameter as well as many.

The **control** argument is a list that can supply any of the following components:

**trace** Non-negative integer. If positive, tracing information on the progress of the optimization is produced. Higher values may produce more tracing information: for method "L-BFGS-B" there are six levels of tracing. (To understand exactly what these do see the source code: higher levels give more detail.)

**fnscale** An overall scaling to be applied to the value of **fn** and **gr** during optimization. If negative, turns the problem into a maximization problem. Optimization is performed on **fn(par)/fnscale**.

**parscale** A vector of scaling values for the parameters. Optimization is performed on **par/parscale** and these should be comparable in the sense that a unit change in any element produces about a unit change in the scaled value.

**ndeps** A vector of step sizes for the finite-difference approximation to the gradient, on **par/parscale** scale. Defaults to **1e-3**.

**maxit** The maximum number of iterations. Defaults to 100 for the derivative-based methods, and 500 for "Nelder-Mead". For "SANN" **maxit** gives the total number of function evaluations. There is no other stopping criterion. Defaults to 10000.

**abstol** The absolute convergence tolerance. Only useful for non-negative functions, as a tolerance for reaching zero.

**reltol** Relative convergence tolerance. The algorithm stops if it is unable to reduce the value by a factor of **reltol \* (abs(val) + reltol)** at a step. Defaults to **sqrt(.Machine\$double.eps)**, typically about **1e-8**.

**alpha**, **beta**, **gamma** Scaling parameters for the "Nelder-Mead" method. **alpha** is the reflection factor (default 1.0), **beta** the contraction factor (0.5) and **gamma** the expansion factor (2.0).

**REPORT** The frequency of reports for the "BFGS" and "L-BFGS-B" methods if **control\$trace** is positive. Defaults to every 10 iterations.

**type** for the conjugate-gradients method. Takes value 1 for the Fletcher-Reeves update, 2 for Polak-Ribiere and 3 for Beale-Sorenson.

**lmm** is an integer giving the number of BFGS updates retained in the "L-BFGS-B" method. It defaults to 5.

**factr** controls the convergence of the "L-BFGS-B" method. Convergence occurs when the reduction in the objective is within this factor of the machine tolerance. Default is `1e7`, that is a tolerance of about `1e-8`.

**pgtol** helps to control the convergence of the "L-BFGS-B" method. It is a tolerance on the projected gradient in the current search direction. This defaults to zero, when the check is suppressed.

**temp** controls the "SANN" method. It is the starting temperature for the cooling schedule. Defaults to 10.

**tmax** is the number of function evaluations at each temperature for the "SANN" method. Defaults to 10.

## Value

A list with components:

<b>par</b>	The best set of parameters found.
<b>value</b>	The value of <b>fn</b> corresponding to <b>par</b> .
<b>counts</b>	A two-element integer vector giving the number of calls to <b>fn</b> and <b>gr</b> respectively. This excludes those calls needed to compute the Hessian, if requested, and any calls to <b>fn</b> to compute a finite-difference approximation to the gradient.
<b>convergence</b>	An integer code. 0 indicates successful convergence. Error codes are <ul style="list-style-type: none"> <li>1 indicates that the iteration limit <b>maxit</b> had been reached.</li> <li>10 indicates degeneracy of the Nelder–Mead simplex.</li> <li>51 indicates a warning from the "L-BFGS-B" method; see component <b>message</b> for further details.</li> <li>52 indicates an error from the "L-BFGS-B" method; see component <b>message</b> for further details.</li> </ul>
<b>message</b>	A character string giving any additional information returned by the optimizer, or <b>NULL</b> .
<b>hessian</b>	Only if argument <b>hessian</b> is true. A symmetric matrix giving an estimate of the Hessian at the solution found. Note that this is the Hessian of the unconstrained problem even if the box constraints are active.

## Note

`optim` will work with one-dimensional `pars`, but the default method does not work well (and will warn). Use `optimize` instead.

The code for methods "Nelder-Mead", "BFGS" and "CG" was based originally on Pascal code in Nash (1990) that was translated by `p2c` and then hand-optimized. Dr Nash has agreed that the code can be made freely available.

The code for method "L-BFGS-B" is based on Fortran code by Zhu, Byrd, Lu-Chen and Nocedal obtained from Netlib (file `'opt/lbfgs_bcm.shar'`: another version is in `'toms/778'`).

The code for method "SANN" was contributed by A. Trapletti.

## References

- Belisle, C. J. P. (1992) Convergence theorems for a class of simulated annealing algorithms on  $R^d$ . *J Applied Probability*, **29**, 885–895.
- Byrd, R. H., Lu, P., Nocedal, J. and Zhu, C. (1995) A limited memory algorithm for bound constrained optimization. *SIAM J. Scientific Computing*, **16**, 1190–1208.
- Fletcher, R. and Reeves, C. M. (1964) Function minimization by conjugate gradients. *Computer Journal* **7**, 148–154.
- Nash, J. C. (1990) *Compact Numerical Methods for Computers. Linear Algebra and Function Minimisation*. Adam Hilger.
- Nelder, J. A. and Mead, R. (1965) A simplex algorithm for function minimization. *Computer Journal* **7**, 308–313.
- Nocedal, J. and Wright, S. J. (1999) *Numerical Optimization*. Springer.

## See Also

`nlm`, `optimize`, `constrOptim`

## Examples

```
fr <- function(x) {    ## Rosenbrock Banana function
  x1 <- x[1]
  x2 <- x[2]
  100 * (x2 - x1 * x1)^2 + (1 - x1)^2
}
grr <- function(x) { ## Gradient of 'fr'
  x1 <- x[1]
  x2 <- x[2]
```



```

      c(-400 * x1 * (x2 - x1 * x1) - 2 * (1 - x1),
        200 *      (x2 - x1 * x1))
    }
  optim(c(-1.2,1), fr)
  optim(c(-1.2,1), fr, grr, method = "BFGS")
  optim(c(-1.2,1), fr, NULL, method = "BFGS", hessian = TRUE)
  optim(c(-1.2,1), fr, grr, method = "CG")
  optim(c(-1.2,1), fr, grr, method = "CG",
        control=list(type=2))
  optim(c(-1.2,1), fr, grr, method = "L-BFGS-B")

  flb <- function(x) {
    p <- length(x);
    sum(c(1, rep(4, p-1)) * (x - c(1, x[-p])^2)^2)
  }
  ## 25-dimensional box constrained
  # par[24] is not at boundary
  optim(rep(3, 25), flb, NULL, "L-BFGS-B",
        lower=rep(2, 25), upper=rep(4, 25))

  ## "wild" function , global minimum at about -15.81515
  fw <- function (x)
    10*sin(0.3*x)*sin(1.3*x^2) + 0.00001*x^4 + 0.2*x+80
  plot(fw, -50, 50, n=1000,
       main = "optim() minimising 'wild function'")

  res <- optim(50, fw, method="SANN",
              control=list(maxit=20000, temp=20, parscale=20))
  res
  ## Now improve locally
  (r2 <- optim(res$par, fw, method="BFGS"))
  points(r2$par, r2$val, pch = 8, col = "red", cex = 2)

  ## Combinatorial optimization: Traveling salesman problem
  library(mva) # normally loaded
  library(ts)  # for embed, normally loaded

  data(eurodist)
  eurodistmat <- as.matrix(eurodist)

  # Target function
  distance <- function(sq) {
    sq2 <- embed(sq, 2)

```

```

    return(sum(eurodistmat[cbind(sq2[,2],sq2[,1])]))
  }

# Generate new candidate sequence
genseq <- function(sq) {
  idx <- seq(2, NROW(eurodistmat)-1, by=1)
  changepoints <- sample(idx, size=2, replace=FALSE)
  tmp <- sq[changepoints[1]]
  sq[changepoints[1]] <- sq[changepoints[2]]
  sq[changepoints[2]] <- tmp
  return(sq)
}

sq <- c(1,2:NROW(eurodistmat),1) # Initial sequence
distance(sq)

set.seed(2222) # chosen to get a good soln quickly
res <- optim(sq, distance, genseq, method="SANN",
  control = list(maxit=6000, temp=2000, trace=TRUE))
res # Near optimum distance around 12842

loc <- cmdscale(eurodist)
rx <- range(x <- loc[,1])
ry <- range(y <- -loc[,2])
tspinit <- loc[sq,]
tspres <- loc[res$par,]
s <- seq(NROW(tspres)-1)

plot(x, y, type="n", asp=1, xlab="", ylab="",
  main="initial solution of traveling salesman problem")
arrows(tspinit[s,1], -tspinit[s,2], tspinit[s+1,1],
  -tspinit[s+1,2], angle=10, col="green")
text(x, y, names(eurodist), cex=0.8)

plot(x, y, type="n", asp=1, xlab="", ylab="",
  main="optim() 'solving' traveling salesman problem")
arrows(tspres[s,1], -tspres[s,2], tspres[s+1,1],
  -tspres[s+1,2], angle=10, col="red")
text(x, y, names(eurodist), cex=0.8)

```

---

**optimize**     *One Dimensional Optimization*

---

**Description**

The function **optimize** searches the interval from **lower** to **upper** for a minimum or maximum of the function **f** with respect to its first argument.

**optimise** is an alias for **optimize**.

**Usage**

```
optimize(f = , interval = , lower = min(interval),
         upper = max(interval), maximum = FALSE,
         tol = .Machine$double.eps^0.25, ...)
optimise(f = , interval = , lower = min(interval),
         upper = max(interval), maximum = FALSE,
         tol = .Machine$double.eps^0.25, ...)
```

**Arguments**

<b>f</b>	the function to be optimized. The function is either minimized or maximized over its first argument depending on the value of <b>maximum</b> .
<b>interval</b>	a vector containing the end-points of the interval to be searched for the minimum.
<b>lower</b>	the lower end point of the interval to be searched.
<b>upper</b>	the upper end point of the interval to be searched.
<b>maximum</b>	logical. Should we maximize or minimize (the default)?
<b>tol</b>	the desired accuracy.
<b>...</b>	additional arguments to <b>f</b> .

**Details**

The method used is a combination of golden section search and successive parabolic interpolation. Convergence is never much slower than that for a Fibonacci search. If **f** has a continuous second derivative which is positive at the minimum (which is not at **lower** or **upper**), then convergence is superlinear, and usually of the order of about 1.324.

The function `f` is never evaluated at two points closer together than  $\epsilon|x_0| + (tol/3)$ , where  $\epsilon$  is approximately `sqrt(.Machine$double.eps)` and  $x_0$  is the final abscissa `optimize()$minimum`.

If `f` is a unimodal function and the computed values of `f` are always unimodal when separated by at least  $\epsilon|x| + (tol/3)$ , then  $x_0$  approximates the abscissa of the global minimum of `f` on the interval `lower, upper` with an error less than  $\epsilon|x_0| + tol$ .

If `f` is not unimodal, then `optimize()` may approximate a local, but perhaps non-global, minimum to the same accuracy.

The first evaluation of `f` is always at  $x_1 = a + (1 - \phi)(b - a)$  where  $(a, b) = (\text{lower}, \text{upper})$  and  $\phi = (\sqrt{5} - 1)/2 = 0.61803..$  is the golden section ratio. Almost always, the second evaluation is at  $x_2 = a + \phi(b - a)$ . Note that a local minimum inside  $[x_1, x_2]$  will be found as solution, even when `f` is constant in there, see the last example.

It uses a C translation of Fortran code (from Netlib) based on the Algol 60 procedure `localmin` given in the reference.

## Value

A list with components `minimum` (or `maximum`) and `objective` which give the location of the minimum (or maximum) and the value of the function at that point.

## References

Brent, R. (1973) *Algorithms for Minimization without Derivatives*. Englewood Cliffs N.J.: Prentice-Hall.

## See Also

`nlm`, `uniroot`.

## Examples

```
f <- function (x,a) (x-a)^2
xmin <- optimize(f, c(0, 1), tol = 0.0001, a = 1/3)
xmin

## See where the function is evaluated:
optimize(function(x) x^2*(print(x)-1), l=0, u=10)

## "wrong" solution with unlucky interval and piecewise
## constant f():
f <- function(x)
```

```
      ifelse(x > -1, ifelse(x < 4, exp(-1/abs(x - 1)), 10), 10)
fp <- function(x) { print(x); f(x) }

plot(f, -2,5, ylim = 0:1, col = 2)
optimize(fp, c(-4, 20)) # doesn't see the minimum
optimize(fp, c(-7, 20)) # ok
```

---

**power**      *Create a Power Link Object*

---

## Description

Creates a link object based on the link function  $\eta = \mu^\lambda$ .

## Usage

```
power(lambda = 1)
```

## Arguments

**lambda**              a real number.

## Details

If **lambda** is non-negative, it is taken as zero, and the log link is obtained. The default **lambda = 1** gives the identity link.

## Value

A list with components **linkfun**, **linkinv**, **mu.eta**, and **valideta**. See **make.link** for information on their meaning.

## References

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

## See Also

**make.link**, **family**

To raise a number to a power, see **Arithmetic**.

To calculate the power of a test, see various functions in the **ctest** package, e.g., **power.t.test**.

## Examples

```
power()
quasi(link=power(1/3))[c("linkfun", "linkinv")]
```

---

**predict.glm**     *Predict Method for GLM Fits*

---

**Description**

Obtains predictions and optionally estimates standard errors of those predictions from a fitted generalized linear model object.

**Usage**

```
## S3 method for class 'glm':  
predict(object, newdata = NULL,  
        type = c("link", "response", "terms"),  
        se.fit = FALSE, dispersion = NULL, terms = NULL,  
        na.action = na.pass, ...)
```

**Arguments**

<b>object</b>	a fitted object of class inheriting from "glm".
<b>newdata</b>	optionally, a new data frame from which to make the predictions. If omitted, the fitted linear predictors are used.
<b>type</b>	the type of prediction required. The default is on the scale of the linear predictors; the alternative <b>"response"</b> is on the scale of the response variable. Thus for a default binomial model the default predictions are of log-odds (probabilities on logit scale) and <b>type = "response"</b> gives the predicted probabilities. The <b>"terms"</b> option returns a matrix giving the fitted values of each term in the model formula on the linear predictor scale. The value of this argument can be abbreviated.
<b>se.fit</b>	logical switch indicating if standard errors are required.
<b>dispersion</b>	the dispersion of the GLM fit to be assumed in computing the standard errors. If omitted, that returned by <b>summary</b> applied to the object is used.
<b>terms</b>	with <b>type="terms"</b> by default all terms are returned. A character vector specifies which terms are to be returned

<code>na.action</code>	function determining what should be done with missing values in <code>newdata</code> . The default is to predict NA.
<code>...</code>	further arguments passed to or from other methods.

## Value

If `se = FALSE`, a vector or matrix of predictions. If `se = TRUE`, a list with components

<code>fit</code>	Predictions
<code>se.fit</code>	Estimated standard errors
<code>residual.scale</code>	A scalar giving the square root of the dispersion used in computing the standard errors.

## See Also

`glm`, `SafePrediction`

## Examples

```
## example from Venables and Ripley (2002, pp. 190-2.)
ldose <- rep(0:5, 2)
numdead <- c(1, 4, 9, 13, 18, 20, 0, 2, 6, 10, 12, 16)
sex <- factor(rep(c("M", "F"), c(6, 6)))
SF <- cbind(numdead, numalive=20-numdead)
budworm.lg <- glm(SF ~ sex*ldose, family=binomial)
summary(budworm.lg)

plot(c(1,32), c(0,1), type = "n", xlab = "dose",
      ylab = "prob", log = "x")
text(2^ldose, numdead/20, as.character(sex))
ld <- seq(0, 5, 0.1)
lines(2^ld, predict(budworm.lg, data.frame(ldose=ld,
      sex=factor(rep("M", length(ld)), levels=levels(sex))),
      type = "response"))
lines(2^ld, predict(budworm.lg, data.frame(ldose=ld,
      sex=factor(rep("F", length(ld)), levels=levels(sex))),
      type = "response"))
```



---

**predict.lm**     *Predict method for Linear Model Fits*

---

**Description**

Predicted values based on linear model object

**Usage**

```
## S3 method for class 'lm':
predict(object, newdata, se.fit = FALSE, scale = NULL,
        df = Inf,
        interval = c("none", "confidence", "prediction"),
        level = 0.95, type = c("response", "terms"),
        terms = NULL, na.action = na.pass, ...)
```

**Arguments**

<b>object</b>	Object of class inheriting from "lm"
<b>newdata</b>	Data frame in which to predict
<b>se.fit</b>	A switch indicating if standard errors are required.
<b>scale</b>	Scale parameter for std.err. calculation
<b>df</b>	Degrees of freedom for scale
<b>interval</b>	Type of interval calculation
<b>level</b>	Tolerance/confidence level
<b>type</b>	Type of prediction (response or model term)
<b>terms</b>	If <b>type="terms"</b> , which terms (default is all terms)
<b>na.action</b>	function determining what should be done with missing values in <b>newdata</b> . The default is to predict NA.
<b>...</b>	further arguments passed to or from other methods.

**Details**

**predict.lm** produces predicted values, obtained by evaluating the regression function in the frame **newdata** (which defaults to **model.frame(object)**). If the logical **se.fit** is **TRUE**, standard errors of the predictions are calculated. If the numeric argument **scale** is set (with optional **df**), it is used as the residual standard deviation in the computation of the standard errors, otherwise this is extracted from the model

fit. Setting **intervals** specifies computation of confidence or prediction (tolerance) intervals at the specified **level**.

If the fit is rank-deficient, some of the columns of the design matrix will have been dropped. Prediction from such a fit only makes sense if **newdata** is contained in the same subspace as the original data. That cannot be checked accurately, so a warning is issued.

## Value

**predict.lm** produces a vector of predictions or a matrix of predictions and bounds with column names **fit**, **lwr**, and **upr** if **interval** is set. If **se.fit** is **TRUE**, a list with the following components is returned:

<b>fit</b>	vector or matrix as above
<b>se.fit</b>	standard error of predictions
<b>residual.scale</b>	residual standard deviations
<b>df</b>	degrees of freedom for residual

## Note

Offsets specified by **offset** in the fit by **lm** will not be included in predictions, whereas those specified by an offset term in the formula will be.

## See Also

The model fitting function **lm**, **predict**, **SafePrediction**

## Examples

```
## Predictions
x <- rnorm(15)
y <- x + rnorm(15)
predict(lm(y ~ x))
new <- data.frame(x = seq(-3, 3, 0.5))
predict(lm(y ~ x), new, se.fit = TRUE)
pred.w.plim <-
  predict(lm(y ~ x), new, interval="prediction")
pred.w.clim <-
  predict(lm(y ~ x), new, interval="confidence")
matplot(new$x, cbind(pred.w.clim, pred.w.plim[, -1]),
        lty=c(1,2,2,3,3), type="l", ylab="predicted y")
```

---

**profile**     *Generic Function for Profiling Models*

---

**Description**

Investigates behavior of objective function near the solution represented by `fitted`.

See documentation on method functions for further details.

**Usage**

```
profile(fitted, ...)
```

**Arguments**

<code>fitted</code>	the original fitted model object.
<code>...</code>	additional parameters. See documentation on individual methods.

**Value**

A list with an element for each parameter being profiled. See the individual methods for further details.

**See Also**

`profile.nls` in package **nls**, `profile.glm` in package **MASS**, ...

For profiling code, see **Rprof**.

---

**proj**     *Projections of Models*

---

**Description**

`proj` returns a matrix or list of matrices giving the projections of the data onto the terms of a linear model. It is most frequently used for `aov` models.

**Usage**

```
proj(object, ...)  
  
## S3 method for class 'aov':  
proj(object, onedf = FALSE, unweighted.scale = FALSE, ...)  
  
## S3 method for class 'aovlist':  
proj(object, onedf = FALSE, unweighted.scale = FALSE, ...)  
  
## Default S3 method:  
proj(object, onedf = TRUE, ...)  
  
## S3 method for class 'lm':  
proj(object, onedf = FALSE, unweighted.scale = FALSE, ...)
```

**Arguments**

<code>object</code>	An object of class <code>"lm"</code> or a class inheriting from it, or an object with a similar structure including in particular components <code>qr</code> and <code>effects</code> .
<code>onedf</code>	A logical flag. If <code>TRUE</code> , a projection is returned for all the columns of the model matrix. If <code>FALSE</code> , the single-column projections are collapsed by terms of the model (as represented in the analysis of variance table).
<code>unweighted.scale</code>	If the fit producing <code>object</code> used weights, this determines if the projections correspond to weighted or unweighted observations.
<code>...</code>	Swallow and ignore any other arguments.

## Details

A projection is given for each stratum of the object, so for `aov` models with an `Error` term the result is a list of projections.

## Value

A projection matrix or (for multi-stratum objects) a list of projection matrices.

Each projection is a matrix with a row for each observations and either a column for each term (`onedf = FALSE`) or for each coefficient (`onedf = TRUE`). Projection matrices from the default method have orthogonal columns representing the projection of the response onto the column space of the `Q` matrix from the QR decomposition. The fitted values are the sum of the projections, and the sum of squares for each column is the reduction in sum of squares from fitting that column (after those to the left of it).

The methods for `lm` and `aov` models add a column to the projection matrix giving the residuals (the projection of the data onto the orthogonal complement of the model space).

Strictly, when `onedf = FALSE` the result is not a projection, but the columns represent sums of projections onto the columns of the model matrix corresponding to that term. In this case the matrix does not depend on the coding used.

## Author(s)

The design was inspired by the `S` function of the same name described in Chambers *et al.* (1992).

## References

Chambers, J. M., Freeny, A and Heiberger, R. M. (1992) *Analysis of variance; designed experiments*. Chapter 5 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`aov`, `lm`, `model.tables`

## Examples

```
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
```

```
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)

npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)
npk.aov <- aov(yield ~ block + N*P*K, npk)
proj(npk.aov)

## as a test, not particularly sensible
options(contrasts=c("contr.helmert", "contr.treatment"))
npk.aovE <- aov(yield ~ N*P*K + Error(block), npk)
proj(npk.aovE)
```

---

**relevel**     *Reorder Levels of Factor*

---

**Description**

The levels of a factor are re-ordered so that the level specified by **ref** is first and the others are moved down. This is useful for **contr.treatment** contrasts which take the first level as the reference.

**Usage**

```
relevel(x, ref, ...)
```

**Arguments**

<b>x</b>	An unordered factor.
<b>ref</b>	The reference level.
<b>...</b>	Additional arguments for future methods.

**Value**

A factor of the same length as **x**.

**See Also**

**factor**, **contr.treatment**

**Examples**

```
data(warpbreaks)
warpbreaks$tension <- relevel(warpbreaks$tension, ref="M")
summary(lm(breaks ~ wool + tension, data=warpbreaks))
```

---

<b>replications</b>	<i>Number of Replications of Terms</i>
---------------------	--

---

## Description

Returns a vector or a list of the number of replicates for each term in the formula.

## Usage

```
replications(formula, data=NULL, na.action)
```

## Arguments

<b>formula</b>	a formula or a terms object or a data frame.
<b>data</b>	a data frame used to find the objects in <b>formula</b> .
<b>na.action</b>	function for handling missing values. Defaults to a <b>na.action</b> attribute of <b>data</b> , then a setting of the option <b>na.action</b> , or <b>na.fail</b> if that is not set.

## Details

If **formula** is a data frame and **data** is missing, **formula** is used for **data** with the formula `~ ..`.

## Value

A vector or list with one entry for each term in the formula giving the number(s) of replications for each level. If all levels are balanced (have the same number of replications) the result is a vector, otherwise it is a list with a component for each terms, as a vector, matrix or array as required.

A test for balance is `!is.list(replications(formula,data))`.

## Author(s)

The design was inspired by the S function of the same name described in Chambers *et al.* (1992).



## References

Chambers, J. M., Freeny, A and Heiberger, R. M. (1992) *Analysis of variance; designed experiments*. Chapter 5 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

## See Also

`model.tables`

## Examples

```
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)

npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)
replications(~ . - yield, npk)
```

---

**residuals**     *Extract Model Residuals*

---

**Description**

**residuals** is a generic function which extracts model residuals from objects returned by modeling functions.

The abbreviated form **resid** is an alias for **residuals**. It is intended to encourage users to access object components through an accessor function rather than by directly referencing an object slot.

All object classes which are returned by model fitting functions should provide a **residuals** method. (Note that the method is ‘**residuals**’ and not ‘**resid**’.)

Methods can make use of **naresid** methods to compensate for the omission of missing values. The default method does.

**Usage**

```
residuals(object, ...)  
resid(object, ...)
```

**Arguments**

<b>object</b>	an object for which the extraction of model residuals is meaningful.
<b>...</b>	other arguments.

**Value**

Residuals extracted from the object **object**.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical Models in S*. Wadsworth & Brooks/Cole.

**See Also**

**coefficients**, **fitted.values**, **glm**, **lm**.

---

**se.contrast**      *Standard Errors for Contrasts in Model Terms*


---

**Description**

Returns the standard errors for one or more contrasts in an **aov** object.

**Usage**

```
se.contrast(object, ...)
## S3 method for class 'aov':
se.contrast(object, contrast.obj,
             coef = contr.helmert(ncol(contrast))[, 1],
             data = NULL, ...)
```

**Arguments**

<b>object</b>	A suitable fit, usually from <b>aov</b> .
<b>contrast.obj</b>	The contrasts for which standard errors are requested. This can be specified via a list or via a matrix. A single contrast can be specified by a list of logical vectors giving the cells to be contrasted. Multiple contrasts should be specified by a matrix, each column of which is a numerical contrast vector (summing to zero).
<b>coef</b>	used when <b>contrast.obj</b> is a list; it should be a vector of the same length as the list with zero sum. The default value is the first Helmert contrast, which contrasts the first and second cell means specified by the list.
<b>data</b>	The data frame used to evaluate <b>contrast.obj</b> .
<b>...</b>	further arguments passed to or from other methods.

**Details**

Contrasts are usually used to test if certain means are significantly different; it can be easier to use **se.contrast** than compute them directly from the coefficients.

In multistratum models, the contrasts can appear in more than one stratum; the contrast and standard error are computed in the lowest stratum and adjusted for efficiencies and comparisons between strata.

Suitable matrices for use with **coef** can be found by calling **contrasts** and indexing the columns by a factor.

**Value**

A vector giving the standard errors for each contrast.

**See Also**

`contrasts`, `model.tables`

**Examples**

```
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)

npk <- data.frame(block = gl(6,4), N = factor(N),
  P = factor(P), K = factor(K), yield = yield)
options(contrasts=c("contr.treatment", "contr.poly"))
npk.aov1 <- aov(yield ~ block + N + K, npk)
se.contrast(npk.aov1, list(N=="0", N=="1"), data=npk)
# or via a matrix
cont <- matrix(c(-1,1), 2, 1, dimnames=list(NULL, "N"))
se.contrast(npk.aov1, cont[N, , drop=FALSE]/12, data=npk)

## test a multi-stratum model
npk.aov2 <- aov(yield ~ N + K + Error(block/(N + K)), npk)
se.contrast(npk.aov2, list(N == "0", N == "1"))
```

---

**stat.anova**     *GLM Anova Statistics*

---

**Description**

This is a utility function, used in `lm` and `glm` methods for `anova(..., test != NULL)` and should not be used by the average user.

**Usage**

```
stat.anova(table, test = c("Chisq", "F", "Cp"), scale,
           df.scale, n)
```

**Arguments**

<code>table</code>	numeric matrix as results from <code>anova.glm(..., test=NULL)</code> .
<code>test</code>	a character string, matching one of "Chisq", "F" or "Cp".
<code>scale</code>	a weighted residual sum of squares.
<code>df.scale</code>	degrees of freedom corresponding to scale.
<code>n</code>	number of observations.

**Value**

A matrix which is the original `table`, augmented by a column of test statistics, depending on the `test` argument.

**References**

Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

`anova.lm`, `anova.glm`.

## Examples

```
## Continued from '?glm':  
print(ag <- anova(glm.D93))  
stat.anova(ag$table, test = "Cp",  
  scale = sum(resid(glm.D93, "pearson")^2)/4, df = 4,  
  n = 9)
```

---

<b>step</b>	<i>Choose a model by AIC in a Stepwise Algorithm</i>
-------------	--

---

## Description

Select a formula-based model by AIC.

## Usage

```
step(object, scope, scale = 0,  
      direction = c("both", "backward", "forward"),  
      trace = 1, keep = NULL, steps = 1000, k = 2, ...)
```

## Arguments

<b>object</b>	an object representing a model of an appropriate class (mainly <b>lm</b> and <b>glm</b> ). This is used as the initial model in the stepwise search.
<b>scope</b>	defines the range of models examined in the stepwise search. This should be either a single formula, or a list containing components <b>upper</b> and <b>lower</b> , both formulae. See the details for how to specify the formulae and how they are used.
<b>scale</b>	used in the definition of the AIC statistic for selecting the models, currently only for <b>lm</b> , <b>aov</b> and <b>glm</b> models.
<b>direction</b>	the mode of stepwise search, can be one of <b>"both"</b> , <b>"backward"</b> , or <b>"forward"</b> , with a default of <b>"both"</b> . If the <b>scope</b> argument is missing the default for <b>direction</b> is <b>"backward"</b> .
<b>trace</b>	if positive, information is printed during the running of <b>step</b> . Larger values may give more detailed information.
<b>keep</b>	a filter function whose input is a fitted model object and the associated AIC statistic, and whose output is arbitrary. Typically <b>keep</b> will select a subset of the components of the object and return them. The default is not to keep anything.
<b>steps</b>	the maximum number of steps to be considered. The default is 1000 (essentially as many as required). It is typically used to stop the process early.

<b>k</b>	the multiple of the number of degrees of freedom used for the penalty. Only <b>k</b> = 2 gives the genuine AIC: <b>k</b> = <b>log(n)</b> is sometimes referred to as BIC or SBC.
<b>...</b>	any additional arguments to <b>extractAIC</b> .

## Details

**step** uses **add1** and **drop1** repeatedly; it will work for any method for which they work, and that is determined by having a valid method for **extractAIC**. When the additive constant can be chosen so that AIC is equal to Mallows'  $C_p$ , this is done and the tables are labelled appropriately.

The set of models searched is determined by the **scope** argument. The right-hand-side of its **lower** component is always included in the model, and right-hand-side of the model is included in the **upper** component. If **scope** is a single formula, it specifies the **upper** component, and the **lower** model is empty. If **scope** is missing, the initial model is used as the **upper** model.

Models specified by **scope** can be templates to update **object** as used by **update.formula**.

There is a potential problem in using **glm** fits with a variable **scale**, as in that case the deviance is not simply related to the maximized log-likelihood. The function **extractAIC.glm** makes the appropriate adjustment for a **gaussian** family, but may need to be amended for other cases. (The **binomial** and **poisson** families have fixed **scale** by default and do not correspond to a particular maximum-likelihood problem for variable **scale**.)

## Value

the stepwise-selected model is returned, with up to two additional components. There is an **"anova"** component corresponding to the steps taken in the search, as well as a **"keep"** component if the **keep=** argument was supplied in the call. The **"Resid. Dev"** column of the analysis of deviance table refers to a constant minus twice the maximized log likelihood: it will be a deviance only in cases where a saturated model is well-defined (thus excluding **lm**, **aov** and **survreg** fits, for example).

## Warning

The model fitting must apply the models to the same dataset. This may be a problem if there are missing values and R's default of **na.action** = **na.omit** is used. We suggest you remove the missing values first.



## Note

This function differs considerably from the function in S, which uses a number of approximations and does not compute the correct AIC.

This is a minimal implementation. Use **stepAIC** for a wider range of object classes.

## Author(s)

B. D. Ripley: **step** is a slightly simplified version of **stepAIC** in package **MASS** (Venables & Ripley, 2002 and earlier editions).

The idea of a **step** function follows that described in Hastie & Pregibon (1992); but the implementation in R is more general.

## References

Hastie, T. J. and Pregibon, D. (1992) *Generalized linear models*. Chapter 6 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

Venables, W. N. and Ripley, B. D. (2002) *Modern Applied Statistics with S*. New York: Springer (4th ed).

## See Also

**stepAIC**, **add1**, **drop1**

## Examples

```
example(lm)
step(lm.D9)

data(swiss)
summary(lm1 <- lm(Fertility ~ ., data = swiss))
slm1 <- step(lm1)
summary(slm1)
slm1$anova
```

---

<code>summary.aov</code>	<i>Summarize an Analysis of Variance Model</i>
--------------------------	--

---

## Description

Summarize an analysis of variance model.

## Usage

```
## S3 method for class 'aov':  
summary(object, intercept = FALSE, split,  
        expand.split = TRUE, keep.zero.df = TRUE, ...)  
  
## S3 method for class 'aovlist':  
summary(object, ...)
```

## Arguments

<code>object</code>	An object of class "aov" or "aovlist".
<code>intercept</code>	logical: should intercept terms be included?
<code>split</code>	an optional named list, with names corresponding to terms in the model. Each component is itself a list with integer components giving contrasts whose contributions are to be summed.
<code>expand.split</code>	logical: should the split apply also to interactions involving the factor?
<code>keep.zero.df</code>	logical: should terms with no degrees of freedom be included?
<code>...</code>	Arguments to be passed to or from other methods, for <code>summary.aovlist</code> including those for <code>summary.aov</code> .

## Value

An object of class `c("summary.aov", "listof")` or `"summary.aovlist"` respectively.

## Note

The use of `expand.split = TRUE` is little tested: it is always possible to set it to `FALSE` and specify exactly all the splits required.

**See Also**

aov, summary, model.tables, TukeyHSD

**Examples**

```
## From Venables and Ripley (2002) p.165.
N <- c(0,1,0,1,1,1,0,0,0,1,1,0,1,1,0,0,1,0,1,0,1,1,0,0)
P <- c(1,1,0,0,0,1,0,1,1,1,0,0,0,1,0,1,1,0,0,1,0,1,1,0)
K <- c(1,0,0,1,0,1,1,0,0,1,0,1,0,1,1,0,0,0,1,1,1,0,1,0)
yield <- c(49.5,62.8,46.8,57.0,59.8,58.5,55.5,56.0,
           62.8,55.8,69.5,55.0,62.0,48.8,45.5,44.2,
           52.0,51.5,49.8,48.8,57.2,59.0,53.2,56.0)
npk <- data.frame(block=gl(6,4), N=factor(N), P=factor(P),
                  K=factor(K), yield=yield)

( npk.aov <- aov(yield ~ block + N*P*K, npk) )
summary(npk.aov)
coefficients(npk.aov)

# Cochran and Cox (1957, p.164)
# 3x3 factorial with ordered factors, each is average
# of 12.
CC <- data.frame(
  y = c(449, 413, 326, 409, 358, 291, 341, 278, 312)/12,
  P = ordered(gl(3, 3)), N = ordered(gl(3, 1, 9))
)
CC.aov <- aov(y ~ N * P, data = CC , weights = rep(12, 9))
summary(CC.aov)

# Split both main effects into linear and quadratic parts.
summary(CC.aov, split = list(N = list(L = 1, Q = 2),
  P = list(L = 1, Q = 2)))

# Split only the interaction
summary(CC.aov, split = list("N:P" = list(L.L = 1,
  Q = 2:4)))

# split on just one var
summary(CC.aov, split = list(P = list(lin = 1, quad = 2)))
summary(CC.aov, split = list(P = list(lin = 1, quad = 2)),
  expand.split=FALSE)
```

---

**summary.glm**      *Summarizing Generalized Linear Model Fits*


---

**Description**

These functions are all **methods** for class `glm` or `summary.glm` objects.

**Usage**

```
## S3 method for class 'glm':
summary(object, dispersion = NULL, correlation = FALSE,
        symbolic.cor = FALSE, ...)

## S3 method for class 'summary.glm':
print(x, digits = max(3, getOption("digits") - 3),
      symbolic.cor = x$symbolic.cor,
      signif.stars = getOption("show.signif.stars"), ...)
```

**Arguments**

<code>object</code>	an object of class <code>"glm"</code> , usually, a result of a call to <code>glm</code> .
<code>x</code>	an object of class <code>"summary.glm"</code> , usually, a result of a call to <code>summary.glm</code> .
<code>dispersion</code>	the dispersion parameter for the fitting family. By default it is obtained from <code>object</code> .
<code>correlation</code>	logical; if <code>TRUE</code> , the correlation matrix of the estimated parameters is returned and printed.
<code>digits</code>	the number of significant digits to use when printing.
<code>symbolic.cor</code>	logical. If <code>TRUE</code> , print the correlations in a symbolic form (see <code>symnum</code> ) rather than as numbers.
<code>signif.stars</code>	logical. If <code>TRUE</code> , “significance stars” are printed for each coefficient.
<code>...</code>	further arguments passed to or from other methods.

**Details**

`print.summary.glm` tries to be smart about formatting the coefficients, standard errors, etc. and additionally gives “significance stars” if `signif.stars` is `TRUE`.

Aliased coefficients are omitted in the returned object but (as from R 1.8.0) restored by the `print` method.

Correlations are printed to two decimal places (or symbolically): to see the actual correlations print `summary(object)$correlation` directly.

## Value

`summary.glm` returns an object of class "`summary.glm`", a list with components

<code>call</code>	the component from <code>object</code> .
<code>family</code>	the component from <code>object</code> .
<code>deviance</code>	the component from <code>object</code> .
<code>contrasts</code>	the component from <code>object</code> .
<code>df.residual</code>	the component from <code>object</code> .
<code>null.deviance</code>	the component from <code>object</code> .
<code>df.null</code>	the component from <code>object</code> .
<code>deviance.resid</code>	the deviance residuals: see <code>residuals.glm</code> .
<code>coefficients</code>	the matrix of coefficients, standard errors, z-values and p-values. Aliased coefficients are omitted.
<code>aliased</code>	named logical vector showing if the original coefficients are aliased.
<code>dispersion</code>	either the supplied argument or the estimated dispersion if the latter in <code>NULL</code>
<code>df</code>	a 3-vector of the rank of the model and the number of residual degrees of freedom, plus number of non-aliased coefficients.
<code>cov.unscaled</code>	the unscaled ( <code>dispersion = 1</code> ) estimated covariance matrix of the estimated coefficients.
<code>cov.scaled</code>	ditto, scaled by <code>dispersion</code> .
<code>correlation</code>	(only if <code>correlation</code> is true.) The estimated correlations of the estimated coefficients.
<code>symbolic.cor</code>	(only if <code>correlation</code> is true.) The value of the argument <code>symbolic.cor</code> .

## See Also

`glm`, `summary`.

**Examples**

```
# Continuing the Example from '?glm':  
summary(glm.D93)
```

---

**summary.lm**      *Summarizing Linear Model Fits*


---

**Description**

`summary` method for class `"lm"`.

**Usage**

```
## S3 method for class 'lm':
summary(object, correlation = FALSE, symbolic.cor = FALSE,
  ...)

## S3 method for class 'summary.lm':
print(x, digits = max(3, getOption("digits") - 3),
  symbolic.cor = x$symbolic.cor,
  signif.stars = getOption("show.signif.stars"), ...)
```

**Arguments**

<code>object</code>	an object of class <code>"lm"</code> , usually, a result of a call to <code>lm</code> .
<code>x</code>	an object of class <code>"summary.lm"</code> , usually, a result of a call to <code>summary.lm</code> .
<code>correlation</code>	logical; if <code>TRUE</code> , the correlation matrix of the estimated parameters is returned and printed.
<code>digits</code>	the number of significant digits to use when printing.
<code>symbolic.cor</code>	logical. If <code>TRUE</code> , print the correlations in a symbolic form (see <code>symnum</code> ) rather than as numbers.
<code>signif.stars</code>	logical. If <code>TRUE</code> , “significance stars” are printed for each coefficient.
<code>...</code>	further arguments passed to or from other methods.

**Details**

`print.summary.lm` tries to be smart about formatting the coefficients, standard errors, etc. and additionally gives “significance stars” if `signif.stars` is `TRUE`.

Correlations are printed to two decimal places (or symbolically): to see the actual correlations print `summary(object)$correlation` directly.

## Value

The function `summary.lm` computes and returns a list of summary statistics of the fitted linear model given in `object`, using the components (list elements) `"call"` and `"terms"` from its argument, plus

<code>residuals</code>	the <i>weighted</i> residuals, the usual residuals rescaled by the square root of the weights specified in the call to <code>lm</code> .
<code>coefficients</code>	a $p \times 4$ matrix with columns for the estimated coefficient, its standard error, t-statistic and corresponding (two-sided) p-value. Aliased coefficients are omitted.
<code>aliased</code>	named logical vector showing if the original coefficients are aliased.
<code>sigma</code>	the square root of the estimated variance of the random error

$$\hat{\sigma}^2 = \frac{1}{n-p} \sum_i R_i^2,$$

where  $R_i$  is the  $i$ -th residual, `residuals[i]`.

<code>df</code>	degrees of freedom, a 3-vector $(p, n-p, p^*)$ , the last being the number of non-aliased coefficients.
<code>fstatistic</code>	(for models including non-intercept terms) a 3-vector with the value of the F-statistic with its numerator and denominator degrees of freedom.
<code>r.squared</code>	$R^2$ , the “fraction of variance explained by the model”,

$$R^2 = 1 - \frac{\sum_i R_i^2}{\sum_i (y_i - y^*)^2},$$

where  $y^*$  is the mean of  $y_i$  if there is an intercept and zero otherwise.

<code>adj.r.squared</code>	the above $R^2$ statistic “ <i>adjusted</i> ”, penalizing for higher $p$ .
<code>cov.unscaled</code>	a $p \times p$ matrix of (unscaled) covariances of the $\hat{\beta}_j$ , $j = 1, \dots, p$ .
<code>correlation</code>	the correlation matrix corresponding to the above <code>cov.unscaled</code> , if <code>correlation = TRUE</code> is specified.
<code>symbolic.cor</code>	(only if <code>correlation</code> is true.) The value of the argument <code>symbolic.cor</code> .

## See Also

The model fitting function `lm`, `summary`.



**Examples**

```
## Continuing the lm(.) example:
coef(lm.D90) # the bare coefficients
# omitting intercept
sld90 <- summary(lm.D90 <- lm(weight ~ group -1))
sld90
coef(sld90) # much more
```

---

<code>summary.manova</code>	<i>Summary Method for Multivariate Analysis of Variance</i>
-----------------------------	---

---

**Description**

A summary method for class "manova".

**Usage**

```
## S3 method for class 'manova':
summary(object,
  test = c("Pillai", "Wilks", "Hotelling-Lawley", "Roy"),
  intercept = FALSE, ...)
```

**Arguments**

<code>object</code>	An object of class "manova" or an aov object with multiple responses.
<code>test</code>	The name of the test statistic to be used. Partial matching is used so the name can be abbreviated.
<code>intercept</code>	logical. If TRUE, the intercept term is included in the table.
<code>...</code>	further arguments passed to or from other methods.

**Details**

The `summary.manova` method uses a multivariate test statistic for the summary table. Wilks' statistic is most popular in the literature, but the default Pillai-Bartlett statistic is recommended by Hand and Taylor (1987).

**Value**

A list with components

<code>SS</code>	A named list of sums of squares and product matrices.
<code>Eigenvalues</code>	A matrix of eigenvalues,
<code>stats</code>	A matrix of the statistics, approximate F value and degrees of freedom.

## References

Krzanowski, W. J. (1988) *Principles of Multivariate Analysis. A User's Perspective*. Oxford.

Hand, D. J. and Taylor, C. C. (1987) *Multivariate Analysis of Variance and Repeated Measures*. Chapman and Hall.

## See Also

manova, aov

## Examples

```
## Example on producing plastic film from Krzanowski (1998,
## p. 381)
tear <- c(6.5, 6.2, 5.8, 6.5, 6.5, 6.9, 7.2, 6.9, 6.1, 6.3,
          6.7, 6.6, 7.2, 7.1, 6.8, 7.1, 7.0, 7.2, 7.5, 7.6)
gloss <- c(9.5, 9.9, 9.6, 9.6, 9.2, 9.1, 10.0, 9.9, 9.5,
          9.4, 9.1, 9.3, 8.3, 8.4, 8.5, 9.2, 8.8, 9.7,
          10.1, 9.2)
opacity <- c(4.4, 6.4, 3.0, 4.1, 0.8, 5.7, 2.0, 3.9, 1.9,
            5.7, 2.8, 4.1, 3.8, 1.6, 3.4, 8.4, 5.2, 6.9,
            2.7, 1.9)
Y <- cbind(tear, gloss, opacity)
rate <- factor(gl(2,10), labels=c("Low", "High"))
additive <- factor(gl(2, 5, len=20),
  labels=c("Low", "High"))

fit <- manova(Y ~ rate * additive)
summary.aov(fit) # univariate ANOVA tables
summary(fit, test="Wilks") # ANOVA table of Wilks' lambda
```

---

**terms**     *Model Terms*

---

**Description**

The function **terms** is a generic function which can be used to extract *terms* objects from various kinds of R data objects.

**Usage**

```
terms(x, ...)
```

**Arguments**

**x**                      object used to select a method to dispatch.  
**...**                    further arguments passed to or from other methods.

**Details**

There are methods for classes "**aovlist**", and "**terms**" "**formula**" (see **terms.formula**): the default method just extracts the **terms** component of the object (if any).

**Value**

An object of class **c("terms", "formula")** which contains the *terms* representation of a symbolic model. See **terms.object** for its structure.

**References**

Chambers, J. M. and Hastie, T. J. (1992) *Statistical models*. Chapter 2 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

**terms.object**, **terms.formula**, **lm**, **glm**, **formula**.

---

<code>terms.formula</code>	<i>Construct a terms Object from a Formula</i>
----------------------------	--

---

## Description

This function takes a formula and some optional arguments and constructs a terms object. The terms object can then be used to construct a `model.matrix`.

## Usage

```
## S3 method for class 'formula':
terms(x, specials = NULL, abb = NULL, data = NULL,
      neg.out = TRUE, keep.order = FALSE, simplify = FALSE,
      ...)
```

## Arguments

<code>x</code>	a formula.
<code>specials</code>	which functions in the formula should be marked as special in the terms object.
<code>abb</code>	Not implemented in R.
<code>data</code>	a data frame from which the meaning of the special symbol <code>.</code> can be inferred. It is unused if there is no <code>.</code> in the formula.
<code>neg.out</code>	Not implemented in R.
<code>keep.order</code>	a logical value indicating whether the terms should keep their positions. If <code>FALSE</code> the terms are reordered so that main effects come first, followed by the interactions, all second-order, all third-order and so on. Effects of a given order are kept in the order specified.
<code>simplify</code>	should the formula be expanded and simplified, the pre-1.7.0 behaviour?
<code>...</code>	further arguments passed to or from other methods.

## Details

Not all of the options work in the same way that they do in S and not all are implemented.

**Value**

A `terms.object` object is returned. The object itself is the re-ordered (unless `keep.order = TRUE`) formula. In all cases variables within an interaction term in the formula are re-ordered by the ordering of the `"variables"` attribute, which is the order in which the variables occur in the formula.

**See Also**

`terms`, `terms.object`

---

<code>terms.object</code>	<i>Description of Terms Objects</i>
---------------------------	-------------------------------------

---

## Description

An object of class **terms** holds information about a model. Usually the model was specified in terms of a **formula** and that formula was used to determine the terms object.

## Value

The object itself is simply the formula supplied to the call of **terms.formula**. The object has a number of attributes and they are used to construct the model frame:

<b>factors</b>	A matrix of variables by terms showing which variables appear in which terms. The entries are 0 if the variable does not occur in the term, 1 if it does occur and should be coded by contrasts, and 2 if it occurs and should be coded via dummy variables for all levels (as when an intercept or lower-order term is missing).
<b>term.labels</b>	A character vector containing the labels for each of the terms in the model. Non-syntactic names will be quoted by backticks.
<b>variables</b>	A call to <b>list</b> of the variables in the model.
<b>intercept</b>	Either 0, indicating no intercept is to be fit, or 1 indicating that an intercept is to be fit.
<b>order</b>	A vector of the same length as <b>term.labels</b> indicating the order of interaction for each term
<b>response</b>	The index of the variable (in <b>variables</b> ) of the response (the left hand side of the formula).
<b>offset</b>	If the model contains <b>offset</b> terms there is an <b>offset</b> attribute indicating which variables are offsets
<b>specials</b>	If the <b>specials</b> argument was given to <b>terms.formula</b> there is a <b>specials</b> attribute, a list of vectors indicating the terms that contain these special functions.

The object has class `c("terms", "formula")`.

**Note**

These objects are different from those found in S. In particular there is no **formula** attribute, instead the object is itself a formula. Thus, the mode of a terms object is different as well.

Examples of the **specials** argument can be seen in the **aov** and **coxph** functions.

**See Also**

**terms**, **formula**.



---

**TukeyHSD**      *Compute Tukey Honest Significant Differences*

---

**Description**

Create a set of confidence intervals on the differences between the means of the levels of a factor with the specified family-wise probability of coverage. The intervals are based on the Studentized range statistic, Tukey's 'Honest Significant Difference' method. There is a `plot` method.

**Usage**

```
TukeyHSD(x, which, ordered = FALSE, conf.level = 0.95, ...)
```

**Arguments**

<code>x</code>	A fitted model object, usually an <code>aov</code> fit.
<code>which</code>	A list of terms in the fitted model for which the intervals should be calculated. Defaults to all the terms.
<code>ordered</code>	A logical value indicating if the levels of the factor should be ordered according to increasing average in the sample before taking differences. If <code>ordered</code> is true then the calculated differences in the means will all be positive. The significant differences will be those for which the <code>lwr</code> end point is positive.
<code>conf.level</code>	A numeric value between zero and one giving the family-wise confidence level to use.
<code>...</code>	Optional additional arguments. None are used at present.

**Details**

When comparing the means for the levels of a factor in an analysis of variance, a simple comparison using t-tests will inflate the probability of declaring a significant difference when it is not in fact present. This is because the intervals are calculated with a given coverage probability for each interval but the interpretation of the coverage is usually with respect to the entire family of intervals.

John Tukey introduced intervals based on the range of the sample means rather than the individual differences. The intervals returned by this function are based on this Studentized range statistics.

Technically the intervals constructed in this way would only apply to balanced designs where there are the same number of observations made at each level of the factor. This function incorporates an adjustment for sample size that produces sensible intervals for mildly unbalanced designs.

## Value

A list with one component for each term requested in `which`. Each component is a matrix with columns `diff` giving the difference in the observed means, `lwr` giving the lower end point of the interval, and `upr` giving the upper end point.

## Author(s)

Douglas Bates

## References

Miller, R. G. (1981) *Simultaneous Statistical Inference*. Springer.  
Yandell, B. S. (1997) *Practical Data Analysis for Designed Experiments*. Chapman & Hall.

## See Also

`aov`, `qtukey`, `model.tables`

## Examples

```
data(warpbreaks)
summary(fm1 <- aov(breaks ~ wool + tension,
  data = warpbreaks))
TukeyHSD(fm1, "tension", ordered = TRUE)
plot(TukeyHSD(fm1, "tension"))
```

---

**uniroot**      *One Dimensional Root (Zero) Finding*

---

**Description**

The function **uniroot** searches the interval from **lower** to **upper** for a root (i.e., zero) of the function **f** with respect to its first argument.

**Usage**

```
uniroot(f, interval, lower = min(interval),  
        upper = max(interval), tol = .Machine$double.eps^0.25,  
        maxiter = 1000, ...)
```

**Arguments**

<b>f</b>	the function for which the root is sought.
<b>interval</b>	a vector containing the end-points of the interval to be searched for the root.
<b>lower</b>	the lower end point of the interval to be searched.
<b>upper</b>	the upper end point of the interval to be searched.
<b>tol</b>	the desired accuracy (convergence tolerance).
<b>maxiter</b>	the maximum number of iterations.
<b>...</b>	additional arguments to <b>f</b> .

**Details**

Either **interval** or both **lower** and **upper** must be specified. The function uses Fortran subroutine “**zeroin**” (from Netlib) based on algorithms given in the reference below.

If the algorithm does not converge in **maxiter** steps, a warning is printed and the current approximation is returned.

**Value**

A list with four components: **root** and **f.root** give the location of the root and the value of the function evaluated at that point. **iter** and **estim.prec** give the number of iterations used and an approximate estimated precision for **root**.

## References

Brent, R. (1973) *Algorithms for Minimization without Derivatives*. Englewood Cliffs, NJ: Prentice-Hall.

## See Also

polyroot for all complex roots of a polynomial; optimize, nlm.

## Examples

```
f <- function(x,a) x - a
str(xmin <- uniroot(f, c(0, 1), tol = 0.0001, a = 1/3))
str(uniroot(function(x) x*(x^2-1) + .5,
           low = -2, up = 2, tol = 0.0001), dig = 10)
str(uniroot(function(x) x*(x^2-1) + .5,
           low = -2, up = 2, tol = 1e-10), dig = 10)

## Find the smallest value x for which exp(x) > 0
## (numerically):
r <- uniroot(function(x) 1e80*exp(x) - 1e-300, , -1000, 0,
           tol=1e-20)
str(r, digits= 15) ## around -745.1332191

exp(r$r)           # = 0, but not for r$r * 0.999...
minexp <- r$r * (1 - .Machine$double.eps)
exp(minexp)        # typically denormalized
```

---

**update**      *Update and Re-fit a Model Call*

---

**Description**

`update` will update and (by default) re-fit a model. It does this by extracting the call stored in the object, updating the call and (by default) evaluating that call. Sometimes it is useful to call `update` with only one argument, for example if the data frame has been corrected.

**Usage**

```
update(object, ...)
```

```
## Default S3 method:
```

```
update(object, formula., ..., evaluate = TRUE)
```

**Arguments**

<code>object</code>	An existing fit from a model function such as <code>lm</code> , <code>glm</code> and many others.
<code>formula.</code>	Changes to the formula – see <code>update.formula</code> for details.
<code>...</code>	Additional arguments to the call, or arguments with changed values. Use <code>name=NULL</code> to remove the argument <code>name</code> .
<code>evaluate</code>	If true evaluate the new call else return the call.

**Value**

If `evaluate = TRUE` the fitted object, otherwise the updated call.

**References**

Chambers, J. M. (1992) *Linear models*. Chapter 4 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

**See Also**

`update.formula`

## Examples

```
oldcon <- options(contrasts = c("contr.treatment",
                                "contr.poly"))
## Annette Dobson (1990) "An Introduction to Generalized
## Linear Models". Page 9: Plant Weight Data.
ctl <- c(4.17,5.58,5.18,6.11,4.50,4.61,5.17,4.53,5.33,5.14)
trt <- c(4.81,4.17,4.41,3.59,5.87,3.83,6.03,4.89,4.32,4.69)
group <- gl(2, 10, 20, labels = c("Ctl", "Trt"))
weight <- c(ctl, trt)
lm.D9 <- lm(weight ~ group)
lm.D9
summary(lm.D90 <- update(lm.D9, . ~ . - 1))
options(contrasts = c("contr.helmert", "contr.poly"))
update(lm.D9)
options(oldcon)
```

---

**update.formula**      *Model Updating*

---

**Description**

`update.formula` is used to update model formulae. This typically involves adding or dropping terms, but updates can be more general.

**Usage**

```
## S3 method for class 'formula':  
update(old, new, ...)
```

**Arguments**

<code>old</code>	a model formula to be updated.
<code>new</code>	a formula giving a template which specifies how to update.
<code>...</code>	further arguments passed to or from other methods.

**Details**

The function works by first identifying the *left-hand side* and *right-hand side* of the `old` formula. It then examines the `new` formula and substitutes the *lhs* of the `old` formula for any occurrence of `"."` on the left of `new`, and substitutes the *rhs* of the `old` formula for any occurrence of `"."` on the right of `new`.

**Value**

The updated formula is returned.

**See Also**

`terms`, `model.matrix`.

**Examples**

```
update(y ~ x,      ~ . + x2) # y ~ x + x2  
update(y ~ x, log(.) ~ . ) # log(y) ~ x
```

---

<b>vcov</b>	<i>Calculate Variance-Covariance Matrix for a Fitted Model Object</i>
-------------	---

---

## Description

Returns the variance-covariance matrix of the main parameters of a fitted model object.

## Usage

```
vcov(object, ...)
```

## Arguments

<b>object</b>	a fitted model object.
<b>...</b>	additional arguments for method functions. For the <b>glm</b> method this can be used to pass a <b>dispersion</b> parameter.

## Details

This is a generic function. Functions with names beginning in **vcov**. will be methods for this function. Classes with methods for this function include: **lm**, **glm**, **nls**, **lme**, **gls**, **coxph** and **survreg**

## Value

A matrix of the estimated covariances between the parameter estimates in the linear or non-linear predictor of the model.



---

<code>weighted.residuals</code>	<i>Compute Weighted Residuals</i>
---------------------------------	-----------------------------------

---

## Description

Computed weighted residuals from a linear model fit.

## Usage

```
weighted.residuals(obj, drop0 = TRUE)
```

## Arguments

<code>obj</code>	R object, typically of class <code>lm</code> or <code>glm</code> .
<code>drop0</code>	logical. If <code>TRUE</code> , drop all cases with <code>weights == 0</code> .

## Details

Weighted residuals are the usual residuals  $R_i$ , multiplied by  $\sqrt{w_i}$ , where  $w_i$  are the `weights` as specified in `lm`'s call.

Dropping cases with weights zero is compatible with `influence` and related functions.

## Value

Numeric vector of length  $n'$ , where  $n'$  is the number of non-0 weights (`drop0 = TRUE`) or the number of observations, otherwise.

## See Also

`residuals`, `lm.influence`, etc.

## Examples

```
example("lm")
all.equal(weighted.residuals(lm.D9),
          residuals(lm.D9))
x <- 1:10
w <- 0:9
y <- rnorm(x)
weighted.residuals(lmxy <- lm(y ~ x, weights = w))
weighted.residuals(lmxy, drop0 = FALSE)
```



## Chapter 5

# Base package — dates, time and time-series

---

**as.POSIX\***      *Date-time Conversion Functions*


---

**Description**

Functions to manipulate objects of classes "POSIXlt" and "POSIXct" representing calendar dates and times (to the nearest second).

**Usage**

```
as.POSIXct(x, tz = "")
as.POSIXlt(x, tz = "")
```

**Arguments**

<b>x</b>	An object to be converted.
<b>tz</b>	A timezone specification to be used for the conversion, <i>if one is required</i> . System-specific, but "" is the current timezone, and "GMT" is UTC (Coordinated Universal Time, in French).

**Details**

The `as.POSIX*` functions convert an object to one of the two classes used to represent date/times (calendar dates plus time to the nearest second). They can convert a wide variety of objects, including objects of the other class and of classes "date" (from package `[date:as.date]date` or `[date:as.date]survival`), "chron" and "dates" (from package `[chron]chron`) to these classes. They can also convert character strings of the formats "2001-02-03" and "2001/02/03" optionally followed by white space and a time in the format "14:52" or "14:52:03". (Formats such as "01/02/03" are ambiguous but can be converted via a format specification by `strptime`.)

Logical NAs can be converted to either of the classes, but no other logical vectors can be.

**Value**

`as.POSIXct` and `as.POSIXlt` return an object of the appropriate class. If `tz` was specified, `as.POSIXlt` will give an appropriate "tzone" attribute.

## Note

If you want to extract specific aspects of a time (such as the day of the week) just convert it to class `"POSIXlt"` and extract the relevant component(s) of the list, or if you want a character representation (such as a named day of the week) use `format.POSIXlt` or `format.POSIXct`.

If a timezone is needed and that specified is invalid on your system, what happens is system-specific but it will probably be ignored.

## See Also

`Date``TimeClasses` for details of the classes; `strptime` for conversion to and from character representations.

## Examples

```
(z <- Sys.time()) # the current date, as class "POSIXct"
unclass(z)        # a large integer
floor(unclass(z)/86400) # number of days since 1970-01-01
(z <- as.POSIXlt(Sys.time())) # current date,
                               # as class "POSIXlt"
unlist(unclass(z)) # a list shown as a named vector

as.POSIXlt(Sys.time(), "GMT") # the current time in GMT
```

---

<code>cut.POSIXt</code>	<i>Convert a Date-Time Object to a Factor</i>
-------------------------	---

---

## Description

Method for `cut` applied to date-time objects.

## Usage

```
## S3 method for class 'POSIXt':  
cut(x, breaks, labels = NULL, start.on.monday = TRUE,  
    right = FALSE, ...)
```

## Arguments

<code>x</code>	an object inheriting from class <code>"POSIXt"</code> .
<code>breaks</code>	a vector of cut points <i>or</i> number giving the number of intervals which <code>x</code> is to be cut into <i>or</i> an interval specification, one of <code>"sec"</code> , <code>"min"</code> , <code>"hour"</code> , <code>"day"</code> , <code>"DSTday"</code> , <code>"week"</code> , <code>"month"</code> or <code>"year"</code> , optionally preceded by an integer and a space, or followed by <code>"s"</code> .
<code>labels</code>	labels for the levels of the resulting category. By default, labels are constructed from the left-hand end of the intervals (which are include for the default value of <code>right</code> ). If <code>labels = FALSE</code> , simple integer codes are returned instead of a factor.
<code>start.on.monday</code>	logical. If <code>breaks = "weeks"</code> , should the week start on Mondays or Sundays?
<code>right, ...</code>	arguments to be passed to or from other methods.

## Value

A factor is returned, unless `labels = FALSE` which returns the integer level codes.

## See Also

`seq.POSIXt`, `cut`

**Examples**

```
## random dates in a 10-week period  
cut(ISOdate(2001, 1, 1) + 70*86400*runif(100), "weeks")
```

---

**DateTimeClasses**      *Date-Time Classes*


---

**Description**

Description of the classes "POSIXlt" and "POSIXct" representing calendar dates and times (to the nearest second).

**Usage**

```
## S3 method for class 'POSIXct':
print(x, ...)

## S3 method for class 'POSIXct':
summary(object, digits = 15, ...)

time + number
time - number
time1 lop time2
```

**Arguments**

<code>x</code> , <code>object</code>	An object to be printed or summarized from one of the date-time classes.
<code>digits</code>	Number of significant digits for the computations: should be high enough to represent the least important time unit exactly.
<code>...</code>	Further arguments to be passed from or to other methods.
<code>time</code> , <code>time1</code> , <code>time2</code>	date-time objects.
<code>number</code>	a numeric object.
<code>lop</code>	One of <code>==</code> , <code>!=</code> , <code>&lt;</code> , <code>&lt;=</code> , <code>&gt;</code> or <code>&gt;=</code> .

**Details**

There are two basic classes of date/times. Class "POSIXct" represents the (signed) number of seconds since the beginning of 1970 as a numeric vector. Class "POSIXlt" is a named list of vectors representing

`sec` 0–61: seconds

`min` 0–59: minutes



**hour** 0–23: hours

**mday** 1–31: day of the month

**mon** 0–11: months after the first of the year.

**year** Years since 1900.

**wday** 0–6 day of the week, starting on Sunday.

**yday** 0–365: day of the year.

**isdst** Daylight savings time flag. Positive if in force, zero if not, negative if unknown.

The classes correspond to the ANSI C constructs of “calendar time” (the `time_t` data type) and “local time” (or broken-down time, the `struct tm` data type), from which they also inherit their names.

"POSIXct" is more convenient for including in data frames, and "POSIXlt" is closer to human-readable forms. A virtual class "POSIXt" inherits from both of the classes: it is used to allow operations such as subtraction to mix the two classes.

Logical comparisons and limited arithmetic are available for both classes. One can add or subtract a number of seconds or a `difftime` object from a date-time object, but not add two date-time objects. Subtraction of two date-time objects is equivalent to using `difftime`. Be aware that "POSIXlt" objects will be interpreted as being in the current timezone for these operations, unless a timezone has been specified.

"POSIXlt" objects will often have an attribute "tzone", a character vector of length 3 giving the timezone name from the TZ environment variable and the names of the base timezone and the alternate (daylight-saving) timezone. Sometimes this may just be of length one, giving the timezone name.

Unfortunately, the conversion is complicated by the operation of time zones and leap seconds (22 days have been 86401 seconds long so far: the times of the extra seconds are in the object `.leap.seconds`). The details of this are entrusted to the OS services where possible. This will usually cover the period 1970–2037, and on Unix machines back to 1902 (when time zones were in their infancy). Outside those ranges we use our own C code. This uses the offset from GMT in use in the timezone in 2000, and uses the alternate (daylight-saving) timezone only if `isdst` is positive.

It seems that some systems use leap seconds but most do not. This is detected and corrected for at build time, so all "POSIXct" times used by R do not include leap seconds. (Conceivably this could be wrong if the system has changed since build time, just possibly by changing locales.)

Using `c` on `"POSIXlt"` objects converts them to the current time zone.

### Warning

Some Unix-like systems (especially GNU/Linux ones) do not have `"TZ"` set, yet have internal code that expects it (as does POSIX). We have tried to work around this, but if you get unexpected results try setting `"TZ"`.

### See Also

`as.POSIXct` and `as.POSIXlt` for conversion between the classes.  
`strptime` for conversion to and from character representations.  
`Sys.time` for clock time as a `"POSIXct"` object.  
`difftime` for time intervals.  
`cut.POSIXt`, `seq.POSIXt`, `round.POSIXt` and `trunc.POSIXt` for methods for these classes.  
`weekdays.POSIXt` for convenience extraction functions.

### Examples

```
(z <- Sys.time()) # the current date, as class "POSIXct"

Sys.time() - 3600 # an hour ago

as.POSIXlt(Sys.time(), "GMT") # the current time in GMT
format(.leap.seconds) # all 22 leapseconds in your timezone
```

---

**diff**      *Lagged Differences*

---

**Description**

Returns suitably lagged and iterated differences.

**Usage**

```
diff(x, ...)

## Default S3 method:
diff(x, lag = 1, differences = 1, ...)

## S3 method for class 'POSIXt':
diff(x, lag = 1, differences = 1, ...)
```

**Arguments**

<b>x</b>	a numeric vector or matrix containing the values to be differenced.
<b>lag</b>	an integer indicating which lag to use.
<b>differences</b>	an integer indicating the order of the difference.
<b>...</b>	further arguments to be passed to or from methods.

**Details**

**diff** is a generic function with a default method and ones for classes "ts" and "POSIXt". NA's propagate.

**Value**

If **x** is a vector of length **n** and **differences=1**, then the computed result is equal to the successive differences **x[(1+lag):n] - x[1:(n-lag)]**.

If **difference** is larger than one this algorithm is applied recursively to **x**. Note that the returned value is a vector which is shorter than **x**.

If **x** is a matrix then the difference operations are carried out on each column separately.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`diff.ts`, `diffinv`.

## Examples

```
diff(1:10, 2)
diff(1:10, 2, 2)
x <- cumsum(cumsum(1:10))
diff(x, lag = 2)
diff(x, differences = 2)

diff(.leap.seconds)
```

---

**difftime**     *Time Intervals*

---

**Description**

Create, print and round time intervals.

**Usage**

```
time1 - time2
difftime(time1, time2, tz = "",
  units = c("auto", "secs", "mins", "hours", "days", "weeks"))
as.difftime(tim, format = "%X")

## S3 method for class 'difftime':
round(x, digits = 0)
```

**Arguments**

<b>time1, time2</b>	date-time objects.
<b>tz</b>	a timezone specification to be used for the conversion. System-specific, but "" is the current time zone, and "GMT" is UTC.
<b>units</b>	character. Units in which the results are desired. Can be abbreviated.
<b>tim</b>	character string specifying a time interval.
<b>format</b>	character specifying the format of <b>tim</b> .
<b>x</b>	an object inheriting from class "difftime".
<b>digits</b>	integer. Number of significant digits to retain.

**Details**

Function **difftime** takes a difference of two date/time objects (of either class) and returns an object of class "**difftime**" with an attribute indicating the units. There is a **round** method for objects of this class, as well as methods for the group-generic (see **Ops**) logical and arithmetic operations.

If **units** = "auto", a suitable set of units is chosen, the largest possible (excluding "weeks") in which all the absolute differences are greater than one.

Subtraction of two date-time objects gives an object of this class, by calling `difftime` with `units="auto"`. Alternatively, `as.difftime()` works on character-coded time intervals.

Limited arithmetic is available on "difftime" objects: they can be added or subtracted, and multiplied or divided by a numeric vector. In addition, adding or subtracting a numeric vector implicitly converts the numeric vector to a "difftime" object with the same units as the "difftime" object.

## See Also

`DateTimeClasses`.

## Examples

```
(z <- Sys.time()) - 3600)
Sys.time() - z                # just over 3600 seconds.

## time interval between releases of 1.2.2 and 1.2.3.
ISOdate(2001, 4, 26) - ISOdate(2001, 2, 26)

as.difftime(c("0:3:20", "11:23:15"))
# 3rd gives NA
as.difftime(c("3:20", "23:15", "2:"), format= "%H:%M")
```

---

**hist.POSIXt**      *Histogram of a Date-Time Object*


---

## Description

Method for `hist` applied to date-time objects.

## Usage

```
## S3 method for class 'POSIXt':
hist(x, breaks, ..., plot = TRUE, freq = FALSE,
      start.on.monday = TRUE, format)
```

## Arguments

<code>x</code>	an object inheriting from class <code>"POSIXt"</code> .
<code>breaks</code>	a vector of cut points <i>or</i> number giving the number of intervals which <code>x</code> is to be cut into <i>or</i> an interval specification, one of <code>"secs"</code> , <code>"mins"</code> , <code>"hours"</code> , <code>"days"</code> , <code>"weeks"</code> , <code>"months"</code> or <code>"years"</code> .
<code>...</code>	graphical parameters, or arguments to <code>hist.default</code> such as <code>include.lowest</code> , <code>right</code> and <code>labels</code> .
<code>plot</code>	logical. If <code>TRUE</code> (default), a histogram is plotted, otherwise a list of breaks and counts is returned.
<code>freq</code>	logical; if <code>TRUE</code> , the histogram graphic is a representation of frequencies, i.e, the <code>counts</code> component of the result; if <code>FALSE</code> , <i>relative</i> frequencies (“probabilities”) are plotted.
<code>start.on.monday</code>	logical. If <code>breaks = "weeks"</code> , should the week start on Mondays or Sundays?
<code>format</code>	for the x-axis labels. See <code>strptime</code> .

## Value

An object of class `"histogram"`: see `hist`.

## See Also

`seq.POSIXt`, `axis.POSIXct`, `hist`

**Examples**

```
hist(.leap.seconds, "years", freq = TRUE)
hist(.leap.seconds,
     seq(ISOdate(1970, 1, 10), ISOdate(2002, 1, 1), "5 years"))

## 100 random dates in a 10-week period
random.dates <- ISOdate(2001, 1, 1) + 70*86400*runif(100)
hist(random.dates, "weeks", format = "%d %b")
```



---

**print.ts**     *Printing Time-Series Objects*

---

**Description**

Print method for time series objects.

**Usage**

```
## S3 method for class 'ts':  
print(x, calendar, ...)
```

**Arguments**

<code>x</code>	a time series object.
<code>calendar</code>	enable/disable the display of information about month names, quarter names or year when printing. The default is <code>TRUE</code> for a frequency of 4 or 12, <code>FALSE</code> otherwise.
<code>...</code>	additional arguments to <code>print</code> .

**Details**

This is the `print` methods for objects inheriting from class `"ts"`.

**See Also**

`print`, `ts`.

**Examples**

```
print(ts(1:10, freq = 7, start = c(12, 2)), calendar=TRUE)
```

---

**rep**     *Replicate Elements of Vectors and Lists*


---

**Description**

**rep** replicates the values in **x**. It is a generic function, and the default method is described here.

**rep.int** is a faster simplified version for the commonest case.

**Usage**

```
rep(x, times, ...)
```

```
## Default S3 method:
```

```
rep(x, times, length.out, each, ...)
```

```
rep.int(x, times)
```

**Arguments**

<b>x</b>	a vector (of any mode including a list) or a pairlist or a <code>POSIXct</code> or <code>POSIXlt</code> object.
<b>times</b>	non-negative integer. A vector giving the number of times to repeat each element if of length <code>length(x)</code> , or to repeat the whole vector if of length 1.
<b>length.out</b>	integer. (Optional.) The desired length of the output vector.
<b>each</b>	optional integer. Each element of <b>x</b> is repeated <b>each</b> times.
<b>...</b>	further arguments to be passed to or from other methods.

**Details**

If **times** consists of a single integer, the result consists of the values in **x** repeated this many times. If **times** is a vector of the same length as **x**, the result consists of **x**[1] repeated **times**[1] times, **x**[2] repeated **times**[2] times and so on.

**length.out** may be given in place of **times**, in which case **x** is repeated as many times as is necessary to create a vector of this length. If both **length.out** and **times** are specified, **times** determines the replication,

and `length.out` can be used to truncate the output vector (or extend it by NAs).

Non-integer values of `times` will be truncated towards zero. If `times` is a computed quantity it is prudent to add a small fuzz.

## Value

A vector of the same class as `x`.

## Note

If the original vector has names, these are also replicated and so will almost always contain duplicates.

If `length.out` is used to extend the vector, the behaviour is different from that of S-PLUS, which recycles the existing vector.

Function `rep.int` is a simple case handled by internal code, and provided as a separate function purely for S compatibility.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`seq`, `sequence`.

## Examples

```
rep(1:4, 2)
rep(1:4, each = 2)      # not the same.
rep(1:4, c(2,2,2,2))    # same as second.
rep(1:4, c(2,1,2,1))
rep(1:4, each = 2, len = 4) # first 4 only.
rep(1:4, each = 2, len = 10) # 8 integers plus two NAs

rep(1, 40*(1-.8)) # length 7 on most platforms
rep(1, 40*(1-.8)+1e-7) # better

## replicate a list
fred <- list(happy = 1:10, name = "squash")
rep(fred, 5)

# date-time objects
```

```
x <- .leap.seconds[1:3]
rep(x, 2)
rep(as.POSIXlt(x), rep(2, 3))
```

---

**round.POSIXt**      *Round / Truncate Data-Time Objects*

---

**Description**

Round or truncate date-time objects.

**Usage**

```
## S3 method for class 'POSIXt':  
round(x, units=c("secs", "mins", "hours", "days"))  
## S3 method for class 'POSIXt':  
trunc(x, units=c("secs", "mins", "hours", "days"))
```

**Arguments**

<b>x</b>	an object inheriting from "POSIXt".
<b>units</b>	one of the units listed. Can be abbreviated.

**Details**

The time is rounded or truncated to the second, minute, hour or day. Timezones are only relevant to days, when midnight in the current time-zone is used.

**Value**

An object of class "POSIXt".

**See Also**

`DateTimeClasses`

**Examples**

```
round(.leap.seconds + 1000, "hour")  
trunc.POSIXt(Sys.time(), "day")
```

---

**seq.POSIXt**      *Generate Regular Sequences of Dates*


---

**Description**

The method for **seq** for data-time classes.

**Usage**

```
## S3 method for class 'POSIXt':
seq(from, to, by, length.out=NULL, along.with=NULL, ...)
```

**Arguments**

<b>from</b>	starting date. Required
<b>to</b>	end date. Optional. If supplied must be after <b>from</b> .
<b>by</b>	increment of the sequence. Optional. See Details.
<b>length.out</b>	integer, optional. desired length of the sequence.
<b>along.with</b>	take the length from the length of this argument.
<b>...</b>	arguments passed to or from other methods.

**Details**

**by** can be specified in several ways.

- A number, taken to be in seconds.
- A object of class **difftime**
- A character string, containing one of "**sec**", "**min**", "**hour**", "**day**", "**DSTday**", "**week**", "**month**" or "**year**". This can optionally be preceded by an integer and a space, or followed by "**s**".

The difference between "**day**" and "**DSTday**" is that the former ignores changes to/from daylight savings time and the latter takes the same clock time each day. ("**week**" ignores DST, but "**7 DSTdays**") can be used as an alternative. "**month**" and "**year**" allow for DST as from R 1.5.0.)

**Value**

A vector of class "**POSIXct**".

**See Also**

`DateTimeClasses`

**Examples**

```
## first days of years
seq(ISOdate(1910,1,1), ISOdate(1999,1,1), "years")
## by month
seq(ISOdate(2000,1,1), by="month", length=12)
## quarters
seq(ISOdate(1990,1,1), ISOdate(2000,1,1), by="3 months")
## days vs DSTdays
seq(ISOdate(2000,3,20), by="day", length = 10)
seq(ISOdate(2000,3,20), by="DSTday", length = 10)
seq(ISOdate(2000,3,20), by="7 DSTdays", length = 4)
```

---

**start**     *Encode the Terminal Times of Time Series*

---

## Description

Extract and encode the times the first and last observations were taken. Provided only for compatibility with S version 2.

## Usage

```
start(x, ...)  
end(x, ...)
```

## Arguments

<b>x</b>	a univariate or multivariate time-series, or a vector or matrix.
<b>...</b>	extra arguments for future methods.

## Details

These are generic functions, which will use the **tsp** attribute of **x** if it exists. Their default methods decode the start time from the original time units, so that for a monthly series **1995.5** is represented as **c(1995, 7)**. For a series of frequency **f**, time **n+i/f** is presented as **c(n, i+1)** (even for **i = 0** and **f = 1**).

## Warning

The representation used by **start** and **end** has no meaning unless the frequency is supplied.

## See Also

**ts**, **time**, **tsp**.



---

**strptime**      *Date-time Conversion Functions to and from Character*


---

**Description**

Functions to convert between character representations and objects of classes "POSIXlt" and "POSIXct" representing calendar dates and times.

**Usage**

```
## S3 method for class 'POSIXct':
format(x, format = "", tz = "", usetz = FALSE, ...)
## S3 method for class 'POSIXlt':
format(x, format = "", usetz = FALSE, ...)

## S3 method for class 'POSIXt':
as.character(x, ...)

strftime(x, format="", usetz = FALSE, ...)
strptime(x, format)

ISOdatetime(year, month, day, hour, min, sec, tz = "")
ISOdate(year, month, day, hour=12, min=0, sec=0, tz="GMT")
```

**Arguments**

<b>x</b>	An object to be converted.
<b>tz</b>	A timezone specification to be used for the conversion. System-specific, but "" is the current time zone, and "GMT" is UTC.
<b>format</b>	A character string. The default is "%Y-%m-%d %H:%M:%S" if any component has a time component which is not midnight, and "%Y-%m-%d" otherwise.
<b>...</b>	Further arguments to be passed from or to other methods.
<b>usetz</b>	logical. Should the timezone be appended to the output? This is used in printing time, and as a workaround for problems with using "%Z" on most GNU/Linux systems.

`year, month, day`

numerical values to specify a day.

`hour, min, sec`

numerical values for a time within a day.

## Details

`strptime` is an alias for `format.POSIXlt`, and `format.POSIXct` first converts to class `"POSIXct"` by calling `as.POSIXct`. Note that only that conversion depends on the time zone.

The usual vector re-cycling rules are applied to `x` and `format` so the answer will be of length that of the longer of the vectors.

Locale-specific conversions to and from character strings are used where appropriate and available. This affects the names of the days and months, the AM/PM indicator (if used) and the separators in formats such as `%x` and `%X`.

The details of the formats are system-specific, but the following are defined by the POSIX standard for `strptime` and are likely to be widely available. Any character in the format string other than the `%` escapes is interpreted literally (and `%%` gives `%`).

`%a` Abbreviated weekday name.

`%A` Full weekday name.

`%b` Abbreviated month name.

`%B` Full month name.

`%c` Date and time, locale-specific.

`%d` Day of the month as decimal number (01–31).

`%H` Hours as decimal number (00–23).

`%I` Hours as decimal number (01–12).

`%j` Day of year as decimal number (001–366).

`%m` Month as decimal number (01–12).

`%M` Minute as decimal number (00–59).

`%p` AM/PM indicator in the locale. Used in conjunction with `%I` and **not** with `%H`.

`%S` Second as decimal number (00–61), allowing for up to two leap-seconds.

`%U` Week of the year as decimal number (00–53) using the first Sunday as day 1 of week 1.

`%w` Weekday as decimal number (0–6, Sunday is 0).

- %W** Week of the year as decimal number (00–53) using the first Monday as day 1 of week 1.
- %x** Date, locale-specific.
- %X** Time, locale-specific.
- %y** Year without century (00–99). If you use this on input, which century you get is system-specific. So don't! Often values up to 69 are prefixed by 20 and 70–99 by 19.
- %Y** Year with century.
- %Z** (output only.) Time zone as a character string (empty if not available). Note: do not use this on GNU/Linux unless the TZ environment variable is set.

Where leading zeros are shown they will be used on output but are optional on input.

`ISOdatetime` and `ISOdate` are convenience wrappers for `strptime`, that differ only in their defaults.

## Value

The `format` methods and `strftime` return character vectors representing the time.

`strptime` turns character representations into an object of class `"POSIXlt"`.

`ISOdatetime` and `ISOdate` return an object of class `"POSIXct"`.

## Note

The default formats follow the rules of the ISO 8601 international standard which expresses a day as `"2001-02-03"` and a time as `"14:01:02"` using leading zeroes as here. The ISO form uses no space to separate dates and times.

If the date string does not specify the date completely, the returned answer may be system-specific. The most common behaviour is to assume that unspecified seconds, minutes or hours are zero, and a missing year, month or day is the current one.

If the timezone specified is invalid on your system, what happens is system-specific but it will probably be ignored.

OS facilities will probably not print years before 1CE (aka 1AD) correctly.

## References

International Organization for Standardization (1988, 1997, ...) *ISO 8601. Data elements and interchange formats – Information interchange – Representation of dates and times*. The 1997 version is available on-line at <ftp://ftp.ql.net/pub/g1smd/8601v03.pdf>

## See Also

`DateTimeClasses` for details of the date-time classes; `locales` to query or set a locale.

Your system's help pages on `strftime` and `strptime` to see how to specify their formats.

## Examples

```
## locale-specific version of date()
format(Sys.time(), "%a %b %d %X %Y")
## we would include the timezone as in format(Sys.time(),
## "%a %b %d %X %Y %Z") but this crashes some GNU/Linux
## systems

## read in date info in format 'ddmmmyyyy' This will give
## NA(s) in some locales; setting the C locale as in the
## commented lines will overcome this on most systems. lct
## <- Sys.getlocale("LC_TIME"); Sys.setlocale("LC_TIME",
## "C")
x <- c("1jan1960", "2jan1960", "31mar1960", "30jul1960")
z <- strptime(x, "%d%b%Y")
## Sys.setlocale("LC_TIME", lct)
z

## read in date/time info in format 'm/d/y h:m:s'
dates <- c("02/27/92", "02/27/92", "01/14/92",
           "02/28/92", "02/01/92")
times <- c("23:03:20", "22:29:56", "01:03:30",
           "18:21:03", "16:56:26")
x <- paste(dates, times)
z <- strptime(x, "%m/%d/%y %H:%M:%S")
z
```

---

**Sys.time**     *Get Current Time and Timezone*

---

**Description**

`Sys.time` returns the system's idea of the current time and `Sys.timezone` returns the current time zone.

**Usage**

```
Sys.time()
Sys.timezone()
```

**Value**

`Sys.time` returns an object of class "POSIXct" (see `DateTimeClasses`).  
`Sys.timezone` returns an OS-specific character string, possibly an empty string.

**See Also**

`date` for the system time in a fixed-format character string.

**Examples**

```
Sys.time()
## locale-specific version of date()
format(Sys.time(), "%a %b %d %X %Y")

Sys.timezone()
```

---

**time**     *Sampling Times of Time Series*

---

**Description**

**time** creates the vector of times at which a time series was sampled.

**cycle** gives the positions in the cycle of each observation.

**frequency** returns the number of samples per unit time and **deltat** the time interval between observations (see **ts**).

**Usage**

```
time(x, ...)
## Default S3 method:
time(x, offset=0, ...)

cycle(x, ...)
frequency(x, ...)
deltat(x, ...)
```

**Arguments**

<b>x</b>	a univariate or multivariate time-series, or a vector or matrix.
<b>offset</b>	can be used to indicate when sampling took place in the time unit. 0 (the default) indicates the start of the unit, 0.5 the middle and 1 the end of the interval.
<b>...</b>	extra arguments for future methods.

**Details**

These are all generic functions, which will use the **tsp** attribute of **x** if it exists. **time** and **cycle** have methods for class **ts** that coerce the result to that class.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`ts`, `start`, `tsp`, `window`.

`date` for clock time, `system.time` for CPU usage.

**Examples**

```
data(presidents)
cycle(presidents)
# a simple series plot: c() makes the x and y arguments
# into vectors
plot(c(time(presidents)), c(presidents), type="l")
```

---

**ts**     *Time-Series Objects*


---

**Description**

The function **ts** is used to create time-series objects.

**as.ts** and **is.ts** coerce an object to a time-series and test whether an object is a time series.

**Usage**

```
ts(data = NA, start = 1, end = numeric(0), frequency = 1,
   deltat = 1, ts.eps = getOption("ts.eps"), class = ,
   names = )
as.ts(x)
is.ts(x)
```

**Arguments**

<b>data</b>	a numeric vector or matrix of the observed time-series values. A data frame will be coerced to a numeric matrix via <b>data.matrix</b> .
<b>start</b>	the time of the first observation. Either a single number or a vector of two integers, which specify a natural time unit and a (1-based) number of samples into the time unit. See the examples for the use of the second form.
<b>end</b>	the time of the last observation, specified in the same way as <b>start</b> .
<b>frequency</b>	the number of observations per unit of time.
<b>deltat</b>	the fraction of the sampling period between successive observations; e.g., 1/12 for monthly data. Only one of <b>frequency</b> or <b>deltat</b> should be provided.
<b>ts.eps</b>	time series comparison tolerance. Frequencies are considered equal if their absolute difference is less than <b>ts.eps</b> .
<b>class</b>	class to be given to the result, or none if <b>NULL</b> or <b>"none"</b> . The default is <b>"ts"</b> for a single series, <b>c("mts", "ts")</b> for multiple series.



<b>names</b>	a character vector of names for the series in a multiple series: defaults to the colnames of <b>data</b> , or <b>Series 1</b> , <b>Series 2</b> , ....
<b>x</b>	an arbitrary R object.

## Details

The function **ts** is used to create time-series objects. These are vector or matrices with class of **"ts"** (and additional attributes) which represent data which has been sampled at equispaced points in time. In the matrix case, each column of the matrix **data** is assumed to contain a single (univariate) time series. Time series must have an least one observation, and although they need not be numeric there is very limited support for non-numeric series.

Class **"ts"** has a number of methods. In particular arithmetic will attempt to align time axes, and subsetting to extract subsets of series can be used (e.g., **EuStockMarkets**[, **"DAX"**]). However, subsetting the first (or only) dimension will return a matrix or vector, as will matrix subsetting.

The value of argument **frequency** is used when the series is sampled an integral number of times in each unit time interval. For example, one could use a value of 7 for **frequency** when the data are sampled daily, and the natural time period is a week, or 12 when the data are sampled monthly and the natural time period is a year. Values of 4 and 12 are assumed in (e.g.) **print** methods to imply a quarterly and monthly series respectively.

**as.ts** will use the **tsp** attribute of the object if it has one to set the start and end times and frequency.

**is.ts** tests if an object is a time series. It is generic: you can write methods to handle specific classes of objects, see **InternalMethods**.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

**tsp**, **frequency**, **start**, **end**, **time**, **window**; **print.ts**, the print method for time series objects; **plot.ts**, the plot method for time series objects. Standard package **ts** for many additional time-series functions.

## Examples

```
# 2nd Quarter of 1959
ts(1:10, frequency = 4, start = c(1959, 2))
# print.ts(.)
print(ts(1:10, freq = 7, start = c(12, 2)), calendar=TRUE)
## Using July 1954 as start date:
gnp <- ts(cumsum(1 + round(rnorm(100), 2)),
          start = c(1954, 7), frequency = 12)
plot(gnp) # using 'plot.ts' for time-series plot

## Multivariate
z <- ts(matrix(rnorm(300), 100, 3), start=c(1961, 1),
            frequency=12)
class(z)
plot(z)
plot(z, plot.type="single", lty=1:3)

## A phase plot:
data(nhtemp)
plot(nhtemp, c(nhtemp[-1], NA), cex = .8, col="blue",
     main = "Lag plot of New Haven temperatures")
## a clearer way to do this would be
library(ts)
plot(nhtemp, lag(nhtemp, 1), cex = .8, col="blue",
     main = "Lag plot of New Haven temperatures")
```

---

**ts-methods**      *Methods for Time Series Objects*

---

**Description**

Methods for objects of class "ts", typically the result of `ts`.

**Usage**

```
## S3 method for class 'ts':  
diff(x, lag=1, differences=1, ...)  
  
## S3 method for class 'ts':  
na.omit(object, ...)
```

**Arguments**

<code>x</code>	an object of class "ts" containing the values to be differenced.
<code>lag</code>	an integer indicating which lag to use.
<code>differences</code>	an integer indicating the order of the difference.
<code>object</code>	a univariate or multivariate time series.
<code>...</code>	further arguments to be passed to or from methods.

**Details**

The `na.omit` method omits initial and final segments with missing values in one or more of the series. 'Internal' missing values will lead to failure.

**Value**

For the `na.omit` method, a time series without missing values. The class of `object` will be preserved.

**See Also**

`diff`; `na.omit`, `na.fail`, `na.contiguous`.

---

**tsp**     *Tsp Attribute of Time-Series-like Objects*

---

**Description**

`tsp` returns the `tsp` attribute (or `NULL`). It is included for compatibility with S version 2. `tsp<-` sets the `tsp` attribute. `hasTsp` ensures `x` has a `tsp` attribute, by adding one if needed.

**Usage**

```
tsp(x)
tsp(x) <- value
hasTsp(x)
```

**Arguments**

<code>x</code>	a vector or matrix or univariate or multivariate time-series.
<code>value</code>	a numeric vector of length 3 or <code>NULL</code> .

**Details**

The `tsp` attribute was previously described here as `c(start(x), end(x), frequency(x))`, but this is incorrect. It gives the start time *in time units*, the end time and the frequency.

Assignments are checked for consistency.

Assigning `NULL` which removes the `tsp` attribute *and* any "`ts`" class of `x`.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**See Also**

`ts`, `time`, `start`.

---

**weekdays**     *Extract Parts of a POSIXt Object*

---

**Description**

Extract the weekday, month or quarter, or the Julian time (days since some origin). These are generic functions: the methods for the internal date-time classes are documented here.

**Usage**

```
weekdays(x, abbreviate)
## S3 method for class 'POSIXt':
weekdays(x, abbreviate = FALSE)

months(x, abbreviate)
## S3 method for class 'POSIXt':
months(x, abbreviate = FALSE)

quarters(x, abbreviate)
## S3 method for class 'POSIXt':
quarters(x, ...)

julian(x, ...)
## S3 method for class 'POSIXt':
julian(x, origin = as.POSIXct("1970-01-01", tz="GMT"), ...)
```

**Arguments**

<b>x</b>	an object inheriting from class "POSIXt".
<b>abbreviate</b>	logical. Should the names be abbreviated?
<b>origin</b>	an length-one object inheriting from class "POSIXt".
<b>...</b>	arguments for other methods.

**Value**

**weekdays** and **months** return a character vector of names in the locale in use.

**quarters** returns a character vector of "Q1" to "Q4".

**julian** returns the number of days (possibly fractional) since the origin, with the origin as a "origin" attribute.

**Note**

Other components such as the day of the month or the year are very easy to compute: just use `as.POSIXlt` and extract the relevant component.

**See Also**

`DateTimeClasses`

**Examples**

```
weekdays(.leap.seconds)
months(.leap.seconds)
quarters(.leap.seconds)
```

---

**window**      *Time Windows*

---

**Description**

**window** is a generic function which extracts the subset of the object **x** observed between the times **start** and **end**. If a frequency is specified, the series is then re-sampled at the new frequency.

**Usage**

```
window(x, ...)  
  
## S3 method for class 'ts':  
window(x, ...)  
  
## Default S3 method:  
window(x, start = NULL, end = NULL,  
       frequency = NULL, deltat = NULL, extend = FALSE, ...)
```

**Arguments**

<b>x</b>	a time-series or other object.
<b>start</b>	the start time of the period of interest.
<b>end</b>	the end time of the period of interest.
<b>frequency, deltat</b>	the new frequency can be specified by either (or both if they are consistent).
<b>extend</b>	logical. If true, the <b>start</b> and <b>end</b> values are allowed to extend the series. If false, attempts to extend the series give a warning and are ignored.
<b>...</b>	further arguments passed to or from other methods.

**Details**

The start and end times can be specified as for **ts**. If there is no observation at the new **start** or **end**, the immediately following (**start**) or preceding (**end**) observation time is used.

## Value

The value depends on the method. `window.default` will return a vector or matrix with an appropriate `tsp` attribute.

`window.ts` differs from `window.default` only in ensuring the result is a `ts` object.

If `extend = TRUE` the series will be padded with `NA` if needed.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## See Also

`time`, `ts`.

## Examples

```
data(presidents)
window(presidents, 1960, c(1969,4)) # values in the 1960's
window(presidents, deltat=1) # All Qtr1s
window(presidents, start=c(1945,3), deltat=1) # All Qtr3s
window(presidents, 1944, c(1979,2), extend=TRUE)
```



## Chapter 6

# Base package — datasets

---

<b>airmiles</b>	<i>Passenger Miles on Commercial US Airlines, 1937–1960</i>
-----------------	---

---

## Description

The revenue passenger miles flown by commercial airlines in the United States for each year from 1937 to 1960.

## Usage

```
data(airmiles)
```

## Format

A time-series of 24 observations; yearly, 1937–1960.

## Source

F.A.A. Statistical Handbook of Aviation.

## References

Brown, R. G. (1963) *Smoothing, Forecasting and Prediction of Discrete Time Series*. Prentice-Hall.

## Examples

```
data(airmiles)
plot(airmiles, main = "airmiles data",
     xlab = "Passenger-miles flown by US commercial airlines",
     col = 4)
```

---

airquality	<i>New York Air Quality Measurements</i>
------------	--

---

**Description**

Daily air quality measurements in New York, May to September 1973.

**Usage**

`data(airquality)`

**Format**

A data frame with 154 observations on 6 variables.

[,1]	Ozone	numeric	Ozone (ppb)
[,2]	Solar.R	numeric	Solar R (lang)
[,3]	Wind	numeric	Wind (mph)
[,4]	Temp	numeric	Temperature (degrees F)
[,5]	Month	numeric	Month (1–12)
[,6]	Day	numeric	Day of month (1–31)

**Details**

Daily readings of the following air quality values for May 1, 1973 (a Tuesday) to September 30, 1973.

- **Ozone:** Mean ozone in parts per billion from 1300 to 1500 hours at Roosevelt Island
- **Solar.R:** Solar radiation in Langleys in the frequency band 4000–7700 Angstroms from 0800 to 1200 hours at Central Park
- **Wind:** Average wind speed in miles per hour at 0700 and 1000 hours at LaGuardia Airport
- **Temp:** Maximum daily temperature in degrees Fahrenheit at La Guardia Airport.

**Source**

The data were obtained from the New York State Department of Conservation (ozone data) and the National Weather Service (meteorological data).

## References

Chambers, J. M., Cleveland, W. S., Kleiner, B. and Tukey, P. A. (1983)  
*Graphical Methods for Data Analysis*. Belmont, CA: Wadsworth.

## Examples

```
data(airquality)
pairs(airquality, panel = panel.smooth,
      main = "airquality data")
```

---

<b>anscombe</b>	<i>Anscombe's Quartet of "Identical" Simple Linear Regressions</i>
-----------------	--

---

## Description

Four  $x$ - $y$  datasets which have the same traditional statistical properties (mean, variance, correlation, regression line, etc.), yet are quite different.

## Usage

```
data(anscombe)
```

## Format

A data frame with 11 observations on 8 variables.

x1 == x2 == x3	the integers 4:14, specially arranged
x4	values 8 and 19
y1, y2, y3, y4	numbers in (3, 12.5) with mean 7.5 and sdev 2.03

## Source

Tufte, Edward R. (1989) *The Visual Display of Quantitative Information*, 13–14. Graphics Press.

## References

Anscombe, Francis J. (1973) Graphs in statistical analysis. *American Statistician*, **27**, 17–21.

## Examples

```
data(anscombe)
summary(anscombe)

## now some "magic" to do the 4 regressions in a loop:
ff <- y ~ x
for(i in 1:4) {
  ff[2:3] <- lapply(paste(c("y","x"), i, sep=""), as.name)
  ## or   ff[[2]] <- as.name(paste("y", i, sep=""))
  ##      ff[[3]] <- as.name(paste("x", i, sep=""))
}
```

```

    assign(paste("lm.",i,sep=""),
           lmi <- lm(ff, data= anscombe))
    print(anova(lmi))
  }

## See how close they are (numerically!)
sapply(objects(pat="lm\.[1-4]$", function(n) coef(get(n)))
lapply(objects(pat="lm\.[1-4]$",
              function(n) summary(get(n))$coef)

## Now, do what you should have done in the first place:
## PLOTS
op <- par(mfrow=c(2,2), mar=.1+c(4,4,1,1), oma= c(0,0,2,0))
for(i in 1:4) {
  ff[2:3] <- lapply(paste(c("y","x"), i, sep=""), as.name)
  plot(ff, data =anscombe, col="red", pch=21, bg="orange",
        cex=1.2, xlim=c(3,19), ylim=c(3,13))
  abline(get(paste("lm.",i,sep="")), col="blue")
}
mtext("Anscombe's 4 Regression data sets", outer = TRUE,
      cex=1.5)
par(op)

```

---

attenu      *The Joyner–Boore Attenuation Data*

---

**Description**

This data gives peak accelerations measured at various observation stations for 23 earthquakes in California. The data have been used by various workers to estimate the attenuating affect of distance on ground acceleration.

**Usage**

`data(attenu)`

**Format**

A data frame with 182 observations on 5 variables.

[,1]	event	numeric	Event Number
[,2]	mag	numeric	Moment Magnitude
[,3]	station	factor	Station Number
[,4]	dist	numeric	Station-hypocenter distance (km)
[,5]	accel	numeric	Peak acceleration (g)

**Source**

Joyner, W.B., D.M. Boore and R.D. Porcella (1981). Peak horizontal acceleration and velocity from strong-motion records including records from the 1979 Imperial Valley, California earthquake. USGS Open File report 81-365. Menlo Park, Ca.

**References**

Boore, D. M. and Joyner, W.B.(1982) The empirical prediction of ground motion, *Bull. Seism. Soc. Am.*, **72**, S269–S268.

Bolt, B. A. and Abrahamson, N. A. (1982) New attenuation relations for peak and expected accelerations of strong ground motion, *Bull. Seism. Soc. Am.*, **72**, 2307–2321.

Bolt B. A. and Abrahamson, N. A. (1983) Reply to W. B. Joyner & D. M. Boore’s “Comments on: New attenuation relations for peak and expected accelerations for peak and expected accelerations of strong ground motion”, *Bull. Seism. Soc. Am.*, **73**, 1481–1483.

Brillinger, D. R. and Preisler, H. K. (1984) An exploratory analysis of the Joyner-Boore attenuation data, *Bull. Seism. Soc. Am.*, **74**, 1441–1449.

Brillinger, D. R. and Preisler, H. K. (1984) *Further analysis of the Joyner-Boore attenuation data*. Manuscript.

## Examples

```
data(attenu)
## check the data class of the variables
sapply(attenu, data.class)
summary(attenu)
pairs(attenu, main = "attenu data")
coplot(accel ~ dist | as.factor(event), data = attenu,
       show = FALSE)
coplot(log(accel) ~ log(dist) | as.factor(event),
       data = attenu, panel = panel.smooth,
       show.given = FALSE)
```



---

attitude	<i>The Chatterjee–Price Attitude Data</i>
----------	---

---

Description

From a survey of the clerical employees of a large financial organization, the data are aggregated from the questionnaires of the approximately 35 employees for each of 30 (randomly selected) departments. The numbers give the percent proportion of favourable responses to seven questions in each department.

Usage

data(attitude)

Format

A dataframe with 30 observations on 7 variables. The first column are the short names from the reference, the second one the variable names in the data frame:

Y	rating	numeric	Overall rating
X[1]	complaints	numeric	Handling of employee complaints
X[2]	privileges	numeric	Does not allow special privileges
X[3]	learning	numeric	Opportunity to learn
X[4]	raises	numeric	Raises based on performance
X[5]	critical	numeric	Too critical
X[6]	advancel	numeric	Advancement

Source

Chatterjee, S. and Price, B. (1977) *Regression Analysis by Example*. New York: Wiley. (Section 3.7, p.68ff of 2nd ed.(1991).)

Examples

```
data(attitude)
pairs(attitude, main = "attitude data")
summary(attitude)
summary(fm1 <- lm(rating ~ ., data = attitude))
opar <- par(mfrow = c(2, 2), oma = c(0, 0, 1.1, 0),
            mar = c(4.1, 4.1, 2.1, 1.1))
plot(fm1)
```

```
summary(fm2 <- lm(rating ~ complaints, data = attitude))  
plot(fm2)  
par(opar)
```

---

**cars**      *Speed and Stopping Distances of Cars*


---

**Description**

The data gives the speed of cars and the distances taken to stop. Note that the data were recorded in the 1920s.

**Usage**

```
data(cars)
```

**Format**

A data frame with 50 observations on 2 variables.

[,1]	speed	numeric	Speed (mph)
[,2]	dist	numeric	Stopping distance (ft)

**Source**

Ezekiel, M. (1930) *Methods of Correlation Analysis*. Wiley.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

**Examples**

```
data(cars)
plot(cars, xlab = "Speed (mph)",
      ylab = "Stopping distance (ft)", las = 1)
lines(lowess(cars$speed, cars$dist, f = 2/3, iter = 3),
      col = "red")
title(main = "cars data")
plot(cars, xlab = "Speed (mph)",
      ylab = "Stopping distance (ft)", las = 1, log = "xy")
title(main = "cars data (logarithmic scales)")
lines(lowess(cars$speed, cars$dist, f = 2/3, iter = 3),
      col = "red")
summary(fm1 <- lm(log(dist) ~ log(speed), data = cars))
opar <- par(mfrow = c(2, 2), oma = c(0, 0, 1.1, 0),
            mar = c(4.1, 4.1, 2.1, 1.1))
```

```
plot(fm1)
par(opar)

## An example of polynomial regression
plot(cars, xlab = "Speed (mph)",
      ylab = "Stopping distance (ft)",
      las = 1, xlim = c(0, 25))
d <- seq(0, 25, len = 200)
for(degree in 1:4) {
  fm <- lm(dist ~ poly(speed, degree), data = cars)
  assign(paste("cars", degree, sep="."), fm)
  lines(d, predict(fm, data.frame(speed=d)), col = degree)
}
anova(cars.1, cars.2, cars.3, cars.4)
```

---

**chickwts**     *Chicken Weights by Feed Type*

---

**Description**

An experiment was conducted to measure and compare the effectiveness of various feed supplements on the growth rate of chickens.

**Usage**

```
data(chickwts)
```

**Format**

A data frame with 71 observations on 2 variables.

**weight** a numeric variable giving the chick weight.

**feed** a factor giving the feed type.

**Details**

Newly hatched chicks were randomly allocated into six groups, and each group was given a different feed supplement. Their weights in grams after six weeks are given along with feed types.

**Source**

Anonymous (1948) *Biometrika*, **35**, 214.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

**Examples**

```
data(chickwts)
boxplot(weight ~ feed, data = chickwts, col = "lightgray",
        varwidth = TRUE, notch = TRUE, main = "chickwt data",
        ylab = "Weight at six weeks (gm)")
anova(fm1 <- lm(weight ~ feed, data = chickwts))
opar <- par(mfrow = c(2, 2), oma = c(0, 0, 1.1, 0),
           mar = c(4.1, 4.1, 2.1, 1.1))
plot(fm1)
par(opar)
```

---

**co2**     *Mauna Loa Atmospheric CO2 Concentration*

---

**Description**

Atmospheric concentrations of CO<sub>2</sub> are expressed in parts per million (ppm) and reported in the preliminary 1997 SIO manometric mole fraction scale.

**Usage**

```
data(co2)
```

**Format**

A time series of 468 observations; monthly from 1959 to 1997.

**Details**

The values for February, March and April of 1964 were missing and have been obtained by interpolating linearly between the values for January and May of 1964.

**Source**

Keeling, C. D. and Whorf, T. P., Scripps Institution of Oceanography (SIO), University of California, La Jolla, California USA 92093-0220.

<ftp://cdiac.esd.ornl.gov/pub/maunaloa-co2/maunaloa.co2>.

**References**

Cleveland, W. S. (1993) *Visualizing Data*. New Jersey: Summit Press.

**Examples**

```
data(co2)
plot(co2,
     ylab = expression("Atmospheric concentration of CO"[2]),
     las = 1)
title(main = "co2 data set")
```

---

**data**     *Data Sets*

---

**Description**

Loads specified data sets, or list the available data sets.

**Usage**

```
data(..., list = character(0), package = .packages(),  
      lib.loc = NULL, verbose = getOption("verbose"),  
      envir = .GlobalEnv)
```

**Arguments**

<code>...</code>	a sequence of names or literal character strings.
<code>list</code>	a character vector.
<code>package</code>	a name or character vector giving the packages to look into for data sets. By default, all packages in the search path are used, then the ‘ <b>data</b> ’ subdirectory (if present) of the current working directory.
<code>lib.loc</code>	a character vector of directory names of R libraries, or <code>NULL</code> . The default value of <code>NULL</code> corresponds to all libraries currently known. If the default is used, the loaded packages are searched before the libraries.
<code>verbose</code>	a logical. If <code>TRUE</code> , additional diagnostics are printed.
<code>envir</code>	the environment where the data should be loaded.

**Details**

Currently, four formats of data files are supported:

1. files ending ‘**.R**’ or ‘**.r**’ are `source()`d in, with the R working directory changed temporarily to the directory containing the respective file.
2. files ending ‘**.RData**’ or ‘**.rda**’ are `load()`ed.
3. files ending ‘**.tab**’, ‘**.txt**’ or ‘**.TXT**’ are read using `read.table(..., header = TRUE)`, and hence result in a data frame.
4. files ending ‘**.csv**’ or ‘**.CSV**’ are read using `read.table(..., header = TRUE, sep = ";")`, and also result in a data frame.

If more than one matching file name is found, the first on this list is used.

The data sets to be loaded can be specified as a sequence of names or character strings, or as the character vector `list`, or as both.

For each given data set, the first two types (`'R'` or `'.r'`, and `'RData'` or `'rda'` files) can create several variables in the load environment, which might all be named differently from the data set. The second two (`'tab'`, `'txt'`, or `'TXT'`, and `'csv'` or `'CSV'` files) will always result in the creation of a single variable with the same name as the data set.

If no data sets are specified, `data` lists the available data sets. It looks for a new-style data index in the `'Meta'` or, if this is not found, an old-style `'00Index'` file in the `'data'` directory of each specified package, and uses these files to prepare a listing. If there is a `'data'` area but no index, available data files for loading are computed and included in the listing, and a warning is given: such packages are incomplete. The information about available data sets is returned in an object of class `"packageIQR"`. The structure of this class is experimental. In earlier versions of R, an empty character vector was returned along with listing available data sets.

If `lib.loc` is not specified, the data sets are searched for amongst those packages already loaded, followed by the `'data'` directory (if any) of the current working directory and then packages in the specified libraries. If `lib.loc` is specified, packages are searched for in the specified libraries, even if they are already loaded from another library.

To just look in the `'data'` directory of the current working directory, set `package = NULL`.

## Value

a character vector of all data sets specified, or information about all available data sets in an object of class `"packageIQR"` if none were specified.

## Note

The data files can be many small files. On some file systems it is desirable to save space, and the files in the `'data'` directory of an installed package can be zipped up as a zip archive `'Rdata.zip'`. You will need to provide a single-column file `'filelist'` of file names in that directory.

One can take advantage of the search order and the fact that a `'R'` file will change directory. If raw data are stored in `'mydata.txt'` then one can set up `'mydata.R'` to read `'mydata.txt'` and pre-process it,



e.g., using **transform**. For instance one can convert numeric vectors to factors with the appropriate labels. Thus, the ‘.R’ file can effectively contain a metadata specification for the plaintext formats.

### See Also

**help** for obtaining documentation on data sets, **save** for *creating* the second (‘.rda’) kind of data, typically the most efficient one.

### Examples

```
# list all available data sets
data()
# list the data sets in the base package
data(package = "base")
# load the data sets 'USArrests' and 'VADeaths'
data(USArrests, "VADeaths")
# give information on data set 'USArrests'
help(USArrests)
```

---

discoveries	<i>Yearly Numbers of Important Discoveries</i>
-------------	--

---

## Description

The numbers of “great” inventions and scientific discoveries in each year from 1860 to 1959.

## Usage

```
data(discoveries)
```

## Format

A time series of 100 values.

## Source

The World Almanac and Book of Facts, 1975 Edition, pages 315–318.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

## Examples

```
data(discoveries)
plot(discoveries, ylab = "Number of important discoveries",
     las = 1)
title(main = "discoveries data set")
```

---

esoph     *Smoking, Alcohol and (O)esophageal Cancer*

---

**Description**

Data from a case-control study of (o)esophageal cancer in Ile-et-Vilaine, France.

**Usage**

`data(esoph)`

**Format**

A data frame with records for 88 age/alcohol/tobacco combinations.

[,1]	"agegp"	Age group	1	25–34 years
			2	35–44
			3	45–54
			4	55–64
			5	65–74
			6	75+
[,2]	"alcgp"	Alcohol consumption	1	0–39 gm/day
			2	40–79
			3	80–119
			4	120+
[,3]	"tobgp"	Tobacco consumption	1	0– 9 gm/day
			2	10–19
			3	20–29
			4	30+
[,4]	"ncases"	Number of cases		
[,5]	"ncontrols"	Number of controls		

**Author(s)**

Thomas Lumley

**Source**

Breslow, N. E. and Day, N. E. (1980) *Statistical Methods in Cancer Research. 1: The Analysis of Case-Control Studies*. IARC Lyon / Oxford University Press.

## Examples

```
data(esoph)
summary(esoph)
## effects of alcohol, tobacco and interaction,
## age-adjusted
model1 <-
  glm(cbind(ncases, ncontrols) ~ agegp + tobgp * alcgp,
      data = esoph, family = binomial())
anova(model1)
## Try a linear effect of alcohol and tobacco
model2 <-
  glm(cbind(ncases, ncontrols) ~ agegp + unclass(tobgp)
      + unclass(alcgp), data = esoph, family = binomial())
summary(model2)
## Re-arrange data for a mosaic plot
ttt <- table(esoph$agegp, esoph$alcgp, esoph$tobgp)
ttt[ttt == 1] <- esoph$ncases
tt1 <- table(esoph$agegp, esoph$alcgp, esoph$tobgp)
tt1[tt1 == 1] <- esoph$ncontrols
tt <- array(c(ttt, tt1), c(dim(ttt),2),
            c(dimnames(ttt), list(c("Cancer", "control"))))
mosaicplot(tt, main = "esoph data set", color = TRUE)
```

---

**euro**      *Conversion Rates of Euro Currencies*

---

**Description**

Conversion rates between the various Euro currencies.

**Usage**

```
data(euro)
```

**Format**

**euro** is a named vector of length 11, **euro.cross** a named matrix of size 11 by 11.

**Details**

The data set **euro** contains the value of 1 Euro in all currencies participating in the European monetary union (Austrian Schilling ATS, Belgian Franc BEF, German Mark DEM, Spanish Peseta ESP, Finnish Markka FIM, French Franc FRF, Irish Punt IEP, Italian Lira ITL, Luxembourg Franc LUF, Dutch Guilder NLG and Portugese Escudo PTE). These conversion rates were fixed by the European Union on December 31, 1998. To convert old prices to Euro prices, divide by the respective rate and round to 2 digits.

The data set **euro.cross** contains conversion rates between the various Euro currencies, i.e., the result of `outer(1 / euro, euro)`.

**Examples**

```
data(euro)
cbind(euro)

## These relations hold:
# [6 digit precision in Euro's definition]
euro == signif(euro,6)
all(euro.cross == outer(1/euro, euro))

## Convert 20 Euro to Belgian Franc
20 * euro["BEF"]
## Convert 20 Austrian Schilling to Euro
20 / euro["ATS"]
```

```
## Convert 20 Spanish Pesetas to Italian Lira
20 * euro.cross["ESP", "ITL"]

dotchart(euro,
  main = "euro data: 1 Euro in currency unit")
dotchart(1/euro,
  main = "euro data: 1 currency unit in Euros")
dotchart(log(euro, 10),
  main = "euro data: log10(1 Euro in currency unit)")
```

---

eurodist	<i>Distances Between European Cities</i>
----------	--

---

## Description

The data give the road distances (in km) between 21 cities in Europe. The data are taken from a table in “The Cambridge Encyclopaedia”.

## Usage

```
data(eurodist)
```

## Format

A `dist` object based on 21 objects. (You must have the **mva** package loaded to have the methods for this kind of object available).

## Source

Crystal, D. Ed. (1990) *The Cambridge Encyclopaedia*. Cambridge: Cambridge University Press,

---

faithful	Old Faithful Geyser Data
----------	--------------------------

---

**Description**

Waiting time between eruptions and the duration of the eruption for the Old Faithful geyser in Yellowstone National Park, Wyoming, USA.

**Usage**

```
data(faithful)
```

**Format**

A data frame with 272 observations on 2 variables.

[,1]	eruptions	numeric	Eruption time in mins
[,2]	waiting	numeric	Waiting time to next eruption

**Details**

A closer look at `faithful$eruptions` reveals that these are heavily rounded times originally in seconds, where multiples of 5 are more frequent than expected under non-human measurement. For a “better” version of the eruptions times, see the example below.

There are many versions of this dataset around: Azzalini and Bowman (1990) use a more complete version.

**Source**

W. Härdle.

**References**

Härdle, W. (1991) *Smoothing Techniques with Implementation in S*. New York: Springer.

Azzalini, A. and Bowman, A. W. (1990). A look at some data on the Old Faithful geyser. *Applied Statistics* **39**, 357–365.

**See Also**

`geyser` in package **MASS** for the Azzalini-Bowman version.



**Examples**

```
data(faithful)
f.tit <- "faithful data: Eruptions of Old Faithful"

ne60 <- round(e60 <- 60 * faithful$eruptions)
all.equal(e60, ne60)           # relative diff. ~ 1/10000
table(zapsmall(abs(e60 - ne60))) # 0, 0.02 or 0.04
faithful$better.eruptions <- ne60 / 60
te <- table(ne60)
te[te >= 4] # (too) many multiples of 5 !
plot(names(te), te, type="h", main = f.tit,
      xlab = "Eruption time (sec)")

plot(faithful[, -3], main = f.tit,
      xlab = "Eruption time (min)",
      ylab = "Waiting time to next eruption (min)")
lines(lowess(faithful$eruptions, faithful$waiting,
             f = 2/3, iter = 3),
      col = "red")
```

---

**Formaldehyde**      *Determination of Formaldehyde*

---

**Description**

These data are from a chemical experiment to prepare a standard curve for the determination of formaldehyde by the addition of chromatropic acid and concentrated sulphuric acid and the reading of the resulting purple color on a spectrophotometer.

**Usage**

```
data(Formaldehyde)
```

**Format**

A data frame with 6 observations on 2 variables.

[,1]	carb	numeric	Carbohydrate (ml)
[,2]	optden	numeric	Optical Density

**Source**

Bennett, N. A. and N. L. Franklin (1954) *Statistical Analysis in Chemistry and the Chemical Industry*. New York: Wiley.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

**Examples**

```
data(Formaldehyde)
plot(optden ~ carb, data = Formaldehyde,
     xlab = "Carbohydrate (ml)", ylab = "Optical Density",
     main = "Formaldehyde data", col = 4, las = 1)
abline(fm1 <- lm(optden ~ carb, data = Formaldehyde))
summary(fm1)
opar <- par(mfrow = c(2,2), oma = c(0, 0, 1.1, 0))
plot(fm1)
par(opar)
```

---

**freeny**     *Freeny's Revenue Data*

---

**Description**

Freeny's data on quarterly revenue and explanatory variables.

**Usage**

```
data(freeny)
```

**Format**

There are three 'freeny' data sets.

**freeny.y** is a time series with 39 observations on quarterly revenue from (1962,2Q) to (1971,4Q).

**freeny.x** is a matrix of explanatory variables. The columns are **freeny.y** lagged 1 quarter, price index, income level, and market potential.

Finally, **freeny** is a data frame with variables **y**, **lag.quarterly.revenue**, **price.index**, **income.level**, and **market.potential** obtained from the above two data objects.

**Source**

A. E. Freeny (1977) *A Portable Linear Regression Package with Test Programs*. Bell Laboratories memorandum.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**Examples**

```
data(freeny)
summary(freeny)
pairs(freeny, main = "freeny data")
summary(fm1 <- lm(y ~ ., data = freeny))
opar <- par(mfrow = c(2, 2), oma = c(0, 0, 1.1, 0),
            mar = c(4.1, 4.1, 2.1, 1.1))
plot(fm1)
par(opar)
```

---

HairEyeColor	<i>Hair and Eye Color of Statistics Students</i>
--------------	--

---

**Description**

Distribution of hair and eye color and sex in 592 statistics students.

**Usage**

```
data(HairEyeColor)
```

**Format**

A 3-dimensional array resulting from cross-tabulating 592 observations on 3 variables. The variables and their levels are as follows:

No	Name	Levels
1	Hair	Black, Brown, Red, Blond
2	Eye	Brown, Blue, Hazel, Green
3	Sex	Male, Female

**Details**

This data set is useful for illustrating various techniques for the analysis of contingency tables, such as the standard chi-squared test or, more generally, log-linear modelling, and graphical methods such as mosaic plots, sieve diagrams or association plots.

**References**

Snee, R. D. (1974), Graphical display of two-way contingency tables. *The American Statistician*, **28**, 9–12.

Friendly, M. (1992), Graphical methods for categorical data. *SAS User Group International Conference Proceedings*, **17**, 190–200. <http://www.math.yorku.ca/SCS/sugi/sugi17-paper.html>

Friendly, M. (1992), Mosaic displays for loglinear models. *Proceedings of the Statistical Graphics Section*, American Statistical Association, pp. 61–68. <http://www.math.yorku.ca/SCS/Papers/asa92.html>

**See Also**

```
chisq.test, loglin, mosaicplot
```

**Examples**

```
data(HairEyeColor)
## Full mosaic
mosaicplot(HairEyeColor)
## Aggregate over sex:
x <- apply(HairEyeColor, c(1, 2), sum)
x
mosaicplot(x, main = "Relation between hair and eye color")
```

---

<code>infert</code>	<i>Infertility after Spontaneous and Induced Abortion</i>
---------------------	---

---

**Description**

This is a matched case-control study dating from before the availability of conditional logistic regression.

**Usage**

`data(infert)`

**Format**

1.	Education	0 = 0-5 years 1 = 6-11 years 2 = 12+ years
2.	age	age in years of case
3.	parity	count
4.	number of prior induced abortions	0 = 0 1 = 1 2 = 2 or more
5.	case status	1 = case 0 = control
6.	number of prior spontaneous abortions	0 = 0 1 = 1 2 = 2 or more
7.	matched set number	1-83
8.	stratum number	1-63

**Note**

One case with two prior spontaneous abortions and two prior induced abortions is omitted.

**Source**

Trichopoulos et al. (1976) *Br. J. of Obst. and Gynaec.* **83**, 645–650.

## Examples

```
data(infert)
model1 <- glm(case ~ spontaneous+induced, data=infert,
              family=binomial())
summary(model1)
## adjusted for other potential confounders:
summary(model2 <-
  glm(case ~ age+parity+education+spontaneous+induced,
       data=infert, family=binomial()))
## Really should be analysed by conditional logistic
## regression which is in the survival package
if(require(survival)){
  model3 <-
    clogit(case~spontaneous+induced+strata(stratum),
           data=infert)
  summary(model3)
  detach() # survival (conflicts)
}
```

---

**InsectSprays**     *Effectiveness of Insect Sprays*

---

**Description**

The counts of insects in agricultural experimental units treated with different insecticides.

**Usage**

```
data(InsectSprays)
```

**Format**

A data frame with 72 observations on 2 variables.

[,1]	count	numeric	Insect count
[,2]	spray	factor	The type of spray

**Source**

Beall, G., (1942) The Transformation of data from entomological field experiments, *Biometrika*, **29**, 243–262.

**References**

McNeil, D. (1977) *Interactive Data Analysis*. New York: Wiley.

**Examples**

```
data(InsectSprays)
boxplot(count ~ spray, data = InsectSprays,
        xlab = "Type of spray", ylab = "Insect count",
        main = "InsectSprays data", varwidth = TRUE,
        col = "lightgray")
fm1 <- aov(count ~ spray, data = InsectSprays)
summary(fm1)
opar <- par(mfrow = c(2,2), oma = c(0, 0, 1.1, 0))
plot(fm1)
fm2 <- aov(sqrt(count) ~ spray, data = InsectSprays)
summary(fm2)
plot(fm2)
par(opar)
```



---

**iris**     *Edgar Anderson's Iris Data*

---

## Description

This famous (Fisher's or Anderson's) iris data set gives the measurements in centimeters of the variables sepal length and width and petal length and width, respectively, for 50 flowers from each of 3 species of iris. The species are *Iris setosa*, *versicolor*, and *virginica*.

## Usage

```
data(iris)
data(iris3)
```

## Format

**iris** is a data frame with 150 cases (rows) and 5 variables (columns) named **Sepal.Length**, **Sepal.Width**, **Petal.Length**, **Petal.Width**, and **Species**.

**iris3** gives the same data arranged as a 3-dimensional array of size 50 by 4 by 3, as represented by S-PLUS. The first dimension gives the case number within the species subsample, the second the measurements with names **Sepal L.**, **Sepal W.**, **Petal L.**, and **Petal W.**, and the third the species.

## Source

Fisher, R. A. (1936) The use of multiple measurements in taxonomic problems. *Annals of Eugenics*, **7**, Part II, 179–188.

The data were collected by Anderson, Edgar (1935). The irises of the Gaspé Peninsula, *Bulletin of the American Iris Society*, **59**, 2–5.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole. (has **iris3** as **iris**.)

## See Also

**matplot** some examples of which use **iris**.

## Examples

```
data(iris3)
dni3 <- dimnames(iris3)
ii <- data.frame(matrix(aperm(iris3, c(1,3,2)), ncol=4,
                        dimnames=list(NULL, sub(" L.", ".Length",
                        sub(" W.", ".Width", dni3[[2]]))),
                        Species = gl(3,50,lab=sub("S", "s",
                        sub("V", "v", dni3[[3]]))))
data(iris)
all.equal(ii, iris) # TRUE
```

---

islands	<i>Areas of the World's Major Landmasses</i>
---------	--

---

## Description

The areas in thousands of square miles of the landmasses which exceed 10,000 square miles.

## Usage

```
data(islands)
```

## Format

A named vector of length 48.

## Source

The World Almanac and Book of Facts, 1975, page 406.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

## Examples

```
data(islands)
dotchart(log(islands, 10),
  main = "islands data: log10(area) (log10(sq. miles))")
dotchart(log(islands[order(islands)], 10),
  main = "islands data: log10(area) (log10(sq. miles))")
```

---

LifeCycleSavings	<i>Intercountry Life-Cycle Savings Data</i>
------------------	---

---

**Description**

Data on the savings ratio 1960–1970.

**Usage**

```
data(LifeCycleSavings)
```

**Format**

A data frame with 50 observations on 5 variables.

[,1]	sr	numeric	aggregate personal savings
[,2]	pop15	numeric	% of population under 15
[,3]	pop75	numeric	% of population over 75
[,4]	dpi	numeric	real per-capita disposable income
[,5]	ddpi	numeric	% growth rate of dpi

**Details**

Under the life-cycle savings hypothesis as developed by Franco Modigliani, the savings ratio (aggregate personal saving divided by disposable income) is explained by per-capita disposable income, the percentage rate of change in per-capita disposable income, and two demographic variables: the percentage of population less than 15 years old and the percentage of the population over 75 years old. The data are averaged over the decade 1960–1970 to remove the business cycle or other short-term fluctuations.

**Source**

The data were obtained from Belsley, Kuh and Welsch (1980). They in turn obtained the data from Sterling (1977).

**References**

Sterling, Arnie (1977) Unpublished BS Thesis. Massachusetts Institute of Technology.

Belsley, D. A., Kuh. E. and Welsch, R. E. (1980) *Regression Diagnostics*. New York: Wiley.

**Examples**

```
data(LifeCycleSavings)
pairs(LifeCycleSavings, panel = panel.smooth,
      main = "LifeCycleSavings data")
fm1 <- lm(sr ~ pop15 + pop75 + dpi + ddpi,
          data = LifeCycleSavings)
summary(fm1)
```

---

longley	<i>Longley's Economic Regression Data</i>
---------	---

---

## Description

A macroeconomic data set which provides a well-known example for a highly collinear regression.

## Usage

```
data(longley)
```

## Format

A data frame with 7 economical variables, observed yearly from 1947 to 1962 ( $n = 16$ ).

**GNP.deflator:** GNP implicit price deflator (1954 = 100)

**GNP:** Gross National Product.

**Unemployed:** number of unemployed.

**Armed.Forces:** number of people in the armed forces.

**Population:** 'noninstitutionalized' population  $\geq 14$  years of age.

**Year:** the year (time).

**Employed:** number of people employed.

The regression `lm(Employed ~ .)` is known to be highly collinear.

## Source

J. W. Longley (1967) An appraisal of least-squares programs from the point of view of the user. *Journal of the American Statistical Association*, **62**, 819–841.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

**Examples**

```
## give the data set in the form it is used in S-PLUS:
data(longley)
longley.x <- data.matrix(longley[, 1:6])
longley.y <- longley[, "Employed"]
pairs(longley, main = "longley data")
summary(fm1 <- lm(Employed ~ ., data = longley))
opar <- par(mfrow = c(2, 2), oma = c(0, 0, 1.1, 0),
            mar = c(4.1, 4.1, 2.1, 1.1))
plot(fm1)
par(opar)
```

---

**morley**     *Michaelson-Morley Speed of Light Data*

---

**Description**

The classical data of Michaelson and Morley on the speed of light. The data consists of five experiments, each consisting of 20 consecutive ‘runs’. The response is the speed of light measurement, suitably coded.

**Usage**

```
data(morley)
```

**Format**

A data frame contains the following components:

**Expt** The experiment number, from 1 to 5.

**Run** The run number within each experiment.

**Speed** Speed-of-light measurement.

**Details**

The data is here viewed as a randomized block experiment with ‘experiment’ and ‘run’ as the factors. ‘run’ may also be considered a quantitative variate to account for linear (or polynomial) changes in the measurement over the course of a single experiment.

**Source**

A. J. Weekes (1986) *A Genstat Primer*. London: Edward Arnold.

**Examples**

```
data(morley)
morley$Expt <- factor(morley$Expt)
morley$Run <- factor(morley$Run)
attach(morley)
plot(Expt, Speed, main = "Speed of Light Data",
     xlab = "Experiment No.")
fm <- aov(Speed ~ Run + Expt, data = morley)
summary(fm)
```



```
fm0 <- update(fm, . ~ . - Run)
anova(fm0, fm)
detach(morley)
```

**mtcars**     *Motor Trend Car Road Tests*

---

## Description

The data was extracted from the 1974 *Motor Trend* US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models).

## Usage

```
data(mtcars)
```

## Format

A data frame with 32 observations on 11 variables.

[, 1]	mpg	Miles/(US) gallon
[, 2]	cyl	Number of cylinders
[, 3]	disp	Displacement (cu.in.)
[, 4]	hp	Gross horsepower
[, 5]	drat	Rear axle ratio
[, 6]	wt	Weight (lb/1000)
[, 7]	qsec	1/4 mile time
[, 8]	vs	V/S
[, 9]	am	Transmission (0 = automatic, 1 = manual)
[,10]	gear	Number of forward gears
[,11]	carb	Number of carburettors

## Source

Henderson and Velleman (1981), Building multiple regression models interactively. *Biometrics*, **37**, 391–411.

## Examples

```
data(mtcars)
pairs(mtcars, main = "mtcars data")
coplot(mpg ~ disp | as.factor(cyl), data = mtcars,
       panel = panel.smooth, rows = 1)
```

---

<b>nhtemp</b>	<i>Average Yearly Temperatures in New Haven</i>
---------------	---

---

## Description

The mean annual temperature in degrees Fahrenheit in New Haven, Connecticut, from 1912 to 1971.

## Usage

```
data(nhtemp)
```

## Format

A time series of 60 observations.

## Source

Vaux, J. E. and Brinker, N. B. (1972) *Cycles*, **1972**, 117–121.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

## Examples

```
data(nhtemp)
plot(nhtemp, main = "nhtemp data",
     ylab = "Mean annual temperature in New Haven (deg F)")
```

---

OrchardSpraysPotency of Orchard Sprays

---

Description

An experiment was conducted to assess the potency of various constituents of orchard sprays in repelling honeybees, using a Latin square design.

Usage

`data(OrchardSprays)`

Format

A data frame with 64 observations on 4 variables.

[,1]	rowpos	numeric	Row of the design
[,2]	colpos	numeric	Column of the design
[,3]	treatment	factor	Treatment level
[,4]	decrease	numeric	Response

Details

Individual cells of dry comb were filled with measured amounts of lime sulphur emulsion in sucrose solution. Seven different concentrations of lime sulphur ranging from a concentration of 1/100 to 1/1,562,500 in successive factors of 1/5 were used as well as a solution containing no lime sulphur.

The responses for the different solutions were obtained by releasing 100 bees into the chamber for two hours, and then measuring the decrease in volume of the solutions in the various cells.

An  $8 \times 8$  Latin square design was used and the treatments were coded as follows:

- Ahighest level of lime sulphur
- Bnext highest level of lime sulphur
- .
- .
- .
- Glowest level of lime sulphur
- Hno lime sulphur

## Source

Finney, D. J. (1947) *Probit Analysis*. Cambridge.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

## Examples

```
data(OrchardSprays)
pairs(OrchardSprays, main = "OrchardSprays data")
```

---

phones      *The World's Telephones*

---

## Description

The number of telephones in various regions of the world (in thousands).

## Usage

```
data(phones)
```

## Format

A matrix with 7 rows and 8 columns. The columns of the matrix give the figures for a given region, and the rows the figures for a year.

The regions are: North America, Europe, Asia, South America, Oceania, Africa, Central America.

The years are: 1951, 1956, 1957, 1958, 1959, 1960, 1961.

## Source

AT&T (1961) *The World's Telephones*.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

## Examples

```
data(phones)
matplot(rownames(phones), phones, type = "b", log = "y",
        xlab = "Year",
        ylab = "Number of telephones (1000's)")
legend(1951.5, 80000, colnames(phones), col = 1:6,
       lty = 1:5, pch = rep(21, 7))
title(main = "phones data: log scale for response")
```

---

**PlantGrowth**     *Results from an Experiment on Plant Growth*

---

**Description**

Results from an experiment to compare yields (as measured by dried weight of plants) obtained under a control and two different treatment conditions.

**Usage**

```
data(PlantGrowth)
```

**Format**

A data frame of 30 cases on 2 variables.

[, 1]	weight	numeric
[, 2]	group	factor

The levels of `group` are 'ctrl', 'trt1', and 'trt2'.

**Source**

Dobson, A. J. (1983) *An Introduction to Statistical Modelling*. London: Chapman and Hall.

**Examples**

```
## One factor ANOVA example from Dobson's book, cf. Table
## 7.4:
data(PlantGrowth)
boxplot(weight ~ group, data = PlantGrowth,
        main = "PlantGrowth data",
        ylab = "Dried weight of plants", col = "lightgray",
        notch = TRUE, varwidth = TRUE)
anova(lm(weight ~ group, data = PlantGrowth))
```

**precip**     *Annual Precipitation in US Cities*

---

### Description

The average amount of precipitation (rainfall) in inches for each of 70 United States (and Puerto Rico) cities.

### Usage

```
data(precip)
```

### Format

A named vector of length 70.

### Source

Statistical Abstracts of the United States, 1975.

### References

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

### Examples

```
data(precip)
dotchart(precip[order(precip)], main = "precip data")
title(sub = "Average annual precipitation (in.)")
```



---

presidents	<i>Quarterly Approval Ratings of US Presidents</i>
------------	--

---

**Description**

The (approximately) quarterly approval rating for the President of the United states from the first quarter of 1945 to the last quarter of 1974.

**Usage**

```
data(presidents)
```

**Format**

A time series of 120 values.

**Details**

The data are actually a fudged version of the approval ratings. See McNeil's book for details.

**Source**

The Gallup Organisation.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

**Examples**

```
data(presidents)
plot(presidents, las = 1, ylab = "Approval rating (%)",
     main = "presidents data")
```

---

<b>pressure</b> <i>ature</i>	<i>Vapor Pressure of Mercury as a Function of Temper-</i>
---------------------------------	---

---

## Description

Data on the relation between temperature in degrees Celsius and vapor pressure of mercury in millimeters (of mercury).

## Usage

```
data(pressure)
```

## Format

A data frame with 19 observations on 2 variables.

[, 1]	temperature	numeric	temperature (deg C)
[, 2]	pressure	numeric	pressure (mm)

## Source

Weast, R. C., ed. (1973) *Handbook of Chemistry and Physics*. CRC Press.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

## Examples

```
data(pressure)
plot(pressure, xlab = "Temperature (deg C)",
     ylab = "Pressure (mm of Hg)",
     main = "pressure data: Vapor Pressure of Mercury")
plot(pressure, xlab = "Temperature (deg C)", log = "y",
     ylab = "Pressure (mm of Hg)",
     main = "pressure data: Vapor Pressure of Mercury")
```

---

quakes	<i>Locations of Earthquakes off Fiji</i>
--------	--

---

**Description**

The data set give the locations of 1000 seismic events of MB > 4.0 The events occurred in a cube near Fiji since 1964.

**Usage**

`data(quakes)`

**Format**

A data frame with 1000 observations on 5 variables.

[,1]	lat	numeric	Latitude of event
[,2]	long	numeric	Longitude
[,3]	depth	numeric	Depth (km)
[,4]	mag	numeric	Richter Magnitude
[,5]	stations	numeric	Number of stations reporting

**Details**

There are two clear planes of seismic activity. One is a major plate junction; the other is the Tonga trench off New Zealand. These data constitute a subsample from a larger dataset of containing 5000 observations.

**Source**

This is one of the Harvard PRIM-H project data sets. They in turn obtained it from Dr. John Woodhouse, Dept. of Geophysics, Harvard University.

**Examples**

```
data(quakes)
pairs(quakes, main = "Fiji Earthquakes, N = 1000",
      cex.main=1.2, pch=".")
```

---

randu	<i>Random Numbers from Congruential Generator RANDU</i>
-------	---

---

## Description

400 triples of successive random numbers were taken from the VAX FORTRAN function RANDU running under VMS 1.5.

## Usage

```
data(randu)
```

## Format

A data frame with 400 observations on 3 variables named **x**, **y** and **z** which give the first, second and third random number in the triple.

## Details

In three dimensional displays it is evident that the triples fall on 15 parallel planes in 3-space. This can be shown theoretically to be true for all triples from the RANDU generator.

These particular 400 triples start 5 apart in the sequence, that is they are  $((U[5i+1], U[5i+2], U[5i+3]), i = 0, \dots, 399)$ , and they are rounded to 6 decimal places.

Under VMS versions 2.0 and higher, this problem has been fixed.

## Source

David Donoho

## Examples

```
## We could re-generate the dataset by the following R code
seed <- as.double(1)
RANDU <- function() {
  seed <- ((2^16 + 3) * seed) %% (2^31)
  seed/(2^31)
}
for(i in 1:400) {
  U <- c(RANDU(), RANDU(), RANDU(), RANDU(), RANDU())
  print(round(U[1:3], 6))
}
```

---

<b>rivers</b>	<i>Lengths of Major North American Rivers</i>
---------------	---

---

**Description**

This data set gives the lengths (in miles) of 141 “major” rivers in North America, as compiled by the US Geological Survey.

**Usage**

`data(rivers)`

**Format**

A vector containing 141 observations.

**Source**

World Almanac and Book of Facts, 1975, page 406.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

---

**sleep**     *Student's Sleep Data*

---

### Description

Data which show the effect of two soporific drugs (increase in hours of sleep) on groups consisting of 10 patients each.

### Usage

```
data(sleep)
```

### Format

A data frame with 20 observations on 2 variables.

[, 1]	extra	numeric	increase in hours of sleep
[, 2]	group	factor	patient group

### Source

Student (1908) The probable error of the mean. *Biometrika*, **6**, 20.

### References

Scheffé, Henry (1959) *The Analysis of Variance*. New York, NY: Wiley.

### Examples

```
data(sleep)
## ANOVA
anova(lm(extra ~ group, data = sleep))
```

---

**stackloss**     *Brownlee's Stack Loss Plant Data*

---

## Description

Operational data of a plant for the oxidation of ammonia to nitric acid.

## Usage

```
data(stackloss)
```

## Format

**stackloss** is a data frame with 21 observations on 4 variables.

[,1]	<b>Air Flow</b>	Flow of cooling air
[,2]	<b>Water Temp</b>	Cooling Water Inlet Temperature
[,3]	<b>Acid Conc.</b>	Concentration of acid [per 1000, minus 500]
[,4]	<b>stack.loss</b>	Stack loss

For compatibility with S-PLUS, the data sets **stack.x**, a matrix with the first three (independent) variables of the data frame, and **stack.loss**, the numeric vector giving the fourth (dependent) variable, are provided as well.

## Details

“Obtained from 21 days of operation of a plant for the oxidation of ammonia ( $\text{NH}_3$ ) to nitric acid ( $\text{HNO}_3$ ). The nitric oxides produced are absorbed in a countercurrent absorption tower”. (Brownlee, cited by Dodge, slightly reformatted by MM.)

**Air Flow** represents the rate of operation of the plant. **Water Temp** is the temperature of cooling water circulated through coils in the absorption tower. **Acid Conc.** is the concentration of the acid circulating, minus 50, times 10: that is, 89 corresponds to 58.9 per cent acid. **stack.loss** (the dependent variable) is 10 times the percentage of the ingoing ammonia to the plant that escapes from the absorption column unabsorbed; that is, an (inverse) measure of the over-all efficiency of the plant.

## Source

Brownlee, K. A. (1960, 2nd ed. 1965) *Statistical Theory and Methodology in Science and Engineering*. New York: Wiley. pp. 491–500.

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

Dodge, Y. (1996) The guinea pig of multiple regression. In: *Robust Statistics, Data Analysis, and Computer Intensive Methods; In Honor of Peter Huber's 60th Birthday, 1996, Lecture Notes in Statistics* **109**, Springer-Verlag, New York.

## Examples

```
data(stackloss)
summary(lm.stack <- lm(stack.loss ~ stack.x))
```



---

**state**     *US State Facts and Figures*


---

**Description**

Data sets related to the 50 states of the United States of America.

**Usage**

```
data(state)
```

**Details**

R currently contains the following “state” data sets. Note that all data are arranged according to alphabetical order of the state names.

**state.abb:** character vector of 2-letter abbreviations for the state names.

**state.area:** numeric vector of state areas (in square miles).

**state.center:** list with components named **x** and **y** giving the approximate geographic center of each state in negative longitude and latitude. Alaska and Hawaii are placed just off the West Coast.

**state.division:** factor giving state divisions (New England, Middle Atlantic, South Atlantic, East South Central, West South Central, East North Central, West North Central, Mountain, and Pacific).

**state.name:** character vector giving the full state names.

**state.region:** factor giving the region (Northeast, South, North Central, West) that each state belongs to.

**state.x77:** matrix with 50 rows and 8 columns giving the following statistics in the respective columns.

**Population:** population estimate as of July 1, 1975

**Income:** per capita income (1974)

**Illiteracy:** illiteracy (1970, percent of population)

**Life Exp:** life expectancy in years (1969–71)

**Murder:** murder and non-negligent manslaughter rate per 100,000 population (1976)

**HS Grad:** percent high-school graduates (1970)

**Frost:** mean number of days with minimum temperature below freezing (1931–1960) in capital or large city

**Area:** land area in square miles

**Source**

U.S. Department of Commerce, Bureau of the Census (1977) *Statistical Abstract of the United States*.

U.S. Department of Commerce, Bureau of the Census (1977) *County and City Data Book*.

**References**

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

---

**sunspots**      *Monthly Sunspot Numbers, 1749–1983*

---

**Description**

Monthly mean relative sunspot numbers from 1749 to 1983. Collected at Swiss Federal Observatory, Zurich until 1960, then Tokyo Astronomical Observatory.

**Usage**

```
data(sunspots)
```

**Format**

A time series of monthly data from 1749 to 1983.

**Source**

Andrews, D. F. and Herzberg, A. M. (1985) *Data: A Collection of Problems from Many Fields for the Student and Research Worker*. New York: Springer-Verlag.

**See Also**

`sunspot.month` (package `ts`) has a longer (and a bit different) series.

**Examples**

```
data(sunspots)
plot(sunspots, main = "sunspots data", xlab = "Year",
     ylab = "Monthly sunspot numbers")
```

---

swiss	<i>Swiss Fertility and Socioeconomic Indicators (1888) Data</i>
-------	---

---

**Description**

Standardized fertility measure and socio-economic indicators for each of 47 French-speaking provinces of Switzerland at about 1888.

**Usage**

`data(swiss)`

**Format**

A data frame with 47 observations on 6 variables, *each* of which is in percent, i.e., in [0,100].

[,1]	Fertility	$I_g$ , “common standardized fertility measure”
[,2]	Agriculture	% of males involved in agricultural occupation
[,3]	Examination	% “draftees” with highest mark on army exam
[,4]	Education	% “draftees” educated beyond primary school
[,5]	Catholic	% catholic (as opposed to “protestant”)
[,6]	Infant.Mortality	live births who live less than 1 year.

All variables but ‘Fertility’ give proportions of the population.

**Details**

(paraphrasing Mosteller and Tukey):

Switzerland, in 1888, was entering a period known as the “demographic transition”; i.e., its fertility was beginning to fall from the high level typical of underdeveloped countries.

The data collected are for 47 French-speaking “provinces” at about 1888.

Here, all variables are scaled to [0,100], where in the original, all but “Catholic” were scaled to [0,1].

**Source**

Project “16P5”, pages 549–551 in

Mosteller, F. and Tukey, J. W. (1977) *Data Analysis and Regression: A Second Course in Statistics*. Addison-Wesley, Reading Mass.

indicating their source as “Data used by permission of Franice van de Walle. Office of Population Research, Princeton University, 1976. Unpublished data assembled under NICHD contract number No 1-HD-O-2077.”

## References

Becker, R. A., Chambers, J. M. and Wilks, A. R. (1988) *The New S Language*. Wadsworth & Brooks/Cole.

## Examples

```
data(swiss)
pairs(swiss, panel = panel.smooth, main = "swiss data",
      col = 3 + (swiss$Catholic > 50))
summary(lm(Fertility ~ . , data = swiss))
```

---

Titanic      *Survival of passengers on the Titanic*

---

**Description**

This data set provides information on the fate of passengers on the fatal maiden voyage of the ocean liner ‘Titanic’, summarized according to economic status (class), sex, age and survival.

**Usage**

```
data(Titanic)
```

**Format**

A 4-dimensional array resulting from cross-tabulating 2201 observations on 4 variables. The variables and their levels are as follows:

No	Name	Levels
1	Class	1st, 2nd, 3rd, Crew
2	Sex	Male, Female
3	Age	Child, Adult
4	Survived	No, Yes

**Details**

The sinking of the Titanic is a famous event, and new books are still being published about it. Many well-known facts—from the proportions of first-class passengers to the “women and children first” policy, and the fact that that policy was not entirely successful in saving the women and children in the third class—are reflected in the survival rates for various classes of passenger.

These data were originally collected by the British Board of Trade in their investigation of the sinking. Note that there is not complete agreement among primary sources as to the exact numbers on board, rescued, or lost.

**Source**

Dawson, Robert J. MacG. (1995), The ‘Unusual Episode’ Data Revisited. *Journal of Statistics Education*, **3**. <http://www.amstat.org/publications/jse/v3n3/datasets.dawson.html>

The source provides a data set recording class, sex, age, and survival status for each person on board of the Titanic, and is based on data originally collected by the British Board of Trade and reprinted in:

British Board of Trade (1990), *Report on the Loss of the 'Titanic' (S.S.)*. British Board of Trade Inquiry Report (reprint). Gloucester, UK: Allan Sutton Publishing.

## Examples

```
data(Titanic)
mosaicplot(Titanic, main = "Survival on the Titanic")
## Higher survival rates in children?
apply(Titanic, c(3, 4), sum)
## Higher survival rates in females?
apply(Titanic, c(2, 4), sum)
## Use loglm() in package 'MASS' for further analysis ...
```

---

**ToothGrowth**     *The Effect of Vitamin C on Tooth Growth in Guinea Pigs*

---

**Description**

The response is the length of odontoblasts (teeth) in each of 10 guinea pigs at each of three dose levels of Vitamin C (0.5, 1, and 2 mg) with each of two delivery methods (orange juice or ascorbic acid).

**Usage**

```
data(ToothGrowth)
```

**Format**

A data frame with 60 observations on 3 variables.

[,1]	len	numeric	Tooth length
[,2]	supp	factor	Supplement type (VC or OJ).
[,3]	dose	numeric	Dose in milligrams.

**Source**

C. I. Bliss (1952) *The Statistics of Bioassay*. Academic Press.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

**Examples**

```
data(ToothGrowth)
coplot(len ~ dose | supp, data = ToothGrowth,
       panel = panel.smooth,
       xlab = "length vs dose, given type of supplement")
```



---

**trees**     *Girth, Height and Volume for Black Cherry Trees*

---

**Description**

This data set provides measurements of the girth, height and volume of timber in 31 felled black cherry trees. Note that girth is the diameter of the tree (in inches) measured at 4 ft 6 in above the ground.

**Usage**

```
data(trees)
```

**Format**

A data frame with 31 observations on 3 variables.

[,1]	Girth	numeric	Tree diameter in inches
[,2]	Height	numeric	Height in ft
[,3]	Volume	numeric	Volume of timber in cubic ft

**Source**

Ryan, T. A., Joiner, B. L. and Ryan, B. F. (1976) *The Minitab Student Handbook*. Duxbury Press.

**References**

Atkinson, A. C. (1985) *Plots, Transformations and Regression*. Oxford University Press.

**Examples**

```
data(trees)
pairs(trees, panel = panel.smooth, main = "trees data")
plot(Volume ~ Girth, data = trees, log = "xy")
coplot(log(Volume) ~ log(Girth) | Height, data = trees,
        panel = panel.smooth)
summary(fm1 <- lm(log(Volume) ~ log(Girth), data=trees))
summary(fm2 <- update(fm1, ~ . + log(Height), data=trees))
step(fm2)
## i.e., Volume ~ c * Height * Girth^2 seems reasonable
```

---

UCBA admissions      *Student Admissions at UC Berkeley*

---

## Description

Aggregate data on applicants to graduate school at Berkeley for the six largest departments in 1973 classified by admission and sex.

## Usage

```
data(UCBA admissions)
```

## Format

A 3-dimensional array resulting from cross-tabulating 4526 observations on 3 variables. The variables and their levels are as follows:

No	Name	Levels
1	Admit	Admitted, Rejected
2	Gender	Male, Female
3	Dept	A, B, C, D, E, F

## Details

This data set is frequently used for illustrating Simpson’s paradox, see Bickel et al. (1975). At issue is whether the data show evidence of sex bias in admission practices. There were 2691 male applicants, of whom 1198 (44.5%) were admitted, compared with 1835 female applicants of whom 557 (30.4%) were admitted. This gives a sample odds ratio of 1.83, indicating that males were almost twice as likely to be admitted. In fact, graphical methods (as in the example below) or log-linear modelling show that the apparent association between admission and sex stems from differences in the tendency of males and females to apply to the individual departments (females used to apply “more” to departments with higher rejection rates).

This data set can also be used for illustrating methods for graphical display of categorical data, such as the general-purpose mosaic plot or the “fourfold display” for 2-by-2-by- $k$  tables. See the home page of Michael Friendly (<http://www.math.yorku.ca/SCS/friendly.html>) for further information.

## References

Bickel, P. J., Hammel, E. A., and O'Connell, J. W. (1975) Sex bias in graduate admissions: Data from Berkeley. *Science*, **187**, 398–403.

## Examples

```
data(UCBAmissions)
## Data aggregated over departments
apply(UCBAmissions, c(1, 2), sum)
mosaicplot(apply(UCBAmissions, c(1, 2), sum),
            main = "Student admissions at UC Berkeley")
## Data for individual departments
opar <- par(mfrow = c(2, 3), oma = c(0, 0, 2, 0))
for(i in 1:6)
  mosaicplot(UCBAmissions[,i],
             xlab = "Admit", ylab = "Sex",
             main = paste("Department", LETTERS[i]))
mtext(expression(bold("Student admissions at Berkeley")),
        outer = TRUE, cex = 1.5)
par(opar)
```

---

USArrests	<i>Violent Crime Rates by US State</i>
-----------	--

---

**Description**

This data set contains statistics, in arrests per 100,000 residents for assault, murder, and rape in each of the 50 US states in 1973. Also given is the percent of the population living in urban areas.

**Usage**

```
data(USArrests)
```

**Format**

A data frame with 50 observations on 4 variables.

[,1]	Murder	numeric	Murder arrests (per 100,000)
[,2]	Assault	numeric	Assault arrests (per 100,000)
[,3]	UrbanPop	numeric	Percent urban population
[,4]	Rape	numeric	Rape arrests (per 100,000)

**Source**

- World Almanac and Book of facts 1975. (Crime rates).
- Statistical Abstracts of the United States 1975. (Urban rates).

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

**See Also**

The `state` data sets.

**Examples**

```
data(USArrests)
pairs(USArrests, panel = panel.smooth,
      main = "USArrests data")
```

---

USJudgeRatings	<i>Lawyers' Ratings of State Judges in the US Superior Court</i>
----------------	--

---

**Description**

Lawyers' ratings of state judges in the US Superior Court.

**Usage**

```
data(USJudgeRatings)
```

**Format**

A data frame containing 43 observations on 12 numeric variables.

[,1]	CONT	Number of contacts of lawyer with judge.
[,2]	INTG	Judicial integrity.
[,3]	DMNR	Demeanor.
[,4]	DILG	Diligence.
[,5]	CFMG	Case flow managing.
[,6]	DECI	Prompt decisions.
[,7]	PREP	Preparation for trial.
[,8]	FAMI	Familiarity with law.
[,9]	ORAL	Sound oral rulings.
[,10]	WRIT	Sound written rulings.
[,11]	PHYS	Physical ability.
[,12]	RTEN	Worthy of retention.

**Source**

New Haven Register, 14 January, 1977 (from John Hartigan).

**Examples**

```
data(USJudgeRatings)
pairs(USJudgeRatings, main = "USJudgeRatings data")
```

---

USPersonalExpenditure	<i>Personal Expenditure Data</i>
-----------------------	----------------------------------

---

## Description

This data set consists of United States personal expenditures (in billions of dollars) in the categories; food and tobacco, household operation, medical and health, personal care, and private education for the years 1940, 1945, 1950, and 1960.

## Usage

```
data(USPersonalExpenditure)
```

## Format

A matrix with 5 rows and 5 columns.

## Source

The World Almanac and Book of Facts, 1962, page 756.

## References

Tukey, J. W. (1977) *Exploratory Data Analysis*. Addison-Wesley.  
McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

## Examples

```
data(USPersonalExpenditure)
USPersonalExpenditure
eda::medpolish(log10(USPersonalExpenditure))
```

---

**uspop**     *Populations Recorded by the US Census*

---

**Description**

This data set gives the population of the United States (in millions) as recorded by the decennial census for the period 1790–1970.

**Usage**

```
data(uspop)
```

**Format**

A time series of 19 values.

**Source**

McNeil, D. R. (1977) *Interactive Data Analysis*. New York: Wiley.

**Examples**

```
data(uspop)
plot(uspop, log = "y", main = "uspop data", xlab = "Year",
      ylab = "U.S. Population (millions)")
```

---

VADeaths	<i>Death Rates in Virginia (1940)</i>
----------	---------------------------------------

---

## Description

Death rates per 100 in Virginia in 1940.

## Usage

```
data(VADeaths)
```

## Format

A matrix with 5 rows and 5 columns.

## Details

The death rates are cross-classified by age group (rows) and population group (columns). The age groups are: 50–54, 55–59, 60–64, 65–69, 70–74 and the population groups are Rural/Male, Rural/Female, Urban/Male and Urban/Female.

This provides a rather nice 3-way analysis of variance example.

## Source

Moyneau, L., Gilliam, S. K., and Florant, L. C.(1947) Differences in Virginia death rates by color, sex, age, and rural or urban residence. *American Sociological Review*, **12**, 525–535.

## References

McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

## Examples

```
data(VADeaths)
n <- length(dr <- c(VADeaths))
nam <- names(VADeaths)
d.VAD <- data.frame(
  Drate = dr,
  age = rep(ordered(rownames(VADeaths)),length=n),
  gender= gl(2,5,n, labels= c("M", "F")),
  site = gl(2,10, labels= c("rural", "urban")))
```



```
coplot(Drate ~ as.numeric(age) | gender * site,  
       data = d.VAD, panel = panel.smooth,  
       xlab = "VADeaths data - Given: gender")  
summary(aov.VAD <- aov(Drate ~ .^2, data = d.VAD))  
opar <- par(mfrow = c(2,2), oma = c(0, 0, 1.1, 0))  
plot(aov.VAD)  
par(opar)
```

---

<b>volcano</b>	<i>Topographic Information on Auckland's Maunga Whau Volcano</i>
----------------	--

---

## Description

Maunga Whau (Mt Eden) is one of about 50 volcanos in the Auckland volcanic field. This data set gives topographic information for Maunga Whau on a 10m by 10m grid.

## Usage

```
data(volcano)
```

## Format

A matrix with 87 rows and 61 columns, rows corresponding to grid lines running east to west and columns to grid lines running south to north.

## Source

Digitized from a topographic map by Ross Ihaka. These data should not be regarded as accurate.

## See Also

`filled.contour` for a nice plot.

## Examples

```
data(volcano)
filled.contour(volcano, color = terrain.colors, asp = 1)
title(main = "volcano data: filled contour map")
```

---

**warpbreaks**      *The Number of Breaks in Yarn during Weaving*


---

**Description**

This data set gives the number of warp breaks per loom, where a loom corresponds to a fixed length of yarn.

**Usage**

```
data(warpbreaks)
```

**Format**

A data frame with 54 observations on 3 variables.

[,1]	<b>breaks</b>	numeric	The number of breaks
[,2]	<b>wool</b>	factor	The type of wool (A or B)
[,3]	<b>tension</b>	factor	The level of tension (L, M, H)

There are measurements on 9 looms for each of the six types of warp (AL, AM, AH, BL, BM, BH).

**Source**

Tippett, L. H. C. (1950) *Technological Applications of Statistics*. Wiley. Page 106.

**References**

Tukey, J. W. (1977) *Exploratory Data Analysis*. Addison-Wesley.  
 McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

**See Also**

`xtabs` for ways to display these data as a table.

**Examples**

```
data(warpbreaks)
summary(warpbreaks)
opar <- par(mfrow = c(1,2), oma = c(0, 0, 1.1, 0))
plot(breaks ~ tension, data=warpbreaks, col="lightgray",
      varwidth=TRUE, subset = wool == "A", main="Wool A")
```

```
plot(breaks ~ tension, data=warpbreaks, col="lightgray",
      varwidth=TRUE, subset = wool == "B", main="Wool B")
mtext("warpbreaks data", side = 3, outer = TRUE)
par(opar)
summary(fm1 <- lm(breaks ~ wool*tension, data=warpbreaks))
anova(fm1)
```

---

women	<i>Average Heights and Weights for American Women</i>
-------	---

---

**Description**

This data set gives the average heights and weights for American women aged 30–39.

**Usage**

```
data(women)
```

**Format**

A data frame with 15 observations on 2 variables.

[,1]	height	numeric	Height (in)
[,2]	weight	numeric	Weight (lbs)

**Details**

The data set appears to have been taken from the American Society of Actuaries *Build and Blood Pressure Study* for some (unknown to us) earlier year.

The World Almanac notes: “The figures represent weights in ordinary indoor clothing and shoes, and heights with shoes”.

**Source**

The World Almanac and Book of Facts, 1975.

**References**

McNeil, D. R. (1977) *Interactive Data Analysis*. Wiley.

**Examples**

```
data(women)
plot(women, xlab = "Height (in)", ylab = "Weight (lb)",
      main = "women data: American women aged 30-39")
```



# Other Books From The Publisher

Network Theory publishes books about free software under free documentation licenses. Our current catalogue includes the following titles:

- **Comparing and Merging Files with GNU diff and patch** by David MacKenzie, Paul Eggert, and Richard Stallman (ISBN 0-9541617-5-0) \$19.95 (£12.95)
- **Version Management with CVS** by Per Cederqvist et al. (ISBN 0-9541617-1-8) \$29.95 (£19.95)
- **GNU Bash Reference Manual** by Chet Ramey and Brian Fox (ISBN 0-9541617-7-7) \$29.95 (£19.95)
- **An Introduction to R** by W.N. Venables, D.M. Smith and the R Development Core Team (ISBN 0-9541617-4-2) \$19.95 (£12.95)
- **GNU Octave Manual** by John W. Eaton (ISBN 0-9541617-2-6) \$29.99 (£19.99)
- **GNU Scientific Library Reference Manual - Second Edition** by M. Galassi, J. Davies, J. Theiler, B. Gough, G. Jungman, M. Booth, F. Rossi (ISBN 0-9541617-3-4) \$39.99 (£24.99)
- **An Introduction to Python** by Guido van Rossum and Fred L. Drake, Jr. (ISBN 0-9541617-6-9) \$19.95 (£12.95)
- **Python Language Reference Manual** by Guido van Rossum and Fred L. Drake, Jr. (ISBN 0-9541617-8-5) \$19.95 (£12.95)

All titles are available for order from bookstores worldwide. Sales of the manuals fund the development of more free software and documentation. For details visit the website <http://www.network-theory.co.uk/>





# Index

## \*Topic **NA**

- naprint, 491
- naresid, 492

## \*Topic **algebra**

- backsolve, 248
- chol, 253
- chol2inv, 256
- colSums, 258
- crossprod, 262
- eigen, 267
- matrix, 286
- qr, 294
- QR.Auxiliaries, 297
- solve, 304
- svd, 315

## \*Topic **aplot**

- abline, 6
- arrows, 8
- axis, 12
- box, 23
- bxp, 30
- contour, 38
- coplot, 42
- filled.contour, 59
- frame, 65
- grid, 69
- Hershey, 72
- image, 84
- Japanese, 90
- legend, 96
- lines, 102
- matplot, 106
- mtext, 114
- persp, 137

- plot.window, 168
- plot.xy, 170
- plotmath, 171
- points, 178
- polygon, 181
- rect, 198
- rug, 202
- screen, 204
- segments, 207
- symbols, 220
- text, 225
- title, 228

## \*Topic **arith**

- all.equal, 241
- approxfun, 243
- Arithmetic, 246
- cumsum, 263
- diff, 561
- Extremes, 270
- findInterval, 274
- gl, 276
- matmult, 285
- ppoints, 190
- prod, 293
- range, 299
- Round, 301
- sign, 303
- sort, 306
- sum, 314
- tabulate, 317

## \*Topic **array**

- backsolve, 248
- chol, 253
- chol2inv, 256

- colSums, 258
- contrast, 413
- crossprod, 262
- eigen, 267
- expand.grid, 425
- lm.fit, 459
- matmult, 285
- matplot, 106
- matrix, 286
- qr, 294
- QR.Auxiliaries, 297
- svd, 315
- \*Topic **category**
  - gl, 276
  - loglin, 471
  - plot.table, 164
- \*Topic **character**
  - strwidth, 215
- \*Topic **chron**
  - as.POSIX\*, 554
  - axis.POSIXct, 15
  - cut.POSIXt, 556
  - DateTimeClasses, 558
  - difftime, 563
  - hist.POSIXt, 565
  - rep, 568
  - round.POSIXt, 571
  - seq.POSIXt, 572
  - strptime, 575
  - Sys.time, 579
  - weekdays, 587
- \*Topic **color**
  - col2rgb, 35
  - colors, 37
  - gray, 68
  - hsv, 80
  - palette, 121
  - Palettes, 123
  - rgb, 201
- \*Topic **datasets**
  - airmiles, 592
  - airquality, 593
  - anscombe, 595
  - attenu, 597
  - attitude, 599
  - cars, 601
  - chickwts, 603
  - co2, 604
  - data, 605
  - discoveries, 608
  - esoph, 609
  - euro, 611
  - eurodist, 613
  - faithful, 614
  - Formaldehyde, 616
  - freeny, 617
  - HairEyeColor, 618
  - infert, 620
  - InsectSprays, 622
  - iris, 623
  - islands, 625
  - LifeCycleSavings, 626
  - longley, 628
  - morley, 630
  - mtcars, 632
  - nhtemp, 633
  - OrchardSprays, 634
  - phones, 636
  - PlantGrowth, 637
  - precip, 638
  - presidents, 639
  - pressure, 640
  - quakes, 641
  - randu, 642
  - rivers, 643
  - sleep, 644
  - stackloss, 645
  - state, 647
  - sunspots, 649
  - swiss, 650
  - Titanic, 652
  - ToothGrowth, 654
  - trees, 655
  - UCBAdmissions, 656
  - USArrests, 658
  - USJudgeRatings, 659

- USPersonalExpenditure, 660
- uspop, 661
- VADeaths, 662
- volcano, 664
- warpbreaks, 665
- women, 667
- \*Topic **design**
  - contrast, 413
  - contrasts, 415
  - TukeyHSD, 543
- \*Topic **device**
  - dev.xxx, 48
  - dev2, 50
  - Devices, 55
  - Gnome, 66
  - gtk, 71
  - pdf, 135
  - pictex, 141
  - png, 175
  - postscript, 184
  - quartz, 196
  - screen, 204
  - x11, 231
  - xfig, 233
- \*Topic **distribution**
  - bandwidth, 320
  - Beta, 322
  - Binomial, 324
  - birthday, 326
  - Cauchy, 328
  - Chisquare, 330
  - density, 333
  - Exponential, 338
  - FDist, 340
  - GammaDist, 342
  - Geometric, 344
  - hist, 76
  - Hypergeometric, 346
  - Logistic, 348
  - Lognormal, 350
  - Multinomial, 352
  - NegBinomial, 354
  - Normal, 357
  - Poisson, 359
  - ppoints, 190
  - qqnorm, 194
  - r2dtable, 361
  - Random, 363
  - Random.user, 368
  - sample, 370
  - SignRank, 372
  - TDist, 374
  - Tukey, 376
  - Uniform, 378
  - Weibull, 380
  - Wilcoxon, 382
- \*Topic **documentation**
  - data, 605
- \*Topic **dplot**
  - approxfun, 243
  - axTicks, 17
  - boxplot.stats, 28
  - col2rgb, 35
  - colors, 37
  - convolve, 260
  - fft, 272
  - hist, 76
  - hist.POSIXt, 565
  - hsv, 80
  - jitter, 91
  - layout, 93
  - n2mfrow, 117
  - Palettes, 123
  - panel.smooth, 125
  - par, 126
  - plot.density, 152
  - ppoints, 190
  - pretty, 192
  - screen, 204
  - splinefun, 311
  - strwidth, 215
  - units, 230
  - xy.coords, 235
  - xyz.coords, 237
- \*Topic **environment**

- layout, 93
- par, 126
- \*Topic **hplot**
  - assocplot, 10
  - barplot, 19
  - boxplot, 24
  - chull, 33
  - contour, 38
  - coplot, 42
  - curve, 46
  - dotchart, 57
  - filled.contour, 59
  - fourfoldplot, 62
  - hist, 76
  - hist.POSIXt, 565
  - image, 84
  - interaction.plot, 87
  - matplot, 106
  - mosaicplot, 110
  - pairs, 118
  - panel.smooth, 125
  - persp, 137
  - pie, 143
  - plot, 145
  - plot.data.frame, 147
  - plot.default, 148
  - plot.design, 153
  - plot.factor, 156
  - plot.formula, 157
  - plot.histogram, 159
  - plot.lm, 161
  - plot.table, 164
  - plot.ts, 166
  - qqnorm, 194
  - stars, 209
  - stripchart, 213
  - sunflowerplot, 217
  - symbols, 220
  - termplot, 223
- \*Topic **iplot**
  - dev.xxx, 48
  - frame, 65
  - identify, 82
  - layout, 93
  - locator, 104
  - par, 126
  - plot.histogram, 159
  - recordPlot, 197
- \*Topic **logic**
  - all.equal, 241
- \*Topic **manip**
  - cut.POSIXt, 556
  - expand.model.frame, 426
  - model.extract, 483
  - rep, 568
  - seq.POSIXt, 572
  - sort, 306
- \*Topic **math**
  - abs, 240
  - Bessel, 250
  - convolve, 260
  - deriv, 264
  - fft, 272
  - Hyperbolic, 277
  - integrate, 278
  - kappa, 281
  - log, 283
  - nextn, 288
  - poly, 289
  - polyroot, 291
  - Special, 309
  - splinefun, 311
  - Trig, 318
- \*Topic **methods**
  - plot.data.frame, 147
- \*Topic **models**
  - add1, 386
  - AIC, 389
  - alias, 391
  - anova, 394
  - anova.glm, 395
  - anova.lm, 397
  - aov, 399
  - AsIs, 401
  - C, 403
  - case/variable.names, 405

coef, 407  
 confint, 408  
 deviance, 417  
 df.residual, 418  
 dummy.coef, 419  
 eff.aovlist, 421  
 effects, 423  
 expand.grid, 425  
 extractAIC, 428  
 factor.scope, 430  
 family, 432  
 fitted, 435  
 formula, 436  
 glm, 439  
 glm.control, 445  
 glm.summaries, 447  
 is.empty.model, 453  
 labels, 454  
 lm.summaries, 464  
 logLik, 466  
 logLik.glm, 468  
 logLik.lm, 469  
 loglin, 471  
 make.link, 479  
 makepredictcall, 480  
 manova, 482  
 model.extract, 483  
 model.frame, 485  
 model.matrix, 487  
 model.tables, 489  
 naprint, 491  
 naresid, 492  
 offset, 497  
 power, 508  
 predict.glm, 509  
 preplot, 191  
 profile, 513  
 proj, 514  
 relevel, 517  
 replications, 518  
 residuals, 520  
 se.contrast, 521  
 stat.anova, 523

step, 525  
 summary.aov, 528  
 summary.glm, 530  
 summary.lm, 533  
 summary.manova, 536  
 terms, 538  
 terms.formula, 539  
 terms.object, 541  
 TukeyHSD, 543  
 update, 547  
 update.formula, 549  
 vcov, 550

\*Topic **multivariate**

stars, 209  
 symbols, 220

\*Topic **nonlinear**

deriv, 264  
 nlm, 493  
 optim, 498  
 vcov, 550

\*Topic **nonparametric**

sunflowerplot, 217

\*Topic **optimize**

constrOptim, 410  
 glm.control, 445  
 nlm, 493  
 optim, 498  
 optimize, 505  
 uniroot, 545

\*Topic **print**

labels, 454

\*Topic **programming**

model.extract, 483

\*Topic **regression**

anova, 394  
 anova.glm, 395  
 anova.lm, 397  
 aov, 399  
 case/variable.names, 405  
 coef, 407  
 contrast, 413  
 contrasts, 415  
 df.residual, 418

- effects, 423
- expand.model.frame, 426
- fitted, 435
- glm, 439
- glm.summaries, 447
- influence.measures, 449
- lm, 455
- lm.fit, 459
- lm.influence, 461
- lm.summaries, 464
- ls.diag, 474
- ls.print, 476
- lsfit, 477
- plot.lm, 161
- predict.glm, 509
- predict.lm, 511
- qr, 294
- residuals, 520
- stat.anova, 523
- summary.aov, 528
- summary.glm, 530
- summary.lm, 533
- termplot, 223
- weighted.residuals, 551
- \*Topic **smooth**
  - bandwidth, 320
  - density, 333
  - sunflowerplot, 217
- \*Topic **sysdata**
  - colors, 37
  - palette, 121
  - Random, 363
  - Random.user, 368
- \*Topic **ts**
  - diff, 561
  - plot.ts, 166
  - print.ts, 567
  - start, 574
  - time, 580
  - ts, 582
  - ts-methods, 585
  - tsp, 586
  - window, 589
- \*Topic **univar**
  - Extremes, 270
  - range, 299
  - sort, 306
- \*Topic **utilities**
  - all.equal, 241
  - as.POSIX\*, 554
  - axis.POSIXct, 15
  - DateTimeClasses, 558
  - dev2bitmap, 53
  - difftime, 563
  - findInterval, 274
  - integrate, 278
  - jitter, 91
  - n2mfrow, 117
  - relevel, 517
  - strptime, 575
  - Sys.time, 579
- \* (*Arithmetic*), 246
- \*.difftime (*difftime*), 563
- + (*Arithmetic*), 246
- +.POSIXt (*DateTimeClasses*), 558
- (*Arithmetic*), 246
- .POSIXt (*DateTimeClasses*), 558
- .Device (*dev.xxx*), 48
- .Devices (*dev.xxx*), 48
- .Pars (*par*), 126
- .PostScript.Options (*postscript*), 184
- .Random.seed (*Random*), 363
- .Device, 196
- .Machine, 506
- .Random.seed, 358, 379
- .leap.seconds (*DateTimeClasses*), 558
- .ps.prolog (*postscript*), 184
- / (*Arithmetic*), 246
- /.difftime (*difftime*), 563
- ==, 242
- [, 554

- [.AsIs (*AsIs*), 401
- [.POSIXct (*DateTimeClasses*), 558
- [.POSIXlt (*DateTimeClasses*), 558
- [.data.frame, 485
- [.difftime (*difftime*), 563
- [<-.POSIXct (*DateTimeClasses*), 558
- [<-.POSIXlt (*DateTimeClasses*), 558
- [ [.POSIXct (*DateTimeClasses*), 558
- %%, 262
- %% (matmult), 285
- %% (Arithmetic), 246
- %% (Arithmetic), 246
- %o%, 262
- ~ (Arithmetic), 246
- ~ (formula), 436
- abline, 6, 69, 182
- abs, 240, 303
- acos, 277
- acos (*Trig*), 318
- acosh (*Hyperbolic*), 277
- adapt, 279
- add.scope (*factor.scope*), 430
- add1, 386, 394, 429, 430, 526, 527
- AIC, 389, 428, 429
- airmiles, 592
- airquality, 593
- alias, 391, 400
- all, 242
- all.vars, 438
- all.equal, 241
- all.equal.POSIXct (*DateTimeClasses*), 558
- anova, 388, 394, 396, 398, 441, 443, 456, 474, 523
- anova.glm, 441, 443, 448, 523
- anova.lm, 457, 465, 523
- anova.glm, 395
- anova.glmlist (*anova.glm*), 395
- anova.lm, 397
- anova.lmlist (*anova.lm*), 397
- anova.mlm (*anova.lm*), 397
- anscombe, 595
- aov, 153, 387, 388, 399, 414, 415, 420, 421, 423, 428, 430, 455, 457, 464, 482, 490, 514, 515, 525, 529, 537, 542-544
- apply, 258, 259
- approx, 275, 312
- approx (*approxfun*), 243
- approxfun, 243, 312
- Arithmetic, 240, 246, 284, 285, 310, 508
- arrows, 8, 207
- as.data.frame, 466
- as.integer, 302
- as.POSIXct, 560, 576
- as.POSIXlt, 560, 588
- as.character.POSIXt (*strptime*), 575
- as.data.frame, 401
- as.data.frame.logLik (*logLik*), 466
- as.data.frame.POSIXct (*DateTimeClasses*), 558
- as.data.frame.POSIXlt (*DateTimeClasses*), 558
- as.difftime (*difftime*), 563
- as.formula (*formula*), 436
- as.matrix (*matrix*), 286

- as.matrix.POSIXlt  
    (*DateTimeClasses*),  
    558
- as.POSIX\*, 554
- as.POSIXct (*as.POSIX\**), 554
- as.POSIXlt (*as.POSIX\**), 554
- as.qr (*qr*), 294
- as.ts (*ts*), 582
- asin, 277
- asin (*Trig*), 318
- asinh (*Hyperbolic*), 277
- AsIs, 401
- assocplot, 10, 112
- atan, 277
- atan (*Trig*), 318
- atan2 (*Trig*), 318
- atanh (*Hyperbolic*), 277
- attenu, 597
- attitude, 599
- attr.all.equal (*all.equal*),  
    241
- attributes, 242, 271
- axis, 12, 15, 17, 18, 21, 31, 72,  
    126, 171, 173, 202
- axis.POSIXct, 565
- axis.POSIXct, 15
- axTicks, 13, 17, 69, 131
- backsolve, 248, 254, 305
- bandwidth, 320
- bandwidth.nrd, 321
- barplot, 19, 98, 156, 199
- bcv, 321
- Bessel, 250, 310
- bessel (*Bessel*), 250
- besselI (*Bessel*), 250
- besselJ (*Bessel*), 250
- besselK (*Bessel*), 250
- besselY (*Bessel*), 250
- Beta, 322
- beta, 251, 323
- beta (*Special*), 309
- Binomial, 324
- binomial, 443
- binomial (*family*), 432
- birthday, 326
- bitmap, 55, 176
- bitmap (*dev2bitmap*), 53
- box, 23, 31, 60, 164, 182, 199,  
    211
- boxplot, 24, 28–30, 156, 158,  
    213
- boxplot.formula, 26
- boxplot.stats, 26
- boxplot.stats, 28
- bquote, 173
- bs, 480
- bw.nrd, 333, 335
- bw.bcv (*bandwidth*), 320
- bw.nrd (*bandwidth*), 320
- bw.nrd0 (*bandwidth*), 320
- bw.SJ (*bandwidth*), 320
- bw.ucv (*bandwidth*), 320
- bxp, 25, 26, 29, 30
- C, 403, 414, 415
- c, 560
- c.POSIXct (*DateTimeClasses*),  
    558
- c.POSIXlt (*DateTimeClasses*),  
    558
- call, 96, 264
- capabilities, 56, 176, 268,  
    316
- cars, 480, 601
- case.names  
    (*case/variable.names*),  
    405
- case/variable.names, 405
- cat, 445
- Cauchy, 328
- cbind.ts (*ts*), 582
- ceiling (*Round*), 301
- character, 178, 226, 228
- check.options, 185, 188
- chickwts, 603
- chisq.test, 11, 618
- Chisquare, 330, 341



- chol, 248, **253**, 256, 268
- chol2inv, 254, **256**
- choose (*Special*), 309
- chull, **33**
- class, 76, 158, 392, 456
- close.screen (*screen*), 204
- cm (*units*), 230
- cm.colors (*Palettes*), 123
- co.intervals (*coplot*), 42
- co2, **604**
- coef, **407**, 424, 448, 465
- coefficients, 394, 435, 441, 458, 520
- coefficients (*coef*), 407
- col2rgb, **35**, 37, 121, 124, 201
- colMeans (*colSums*), 258
- colors, 35, **37**, 121, 124, 132, 133, 170
- colours (*colors*), 37
- colSums, **258**
- Comparison, 306
- complex, 291
- confint, **408**
- confint.nls, 408
- constrOptim, **410**, 502
- contour, **38**, 61, 72, 75, 85, 90, 139
- contr.helmert, 415
- contr.poly, 290, 415
- contr.sum, 403, 415
- contr.treatment, 415, 517
- contr.helmert (*contrast*), 413
- contr.poly (*contrast*), 413
- contr.sum (*contrast*), 413
- contr.treatment (*contrast*), 413
- contrast, **413**
- contrasts, 403, 414, **415**, 487, 521, 522
- contrasts<- (*contrasts*), 415
- convolve, **260**, 273, 288
- cooks.distance, 162, 462
- cooks.distance  
(*influence.measures*), 449
- coplot, **42**, 125, 438
- cos, 277
- cos (*Trig*), 318
- cosh (*Hyperbolic*), 277
- covratio, 462
- covratio  
(*influence.measures*), 449
- coxph, 224, 542
- crossprod, **262**
- cummax (*cumsum*), 263
- cummin (*cumsum*), 263
- cumprod, 293
- cumprod (*cumsum*), 263
- cumsum, **263**, 293
- curve, **46**
- cut, 85, 556
- cut.POSIXt, 560
- cut.POSIXt, **556**
- cycle (*time*), 580
- D (*deriv*), 264
- data, **605**
- data.frame, 146, 147, 286, 402, 485, 486
- data.matrix, 287
- data.frame, 401
- date, 579, 581
- DateTimeClasses, 15, 555, **558**, 564, 571, 573, 578, 579, 588
- dbeta, 343
- dbeta (*Beta*), 322
- dbinom, 355, 359, 360
- dbinom (*Binomial*), 324
- dcauchy (*Cauchy*), 328
- dchisq, 341, 343
- dchisq (*Chisquare*), 330
- deltat (*time*), 580
- density, 79, 146, 152, 160, 219, 320, 321, **333**

- deparse, 221
- deriv, 264, 493, 495
- deriv3 (*deriv*), 264
- det, 268, 296
- dev.cur, 52, 56
- dev.print, 56, 176
- dev.control (*dev2*), 50
- dev.copy (*dev2*), 50
- dev.copy2eps (*dev2*), 50
- dev.cur (*dev.xxx*), 48
- dev.interactive (*Devices*), 55
- dev.list (*dev.xxx*), 48
- dev.next (*dev.xxx*), 48
- dev.off (*dev.xxx*), 48
- dev.prev (*dev.xxx*), 48
- dev.print (*dev2*), 50
- dev.set (*dev.xxx*), 48
- dev.xxx, 48
- dev2, 50
- dev2bitmap, 53, 56
- deviance, 417, 418, 429, 448, 465
- device (*Devices*), 55
- Devices, 49, 55, 66, 71, 136, 142, 176, 188, 196, 205, 232, 234
- dexp, 381
- dexp (*Exponential*), 338
- df, 375
- df (*FDist*), 340
- df.residual, 417, 448, 465
- df.residual, 418
- dfbeta (*influence.measures*), 449
- dfbetas, 462
- dfbetas (*influence.measures*), 449
- dffits, 462
- dffits (*influence.measures*), 449
- dgamma, 323, 331, 339
- dgamma (*GammaDist*), 342
- dgeom, 355
- dgeom (*Geometric*), 344
- dhyper (*Hypergeometric*), 346
- diag, 285
- diff, 561, 585
- diff.ts, 562
- diff.ts (*ts-methods*), 585
- diffinv, 562
- difftime, 559, 560, 563, 572
- digamma (*Special*), 309
- dim, 271, 286
- dimnames, 164, 286
- discoveries, 608
- dlnorm, 358
- dlnorm (*Lognormal*), 350
- dlogis (*Logistic*), 348
- dmultinom (*Multinomial*), 352
- dnbinom, 325, 345, 360
- dnbinom (*NegBinomial*), 354
- dnorm, 351
- dnorm (*Normal*), 357
- dotchart, 21, 57, 144
- double, 237, 270
- dpois, 325, 355
- dpois (*Poisson*), 359
- drop, 285
- drop.scope (*factor.scope*), 430
- drop1, 394, 396, 398, 429, 430, 526, 527
- drop1 (*add1*), 386
- dsignrank, 383
- dsignrank (*SignRank*), 372
- dt, 329, 341
- dt (*TDist*), 374
- dummy.coef, 419
- dunif (*Uniform*), 378
- dweibull, 339
- dweibull (*Weibull*), 380
- dwilcox, 373
- dwilcox (*Wilcoxon*), 382
- ecdf, 275

- eff.aovlist, **421**
- effects, 394, **423**, 443, 448, 457, 458, 465
- eigen, **267**, 296, 316
- end, 583
- end (*start*), 574
- environment, 605
- erase.screen (*screen*), 204
- Error (*aov*), 399
- esoph, 443, **609**
- euro, **611**
- eurodist, **613**
- exp, 339
- exp (*log*), 283
- expand.model.frame, 486
- expand.grid, **425**
- expand.model.frame, **426**
- expm1 (*log*), 283
- Exponential, **338**
- expression, 96, 215, 226, 228, 264, 265
- extractAIC, 388, 390, 417, **428**, 526
- extractAIC.glm, 526
- Extremes, **270**
- factor, 43, 153, 156, 276, 317, 441, 487, 517
- factor.scope, **430**
- faithful, **614**
- family, **432**, 439, 441, 468, 479, 508
- family.glm (*glm.summaries*), 447
- family.lm (*lm.summaries*), 464
- FDist, **340**
- fft, 260, 261, **272**, 288, 334
- filled.contour, 40, 85, 664
- filled.contour, **59**
- filter, 261
- findInterval, **274**
- fitted, **435**, 448, 465
- fitted.values, 394, 407, 443, 458, 520
- fivenum, 29
- floor (*Round*), 301
- Formaldehyde, **616**
- format, 286
- format.POSIXct, 555
- format.POSIXlt, 555
- format.POSIXct (*strptime*), 575
- format.POSIXlt (*strptime*), 575
- formula, 153, 157, 264, 401, 402, **436**, 486, 538, 541, 542
- formula.lm (*lm.summaries*), 464
- forwardsolve (*backsolve*), 248
- fourfoldplot, **62**
- frame, **65**
- freeny, **617**
- frequency, 583
- frequency (*time*), 580
- function, 43, 146
- Gamma, 468
- Gamma (*family*), 432
- gamma, 250, 251, 343
- gamma (*Special*), 309
- gammaCody (*Bessel*), 250
- GammaDist, **342**
- gaussian, 468
- gaussian (*family*), 432
- Geometric, **344**
- gl, **276**
- glm, 161, 224, 387, 395, 396, 407, 414, 415, 417, 418, 432, 433, 435, 436, 438, **439**, 445, 447, 448, 450, 453, 458, 465, 468, 479, 497, 510, 520, 525,

- 526, 530, 531, 538,  
550, 551
- glm.control, 440
- glm.fit, 445
- glm.control, 445
- glm.summaries, 447
- GNOME, 55
- GNOME (*Gnome*), 66
- Gnome, 66
- gnome (*Gnome*), 66
- graphics.off, 56
- graphics.off (*dev.xxx*), 48
- gray, 37, 68, 80, 121, 124, 133,  
201
- grey (*gray*), 68
- grid, 69
- GTK, 55
- GTK (*gtk*), 71
- gtk, 71
- HairEyeColor, 618
- hasTsp (*tsp*), 586
- hat, 162, 462, 474
- hat (*influence.measures*),  
449
- hatvalues  
(*influence.measures*),  
449
- heat.colors, 37, 84, 85
- heat.colors (*Palettes*), 123
- heatmap, 85
- help, 607
- Hershey, 39, 72, 90, 226
- hist, 21, 76, 159, 160, 168,  
199, 335, 565
- hist.default, 565
- hist.POSIXt, 565
- hsv, 37, 68, 80, 85, 121, 123,  
124, 129, 201
- Hyperbolic, 277
- Hypergeometric, 346
- I, 437
- I (*AsIs*), 401
- identical, 241
- identify, 82, 105
- image, 40, 54, 56, 61, 84, 139,  
168
- Inf, 246
- infer, 443, 620
- influence, 162, 450–452, 465,  
551
- influence (*lm.influence*),  
461
- influence.measures, 461, 462
- influence.measures, 449
- InsectSprays, 622
- Insurance, 497
- integer, 270, 365
- integrate, 278
- interaction.plot, 154
- interaction.plot, 87
- InternalMethods, 286, 583
- interpSpline, 312
- inverse.gaussian, 468
- inverse.gaussian (*family*),  
432
- invisible, 121
- iris, 623
- iris3 (*iris*), 623
- is.empty.model, 453
- is.matrix (*matrix*), 286
- is.mts (*ts*), 582
- is.na.POSIXlt  
(*DateTimeClasses*),  
558
- is.qr (*qr*), 294
- is.ts (*ts*), 582
- is.unsorted (*sort*), 306
- islands, 625
- ISOdate (*strptime*), 575
- ISOdatetime (*strptime*), 575
- Japanese, 75, 90
- jitter, 91, 203, 218
- jpeg, 54–56
- jpeg (*png*), 175
- julian (*weekdays*), 587

kappa, **281**

La.chol (*chol*), 253

La.chol2inv (*chol2inv*), 256

La.eigen (*eigen*), 267

La.svd (*svd*), 315

labels, **454**

layout, 49, **93**, 117, 129, 133,  
205

lbeta (*Special*), 309

lchoose (*Special*), 309

lcm (*layout*), 93

legend, 88, **96**, 199

levelplot, 60, 61

lgamma (*Special*), 309

LifeCycleSavings, **626**

lines, 7, 47, 69, **102**, 107, 125,  
126, 139, 146, 159,  
166, 167, 170, 179,  
182, 207, 223, 236

lines.formula  
(*plot.formula*), 157

lines.histogram  
(*plot.histogram*),  
159

lines.ts (*plot.ts*), 166

list, 295

lm, 161, 224, 387, 388, 397, 398,  
400, 405, 407, 414,  
415, 417, 418, 423,  
428, 430, 435, 436,  
438, 443, 450, 453,  
**455**, 459–462, 464,  
465, 469, 476, 478,  
512, 515, 520, 525,  
533, 534, 538, 551

lm.fit, 456, 458

lm.influence, 162, 450–452,  
458, 474, 476, 551

lm.wfit, 458

lm.fit, **459**

lm.influence, **461**

lm.summaries, **464**

lm.wfit (*lm.fit*), 459

load, 605

locales, 578

locator, 83, 96, **104**

log, 240, **283**

log10 (*log*), 283

log1p (*log*), 283

log2 (*log*), 283

logb (*log*), 283

Logistic, **348**

logLik, 389, 390, **466**

logLik.glm, 467

logLik.gls, 467

logLik.lm, 467, 468

logLik.lme, 467

logLik.glm, **468**

logLik.lm, **469**

loglin, 111, 112, **471**, 618

Lognormal, **350**

longley, **628**

lowess, 125, 236

ls.diag, 476, 478

ls.print, 474, 478

ls.diag, **474**

ls.print, **476**

lsfit, 296, 297, 474, 476, **477**

make.link, 508

make.link, **479**

makepredictcall, **480**

makepredictcall.poly (*poly*),  
289

manova, **482**, 537

Math.difftime (*difftime*),  
563

Math.POSIXlt  
(*DateTimeClasses*),  
558

Math.POSIXt  
(*DateTimeClasses*),  
558

matlines (*matplot*), 106

matmult, **285**

matplot, **106**, 623

matpoints (*matplot*), 106

- matrix, 44, 107, 285, **286**
- max, 244, 299
- max (*Extremes*), 270
- mean, 244, 258, 375
- mean.POSIXct
  - (*DateTimeClasses*), 558
- mean.POSIXlt
  - (*DateTimeClasses*), 558
- Methods, 299
- methods, 241, 447, 464, 530
- min, 244, 299
- min (*Extremes*), 270
- Mod, 242
- mode, 242
- model.extract, 488
- model.frame, 426, 437, 480, 483, 487, 488, 497
- model.frame.default, 480
- model.matrix, 456, 486, 539, 549
- model.offset, 497
- model.tables, 399, 400, 420, 515, 519, 522, 529, 544
- model.tables.aovlist, 421
- model.extract, **483**
- model.frame, **485**
- model.matrix, **487**
- model.offset
  - (*model.extract*), 483
- model.response
  - (*model.extract*), 483
- model.tables, **489**
- model.weights
  - (*model.extract*), 483
- months (*weekdays*), 587
- morley, **630**
- mosaicplot, 11, 63, **110**, 164, 618
- mosaicplot.default, 111
- mosaicplot.formula, 111
- mtcars, **632**
- mtext, 72, **114**, 126, 171, 173, 226, 229
- Multinomial, **352**
- mvfft (*fft*), 272
- n2mfrow, **117**
- NA, 24, 28, 35, 69, 194, 198, 235, 237, 246, 299, 561
- na.action, 461
- na.contiguous, 585
- na.exclude, 461, 462
- na.fail, 426, 440, 455, 485, 585
- na.omit, 426, 440, 455, 485, 585
- na.omit.ts (*ts-methods*), 585
- names, 271
- NaN, 28, 246
- napredict, 435
- napredict (*naresid*), 492
- naprint, **491**
- naresid, **492**, 520
- nchar, 215
- nclass.FD, 78
- nclass.scott, 78
- nclass.Sturges, 78, 79
- NegBinomial, **354**
- nextn, 261, 273, **288**
- nhtemp, **633**
- nlm, 265, **493**, 502, 506, 546
- nls, 265, 495
- Normal, **357**
- ns, 480
- numeric, 299
- offset, 456, 483, **497**, 541
- Ops, 563
- Ops.difftime (*difftime*), 563
- Ops.POSIXt
  - (*DateTimeClasses*), 558
- Ops.ts (*ts*), 582
- optim, 265, 410, 411, 495, **498**
- optimise (*optimize*), 505

- optimize, 495, 502, **505**, 546
- options, 51, 56, 83, 104, 133, 415, 440, 445, 455, 485
- OrchardSprays, **634**
- order, 307
- pairs, 44, **118**, 125, 147
- palette, 35, 37, 61, **121**, 124, 132, 170
- Palettes, **123**
- panel.smooth, 44, 161
- panel.smooth, **125**
- par, 7, 8, 13, 17, 18, 21, 23, 39, 57, 68, 85, 88, 94, 102, 103, 105, 107, 111, 114, 115, 117, 125, **126**, 139, 145, 146, 149, 153, 156, 157, 161, 166, 168, 179, 181, 182, 198, 199, 201, 205, 207, 210, 211, 218, 224, 226, 228, 230
- pbeta (*Beta*), 322
- pbinom (*Binomial*), 324
- pbirthday (*birthday*), 326
- pcauchy (*Cauchy*), 328
- pchisq, 374, 376
- pchisq (*Chisquare*), 330
- pdf, 54, 55, **135**
- pentagamma (*Special*), 309
- periodicSpline, 312
- persp, **137**
- pexp (*Exponential*), 338
- pf (*FDist*), 340
- pgamma, 355
- pgamma (*GammaDist*), 342
- pgeom (*Geometric*), 344
- phones, **636**
- phyper (*Hypergeometric*), 346
- pictex, 55, 141
- pie, **143**
- PlantGrowth, **637**
- plnorm (*Lognormal*), 350
- plogis (*Logistic*), 348
- plot, 21, 69, 85, 98, 102, 103, 106, 107, 115, 126, **145**, 147, 148, 150, 153, 156, 164, 166, 168, 170, 178, 179, 218, 235
- plot.default, 15, 31, 39, 60, 65, 107, 133, 146, 147, 153, 156, 158, 164, 166, 168, 170, 211, 217, 236
- plot.density, 335
- plot.factor, 158, 164
- plot.formula, 146, 156
- plot.histogram, 76, 77
- plot.lm, 224
- plot.new, 168
- plot.ts, 583
- plot.window, 21, 58, 60, 65, 149, 150
- plot.xy, 103, 168, 179
- plot.data.frame, **147**
- plot.default, 131, **148**
- plot.density, **152**
- plot.design, **153**
- plot.factor, **156**
- plot.formula, **157**
- plot.function (*curve*), 46
- plot.histogram, **159**
- plot.lm, **161**
- plot.mlm (*plot.lm*), 161
- plot.new, 130
- plot.new (*frame*), 65
- plot.POSIXct (*axis.POSIXct*), 15
- plot.POSIXlt (*axis.POSIXct*), 15
- plot.table, **164**
- plot.ts, **166**
- plot.TukeyHSD (*TukeyHSD*), 543
- plot.window, **168**

- plot.xy, 170
- plotmath, 72, 97, 115, 141, 171, 226, 229
- pmax (*Extremes*), 270
- pmin (*Extremes*), 270
- pnbinom (*NegBinomial*), 354
- png, 54–56, 175
- pnorm, 377
- pnorm (*Normal*), 357
- points, 31, 44, 69, 97, 103, 106, 107, 125, 126, 139, 146, 149, 161, 170, 178, 224, 236
- points.default, 170
- points.formula (*plot.formula*), 157
- Poisson, 359
- poisson (*family*), 432
- poly, 289, 480
- polygon, 33, 143, 181, 199, 207
- polym (*poly*), 289
- polyroot, 291, 546
- POSIXct (*DateTimeClasses*), 558
- POSIXlt (*DateTimeClasses*), 558
- POSIXt, 561
- POSIXt (*DateTimeClasses*), 558
- postscript, 49, 50, 53–56, 133, 135, 136, 142, 178, 184
- power, 432, 433, 508
- power.t.test, 508
- ppoints, 190, 195
- ppois (*Poisson*), 359
- precip, 638
- predict, 426, 458, 492, 512
- predict.glm, 224, 443
- predict.lm, 457, 458
- predict.glm, 509
- predict.lm, 511
- predict.mlm (*predict.lm*), 511
- predict.poly (*poly*), 289
- preplot, 191
- presidents, 639
- pressure, 640
- pretty, 13, 18, 192
- print, 392, 400, 567
- print.ts, 583
- print.anova (*anova*), 394
- print.aov (*aov*), 399
- print.aovlist (*aov*), 399
- print.AsIs (*AsIs*), 401
- print.density (*density*), 333
- print.difftime (*difftime*), 563
- print.dummy.coef (*dummy.coef*), 419
- print.family (*family*), 432
- print.formula (*formula*), 436
- print.glm (*glm*), 439
- print.infl (*influence.measures*), 449
- print.integrate (*integrate*), 278
- print.lm (*lm*), 455
- print.logLik (*logLik*), 466
- print.mtable (*alias*), 391
- print.packageIQR (*data*), 605
- print.POSIXct (*DateTimeClasses*), 558
- print.POSIXlt (*DateTimeClasses*), 558
- print.recordedplot (*recordPlot*), 197
- print.summary.aov (*summary.aov*), 528
- print.summary.aovlist (*summary.aov*), 528



- print.summary.glm  
    (*summary.glm*), 530
- print.summary.lm  
    (*summary.lm*), 533
- print.summary.manova  
    (*summary.manova*),  
    536
- print.tables.aov  
    (*model.tables*), 489
- print.terms (*terms*), 538
- print.ts, **567**
- print.TukeyHSD (*TukeyHSD*),  
    543
- prod, **293**
- profile, **513**
- profile.glm, 513
- profile.nls, 513
- proj, 400, 490, **514**
- ps.options, 233, 234
- ps.options (*postscript*), 184
- psignrank (*SignRank*), 372
- pt (*TDist*), 374
- ptukey (*Tukey*), 376
- punif (*Uniform*), 378
- pweibull (*Weibull*), 380
- pwilcox (*Wilcoxon*), 382
  
- qbeta (*Beta*), 322
- qbinom (*Binomial*), 324
- qbirthday (*birthday*), 326
- qcauchy (*Cauchy*), 328
- qchisq (*Chisquare*), 330
- qexp (*Exponential*), 338
- qf (*FDist*), 340
- qgamma (*GammaDist*), 342
- qgeom (*Geometric*), 344
- qhyper (*Hypergeometric*), 346
- qlnorm (*Lognormal*), 350
- qlogis (*Logistic*), 348
- qnbinom (*NegBinomial*), 354
- qnorm, 365, 377
- qnorm (*Normal*), 357
- qpois (*Poisson*), 359
- qqline (*qqnorm*), 194
- qqnorm, 190, **194**
- qqplot, 190
- qqplot (*qqnorm*), 194
- qr, 248, 254, 268, 281, 282,  
    **294**, 297, 298, 316,  
    459, 460
- qr.Q, 296
- qr.qy, 298
- qr.R, 296
- qr.solve, 305
- qr.X, 296
- QR.Auxiliaries, **297**
- qr.Q (*QR.Auxiliaries*), 297
- qr.R (*QR.Auxiliaries*), 297
- qr.X (*QR.Auxiliaries*), 297
- qsignrank (*SignRank*), 372
- qt (*TDist*), 374
- qtukey, 544
- qtukey (*Tukey*), 376
- quakes, **641**
- quantile, 28
- quarters (*weekdays*), 587
- quartz, **196**
- quasi (*family*), 432
- quasibinomial (*family*), 432
- quasipoisson (*family*), 432
- qunif (*Uniform*), 378
- quote, 173
- qweibull (*Weibull*), 380
- qwilcox (*Wilcoxon*), 382
  
- R CMD BATCH, 55
- r2dtable, **361**
- rainbow, 37, 68, 80, 85, 121,  
    133, 201
- rainbow (*Palettes*), 123
- Random, **363**
- Random.user, 365, **368**
- randu, **642**
- range, 44, 271, **299**
- range.default, 299
- rank, 307
- rbeta (*Beta*), 322
- rbinom, 353

- rbinom (*Binomial*), 324
- rcauchy (*Cauchy*), 328
- rchisq (*Chisquare*), 330
- read.table, 605
- recordPlot, 197
- rect, 23, 159, 182, **198**
- relevel, **517**
- rep, 235, 237, **568**
- replayPlot (*recordPlot*), 197
- replications, 490, **518**
- resid, 492
- resid (*residuals*), 520
- residuals, 224, 394, 407, 423, 435, 442, 443, 448, 458, 465, **520**, 551
- residuals.glm, 465, 531
- residuals.glm (*glm.summaries*), 447
- residuals.lm (*lm.summaries*), 464
- rexp (*Exponential*), 338
- rf (*FDist*), 340
- rgamma (*GammaDist*), 342
- rgb, 35, 37, 68, 80, 124, 133, **201**
- rgeom (*Geometric*), 344
- rhypcr (*Hypergeometric*), 346
- rivers, **643**
- rlnorm (*Lognormal*), 350
- rlogis (*Logistic*), 348
- rmultinom (*Multinomial*), 352
- rnbinom (*NegBinomial*), 354
- RNG (*Random*), 363
- RNGkind, 368
- RNGkind (*Random*), 363
- RNGversion (*Random*), 363
- rnrm, 367, 379
- rnrm (*Normal*), 357
- Round, **301**
- round, 563
- round (*Round*), 301
- round.POSIXt, 560
- round.difftime (*difftime*), 563
- round.POSIXt, **571**
- rowMeans (*colSums*), 258
- rowsum, 259
- rowSums (*colSums*), 258
- rpois (*Poisson*), 359
- Rprof, 513
- rsignrank (*SignRank*), 372
- rstandard (*influence.measures*), 449
- rstudent, 464
- rstudent (*influence.measures*), 449
- rt (*TDist*), 374
- rug, 92, **202**, 223
- runif, 358, 367
- runif (*Uniform*), 378
- rweibull (*Weibull*), 380
- rwilcox (*Wilcoxon*), 382
- SafePrediction, 510, 512
- SafePrediction (*makepredictcall*), 480
- sample, **370**
- save, 607
- scale, 480
- screen, **204**
- sd, 375
- se.contrast, 490
- se.contrast.aovlist, 421
- se.contrast, **521**
- segments, 7, 9, 182, 199, **207**
- seq, 569, 572
- seq.POSIXt, 556, 560, 565
- seq.POSIXt, **572**
- sequence, 569
- set.seed (*Random*), 363
- sign, **303**
- signif (*Round*), 301
- SignRank, **372**

- sin, 240, 277
- sin (*Trig*), 318
- sinh (*Hyperbolic*), 277
- sleep, 644
- smooth.spline, 312
- solve, 248, 256, **304**
- solve.qr, 296, 305
- solve.qr (*qr*), 294
- sort, **306**
- source, 605
- Special, 240, 247, **309**
- spline, 244
- spline (*splinefun*), 311
- splinefun, 47, 244, **311**
- split.screen, 133
- split.screen, 129
- split.screen (*screen*), 204
- sqrt, 247, 284, 310
- sqrt (*abs*), 240
- stack.loss (*stackloss*), 645
- stack.x (*stackloss*), 645
- stackloss, **645**
- stars, **209**, 221
- start, **574**, 581, 583, 586
- stat.anova, 395
- stat.anova, **523**
- state, **647**, 658
- stem, 79, 160
- step, 388, 428, 429, **525**
- stepAIC, 527
- str.logLik (*logLik*), 466
- str.POSIXt
  - (*DateTimeClasses*), 558
- strftime (*strftime*), 575
- strheight (*strwidth*), 215
- stripchart, 26, 147, **213**
- strftime, 15, 554, 555, 560, 565, **575**
- strwidth, 97, 128, **215**
- substitute, 96, 173
- sum, 258, 293, **314**
- summary, 394, 400, 441, 443, 462, 529, 531, 534
- summary.aov, 400
- summary.glm, 441, 443, 448
- summary.lm, 457, 462, 465, 474
- summary.manova, 482
- summary.aov, **528**
- summary.aovlist
  - (*summary.aov*), 528
- Summary.difftime (*difftime*), 563
- summary.glm, **530**
- summary.infl
  - (*influence.measures*), 449
- summary.lm, **533**
- summary.manova, **536**
- summary.mlm (*summary.lm*), 533
- Summary.POSIXct
  - (*DateTimeClasses*), 558
- summary.POSIXct
  - (*DateTimeClasses*), 558
- Summary.POSIXlt
  - (*DateTimeClasses*), 558
- summary.POSIXlt
  - (*DateTimeClasses*), 558
- sunflowerplot, **217**, 221
- sunspot.month, 649
- sunspots, **649**
- survreg, 224
- svd, 254, 268, 282, 296, **315**
- swiss, **650**
- symbols, **220**
- symnum, 530, 533
- Syntax, 247
- Sys.time, 560
- Sys.time, **579**
- Sys.timezone (*Sys.time*), 579

- system.time, 581
- table, 164, 317, 473
- tabulate, 317
- tan, 277
- tan (*Trig*), 318
- tanh (*Hyperbolic*), 277
- TDist, 374
- termplot, 162, 223
- terms, 153, 438, 442, 457, 488, 538, 540–542, 549
- terms.formula, 538, 541
- terms.object, 538, 540
- terms.formula, 539
- terms.object, 541
- terrain.colors, 84, 85, 121
- terrain.colors (*Palettes*), 123
- tetragamma (*Special*), 309
- text, 39, 72, 75, 90, 98, 114, 115, 126, 127, 166, 167, 171, 173, 187, 215, 225, 229
- time, 235, 574, 580, 583, 586, 590
- Titanic, 652
- title, 21, 31, 39, 58, 60, 72, 107, 115, 126, 137, 145, 153, 161, 173, 226, 228
- ToothGrowth, 654
- topo.colors, 37, 84, 85
- topo.colors (*Palettes*), 123
- trees, 655
- Trig, 284, 318
- trigamma (*Special*), 309
- TRUE, 167
- truehist, 79
- trunc (*Round*), 301
- trunc.POSIXt, 560
- trunc.POSIXt (*round.POSIXt*), 571
- ts, 167, 561, 567, 574, 580, 581, 582, 585, 586, 589, 590
- ts-methods, 585
- tsp, 574, 580, 581, 583, 586, 590
- tsp<- (*tsp*), 586
- Tukey, 376
- TukeyHSD, 400, 490, 529, 543
- typeof, 246
- UCBAdmissions, 656
- ucv, 321
- Uniform, 378
- uniroot, 291, 495, 506, 545
- units, 230
- update, 547
- update.formula, 526, 547
- update.formula, 549
- USArrests, 658
- USJudgeRatings, 659
- USPersonalExpenditure, 660
- uspop, 661
- VADeaths, 662
- variable.names  
(*case/variable.names*), 405
- vcov, 550
- volcano, 664
- warning, 246
- warpbreaks, 665
- weekdays, 587
- weekdays.POSIXt, 560
- Weibull, 380
- weighted.residuals, 465
- weighted.residuals, 551
- weights, 551
- weights (*lm.summaries*), 464
- weights.glm (*glm*), 439
- which.min, 271
- width.SJ, 321
- Wilcoxon, 382
- window, 581, 583, 589
- women, 667

X11, 55, 56, 175, 176  
X11 (*x11*), 231  
x11, 66, 71, 133, **231**  
xfig, 55, **233**  
xinch (*units*), 230  
xtabs, 665  
xy.coords, 33, 96, 97, 102, 148,  
150, 168, 170, 179,  
182, 217, 220, 225,  
238, 243, 311  
xy.coords, **235**  
xyinch (*units*), 230  
xyz.coords, **237**  
  
yinch (*units*), 230  
  
zapsmall (*Round*), 301





